

9. Gaussian Mixture Models


**EECE454 Introduction to
Machine Learning Systems**

2023 Fall, Jaeho Lee

Recap: Clustering by K-means

- **K-means.** Each cluster is represented by the centroid.
 - A datum belongs to the cluster with nearest centroid.

- **Limitations.** Plenty, e.g., cannot handle...
 - overlapping clusters
 - “wider” clusters
 - Example. Non-local residents in Pohang
 - POSCO or POSTECH?




needs a probabilistic approach!

Mixture Models

Mixture models

- **Idea.** Take a **generative approach**, and fit parameters!
 - Example. the previous POSCO vs POSTECH.
 - We draw $Y \in \{0,1\} \sim \text{Bern}(p)$. (0: POSCO, 1: POSTECH)
 - Model the conditional distribution:
 - If $Y = 0$, draw X from $\mathcal{N}(\mu_0, \sigma_0^2)$
 - If $Y = 1$, draw X from $\mathcal{N}(\mu_1, \sigma_1^2)$
 - Allows overlap & can account for wideness.



Mixture models

- **Perk.** If you have “learned” a nice probabilistic model from data, you can not only cluster, but also **generate a new data**.

(Note: Example below requires additional text conditioning...)

a nendoroid
of a cute boy



a nendoroid
of a cute girl



a penguin



a potted
cactus plant



a 3D model
of a fox



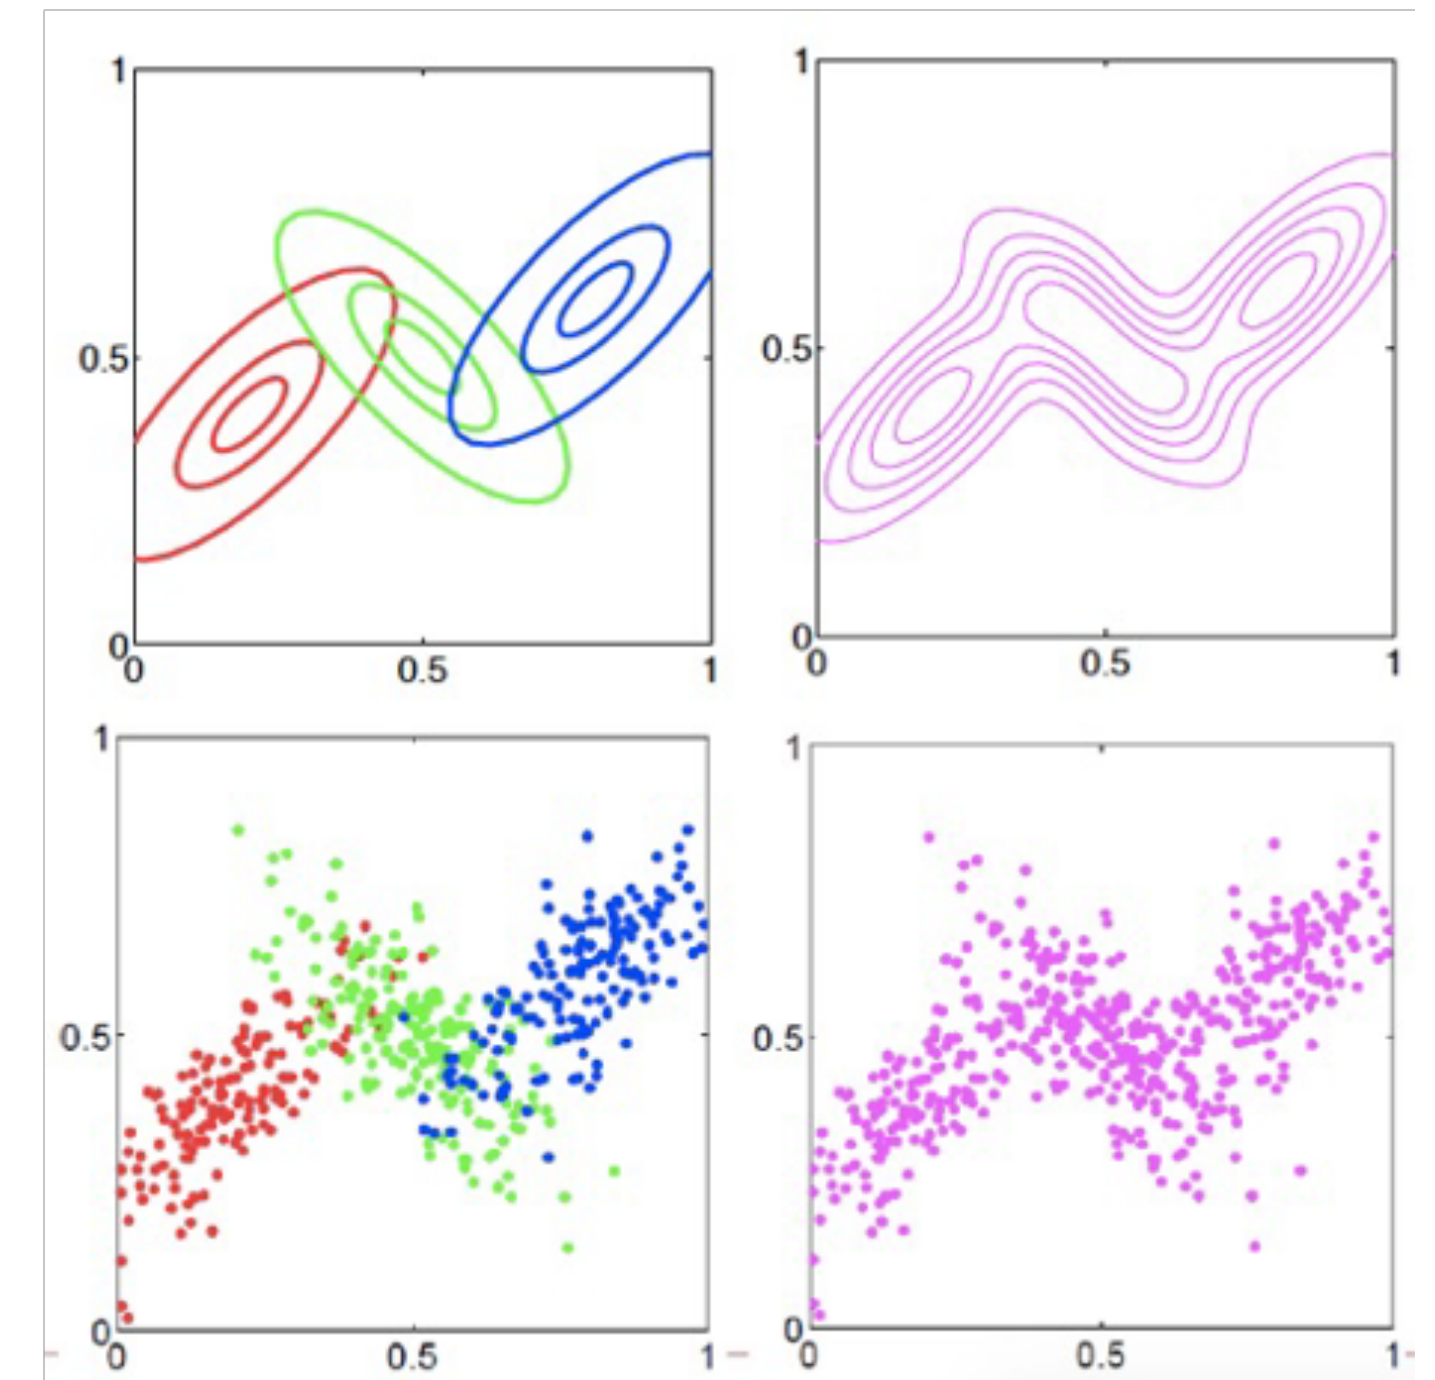

a 3D model
of a soldier



(finite) Mixture models

- More generally we model the data-generating pdf with

$$p(\mathbf{x}) = \sum_{k=1}^K \pi_k \cdot p_k(\mathbf{x}), \quad \pi_k \in [0,1], \sum \pi_k = 1.$$




Gaussian mixture models

- Each base distribution is a Gaussian distribution:

$$p(\mathbf{x} | \theta) = \sum_{k=1}^K \pi_k \cdot \mathcal{N}(\mathbf{x} | \mu_k, \Sigma_k),$$

where $\theta = (\mu_1, \Sigma_1, \dots, \mu_K, \Sigma_K, \pi_1, \dots, \pi_K)$ is the total parameter set.



$$p(x \mid \boldsymbol{\theta}) = 0.5\mathcal{N}(x \mid -2, \tfrac{1}{2}) + 0.2\mathcal{N}(x \mid 1, 2) + 0.3\mathcal{N}(x \mid 4, 1)$$

Gaussian mixture models

- Each base distribution is a Gaussian distribution:

$$p(\mathbf{x} | \theta) = \sum_{k=1}^K \pi_k \cdot \mathcal{N}(\mathbf{x} | \mu_k, \Sigma_k),$$

where $\theta = (\mu_1, \Sigma_1, \dots, \mu_K, \Sigma_K, \pi_1, \dots, \pi_K)$ is the total parameter set.

- **Question.** How do we fit the parameters, given $\{\mathbf{x}_1, \dots, \mathbf{x}_n\}$?
 - **Challenge.** We do not know the true labels!

Maximum Likelihood

- Similar to what we learned in naïve Bayes, what we want to try is the **maximum likelihood**.

$$\begin{aligned} p(\mathbf{x}_{1:n} | \theta) &= \prod_{i=1}^n p(\mathbf{x}_i | \theta) \\ &= \prod_{i=1}^n \sum_{k=1}^K \pi_k \cdot \mathcal{N}(\mathbf{x}_i | \mu_k, \Sigma_k) \end{aligned}$$

⇒ maximize this quantity by tuning $\theta = \{\mu_k, \Sigma_k, \pi_k \mid k \in [K]\}$

Maximum **Log**-Likelihood

- We do the usual log trick to make everything summation...

$$\mathcal{L} := \log p(\mathbf{x}_{1:n} | \theta) = \sum_{i=1}^n \log \left(\sum_{k=1}^K \pi_k \cdot \mathcal{N}(\mathbf{x}_i | \mu_k, \Sigma_k) \right)$$

- Normally, you would try to find the optimum by locating the critical point (i.e., gradient = 0)
 - Give it a try! (let me know if you succeed)

Expectation-Maximization

- **Idea.** Fix some variables and optimize others.
Fix the optimized variables, and optimize the previously fixed.
Repeat ...
- Generally, we call it **expectation-maximization (EM)** algorithm.
- Similar to what we did in K-means!

Algorithm 1 *k*-means algorithm

- 1: Specify the number k of clusters to assign.
 - 2: Randomly initialize k centroids.
 - 3: **repeat**
 - 4: **expectation:** Assign each point to its closest centroid.
 - 5: **maximization:** Compute the new centroid (mean) of each cluster.
 - 6: **until** The centroid positions do not change.
-

Expectation-Maximization

- Recall that, in hard K-means...
 - Randomly initialize **centroids** $\{\mu_k\}$.
 - Fix the **centroids** $\{\mu_k\}$ and optimize the **assignment** $\{r_{ik}\}$.
 - Optimal, if **nearest neighbor**.
 - Fix the **assignment** $\{r_{ik}\}$ and optimize the **centroid** $\{\mu_k\}$.
 - Optimal, if **mean of the assigned data**.
- Repeat.

Expectation-Maximization

- Similarly, what we want to do is...
 - Randomly initialize **parameters** $\theta = \{\mu_k, \Sigma_k, \pi_k\}$.
 - Non-binary, as in soft K-means
↓
 - Fix the **parameters** θ and optimized the **responsibility** $\{r_{ik}\}$.
 - Optimal, if?
 - Fixed the **responsibility** $\{r_{ik}\}$ and optimized the **parameters** θ .
 - Optimal, if?
- Let's think about the optimal conditions...

Recall: Multivariate Gaussian

- Multivariate Gaussians:

$$\mathcal{N}(\mathbf{x} | \mu, \Sigma) = \frac{1}{\sqrt{(2\pi)^d |\Sigma|}} \cdot \exp \left(-\frac{1}{2}(\mathbf{x} - \mu)^\top \Sigma^{-1}(\mathbf{x} - \mu) \right)$$

- Take log, you get:

$$\log \mathcal{N}(\mathbf{x} | \mu, \Sigma) = -\frac{1}{2} \cdot \left(d \log(2\pi) + \log |\Sigma| + (\mathbf{x} - \mu)^\top \Sigma^{-1}(\mathbf{x} - \mu) \right)$$

Recall: Responsibilities

- **Soft K-means.** The softmax value

$$r_{ik} = \frac{\exp(-\beta \|\mathbf{x}_i - \mu_k\|_2^2)}{\sum_j \exp(-\beta \|\mathbf{x}_i - \mu_j\|_2^2)}$$

- **GMM.** We use

$$r_{ik} = \frac{\pi_k \mathcal{N}(\mathbf{x} | \mu_k, \Sigma_k)}{\sum_j \pi_j \mathcal{N}(\mathbf{x} | \mu_j, \Sigma_j)}$$

Recall: Responsibilities

- **Soft K-means.** The softmax value

$$r_{ik} = \frac{\exp(-\beta \|\mathbf{x}_i - \mu_k\|_2^2)}{\sum_j \exp(-\beta \|\mathbf{x}_i - \mu_j\|_2^2)}$$

- **GMM.** We use

$$r_{ik} = \frac{\pi_k \mathcal{N}(\mathbf{x} | \mu_k, \Sigma_k)}{\sum_j \pi_j \mathcal{N}(\mathbf{x} | \mu_j, \Sigma_j)}$$

$$p(y = k | \mathbf{x}) = \frac{p(\mathbf{x}, y = k)}{p(\mathbf{x})}$$

Recall: Responsibilities

- **Soft K-means.** The softmax value

$$r_{ik} = \frac{\exp(-\beta \|\mathbf{x}_i - \mu_k\|_2^2)}{\sum_j \exp(-\beta \|\mathbf{x}_i - \mu_j\|_2^2)}$$

- **GMM.** We use

$$r_{ik} = \frac{\pi_k \mathcal{N}(\mathbf{x} \mid \mu_k, \Sigma_k)}{\sum_j \pi_j \mathcal{N}(\mathbf{x} \mid \mu_j, \Sigma_j)}$$

Note. If $\pi_k = 1/K$, $\Sigma_k = \mathbf{I}/\beta$, then this is identical to soft K-means.

Optimality Condition: Mean

- Recall that

$$\mathcal{L} := \log p(\mathbf{x}_{1:n} | \theta) = \sum_{i=1}^n \log \left(\sum_{k=1}^K \pi_k \cdot \mathcal{N}(\mathbf{x}_i | \mu_k, \Sigma_k) \right)$$

- Partial derivative w.r.t. μ_k is...

$$\nabla_{\mu_k} \mathcal{L} = \sum_{i=1}^n \frac{\pi_k \cdot \nabla_{\mu_k} \mathcal{N}(\mathbf{x} | \mu_k, \Sigma_k)}{\sum \pi_j p(\mathbf{x}_i | \mu_j, \Sigma_j)} = \sum_{i=1}^n r_{ik} (\mathbf{x}_i - \mu_k)^\top \Sigma_k^{-1} = \mathbf{0}$$
$$\Rightarrow \mu_k = \frac{\sum_i r_{ik} \mathbf{x}_i}{\sum_i r_{ik}}$$

Optimality Condition: Variance

- Do the similar thing, and you get

$$\Sigma_k = \frac{1}{n_k} \sum_{i=1}^n r_{ik} (\mathbf{x}_i - \mu_k)(\mathbf{x}_i - \mu_k)^\top$$

where we use the shorthand $n_k = \sum_{i=1}^n r_{ik}$.

see section 11.2.3 of the main textbook

Optimality Condition: Mixture Weights

- Do the similar thing, and you get

$$\pi_k = \frac{n_k}{n}$$

see section 11.2.4 of the main textbook;

this one is trickier as it's constrained—use Lagrange multipliers!

The full E-M

- Do the similar thing, and you get

1. Initialize $\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k, \pi_k$.
2. *E-step*: Evaluate responsibilities r_{nk} for every data point \mathbf{x}_n using current parameters $\pi_k, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k$:

$$r_{nk} = \frac{\pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{\sum_j \pi_j \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)} . \quad (11.53)$$


3. *M-step*: Reestimate parameters $\pi_k, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k$ using the current responsibilities r_{nk} (from E-step):

$$\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{n=1}^N r_{nk} \mathbf{x}_n , \quad (11.54)$$

$$\boldsymbol{\Sigma}_k = \frac{1}{N_k} \sum_{n=1}^N r_{nk} (\mathbf{x}_n - \boldsymbol{\mu}_k)(\mathbf{x}_n - \boldsymbol{\mu}_k)^\top , \quad (11.55)$$

$$\pi_k = \frac{N_k}{N} . \quad (11.56)$$

The full E-M



The full E-M



Cheers

- Next up. Trees, Random Forest, and Boosting