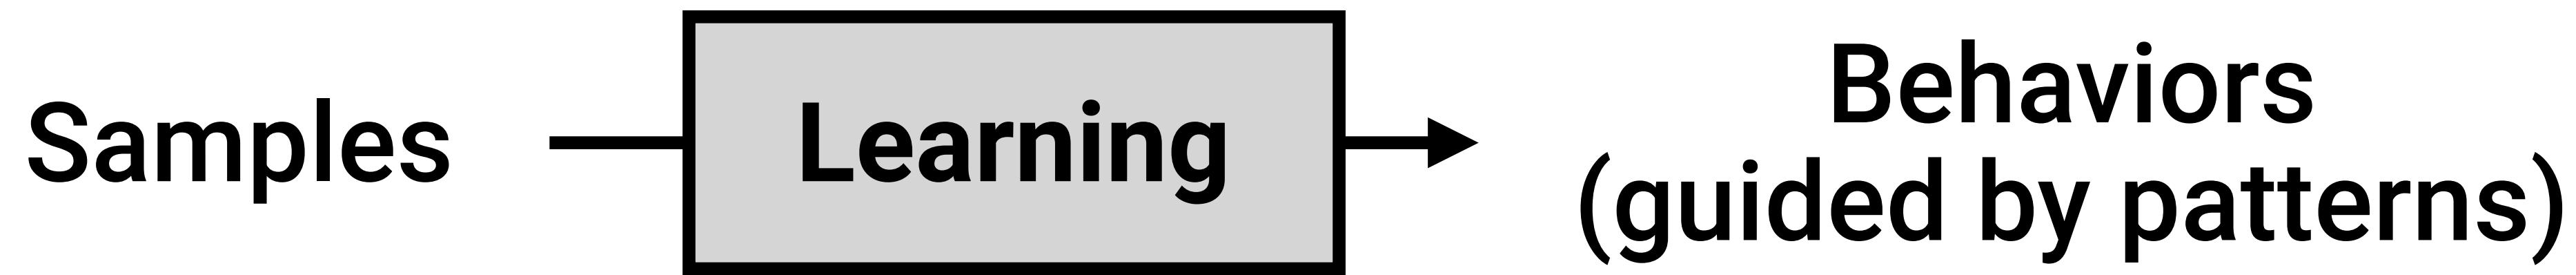


# **Elements of ML**

# Recap: What is learning?

- The process of extracting and utilizing **patterns** from the samples



# Recap: What is learning?

- **Today.** We formalize this concept:
  - What exactly is a pattern?
  - How can we program a machine to find one?
- In particular, we provide some unified but hand-wavy perspectives:
  - starting next week, we look at individual ML algorithms

# **Patterns**

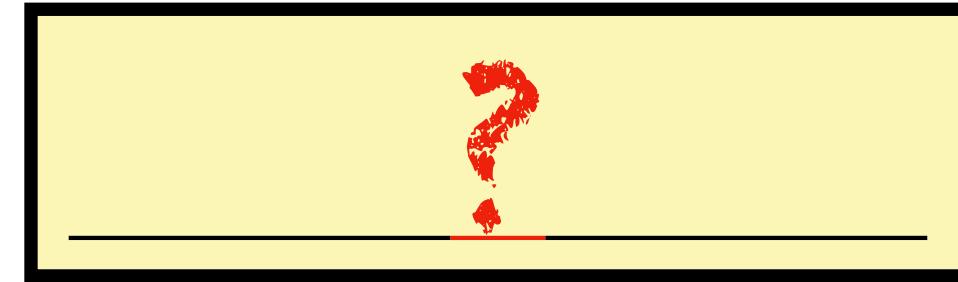
# Patterns

- Associations of distinct variables
- **Example.** “Green pixels” are associated with “another green pixel”



# Patterns

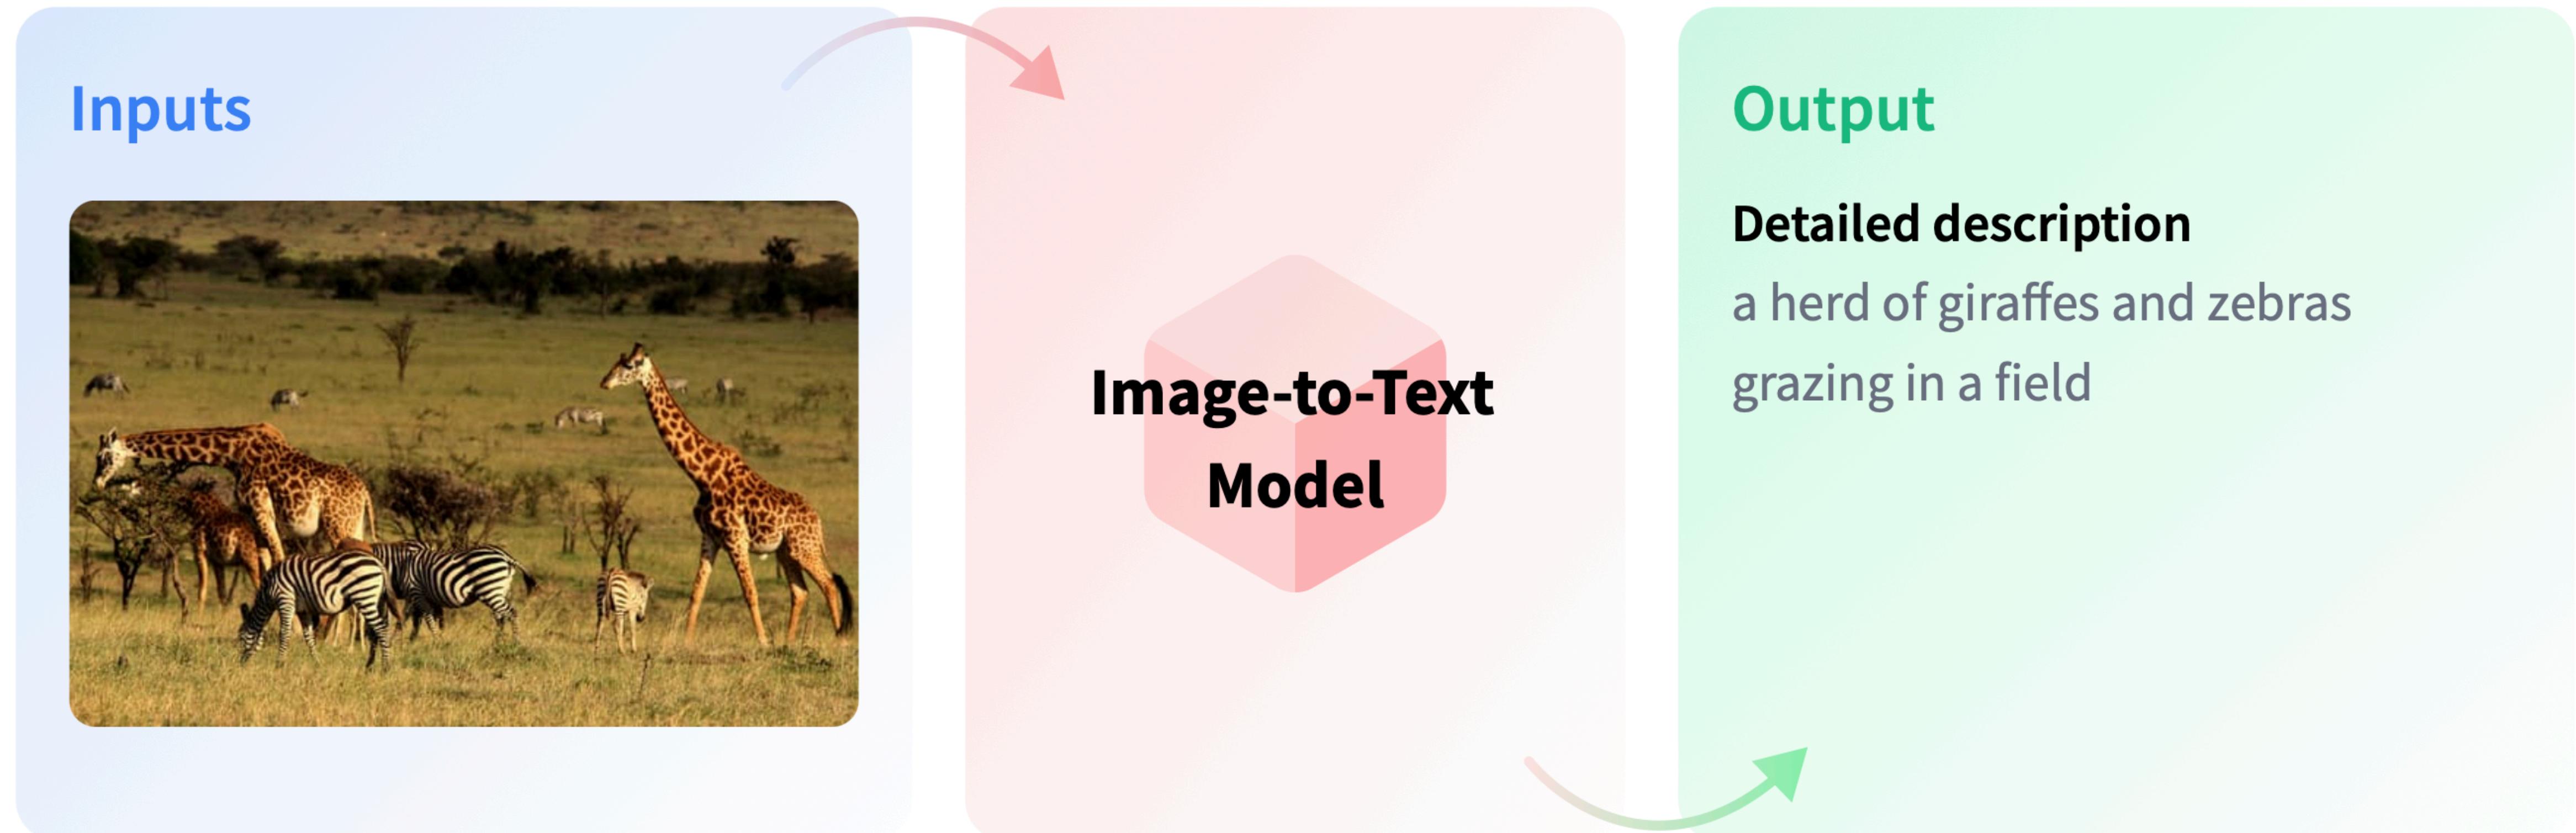
- Associations of distinct variables
- **Example.** The text “A dog is” is associated with the word “cute”

“A dog is  ”

(Next word prediction – GPT is trained this way)

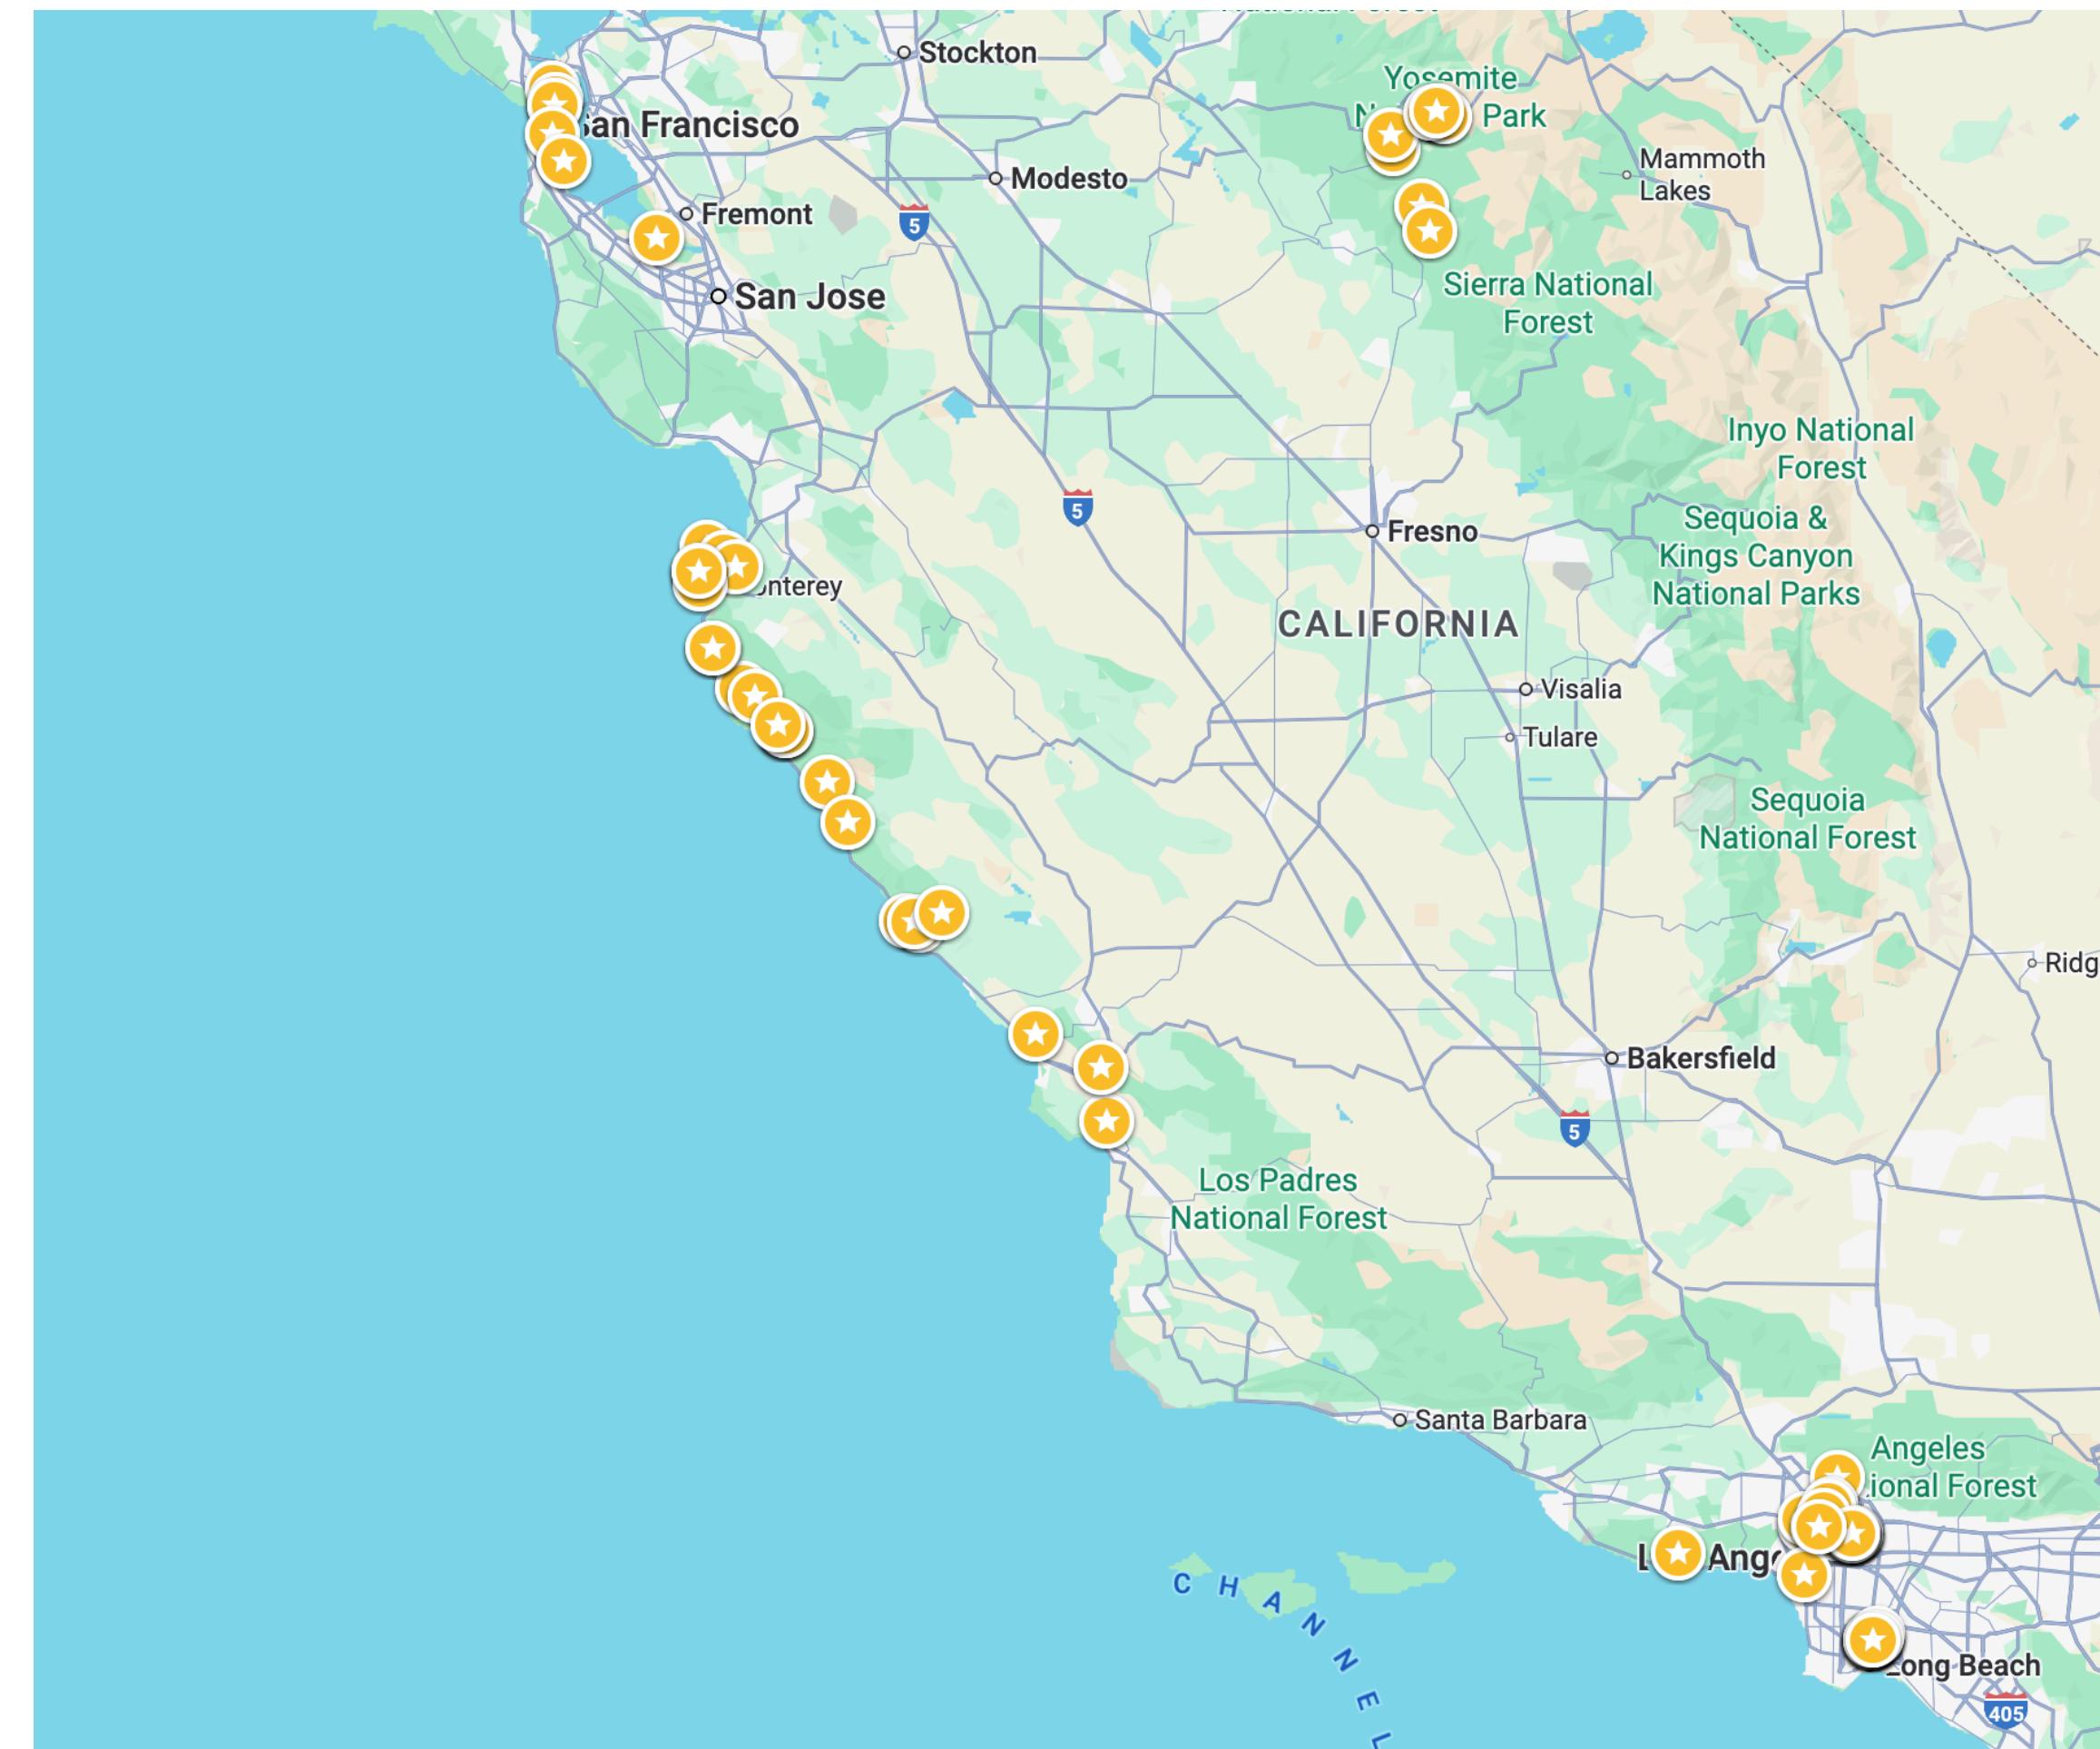
# Patterns

- The variables need not be of same modality
- **Example.** An image is associated with its textual description



# Patterns

- We can make associations with **imaginary variables**
- **Example.** “Locations” are associated with “Categories (imaginary)”



# Categories

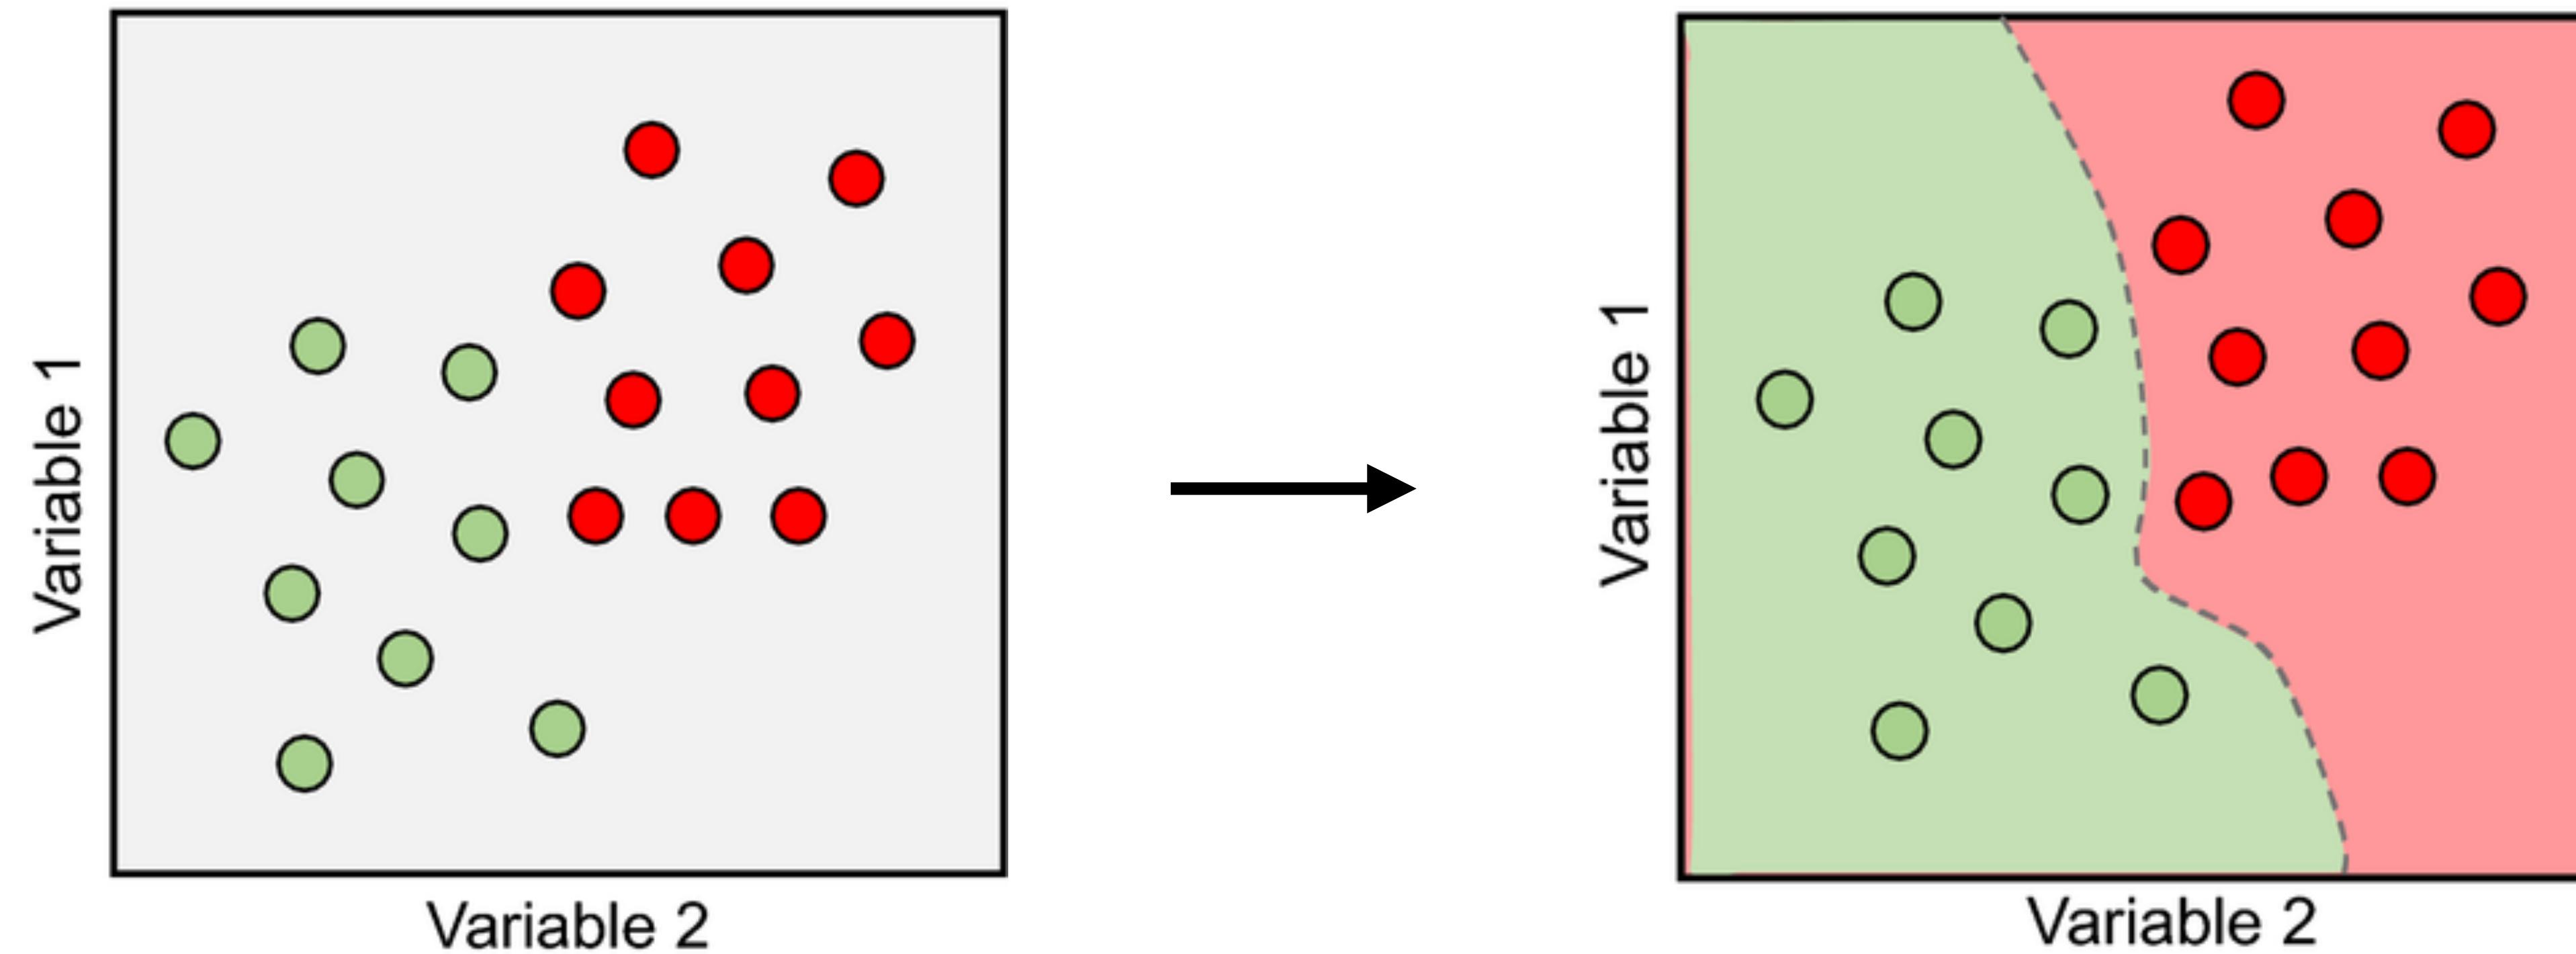
- Roughly, learning is about associating different random variables:  $X, Y$ 
  - Jointly distributed as some probability distribution  $P_{XY}(x, y)$
  - However,  $P_{XY}$  is not known to the learner
  - Instead, we have **training data**

# Categories

- Depending on the **type of data** available, learning can be categorized into:
  - Supervised learning
  - Unsupervised learning
  - Reinforcement learning
- Note: Of course, there are many other terminologies
  - semi-supervised, self-supervised, active, ...
- Note: In this course, we focus on supervised & unsupervised
  - Reinforcement learning as a special session

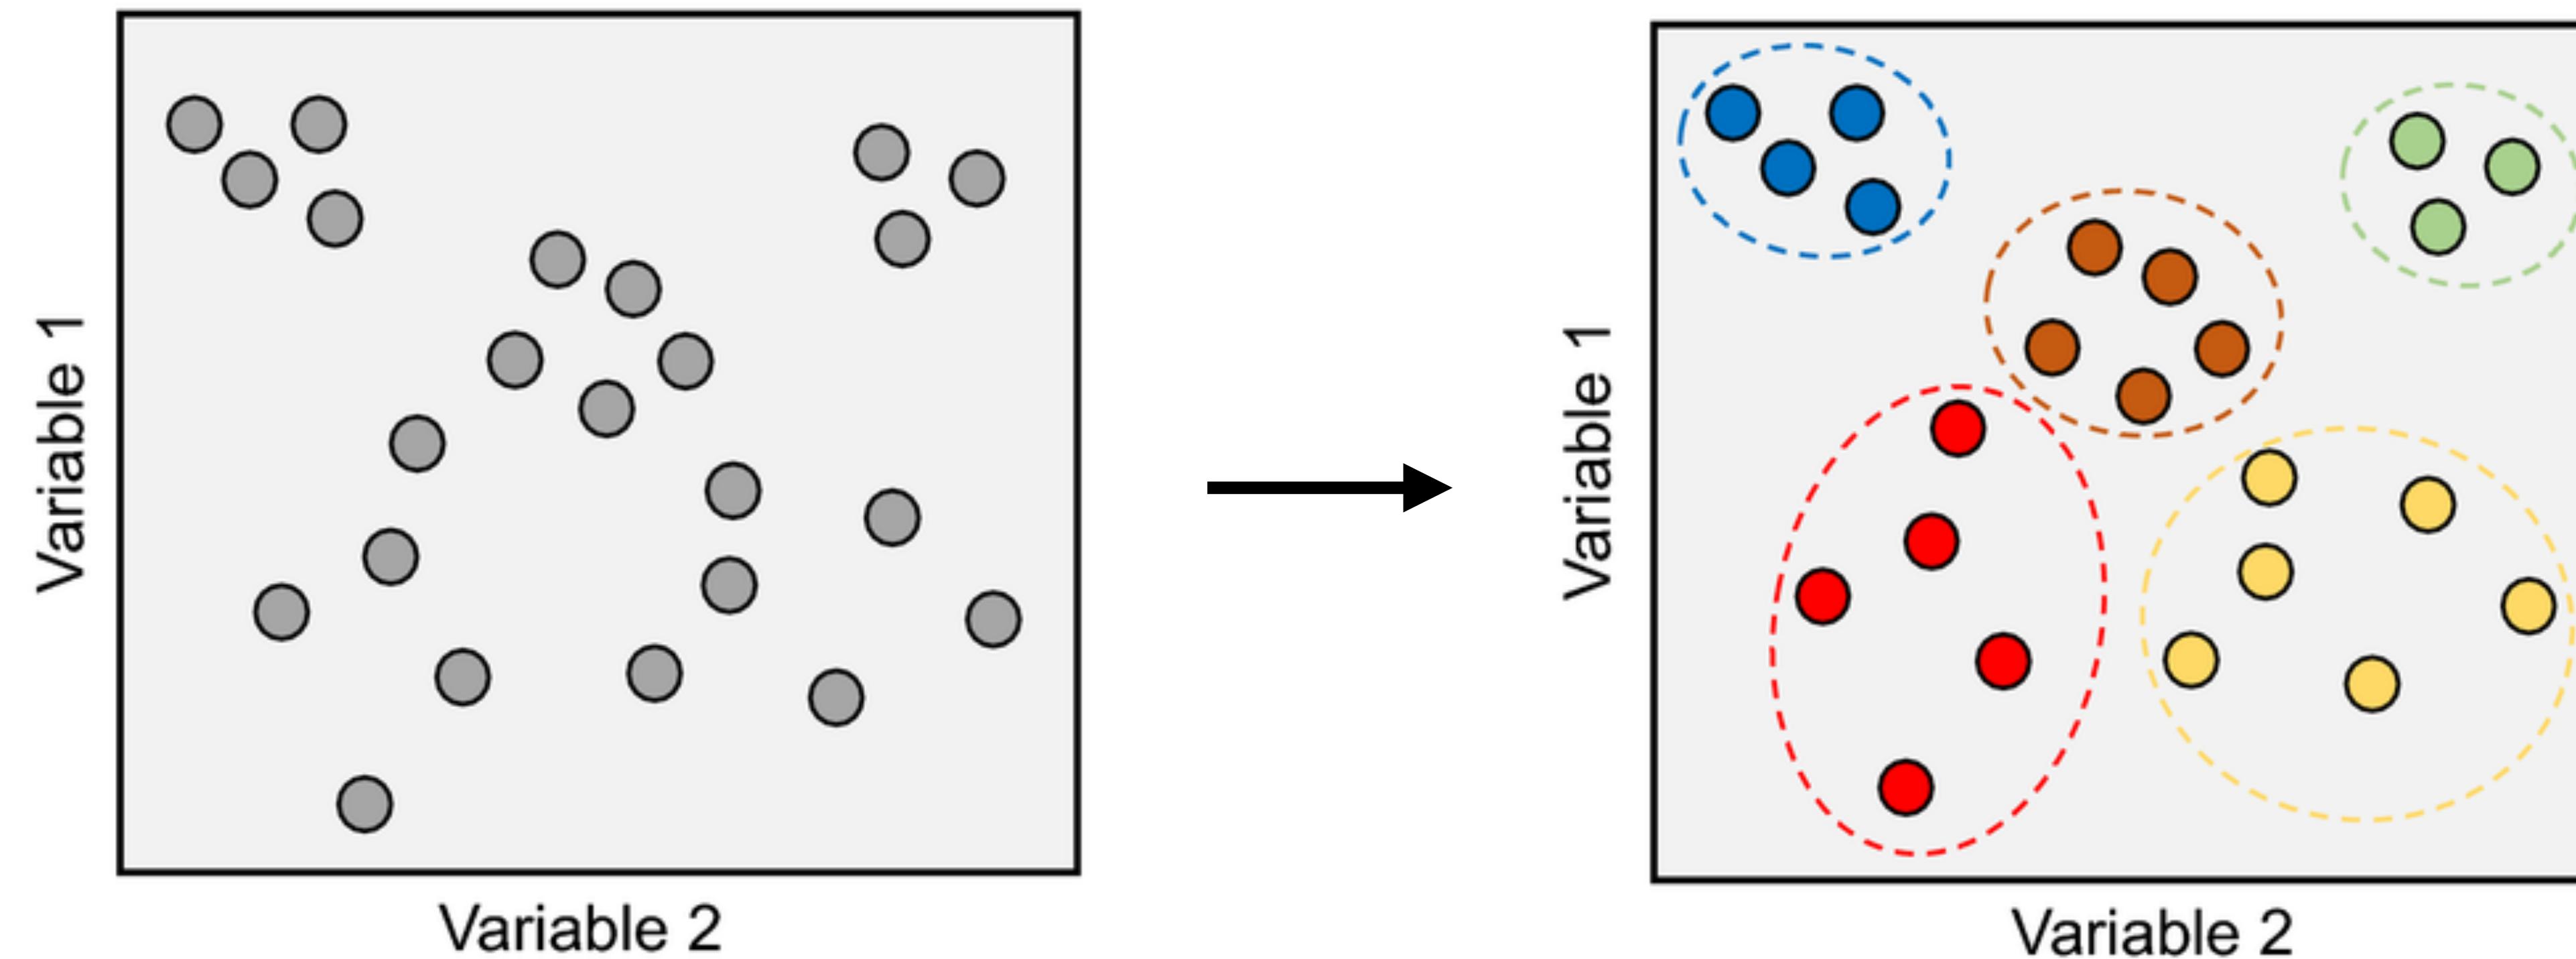
# Categories: Supervised Learning

- We have many **input-output pairs**  $D = \{(X_i, Y_i)\}_{i=1}^n, (X_i, Y_i) \sim P_{XY}$ 
  - Learn the **input-to-output mapping**  
(e.g., learning to predict the color of points)



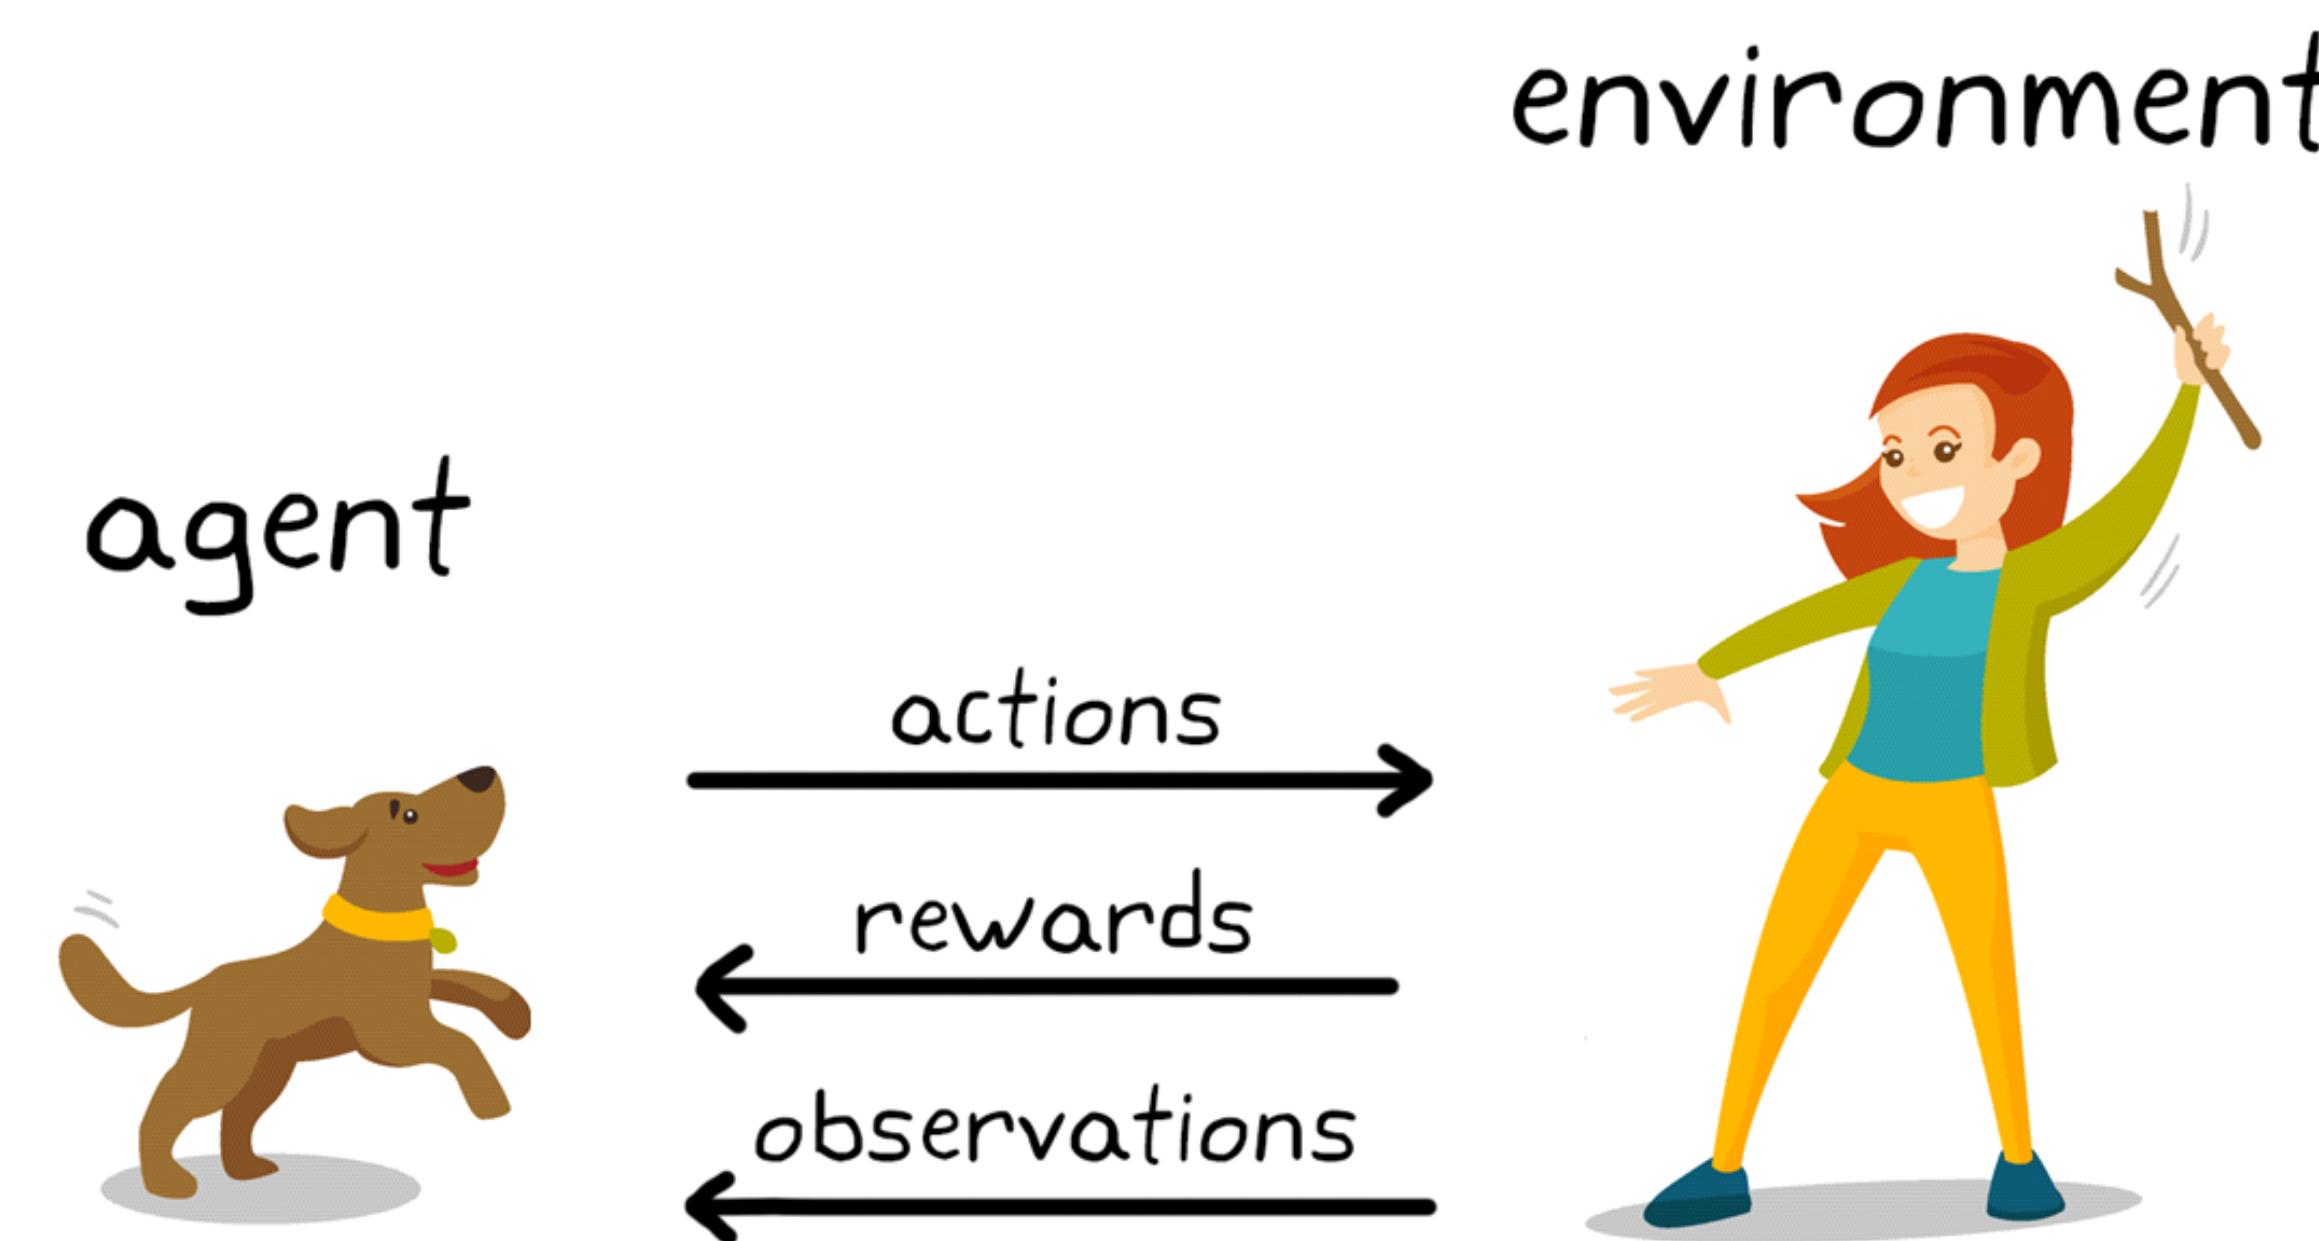
# Categories: Unsupervised Learning

- We have many **unlabeled input data**  $D = \{X_i\}_{i=1}^n, X_i \sim P_X$ 
  - Learn useful **structures of data** – or virtual labels (e.g., identifying clusters of data)
  - The structures may be useful for downstream tasks



# Categories: Reinforcement Learning

- We have an **environment** that we can interact with:
  - Can collect **sequence of interactions** with the environment
    - Actions, Rewards, States
  - Learn an **interaction policy (i.e., agent)** that maximizes the reward
    - Somewhat specialized; we'll discuss this as a special topic later



# **Supervised Learning**

# Supervised Learning

- For now, let's focus on the supervised learning scenario:
  - We have two random variables:  $X, Y$ 
    - Jointly distributed as some probability distribution  $P_{XY}(x, y)$
- Consider a simple **prediction** task, where
  - $X$  is easy to collect (called **features**)
    - e.g., natural images
  - $Y$  is costly to acquire (called **labels**)
    - e.g., human-written labels



# Supervised Learning

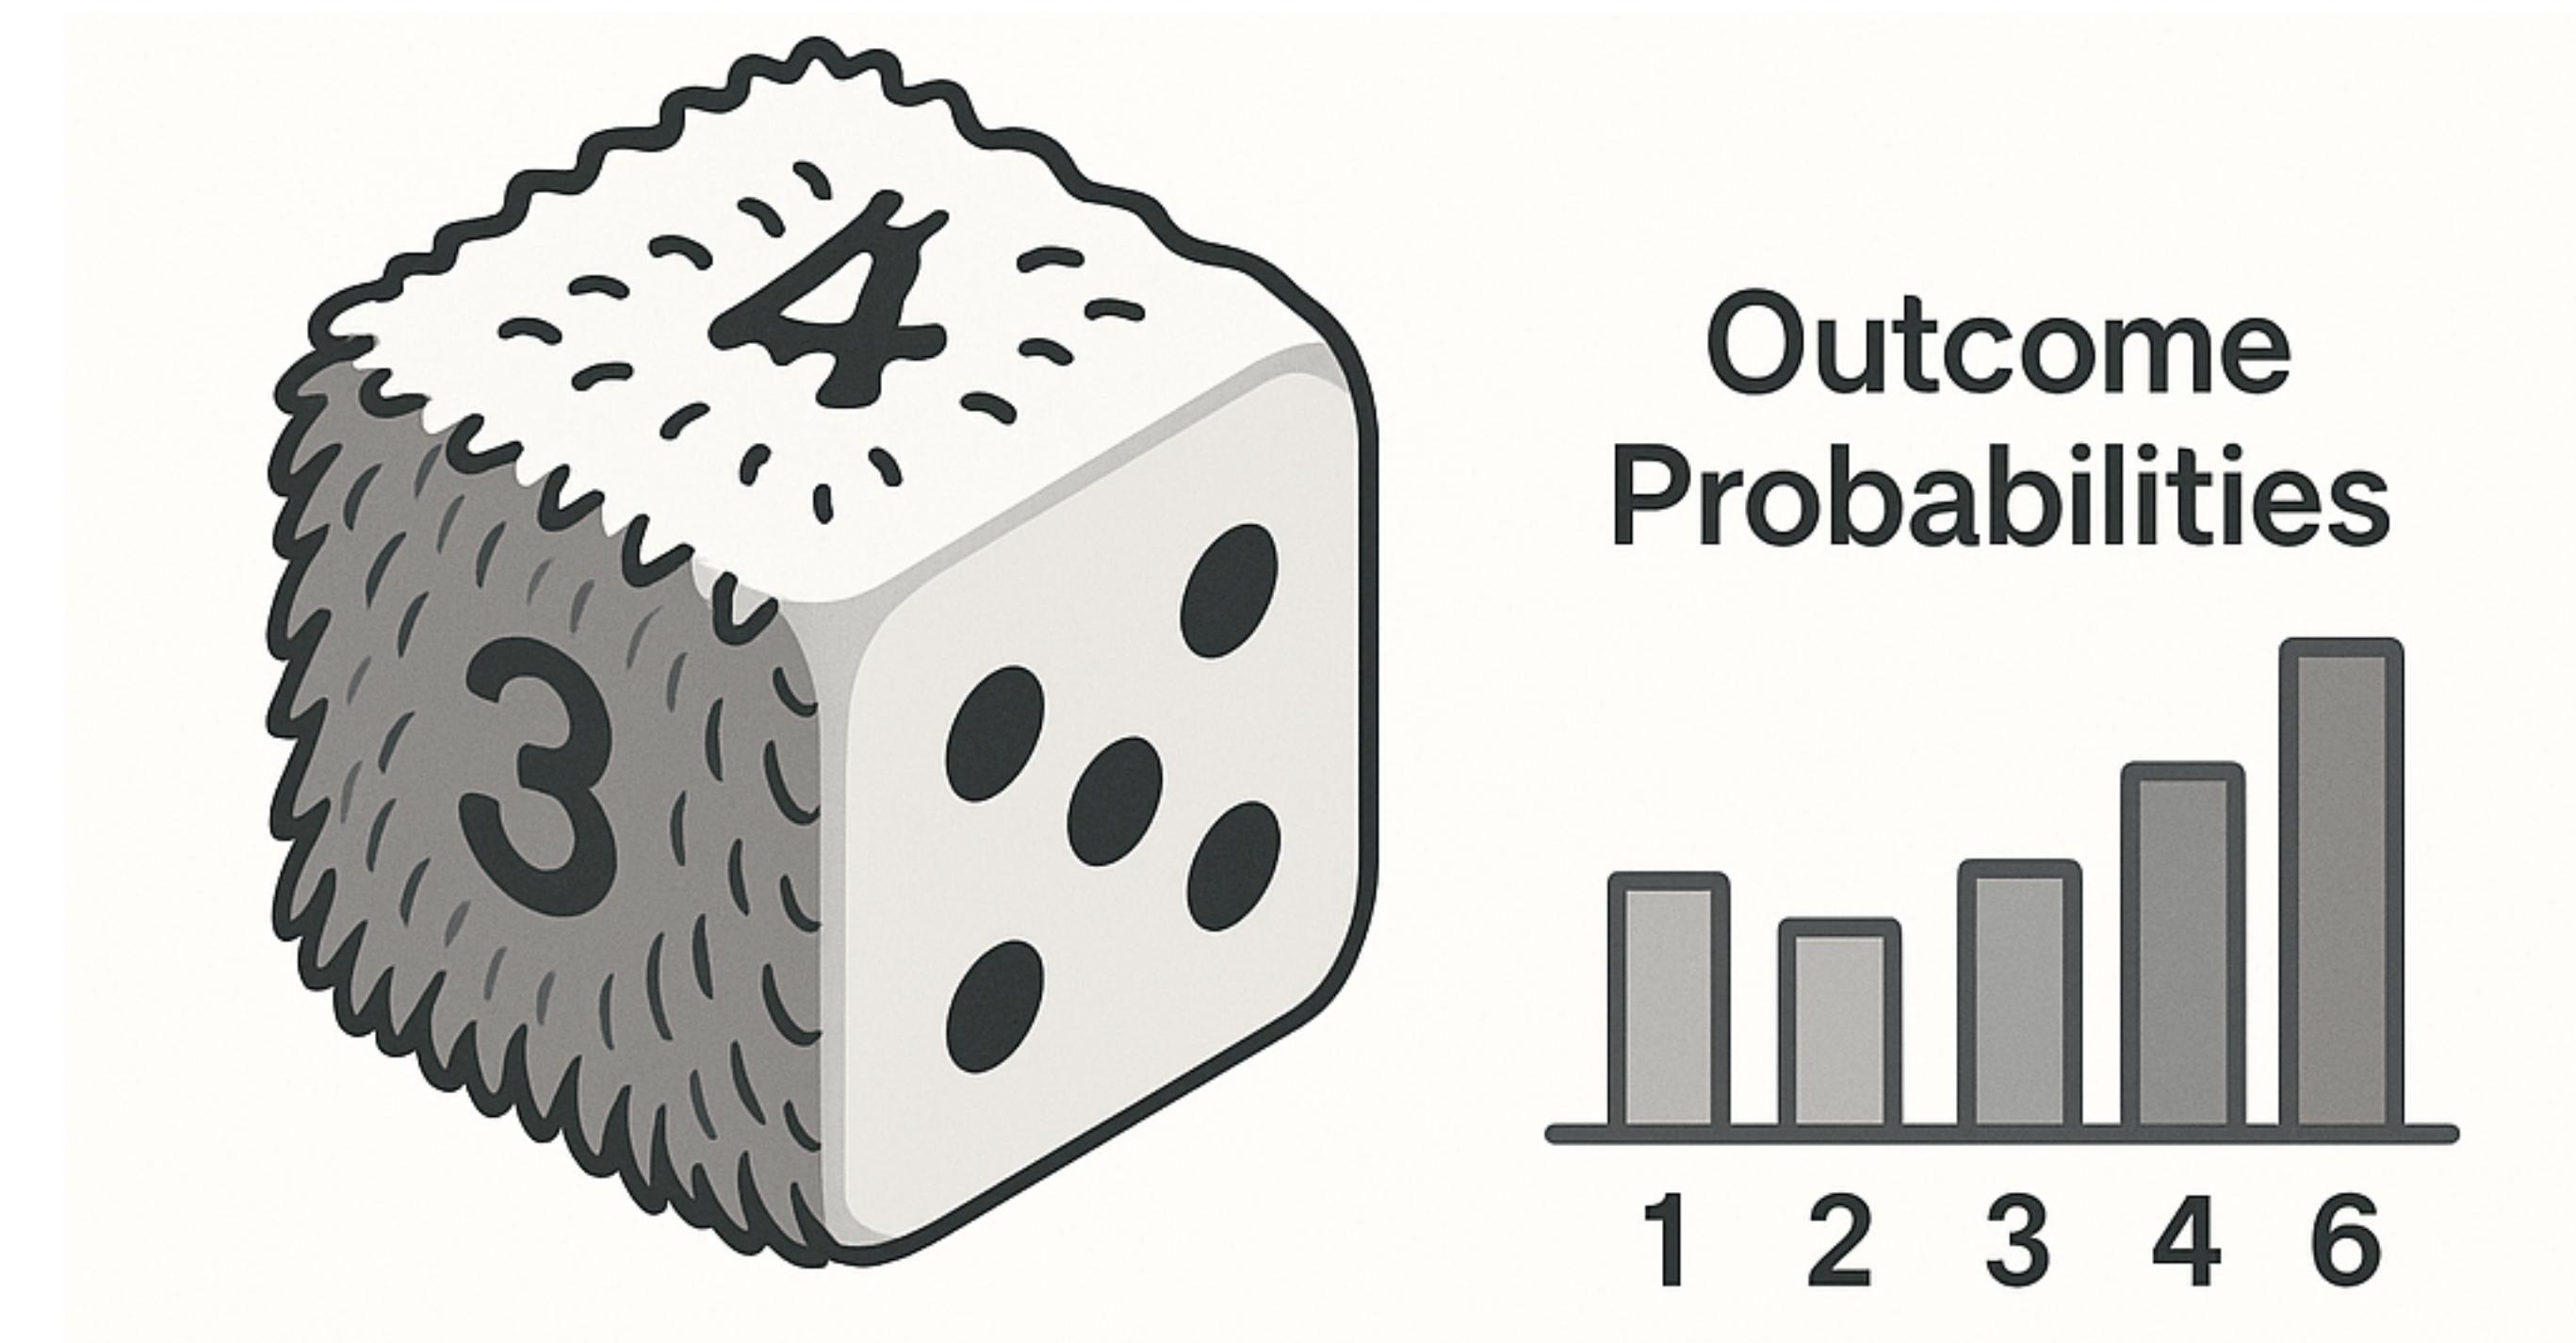
- Goal. Given some  $X$ , predict the associated label  $Y$
- Challenge. We do not know the joint distribution  $P_{XY}(x, y)$ 
  - Instead, we only have access to some data
  - If we knew: we can get the **posterior distribution**...

$$\begin{aligned} P_{Y|X}(y | x) &= \frac{P_{XY}(x, y)}{P_X(x)} \\ &= \frac{P_{XY}(x, y)}{\int_y P_{XY}(x, y) dy} \end{aligned}$$

Warm-up quiz. Given  $P_{Y|X}$ , do you know what to do?

# Quiz: Estimation Basics

- Imagine a **biased die**
  - Its face probabilities depend on the table condition,  $X \in \mathbb{R}^d$
  - The event of each face coming up is represented by  $Y \in \{1, 2, \dots, 6\}$



# Quiz: Estimation Basics

- We are given some weather ( $X = x$ )
  - Then, we can compute the probability of each face:

$$P_{Y|X}(1 | x) = p_1, \quad P_{Y|X}(2 | x) = p_2, \quad \dots, \quad P_{Y|X}(6 | x) = p_6$$

- **Question.** Given these, how will you predict the outcome  $\hat{Y}$ , if you want to:
  - Maximize the probability of being wrong?
  - Minimize the expected error  $\mathbb{E}[(\hat{Y} - Y)^2]$ ?
  - Simulate the (random) outcome  $Y$  of the die?

# Quiz: Estimation Basics

- **Answer.** we can construct, e.g.,

- Maximum a posteriori estimate (MAP), for discrete  $Y$

$$\hat{y} = \arg \max_y P_{Y|X}(y | x)$$

- Minimum Mean-Squared Error (MMSE), for continuous  $Y$

$$\hat{y} = \mathbb{E}[Y | X = x] = \int_y y \cdot P_{Y|X}(y | x) \, dy$$

- Sampling a solution, for diverse generation / prediction

$$\hat{y} \sim P_{Y|X}(\cdot | x)$$

- You can construct similar estimates, e.g., to minimize  $\mathbb{E}[|\hat{Y} - Y|]$ .

# Two approaches in ML

- Now let's get back on track – in learning, we do not know  $P_{XY}$
- Instead, we want to use **training data** to build an estimate
$$\hat{Y} = f(X)$$
- **Question.** How can we do this in a principled way?
  - Generative
  - Discriminative

# Two approaches in ML: Generative

- **Generative** approach aims to directly model  $P_{XY}$ 
    - i.e., capturing the data generation process itself
    - Can be used to construct  $\hat{P}_{Y|X}$
    - Examples: Naïve Bayes, VAE, GAN, Diffusion
- 👍 Once it works well, many other perks
- Generation ( $P_X$ ), Conditional Generation ( $P_{X|Y}$ ), quantify the uncertainty of prediction  $P_{Y|X}(\hat{y} | x)$ , ...
- 👎 Very difficult to achieve – a lot of data, or heavy assumption

# Two approaches in ML: Discriminative

- **Discriminative** approach aims to model  $P_{Y|X}$ 
    - Often model the estimates based on  $P_{Y|X}$  (e.g., MAP), not itself
    - Example: Logistic regression, SVM, neural net classifiers
- 👍 Can learn with relatively less samples
  - Usually better accuracy on the target task
- 👎 Cannot generate data, potentially poor calibration, limited use case

Note: In DL, people used to work on D (~2019), then moved onto G

# **Learning as an optimization**

# Learning as an optimization

- Now, let's focus on the **discriminative** case for **supervised learning**
  - That is, we have a bunch of data

$$\{(X_i, Y_i)\}_{i=1}^n \sim P_{XY}$$

- And our goal is to find a nice model

$$\hat{P}_{Y|X}(y|x)$$

which fits this data the best.

- This optimization is what **learning algorithm** does

# Learning as an optimization

- **Question.** How exactly do we do this?
  - **Answer.** Differs from algorithm to algorithm (sadly)
- There are two unified perspectives toward various ML algorithms
  - Statistical learning
  - Bayesian approach
- Note: These two are – to some degree – interchangeable

# **Learning as an optimization: Statistical learning**

# Statistical Learning

- Under the statistical learning paradigm, each learning algorithm is characterized by three elements:
  - Hypothesis space
  - Loss function
  - Search algorithm



# Statistical Learning

- **Hypothesis space.** A bag of models

$$\mathcal{F} = \left\{ f_{\theta}(\cdot) \mid f_{\theta} : \mathcal{X} \rightarrow \mathcal{Y}, \quad \theta \in \Theta \right\}$$

- $\mathcal{X}$ : set of all possible  $X$  (e.g., set of all 256 x 256 images)
- $\mathcal{Y}$ : set of all possible  $Y$  (e.g., set of all labels)
- $\theta$ : parameters (which we optimize for)
- **Example.** Set of all affine models

$$f_{\theta}(x) = Wx + b, \quad \theta = (W, b), W \in \mathbb{R}^{m \times d}, b \in \mathbb{R}^m$$

# Statistical Learning

- **Loss function.** A measure of “wrongness” of the model prediction:

$$\ell(\cdot, \cdot) : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$$

- If we get a sample  $(X^*, Y^*)$ , the loss of a predictor  $f_\theta$  is:

$$\ell(f_\theta(X), Y)$$

- **Example.** Squared loss

$$\ell(\hat{Y}, Y) = \|\hat{Y} - Y\|_2^2$$

Zero-one loss

$$\ell(\hat{Y}, Y) = \mathbf{1}\{\hat{Y} = Y\}$$

# Statistical Learning

- Before we describe the search algo, let us first formalize our final goal:
- **Objective.** Given the hypothesis space and the loss, our goal is to solve:

$$\min_{\theta \in \Theta} \mathbb{E}_{(X,Y) \sim P_{XY}} [\ell(f_\theta(X), Y)]$$

- That is, we want to find the function  $f_\theta$  which has the smallest average loss on the test sample  $(X, Y)$

# Statistical Learning

- Problem. We don't know  $P_{XY}$
- Idea. We conduct **Empirical Risk Minimization (ERM)**:
  - That is, we find the model which achieves the minimum **average loss** on training dataset:

$$\min_{\theta \in \Theta} \mathbb{E}_{(X,Y) \sim D} [\ell(f_\theta(X), Y)] = \min_{\theta \in \Theta} \left[ \frac{1}{n} \sum_{i=1}^n \ell(f_\theta(X_i), Y_i) \right]$$

# Statistical Learning

- Rationale. If we have enough samples, the law of large numbers say that

$$\frac{1}{n} \sum_{i=1}^n \ell(f_\theta(X_i), Y_i) \xrightarrow{n \rightarrow \infty} \mathbb{E}[\ell(f_\theta(X), Y)]$$

for any fixed  $\theta$

- Thus, the empirical loss can be a good proxy of the population loss
  - Caveat: LLN requires independent(-ish) draws of the samples!

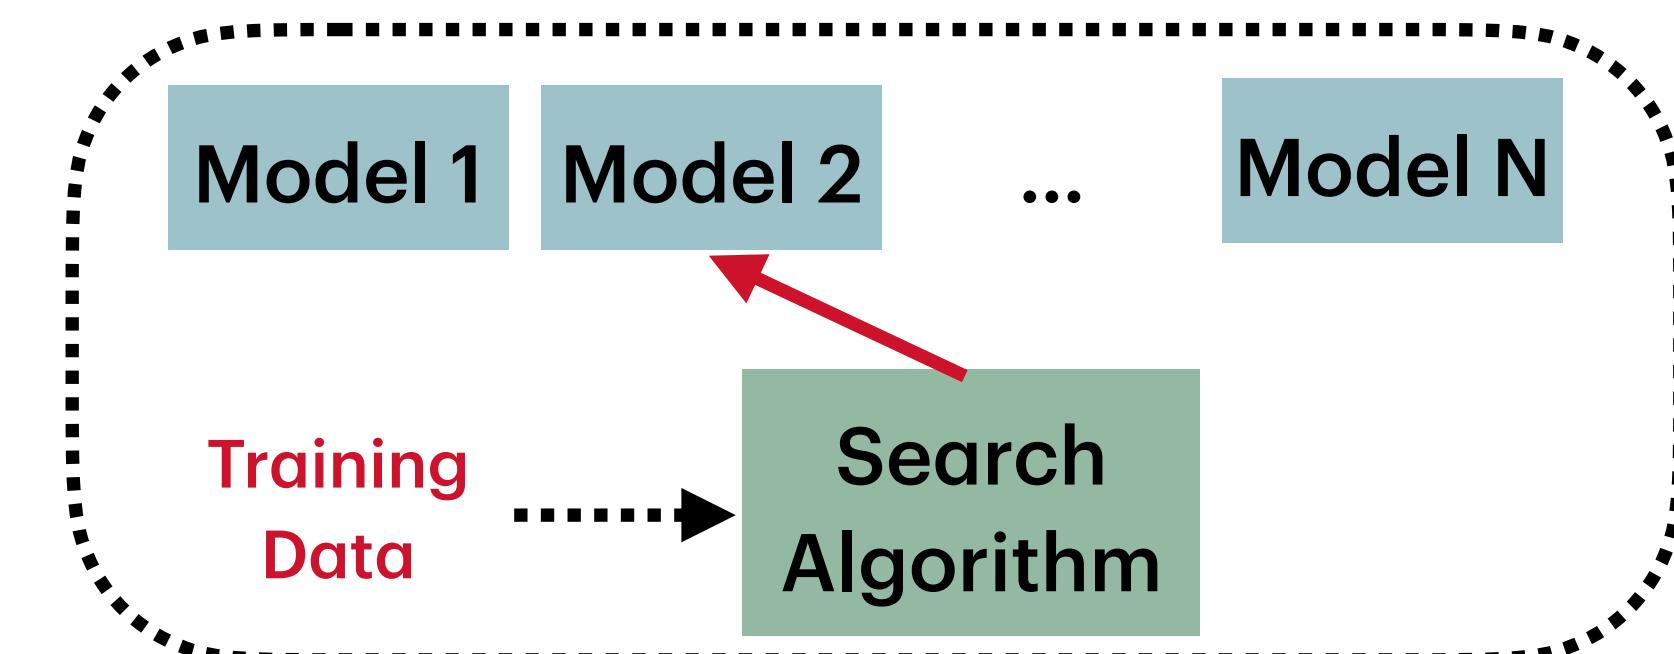
# Statistical Learning

- **Search algorithm.** How we solve this ERM optimization

$$\min_{\theta \in \Theta} \mathbb{E}_{(X,Y) \sim D} [\ell(f_\theta(X), Y)] = \min_{\theta \in \Theta} \left[ \frac{1}{n} \sum_{i=1}^n \ell(f_\theta(X_i), Y_i) \right]$$

- Example.
  - Analytical solution
  - Solve by iterative optimization (e.g., SGD)

(more on this later)

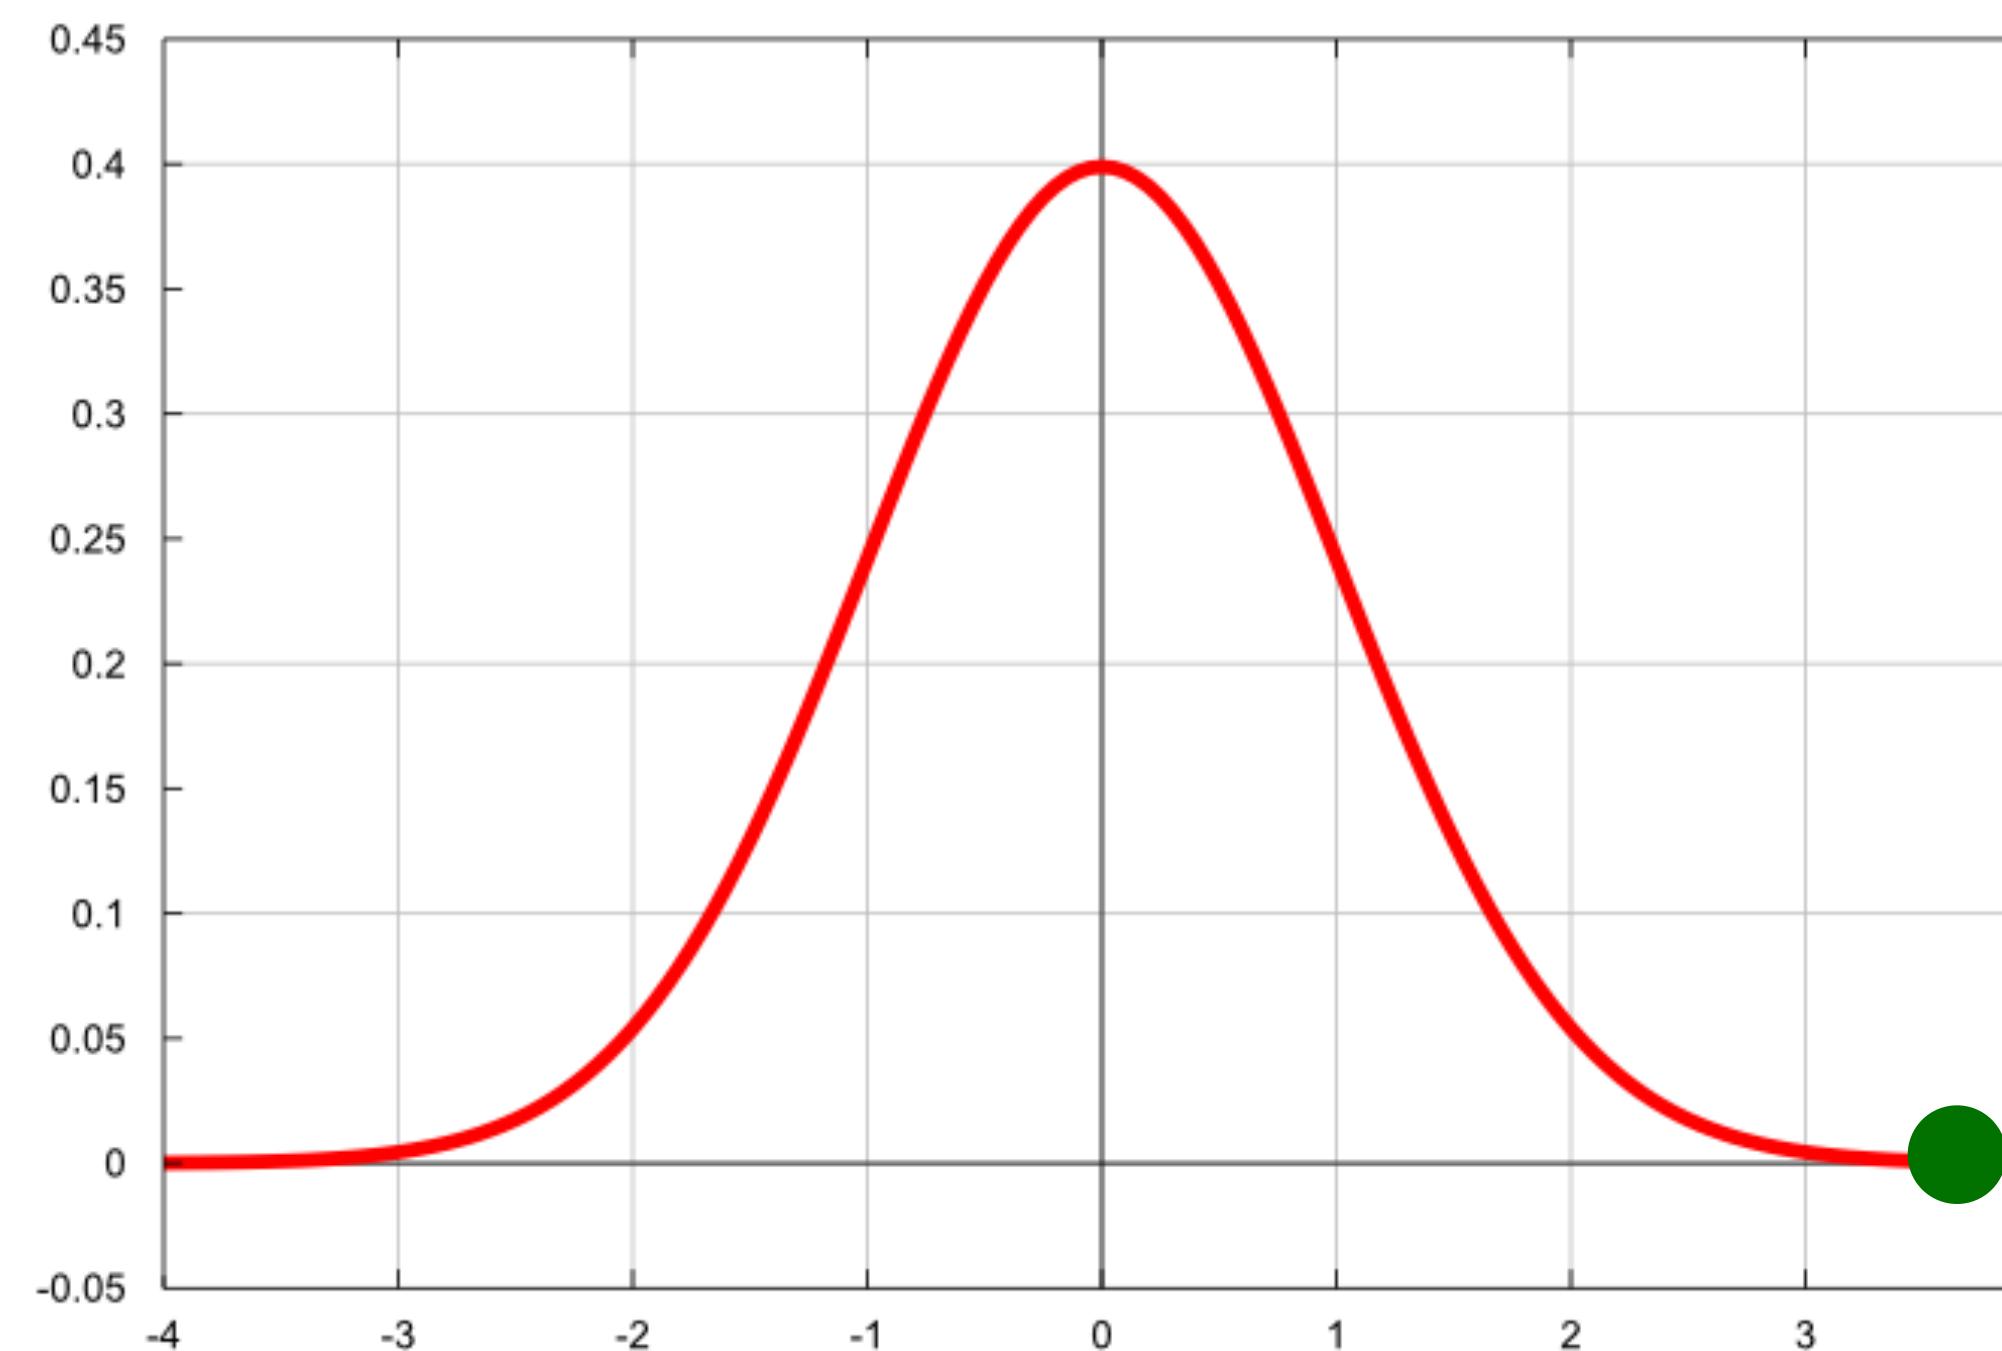


# **Learning as an optimization: Bayesian perspective**

# Bayesian approach

- Bayesians prefer a generative explanation:  
“If our model is correct, the probability of our model generating the data would be high.”

(called the “maximum likelihood principle”)



If we see this data,  
maybe our model is wrong...

# Bayesian approach

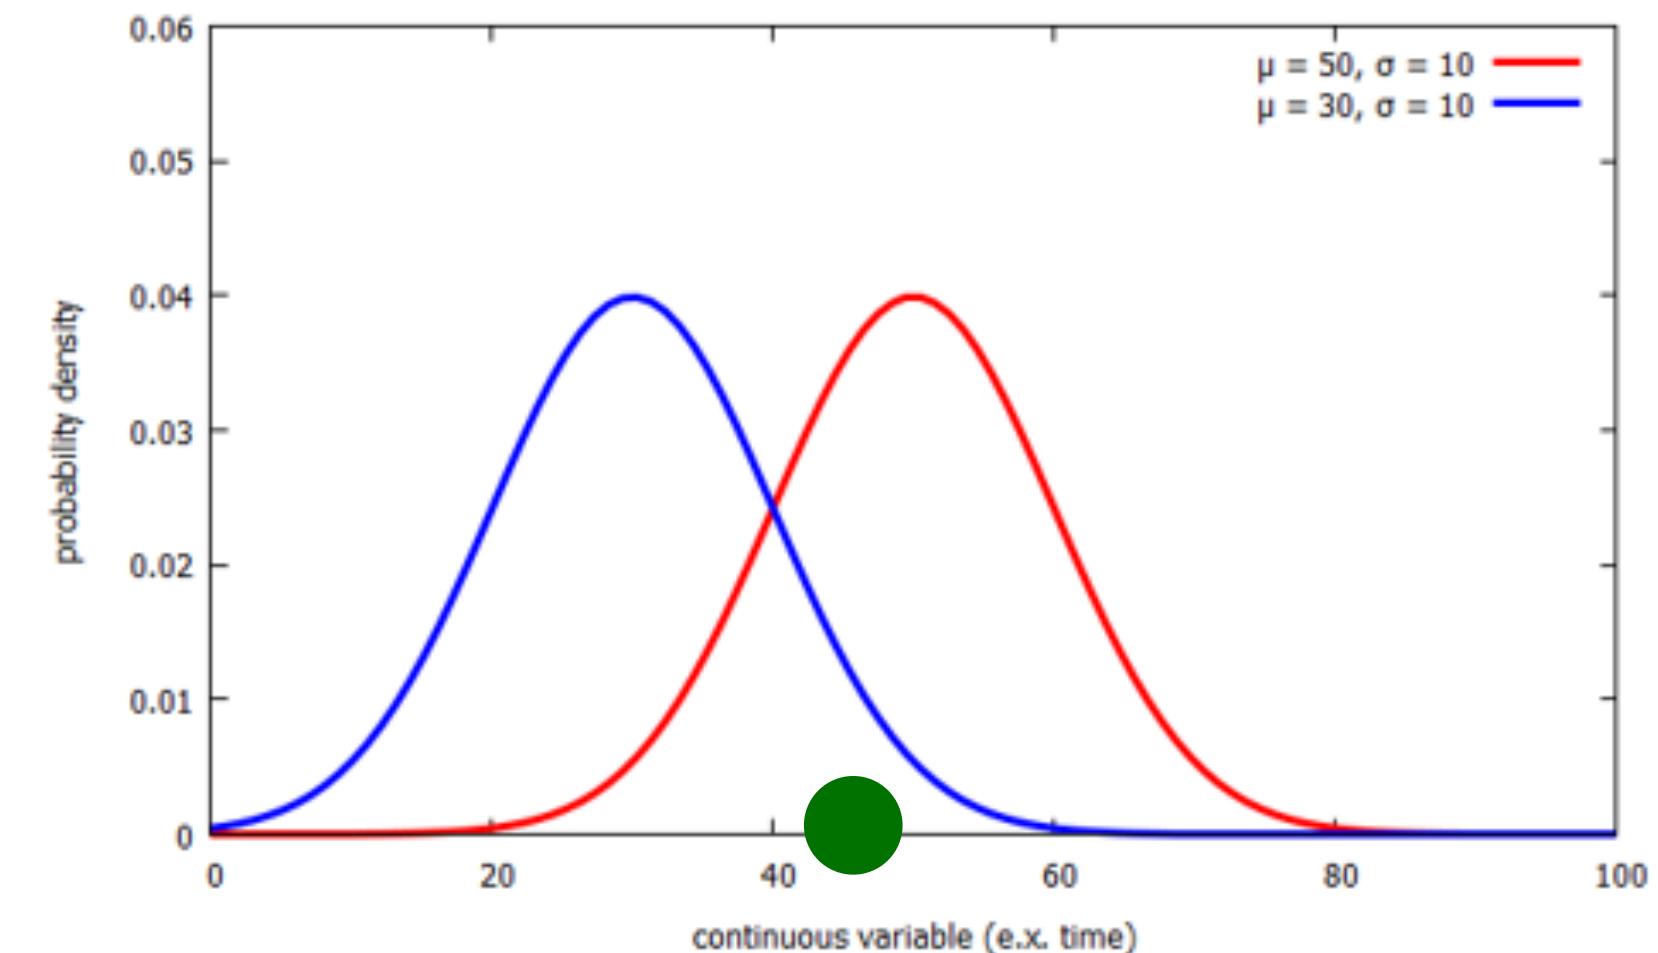
- This principle provides a mean to compare two models:
- **Example.** Suppose that we have two “models”

$$P_{XY}^{(1)}(x, y), \quad P_{XY}^{(2)}(x, y)$$

- Suppose that we are given one sample:  $(X^*, Y^*)$ 
  - If we have

$$P_{XY}^{(1)}(X^*, Y^*) > P_{XY}^{(2)}(X^*, Y^*)$$

then  $P^{(1)}$  is more likelier to be correct!



# Bayesian approach

- Suppose that we have a family of **parametrized joint distributions**

$$P_\theta(x, y) = P_\theta(x)P_\theta(y | x), \quad \theta \in \Theta$$

- **Goal.** Find  $\theta$  maximizing the probability of generating all training data:

$$\max_{\theta} P_\theta((X_1, Y_1), (X_2, Y_2), \dots, (X_n, Y_n))$$

- If all training data are independently drawn, we know that this is:

$$\max_{\theta} \left( \prod_{i=1}^n P_\theta(X_i, Y_i) \right)$$

forgive me for the abuse of notation ;)

# Bayesian approach

- We can apply  $\log(\cdot)$  to make things look simpler:

$$\max_{\theta} \sum_{i=1}^n \log P_{\theta}(X_i, Y_i)$$

- This is what we call the **maximum log-likelihood solution**
- We can break down  $P_{\theta}$  and write:

$$\max_{\theta} \left( \sum_{i=1}^n \log P_{\theta}(Y_i | X_i) + \sum_{i=1}^n \log P_{\theta}(X_i) \right)$$

- For simplicity, ignore the second term

# Bayesian approach

- Notice that this is **similar to doing ERM:**

$$\max_{\theta} \left( \sum_{i=1}^n \log P_{\theta}(Y_i | X_i) \right) \Leftrightarrow \min_{\theta} \left( \frac{1}{n} \sum_{i=1}^n \log \frac{1}{P_{\theta}(Y_i | X_i)} \right)$$

- If we have a nice loss and  $f_{\theta}$  such that

$$\log \frac{1}{P_{\theta}(y | x)} = \ell(f_{\theta}(x), y)$$

then Bayesian approach reduces to ERM!

(many loss functions – e.g., cross-entropy – have this origin)

# Summing up

- Most ML algorithms are **ERM**, with different choice of
  - Hypothesis space
  - Loss
  - Search algorithm
- But why are there so many algorithms?

# **Considerations of building an ML algorithm**

# Which algorithm should we use?

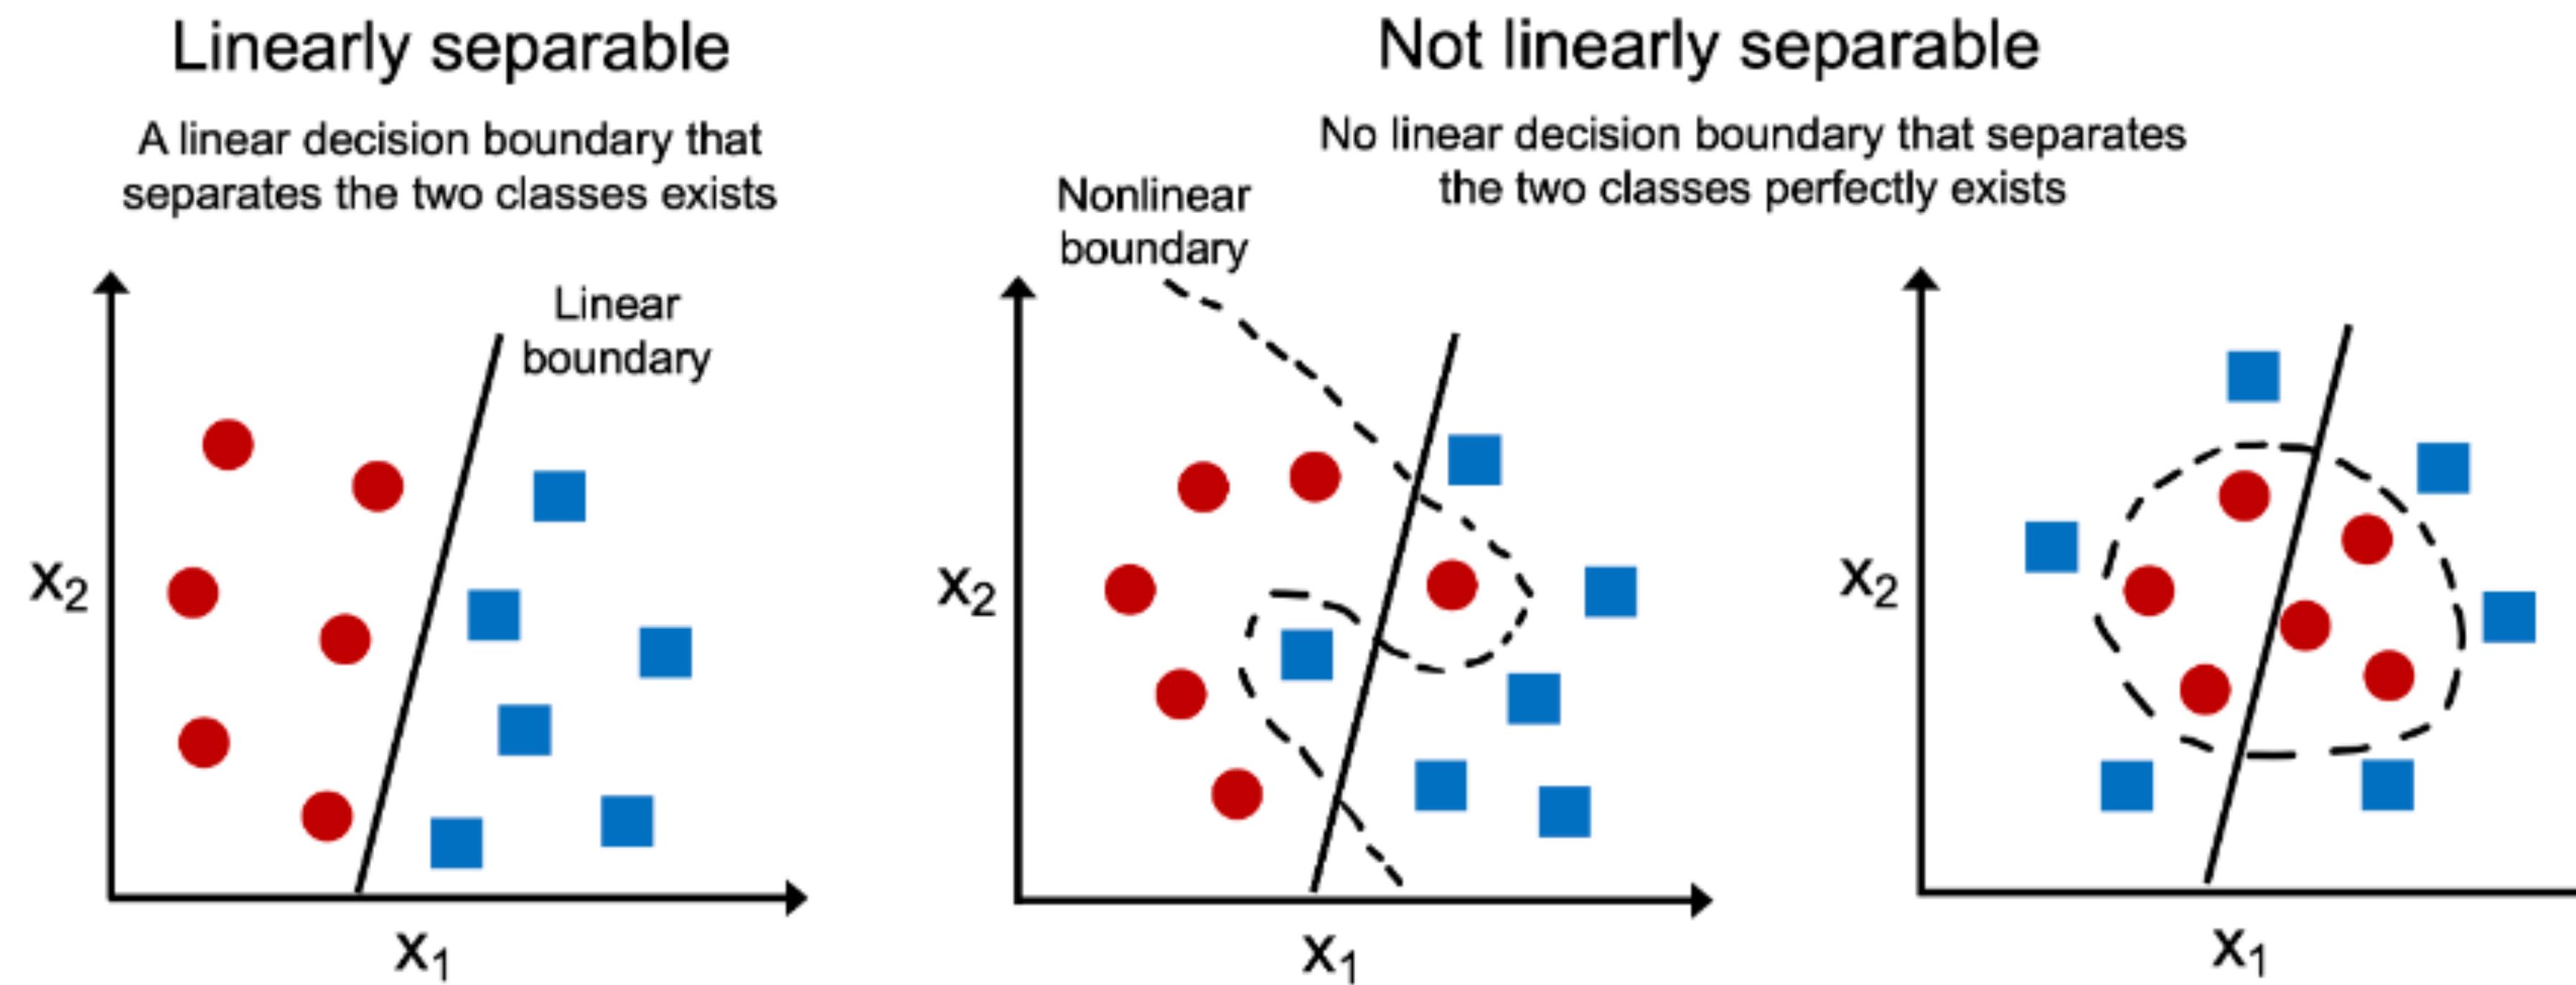
$$\min_{\theta \in \Theta} \frac{1}{n} \sum_{i=1}^n \ell(f_\theta(X_i), Y_i) \quad (+ \text{ regularizers})$$

- Basically, designing the components of this optimization formula

# Which algorithm should we use?

$$\min_{\theta \in \Theta} \frac{1}{n} \sum_{i=1}^n \ell(f_\theta(X_i), Y_i) \quad (+ \text{ regularizers})$$

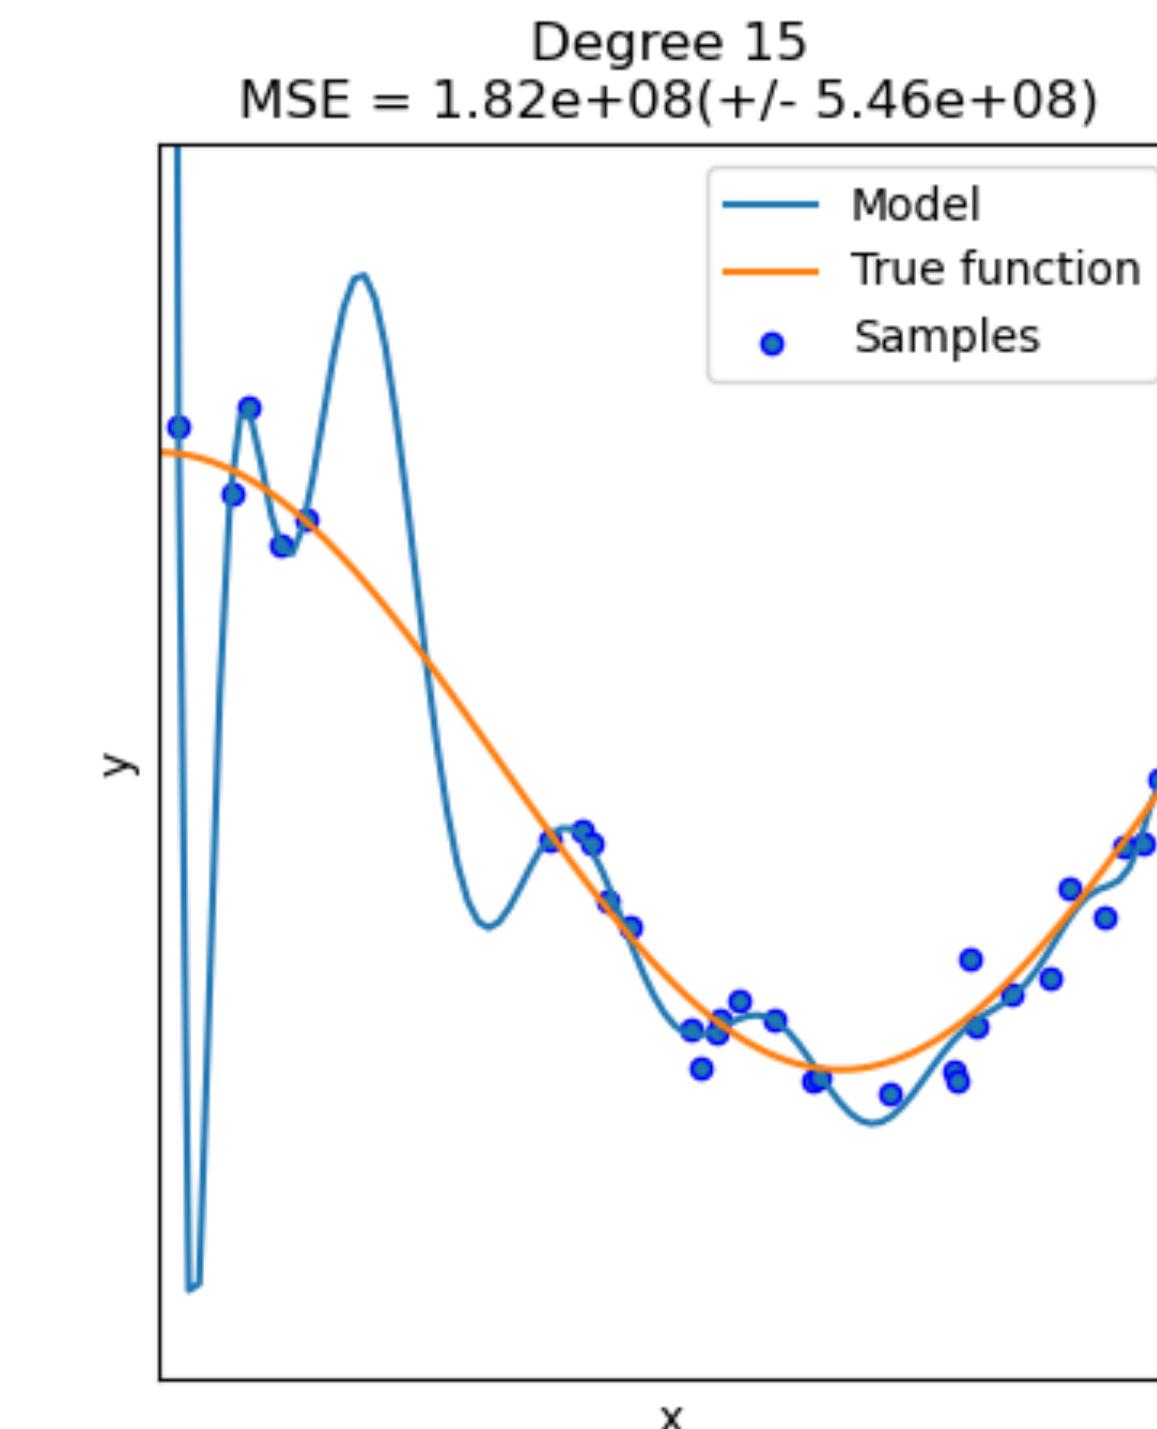
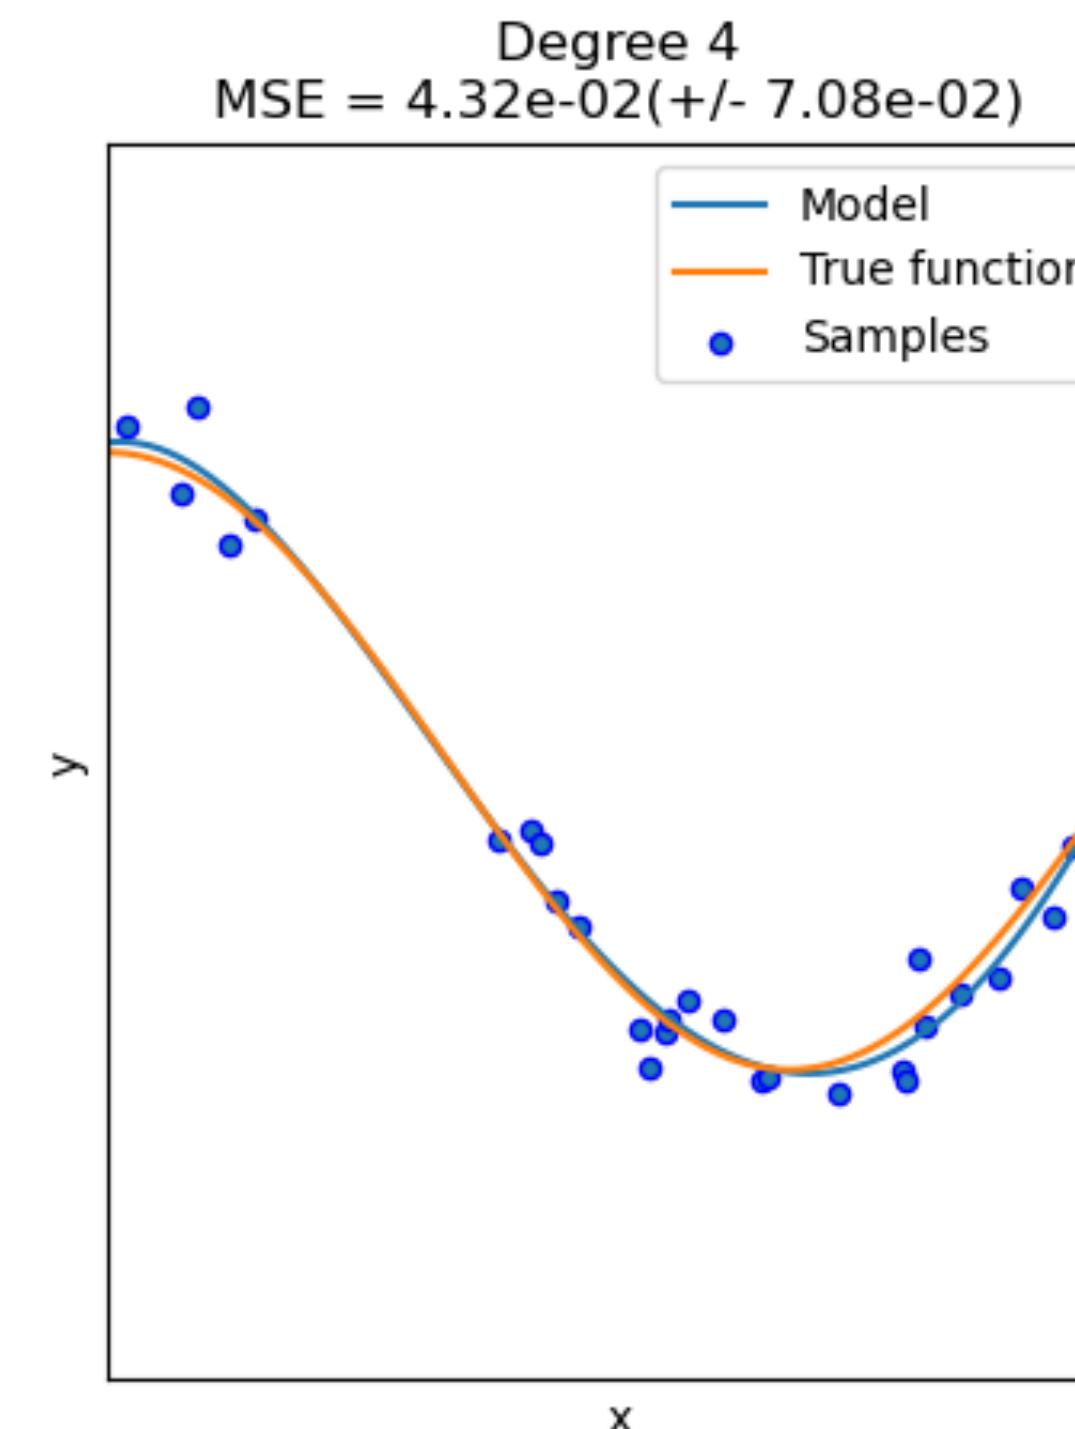
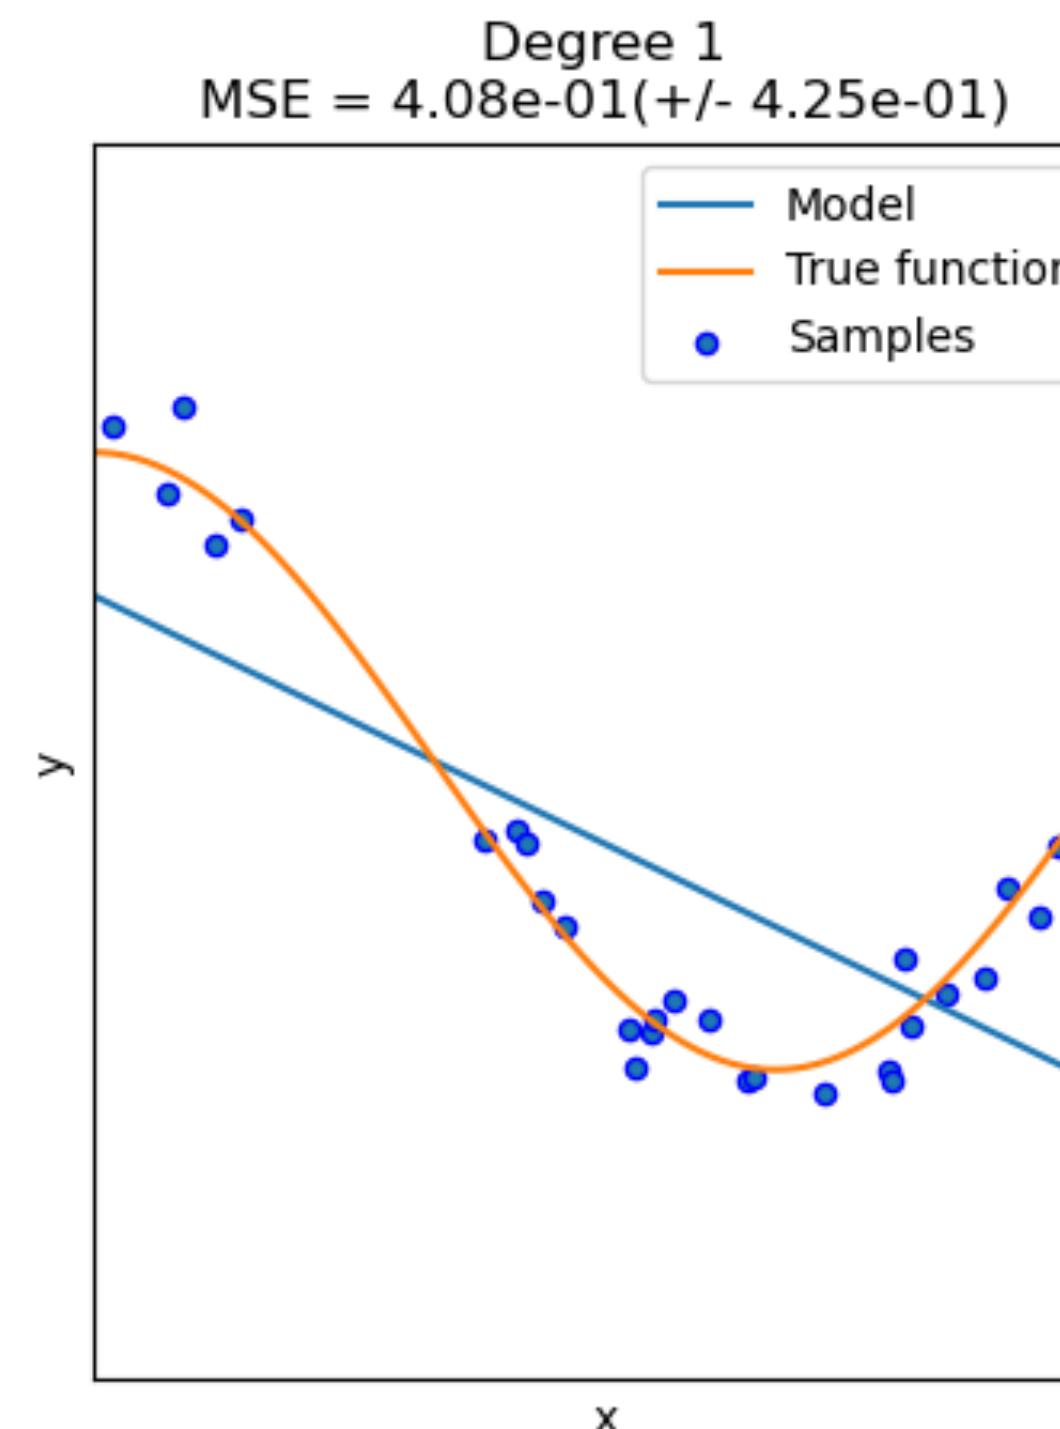
- **Model Size (= Richness of  $\mathcal{F}$ )**
  - If too small, even the best  $f_\theta$  cannot fit the training data



# Which algorithm should we use?

$$\min_{\theta \in \Theta} \frac{1}{n} \sum_{i=1}^n \ell(f_\theta(X_i), Y_i) \quad (+ \text{ regularizers})$$

- **Model Size (= Richness of  $\mathcal{F}$ )**
  - If too large, can overfit the training data + large inference cost

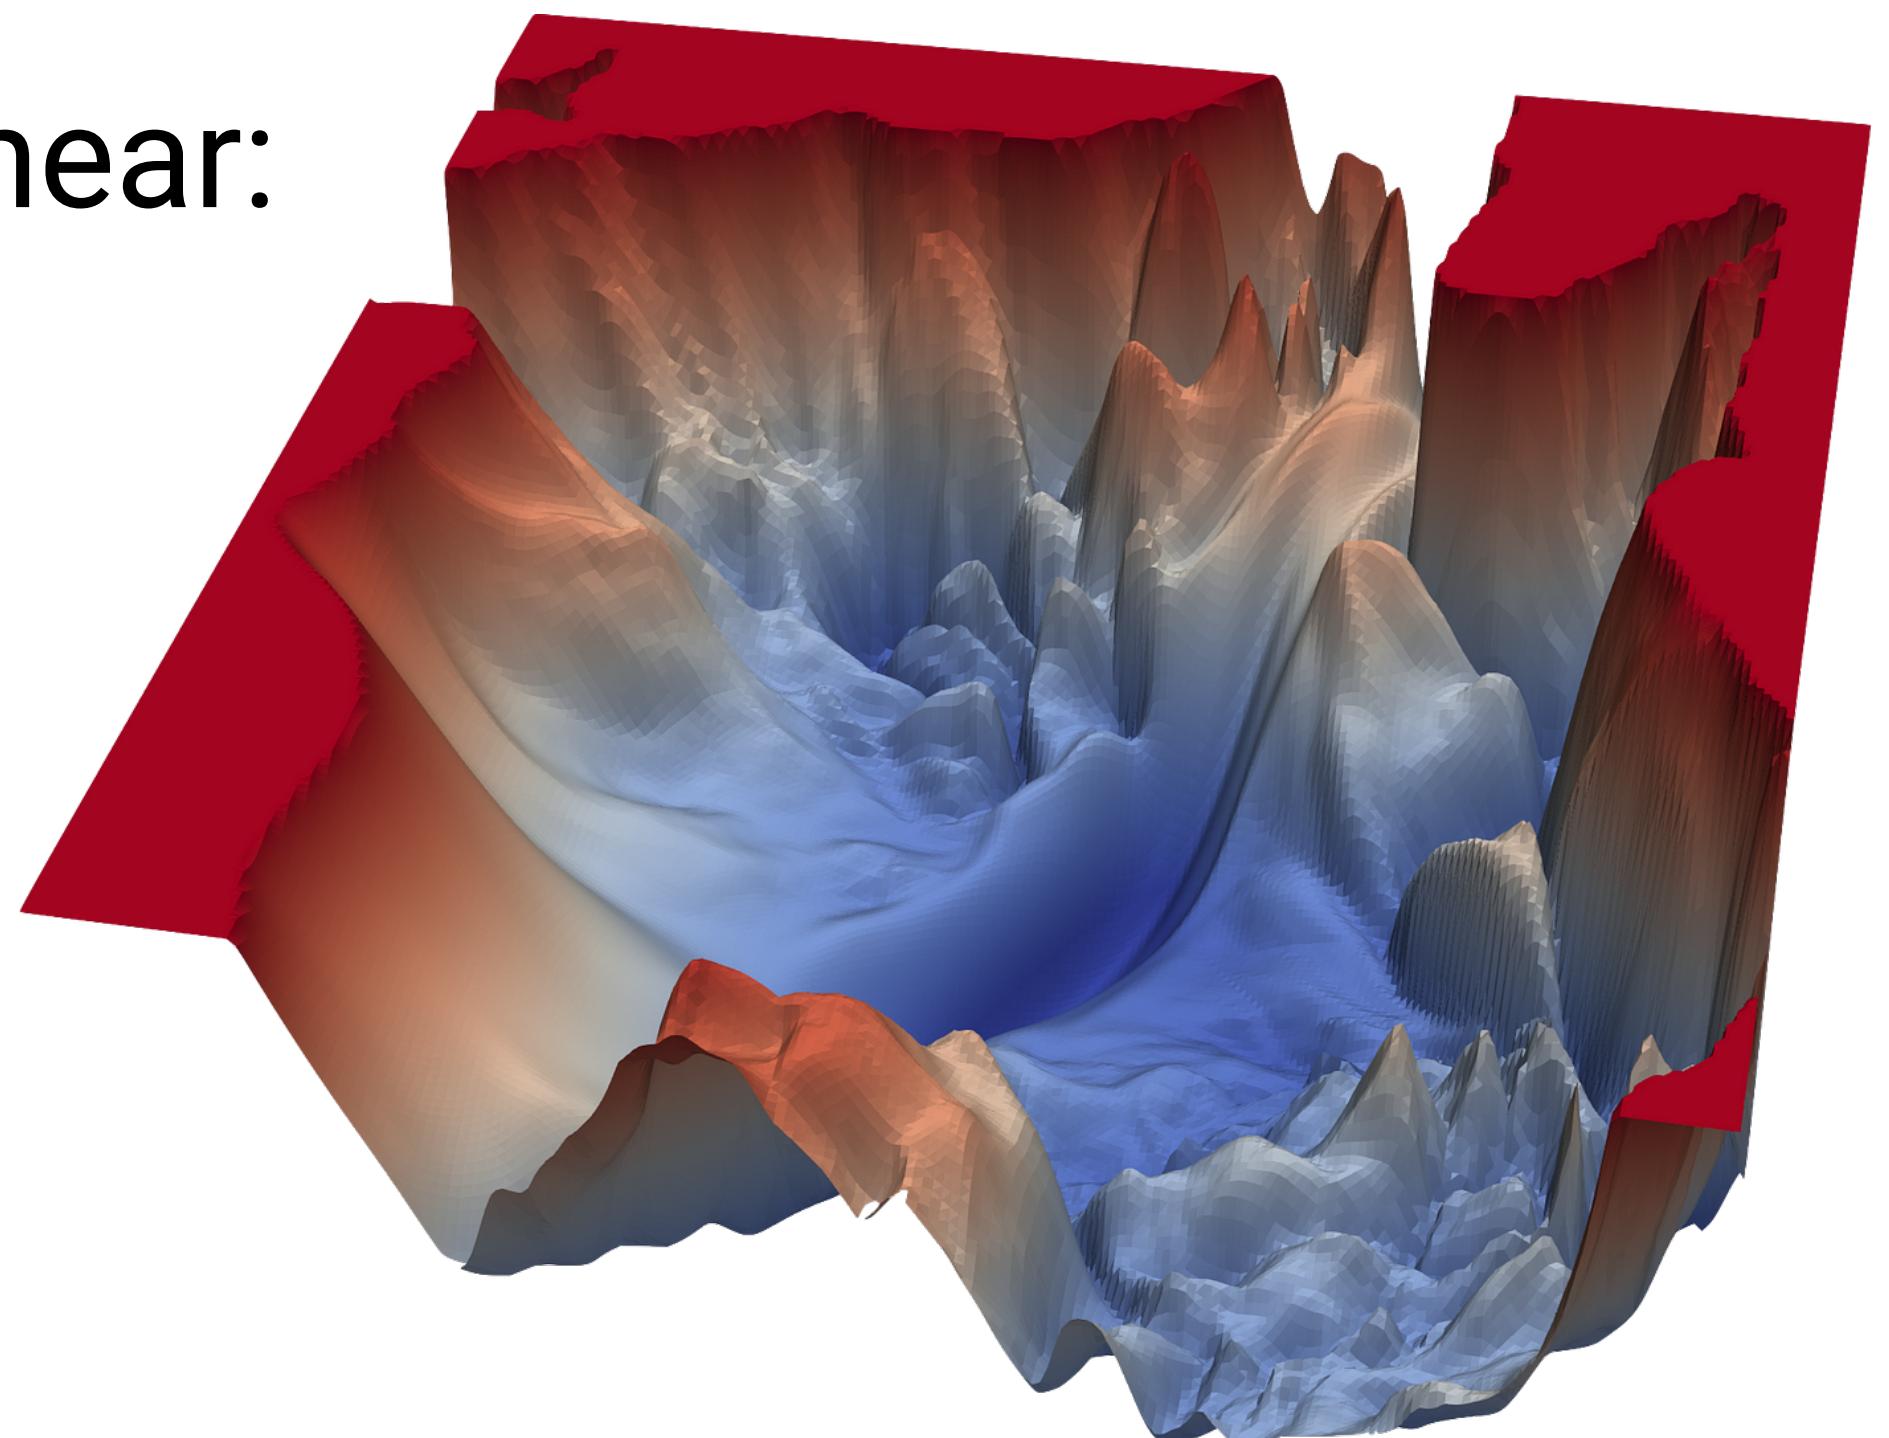


# Which algorithm should we use?

$$\min_{\theta \in \Theta} \frac{1}{n} \sum_{i=1}^n \ell(f_\theta(X_i), Y_i) \quad (+ \text{ regularizers})$$

- **Optimization** (= Difficulty of solving ERM)

- Need tailoring for each model class
- If models are highly complicated & nonlinear:
  - Analytical solution unavailable
  - Takes a long time to solve



# Which algorithm should we use?

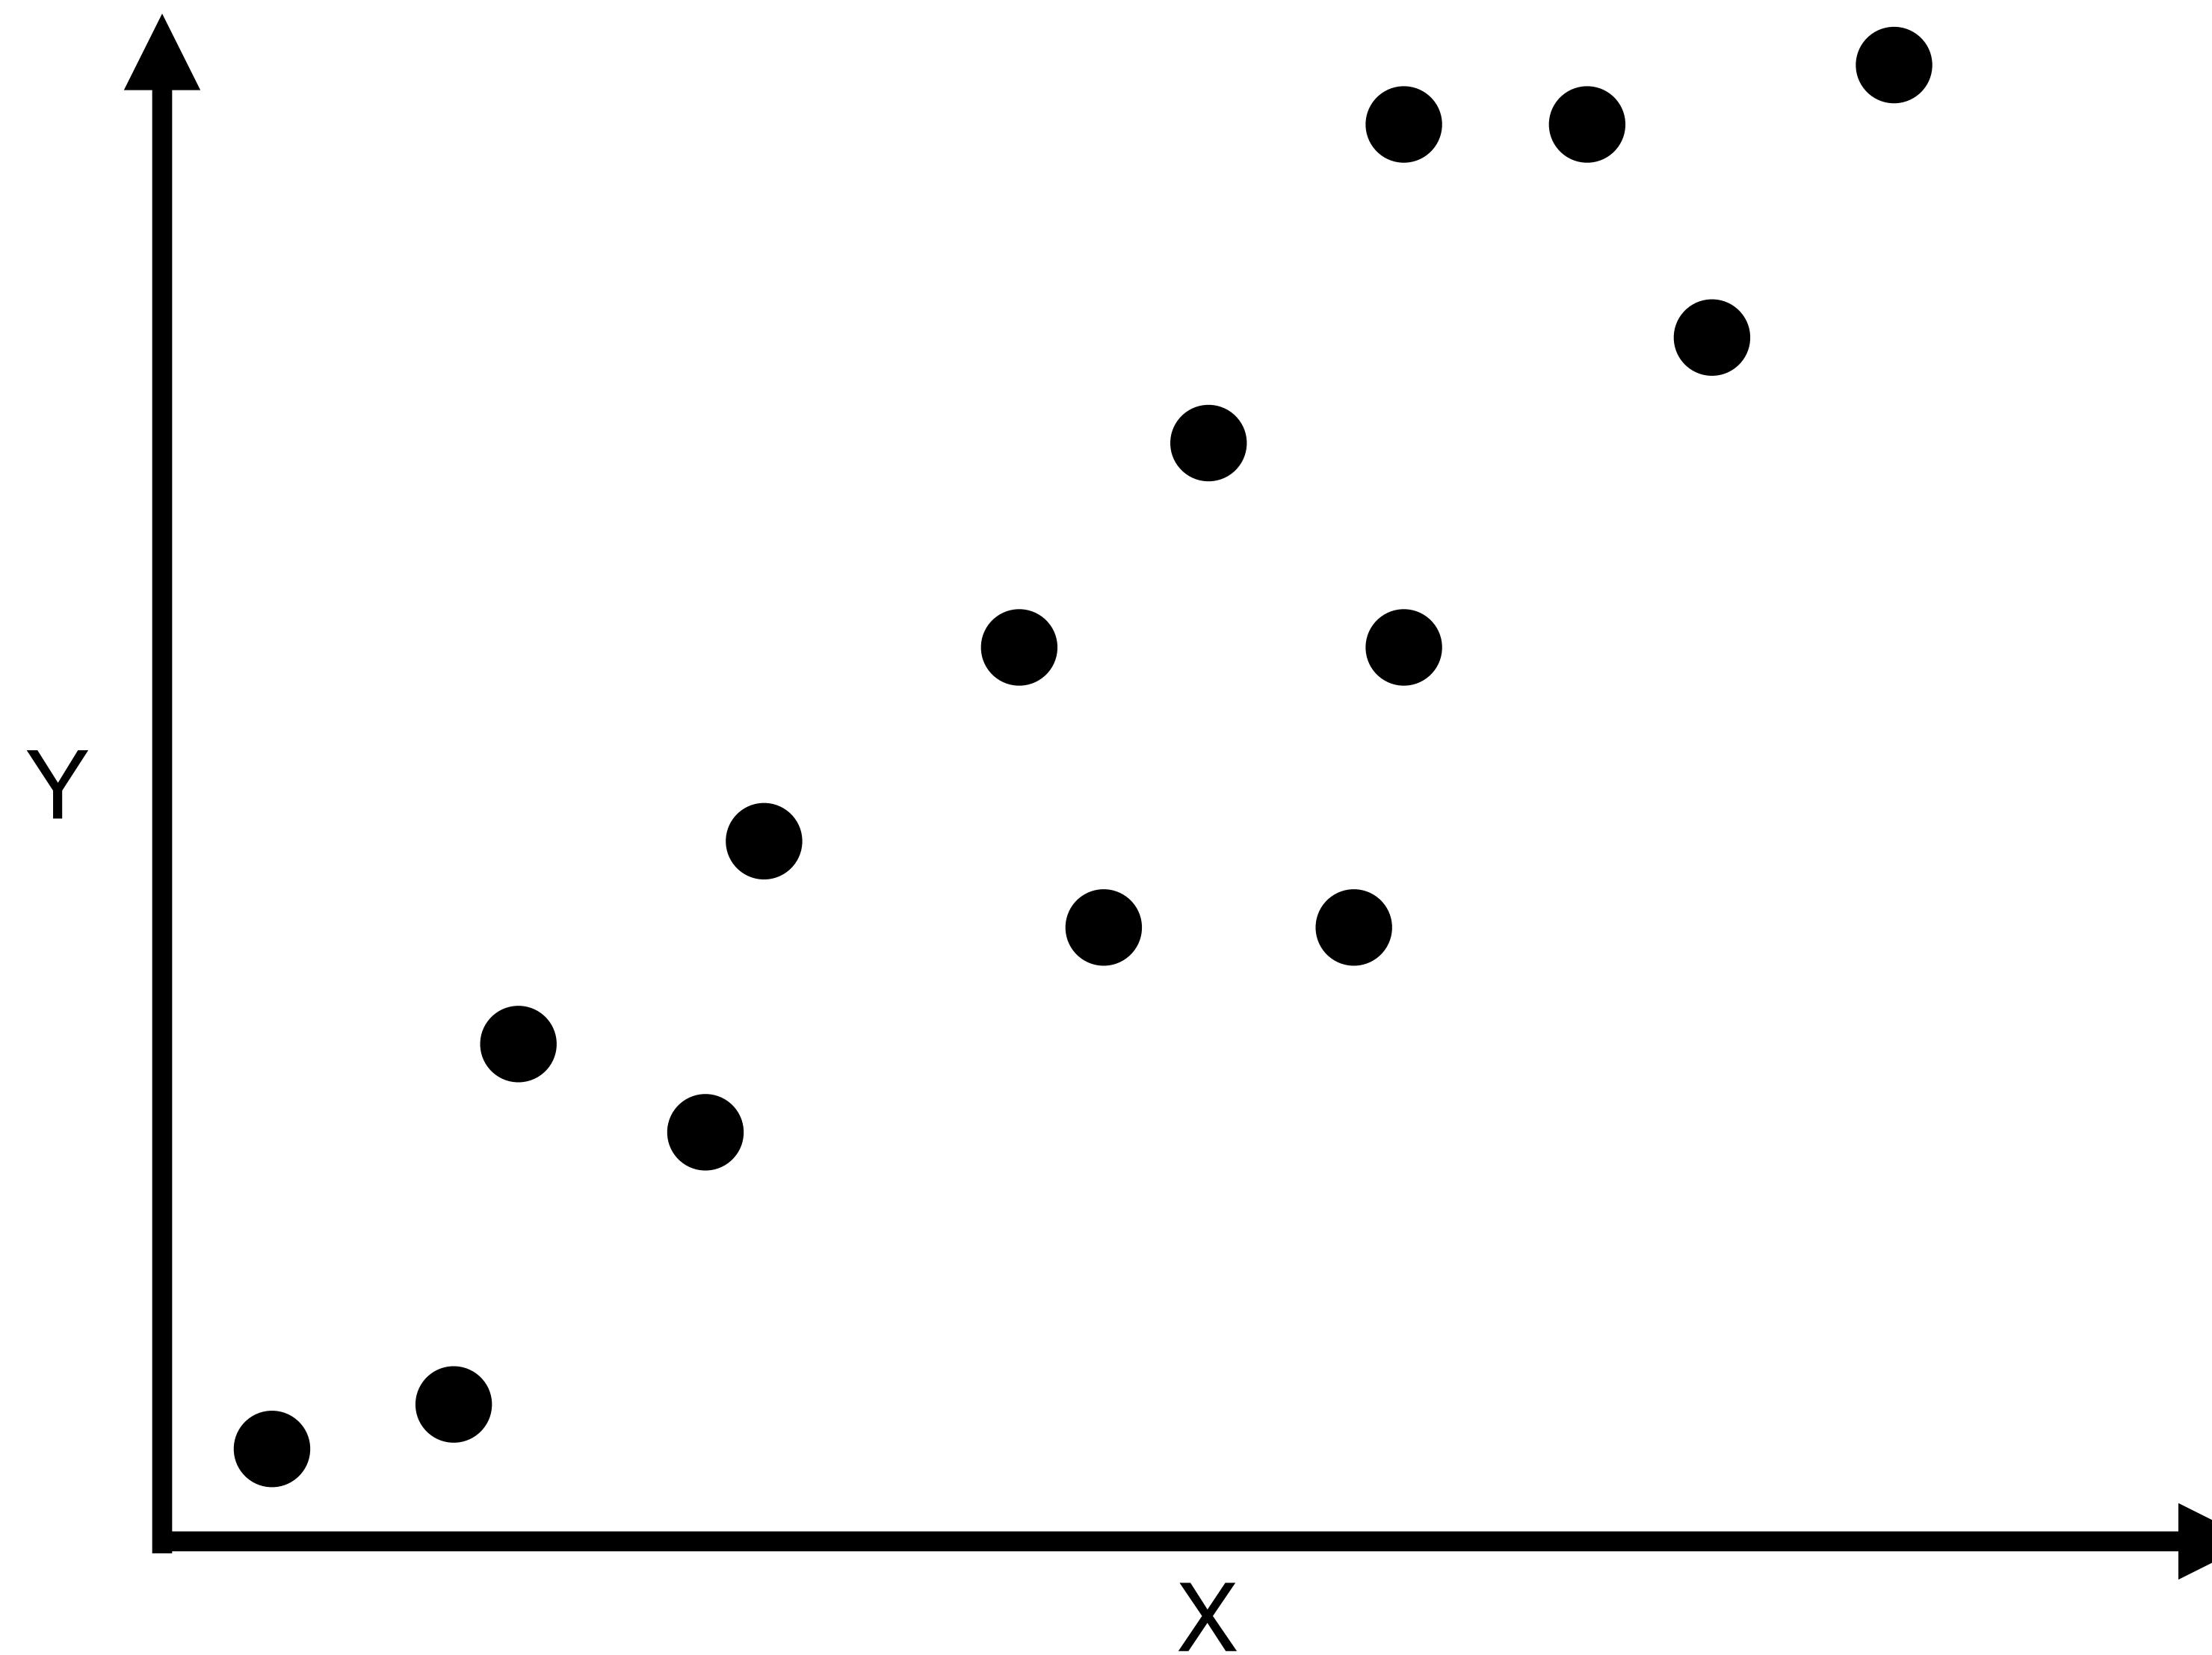
$$\min_{\theta \in \Theta} \frac{1}{n} \sum_{i=1}^n \ell(f_\theta(X_i), Y_i) \quad (+ \text{ regularizers})$$

- **Loss & Regularizer**

- Affects the difficulty of optimization
  - e.g., non-continuous loss
- Affects overfitting
  - e.g., penalizing model complexity

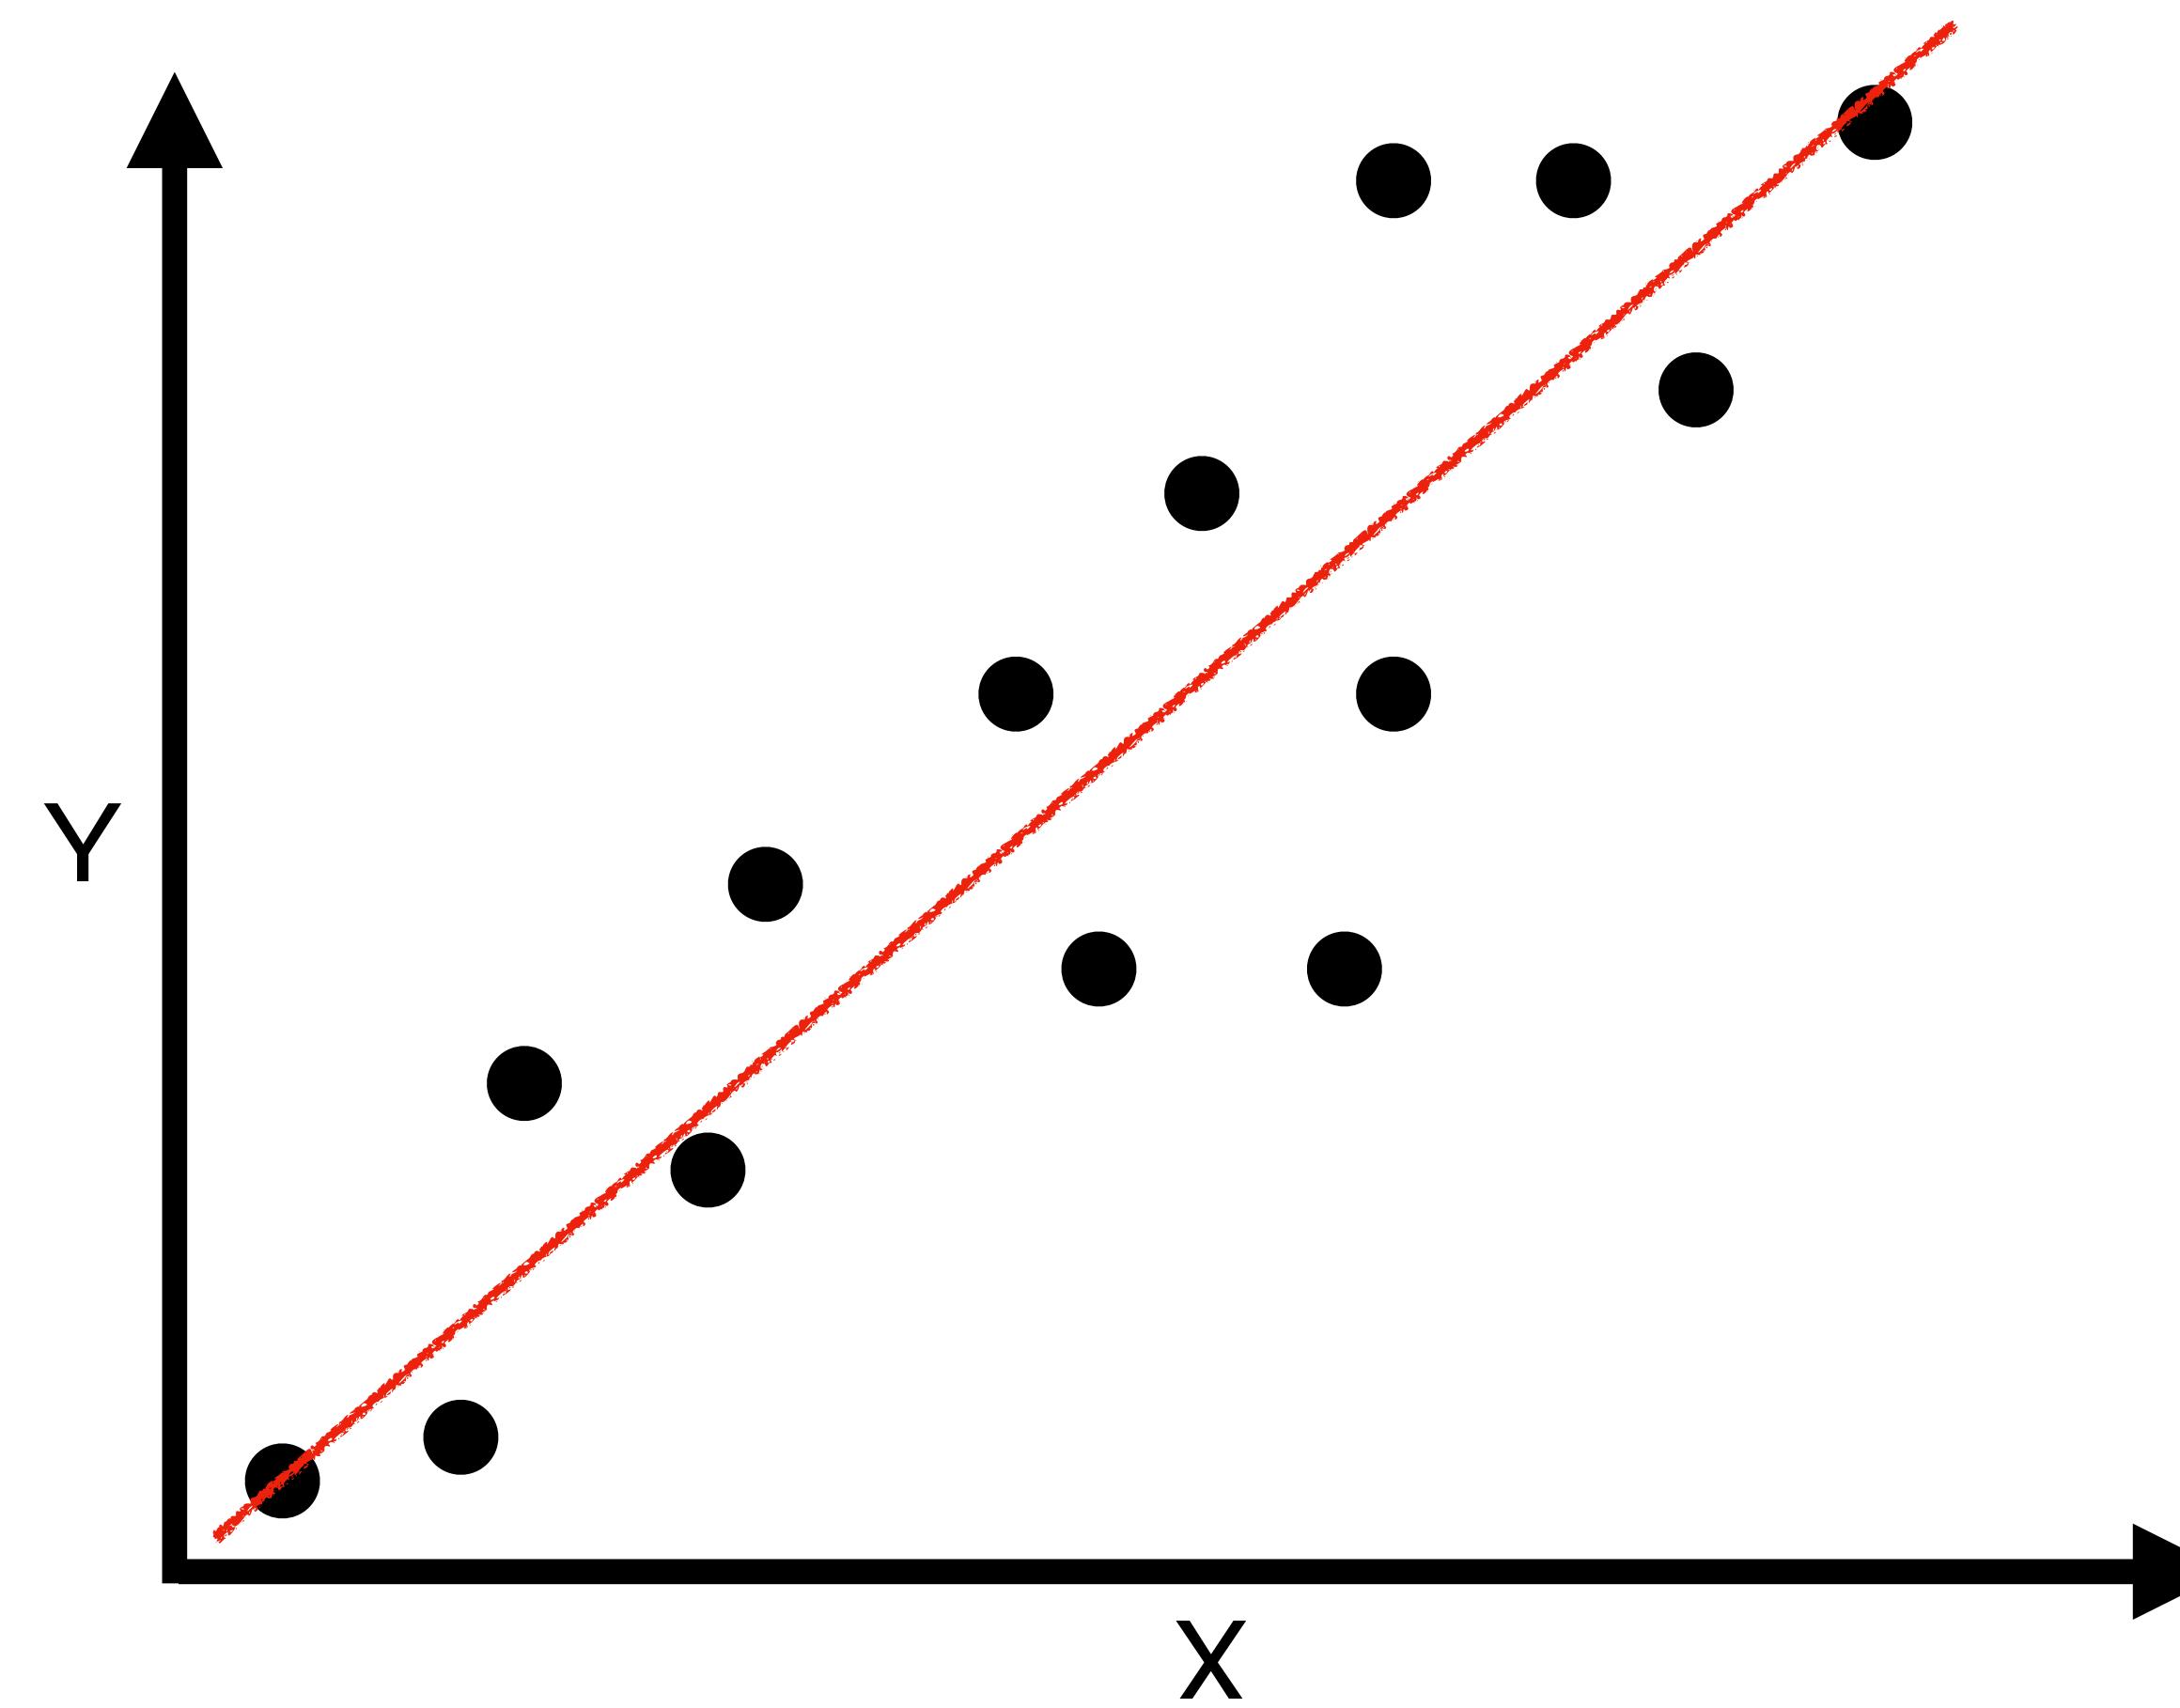
# Designing the right model class

- Think about this data:



# Designing the right model class

- We, as a human, may believe that this is a **straight line + noise**
  - This is due to our **inductive bias** – simpler solutions (in some sense) are more desirable



# From the next class

- We study popular ML algorithms one-by-one
  - Each designed with different inductive bias
    - Different hypothesis space
    - Different optimization mechanism
    - Different loss / regularizer
- Note. Many of these choices heavily depend on tasks:
  - e.g., image vs text vs tabular