


Simple Classifiers

Today

- Various classification algorithms
 - Nearest neighbors
 - Naïve Bayes
 - Linear classifiers
 - Perceptron
 - Logistic regression


Goal

- Modeling the relationship between
 - continuous input $X \in \mathbb{R}^d$ (or discrete)
 - **discrete** output $Y \in \{1, \dots, K\}$
 - called “class”



Binary Classification

- For simplicity, we mostly consider the case of **binary classification**
 - $Y \in \{0,1\}$




Binary Classification

- In binary classification, any classifier can be viewed as selecting a **subset of the input space**


$$f(x) = \begin{cases} 0 & \dots x \in \mathcal{R}_0 \\ 1 & \dots x \in \mathcal{R}_1 \end{cases}$$

- Decision regions $\mathcal{R}_0, \mathcal{R}_1$ are separated using some **decision boundary**




Classification vs. Regression

- Fun fact. Technically, we can use linear regression for classification
 - Simply view 0/1 class labels as outputs to predict



Classification vs. Regression



- **However.** This is not a good idea...
 - Very sensitive to “outliers,” e.g., extremely large yet benign tumor
 - Thus we want better tools



Nearest neighbors


Historical bits

- Can be traced back to a book in 1021
• كتاب المناظر (“the book of optics”) by Ibn al-Haytham



Historical bits

- Viewed human visual recognition as a nearest neighbor
 - “Recognition is the **perception of similarity** between two forms—i.e., of the form
 - (1) sight perceives at the moment of recognition,
 - (2) and the form of that visible object, or its like, that it has perceived one or more times before.”



Setup

- We have a labeled dataset

$$D = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$$

- **Features.** $\mathbf{x}_i \in \mathcal{X}$ (continuous, discrete, mixed, ...)
- **Label.** $y_i \in \{1, \dots, K\}$

Training

- A cool aspect of KNN is that it is **training-free**
 - All we need to do is to **store data** in some database, in a form that we can retrieve them easily



Inference



- Suppose that we are given some test sample $\mathbf{x}^{(\text{new})}$
- Pick k samples with the **highest similarity**:
 - Equivalently, find the training samples with bottom- k distance:

$$\min_i \text{dist}(\mathbf{x}^{(\text{new})}, \mathbf{x}_{(i)})$$

- Then, predict with **majority vote**
(we can also do regression, via weighted averaging)


Properties

- KNN predictor is **nonlinear**
 - Example. $k = 3$



Hyperparameter

- The **neighbor set size k** has a big impact on the predictor
 - Small k : Flexibility Larger k : Smooth decision boundary



Properties



- KNN predictor is **nonparametric**
 - Nonparametric. Using flexible number of (or infinite) parameters
 - e.g., k-NN, Decision trees
 - Parametric. Parameters are finite-dimensional
 - e.g., linear regression, deep learning

Properties

- Computation. K-NN is **difficult to scale up** to large datasets
 - Pros. No training cost
 - Cons. High inference cost
 - For testing, we need to conduct n comparisons
 - Fortunately, there are many techniques to relieve this
 - Used in modern LLMs with RAG

Limitation

- The success depends critically on the **similarity metric**
 - The similarity should represent some semantic knowledge
 - From human
 - From data
 - We'll see later how neural nets can do this



Naïve Bayes

Setup

- Suppose that we have a labeled dataset

$$\{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^n$$

- $\mathbf{x}^{(i)} \in \mathbb{R}^d$
- $y^{(i)} \in \{0, 1\}$
- The data is assumed to have been independently drawn from P_{XY}

Setup


- We assume that entries of each \mathbf{x} are **conditionally independent** given y

$$p(\mathbf{x} \mid y) = \prod_{i=1}^d p(x_i \mid y)$$

- Can be true for tabular data, but not for images (thus naïve)
- From now on, we let $d = 1$, without loss of generality

Bayesian approach

- Based on some human knowledge, we **manually design** two things:
 - Likelihood model $p(x | y)$
 - Prior $p(y)$
- Example. We may have a good physical model of the channel output (x) given the channel input (y)



Training

- Estimating parameters of $p(x | y), p(y)$ from data
- Example. Gaussian likelihood has **four parameters**
 - Mean and variance, for each y

$$p(x | y) = \frac{1}{\sigma_y \sqrt{2\pi}} \exp\left(-\frac{(x - \mu_y)^2}{2\sigma_y^2}\right)$$
$$\theta_l = (\mu_0, \mu_1, \sigma_0, \sigma_1) \in \mathbb{R}^4$$

- Example. Bernoulli prior has **one parameter**

$$\theta_p = p(1) \in [0, 1]$$

Training

- To fit the parameters, we **maximize the joint probability** of the training data given the parameters

$$\begin{aligned} & \max_{\theta} p_{\theta}(\mathbf{x}_1, \dots, \mathbf{x}_n, y_1, \dots, y_n) \\ = & \max_{\theta_\ell, \theta_p} \prod_{i=1}^n p_{\theta_\ell}(\mathbf{x}_i | y_i) p_{\theta_p}(y_i) \end{aligned}$$

- Note. As we have seen last week, this has an ERM interpretation

Training

- We can solve two sub-problems separately

$$\min_{\theta_\ell} \sum_{i=1}^n \left(-\log p_{\theta_\ell}(\mathbf{x}_i | y_i) \right)$$
$$\min_{\theta_p} \sum_{i=1}^n \left(-\log p_{\theta_p}(y_i) \right)$$

- The solution to the upper optimization problem is what we call the **maximum-likelihood estimate (MLE)**

Training

- Example. Consider the subproblem for Gaussian likelihood:

$$\begin{aligned} & \min_{\theta_\ell} \sum_{i=1}^n \left(-\log p_{\theta_\ell}(\mathbf{x}_i | y_i) \right) \\ \Leftrightarrow & \min_{\theta_\ell} \left(\sum_{i=1}^n \frac{\|\mathbf{x}_i - \mu_{(y_i)}\|^2}{2\sigma_{(y_i)}^2} + \log(\sigma_{(y_i)}) \right) \end{aligned}$$

- Solving this optimization will give **class-wise sample mean** and **class-wise sample variance** (check!)

Training

- Example. Consider the subproblem for Bernoulli prior

$$\begin{aligned} & \min_{\theta_p} \sum_{i=1}^n \left(-\log p_{\theta_p}(y_i) \right) \\ \Leftrightarrow & \min_{\theta_p} \left(\sum_{i:y_i=1} -\log(\theta_p) + \sum_{i:y_i=0} -\log(1 - \theta_p) \right) \end{aligned}$$

- Solving this optimization will give the **sample frequency**

$$\theta_p = \frac{\#1s \text{ in dataset}}{n}$$

Inference

- We conduct MAP estimation

$$\begin{aligned} f(\mathbf{x}) &= \arg \max_y p(y \mid \mathbf{x}) \\ &= \arg \max_y p(y) p(\mathbf{x} \mid y) \\ &= \arg \max_y \left(p(y) \prod_{i=1}^d p(x_i \mid y) \right) \end{aligned}$$

Properties

- **Computation.** Quite simple for popular choice of $p(\mathbf{x} | y)$ and $p(y)$
 - Training. Already known, explicit formula
 - Inference. Simply compute $p(y | \mathbf{x})$
- However, these can be very messy for atypical models & priors
 - or if there is any dependency structure

Limitation

- Requires a well-designed prior and likelihood
 - We expect very complicated $p(\mathbf{x} \mid y)$ for, e.g., visual data
 - We want an automated mechanism to design these as well

Perceptrons

Historical bits

- The first “neural network” designed by Rosenblatt (1958)






FIGURE 5
DESIGN OF TYPICAL UNITS




Linear model

- Perceptron is a method to train a **linear classifier**
 - Linear classifier is about drawing a linear decision boundary

$$\mathbf{w}^T \mathbf{x} + b = 0$$

- This divides two regions:
 - $\{\mathbf{x} \mid \mathbf{w}^T \mathbf{x} + b > 0\}$
 - $\{\mathbf{x} \mid \mathbf{w}^T \mathbf{x} + b < 0\}$




Inference

- For inference, we use the **sign** of linear models

$$f_{\theta}(\mathbf{x}) = \mathbf{1}\{\mathbf{w}^T \mathbf{x} + b > 0\}$$

- Again, by stacking, we can write more neatly as

$$f_{\theta}(\mathbf{x}) = \mathbf{1}\{\boldsymbol{\theta}^T \tilde{\mathbf{x}} > 0\}$$



Training

- The most standard way to find a linear classifier would be to solve:

$$\min_{\theta} \frac{1}{n} \sum_{i=1}^n \mathbf{1}\{f_{\theta}(\mathbf{x}_i) \neq y_i\}$$

- Or more neatly, we can write as:

$$\min_{\theta} \frac{1}{n} \sum_{i=1}^n \left(f_{\theta}(\mathbf{x}_i)(1 - y_i) + (1 - f_{\theta}(\mathbf{x}_i))y_i \right)$$

Training

$$\min_{\theta} \frac{1}{n} \sum_{i=1}^n \left(f_{\theta}(\mathbf{x}_i)(1 - y_i) + (1 - f_{\theta}(\mathbf{x}_i))y_i \right)$$

- **Problem.** Difficult to optimize
 - Explicit solution – not available
 - Gradient descent – difficult to evaluate gradient
 - $f_{\theta}(\cdot)$ contains $\mathbf{1}\{\cdot\}$ – gradient is zero almost everywhere

Training

Rosenblatt's solution.


- Instead of the loss

$$\ell(y, f_\theta(\mathbf{x})) = f_\theta(\mathbf{x})(1 - y) + (1 - f_\theta(\mathbf{x}))y$$

use this loss instead:

$$\ell(y, f_\theta(\mathbf{x})) = (f_\theta(\mathbf{x}) - y) \cdot \theta^\top \tilde{\mathbf{x}}$$

- When wrong, the loss is: $|\theta^\top \tilde{\mathbf{x}}|$
- When correct, the loss is: 0
- Intuition. We penalize the “confidence” of misprediction



Training

$$\ell(y, f_\theta(\mathbf{x})) = (f_\theta(\mathbf{x}) - y) \cdot \theta^\top \tilde{\mathbf{x}}$$

- With this new loss, suddenly the gradient is non-zero

$$\nabla_\theta \ell(y, f_\theta(\mathbf{x})) = (f_\theta(\mathbf{x}) - y) \tilde{\mathbf{x}}$$

- The loss like this – not a true loss but helps optimization – is called **surrogate loss**

Optimization

- The original perceptron paper assumes that:
 - the data comes one-by-one
 - we cannot re-use the past data
- Such scenario is called **online learning**

Optimization

- Given a sample, the gradient is:

$$\nabla_{\theta} \ell(y, f_{\theta}(\mathbf{x})) = (f_{\theta}(\mathbf{x}) - y) \tilde{\mathbf{x}}$$

- If wrong for a sample with $y = 1$:

$$\theta^{(i+1)} = \theta^{(i)} + \eta \cdot \tilde{\mathbf{x}}$$


- If wrong for a sample with $y = 0$:

$$\theta^{(i+1)} = \theta^{(i)} - \eta \cdot \tilde{\mathbf{x}}$$

- If correct, no change


Properties

- Computation. Quite easy
 - Training. Simply add or subtract data x
 - Also, provably converges whenever the data is separable
 - Inference. Simply do a dot product



Limitations

- Cannot achieve low training loss on not linearly separable data



Logistic Regression

Logistic Regression

- Another popular version of the linear classifier

$$f_{\theta}(\mathbf{x}) = \mathbf{1}\{\boldsymbol{\theta}^T \tilde{\mathbf{x}} \geq 0\}$$

- Unlike Rosenblatt, logistic regression interprets $\boldsymbol{\theta}^T \tilde{\mathbf{x}}$ as a **log-likelihood ratio** of the model's internal probability estimate

$$\log \left(\frac{p(y=1 | \mathbf{x})}{p(y=0 | \mathbf{x})} \right) \approx \boldsymbol{\theta}^T \tilde{\mathbf{x}}$$

- Brainteaser. Why not interpret as $p(y=1 | \mathbf{x})$?


Logistic Regression

$$\log \left(\frac{p(y = 1 | \mathbf{x})}{p(y = 0 | \mathbf{x})} \right) \approx \boldsymbol{\theta}^\top \tilde{\mathbf{x}}$$

- In other words, we are modeling the posterior distribution as

$$p(y = 1 | \mathbf{x}) = \frac{1}{1 + \exp(-\boldsymbol{\theta}^\top \tilde{\mathbf{x}})}$$

- The function $\sigma(t) = 1/(1 + \exp(-t))$ is the **logistic function**
(a.k.a. sigmoid)



Training

- Given the data, maximize the log-likelihood

$$\max_{\theta} \frac{1}{n} \sum_{i=1}^n \log p(y_i \mid \mathbf{x}_i)$$

- Equivalently, minimize the NLL loss:

$$\min_{\theta} -\frac{1}{n} \sum_{i=1}^n \log \left(\frac{1}{p(y_i \mid \mathbf{x}_i)} \right)$$

Training

- Equivalently again, we are solving:

$$\min_{\theta} \frac{1}{n} \sum_{i=1}^n \ell(y_i, f_{\theta}(\mathbf{x}_i))$$

where

- $f_{\theta}(\cdot)$ is the **sigmoid** of the prediction

$$f_{\theta}(\mathbf{x}) = \sigma(\boldsymbol{\theta}^\top \tilde{\mathbf{x}})$$

- $\ell(\cdot, \cdot)$ is the **cross-entropy**

$$\ell(y, t) = \text{CE}(\mathbf{1}_y, [t, 1-t]) = \log(t)^{-y} + \log(1-t)^{y-1}$$

Training

- More tediously, this can be written as

$$\min_{\theta} \frac{1}{n} \sum_{i=1}^n (-y_i) \log(\sigma(\theta^\top \tilde{\mathbf{x}}_i)) + (y_i - 1) \log(1 - \sigma(\theta^\top \tilde{\mathbf{x}}_i))$$

- No analytic solution, but is convex and can use GD


$$\theta^{(\text{new})} = \theta + \eta \cdot \frac{1}{n} \sum_{i=1}^n (y_i - \sigma(\theta^\top \tilde{\mathbf{x}}_i)) \tilde{\mathbf{x}}_i$$

Properties

- **Computation.** Relatively easy
 - Training. Requires GD, but is convex
 - Inference. Easy – Dot product and apply threshold

Limitation

- Again, cannot fit not-linearly-separable data



Next class

- Sophisticated versions of linear classifiers

</lecture 4>