

# **Dimensionality Reduction**

# Recap

- **Unsupervised learning**
  - Learning from unlabeled data  $\{\mathbf{x}_1, \dots, \mathbf{x}_n\}_{i=1}^m \subseteq \mathbb{R}^d$
  - Easy to scale up – necessary for large-scale training
- **Clustering**
  - Learning a mapping  $\Phi(\cdot) : \mathbb{R}^d \rightarrow \{1, \dots, k\}$
  - Each  $k$  may be represented by some mean  $\mu_i \in \mathbb{R}^d$   
(and variance, and so on ...)
  - K-Means
  - Gaussian Mixture Models

# Today

- **Dimensionality Reduction**
  - Learning a mapping  $\Phi(\cdot) : \mathbb{R}^d \rightarrow \mathbb{R}^k$  ( $k < d$ )
  - In particular, we focus on the case of **linear**  $\Phi(\cdot)$ 
    - Precisely, we discuss **Principal Component Analysis (PCA)**
  - Other examples
    - ICA (Independent Component Analysis)
    - Autoencoders

# Motivations


# Dealing with high-dimensional data

- Many datasets are extremely high-dimensional, in its raw form
- **Example.** Suppose you are an ML engineer at Google
  - Goal. A model that detect **copyrighted clips** from **Youtube shorts**
  - The dimensionality of Youtube shorts  $\mathbf{x} \in \mathbb{R}^d$  are:  
 $1920 \times 1080 \times \text{RGB} \times 60\text{FPS} \times 60 \text{ Seconds}$   
=22.4 Billion dimension




# Curse of dimensionality

- Learning from high-dimensional data is challenging
  - Computation
  - Higher chance of noise
  - Difficult to visualize – for human insights
  - Difficult to find generalizable patterns (**important**)




# Nominal dimensionality vs. True

- But do we really need all these dimensions?
- **Example.** Handwritten digit recognition (MNIST, 28x28)



only looks like this




... and not like this

- That is, we are not **fully utilizing**  $\mathbb{R}^{28 \times 28} = \mathbb{R}^{784}$

# Nominal dimensionality vs. True


- **Hypothesis.**  
There exists some **low-dim. subspace** (or submanifold) in the high-dim. space where the real data lies in
- **Dimensionality Reduction**  
Using **unlabeled data** to find the right mapping b/w high-dim & low-dim spaces
  - Caveat. Data could be noisy



# **Principal Component Analysis**


# Overview

- A dimensionality reduction technique, invented by Karl Pearson (1909)
  - Uses an **affine subspace** of the original space
  - Many aliases – e.g., Karhunen-Loève Transform




# Motivating PCA: Toy Example

- Suppose that we are given a 2D dataset
- **Goal.** Find a nice **1d subspace** and the corresponding **mappings**, such that the mapped data have **desirable properties**




# Motivating PCA: Toy Example

- Let's simplify a bit
  - We confine the mapping to be an **orthogonal projection**
  - Given a **subspace**, the mapping is uniquely determined.




# Motivating PCA: Toy Example

- Goal (restated). Find a nice **1D subspace** such that the projected data have **desirable properties**
  - Exactly what properties do we need?




# Motivating PCA: Toy Example

- **Answer.** Preserve **task-relevant information** as much as possible
  - However, this is a difficult task
    - task-relevance: no label given to us!
    - information: usual metrics, e.g., entropy is hard to estimate
  - **Simpler approach.** Which projection is more informative?



# Motivating PCA: Toy Example

- **Answer.** Left is considered informative, for two reasons
  - (A) Projected points are more **well-spread**
    - Does not ignore differences b/w points
    - Noise-robust
  - (B) Projected points (●) are **closer** to their original data (○)
    - That is, more accurate reconstruction is possible




# Motivating PCA: Toy Example

- **Answer.** Left is considered informative, for two reasons
  - (A) Projected points are more well-spread
    - Does not ignore differences b/w points
    - Noise-robust
  - (B) Projected points (●) are closer to their original data (●)
    - That is, more accurate reconstruction is possible
- Interestingly, these two criteria are **equivalent!**

# Key Result

- We are given a dataset  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$
- **Goal.** Find a  $k$ -dimensional subset  $\mathbf{U} \subseteq \mathbb{R}^d$  with
  - (A) Maximum **variance** of projected points
$$\max_{\mathbf{U}} \text{Var}(\pi_{\mathbf{U}}(\mathbf{x}_1), \dots, \pi_{\mathbf{U}}(\mathbf{x}_n))$$
  - (B) Minimum  $\ell^2$  **distortion** from projection
$$\min_{\mathbf{U}} \sum_{i=1}^n \|\mathbf{x}_i - \pi_{\mathbf{U}}(\mathbf{x}_i)\|_2^2$$
- But first, let's formally define what “projection” is...




# **Formalisms: Projection**

# Formalisms


- A  $k$ -dimensional affine subspace  $U \subset \mathbb{R}^d$  can be characterized by:
  - Orthonormal basis  $\mathbf{u}_1, \dots, \mathbf{u}_k \in \mathbb{R}^d$
  - Orthogonal bias  $\mathbf{b} \in \mathbb{R}^d$

$$U = \{a_1\mathbf{u}_1 + \cdots + a_k\mathbf{u}_k + \mathbf{b} : a_i \in \mathbb{R}\}$$



# Formalisms

- Any element on  $U$  can be represented in two ways:
  - A  $d$ -dimensional vector  $\mathbf{u} \in U$
  - A  $k$ -dimensional vector  $\mathbf{a} = (a_1, a_2, \dots, a_k)$ 
    - where  $\mathbf{u} = a_1\mathbf{u}_1 + \dots + a_k\mathbf{u}_k + \mathbf{b}$  holds




# Formalisms

- A **projection** of a vector  $\mathbf{x} \in \mathbb{R}^d$  to the affine subspace  $U$  is:

$$\pi_U(\mathbf{x}) = \sum_{i=1}^k (\mathbf{u}_i^\top \mathbf{x}) \cdot \mathbf{u}_i + \mathbf{b}$$

- This is a d-dimensional quantity, with an alternative representation:

$$\mathbf{a} = (\mathbf{u}_1^\top \mathbf{x}, \dots, \mathbf{u}_k^\top \mathbf{x}) \in \mathbb{R}^k$$



# Formalisms

- The projection admits a matrix form:

$$\begin{aligned}\pi_U(\mathbf{x}) &= \left( \sum_{i=1}^k \mathbf{u}_i \mathbf{u}_i^\top \right) \mathbf{x} + \mathbf{b} \\ &=: \mathbf{U}\mathbf{x} + \mathbf{b}\end{aligned}$$

- Here, the projection matrix  $\mathbf{U}$  is:

- $d \times d$  matrix with rank  $k$
- $\mathbf{U}^\top = \mathbf{U}$
- $\mathbf{U}^\top \mathbf{U} = \mathbf{U}$
- Conversely, called projection matrix if these are satisfied

# Formalisms

- In a sense, projection consists of two operations
- **Compression**  $\mathbb{R}^d \rightarrow \mathbb{R}^k$ 
  - Also known as “encoding”

$$\mathbf{z} = \mathbf{U}_{\text{enc}} \mathbf{x},$$


where

$$\mathbf{U}_{\text{enc}} = \begin{bmatrix} \leftarrow & \mathbf{u}_1^\top & \rightarrow \\ & \cdots & \\ \leftarrow & \mathbf{u}_k^\top & \rightarrow \end{bmatrix} \in \mathbb{R}^{k \times d}$$

- **Reconstruction**  $\mathbb{R}^k \rightarrow \mathbb{R}^d$ 
  - Also known as “decoding”

$$\hat{\mathbf{x}} = \mathbf{U}_{\text{dec}} \mathbf{z} + \mathbf{b},$$

where  $\mathbf{U}_{\text{dec}} = \mathbf{U}_{\text{enc}}^\top \in \mathbb{R}^{d \times k}$



# **PCA: Variance Maximization**

# Variance Maximization

- In PCA, we want to find a nice  $(\mathbf{U}, \mathbf{b})$  which solves

$$\max_{\mathbf{U}, \mathbf{b}} \text{Var}(\mathbf{Ux}_1 + \mathbf{b}, \dots, \mathbf{Ux}_n + \mathbf{b})$$

- As the constant term does not affect the variance, this is equivalent to

$$\max_{\mathbf{U}} \text{Var}(\mathbf{Ux}_1, \dots, \mathbf{Ux}_n)$$

# Variance Maximization

- Define  $\bar{\mathbf{x}}$  as the mean of  $\{\mathbf{x}_i\}_{i=1}^n$
- Then, the variance can be written as:

$$\begin{aligned}\text{Var}(\mathbf{Ux}_1, \dots, \mathbf{Ux}_n) &= \frac{1}{n} \sum_{i=1}^n \|\mathbf{U}(\mathbf{x}_i - \bar{\mathbf{x}})\|_2^2 \\ &= \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \bar{\mathbf{x}})^T \mathbf{U}^T \mathbf{U} (\mathbf{x}_i - \bar{\mathbf{x}}) \\ &= \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \bar{\mathbf{x}})^T \mathbf{U} (\mathbf{x}_i - \bar{\mathbf{x}})\end{aligned}$$

# Variance Maximization

$$\max_{\mathbf{U}} \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \bar{\mathbf{x}})^\top \mathbf{U} (\mathbf{x}_i - \bar{\mathbf{x}})$$

- By the definition of  $\mathbf{U}$ , we can re-write the above as

$$\begin{aligned} & \max_{\mathbf{U}} \frac{1}{n} \sum_{i=1}^n \sum_{j=1}^k (\mathbf{x}_i - \bar{\mathbf{x}})^\top \mathbf{u}_j \mathbf{u}_j^\top (\mathbf{x}_i - \bar{\mathbf{x}}) \\ &= \max_{\mathbf{U}} \sum_{j=1}^k \mathbf{u}_j^\top \left( \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \bar{\mathbf{x}})(\mathbf{x}_i - \bar{\mathbf{x}})^\top \right) \mathbf{u}_j \end{aligned}$$

= sample covariance matrix  $\mathbf{S}$   
(positive-semidefinite)

# Variance Maximization

- Thus, PCA is about solving the **constrained quadratic optimization**

$$\max_{\mathbf{u}_1, \dots, \mathbf{u}_k} \sum_{j=1}^k \mathbf{u}_j^\top \mathbf{S} \mathbf{u}_j, \quad \text{subject to} \quad \mathbf{u}_i^\top \mathbf{u}_j = \begin{cases} 1 & \dots \quad i = j \\ 0 & \dots \quad i \neq j \end{cases}$$

- **Question.** How do we solve this?

# Solving the quadratic problem

$$\max_{\mathbf{u}_1, \dots, \mathbf{u}_k} \sum_{j=1}^k \mathbf{u}_j^\top \mathbf{S} \mathbf{u}_j, \quad \text{subject to} \quad \mathbf{u}_i^\top \mathbf{u}_j = 1 \{i = j\}$$

- **Answer.** Of course, the method of Lagrangian multipliers
  - Standard derivation requires complicated matrix derivatives – instead, will give you a simplified proof idea.
- **Strategy.** Conduct a greedy optimization
  - Select a nice  $\mathbf{u}_1$  that maximizes  $\mathbf{u}_1^\top \mathbf{S} \mathbf{u}_1$  s.t.  $\mathbf{u}_1^\top \mathbf{u}_1 = 1$
  - Select a nice  $\mathbf{u}_2$  that maximizes  $\mathbf{u}_2^\top \mathbf{S} \mathbf{u}_2$  s.t.  $\mathbf{u}_2^\top \mathbf{u}_2 = 1, \mathbf{u}_2^\top \mathbf{u}_1 = 0$
  - ...

# Solving the quadratic problem

- First step is to determine  $\mathbf{u}_1$

$$\max_{\mathbf{u}} \mathbf{u}^\top \mathbf{S} \mathbf{u}, \quad \text{subject to} \quad \mathbf{u}^\top \mathbf{u} = 1$$

- To solve this, consider the Lagrangian relaxation

$$\max_{\mathbf{u}} \mathbf{u}^\top \mathbf{S} \mathbf{u} + \alpha(1 - \mathbf{u}^\top \mathbf{u})$$

- Critical point is where  $\mathbf{S} \mathbf{u} = \alpha \mathbf{u}$  holds
  - i.e., eigenvectors
  - Choose the **principal component** – i.e., eigenvector w/ maximum eigenvalue – to maximize the value of  $\mathbf{u}^\top \mathbf{S} \mathbf{u}$

# Solving the quadratic problem

- Next, we determine  $\mathbf{u}_2$

$$\max_{\mathbf{u}} \mathbf{u}^\top \mathbf{S} \mathbf{u}, \quad \text{subject to} \quad \mathbf{u}^\top \mathbf{u} = 1, \mathbf{u}^\top \mathbf{u}_1 = 0$$

- Lagrangian relaxation becomes

$$\mathbf{u}^\top \mathbf{S} \mathbf{u} + \alpha(1 - \mathbf{u}^\top \mathbf{u}) - \beta(\mathbf{u}^\top \mathbf{u}_1)$$

- The critical point condition is:

$$\mathbf{S} \mathbf{u} = \alpha \mathbf{u} + \frac{\beta}{2} \mathbf{u}_1$$

# Solving the quadratic problem

$$\mathbf{S}\mathbf{u} = \alpha\mathbf{u} + \frac{\beta}{2}\mathbf{u}_1$$

- Multiplying  $\mathbf{u}_1^\top$  on both sides, we get:

$$0 = 0 + \frac{\beta}{2}$$

- Thus, we have  $\beta = 0$
- Then, the Lagrangian becomes

$$\mathbf{u}^\top \mathbf{S}\mathbf{u} + \alpha(1 - \mathbf{u}^\top \mathbf{u})$$

- Thus the things are the same as in the derivation of  $\mathbf{u}_1$ 
  - Thus, choose the eigenvector for **2nd largest eigenvalue**

# Solving the quadratic problem

- Repeat this, the solution is to let  $\mathbf{u}_1, \dots, \mathbf{u}_k$  be the **top-k principal components** of our sample covariance matrix

- This can be done by performing SVD on the **data matrix**

$$\mathbf{X} = [\mathbf{x}_1 - \bar{\mathbf{x}} \mid \dots \mid \mathbf{x}_n - \bar{\mathbf{x}}] = \mathbf{U}\Sigma\mathbf{V}^T$$

and then selecting the columns of  $\mathbf{U}$  corresponding to top-k singular values


- **Note.** Did not cover determining  $\mathbf{b}$  – will be covered soone

# **PCA: Distortion Minimization**

# Distortion Minimization

- Here is the spirit:

“If the projected point is close to the original point,  
then we did not lose too much information”
- We’ll show that this distortion minimization = variance maximization




# Distortion Minimization

- Formally, we try to find an **affine subspace**

$$U = \{a_1\mathbf{u}_1 + \cdots + a_k\mathbf{u}_k + \mathbf{b} : a_i \in \mathbb{R}\}$$

such that the **mean squared error** of data from projection is minimized

$$\min_U \frac{1}{n} \sum_{i=1}^n \|\mathbf{x}_i - \pi_U(\mathbf{x}_i)\|^2$$



# Distortion Minimization

- Using the definition of projection, we know that

$$\begin{aligned} & \frac{1}{n} \sum_{i=1}^n \|\mathbf{x}_i - \pi_U(\mathbf{x}_i)\|^2 \\ &= \frac{1}{n} \sum_{i=1}^n \|\mathbf{x}_i - \mathbf{U}\mathbf{x}_i - \mathbf{b}\|^2 \\ &= \frac{1}{n} \sum_{i=1}^n (\|\mathbf{x}_i\|^2 + \|\mathbf{b}\|^2 - \mathbf{x}_i^\top \mathbf{U} \mathbf{x}_i - 2\mathbf{b}^\top \mathbf{x}_i + 2\mathbf{b}^\top \mathbf{U} \mathbf{x}_i) \\ &= \frac{1}{n} \left( \sum_{i=1}^n \|\mathbf{x}_i\|^2 \right) + \|\mathbf{b}\|^2 - \frac{1}{n} \left( \sum_{i=1}^n \mathbf{x}_i^\top \mathbf{U} \mathbf{x}_i \right) - 2\mathbf{b}^\top \bar{\mathbf{x}} + 2\mathbf{b}^\top \mathbf{U} \bar{\mathbf{x}} \end{aligned}$$

# Distortion Minimization

- Removing the irrelevant terms, we are solving:

$$\min_{\mathbf{U}, \mathbf{b}} \left( \|\mathbf{b}\|^2 - \frac{1}{n} \sum \mathbf{x}_i^\top \mathbf{U} \mathbf{x}_i - 2\mathbf{b}^\top \bar{\mathbf{x}} + 2\mathbf{b}^\top \mathbf{U} \bar{\mathbf{x}} \right)$$

- For any fixed  $\mathbf{U}$ , we have

$$\mathbf{b}^* = \bar{\mathbf{x}} - \mathbf{U} \bar{\mathbf{x}}$$

- Plugging in and removing constant terms again, we get:

$$\min_{\mathbf{U}} \left( \bar{\mathbf{x}}^\top \mathbf{U} \bar{\mathbf{x}} - \frac{1}{n} \sum \mathbf{x}_i^\top \mathbf{U} \mathbf{x}_i \right) = - \max_{\mathbf{U}} \left( \sum_{j=1}^k \mathbf{u}_j^\top \mathbf{S} \mathbf{u}_j \right)$$

# **Applications & Limitations**

# Face Recognition

- Many applications, but here's an interesting one: **Eigenface (1991)**
- **Goal.** Identify specific person, based on facial image
  - Robust to glass, lightning, ...
  - Using  $256 \times 256$  is difficult!



# Face Recognition

- **Idea.** Build a PCA database for whole dataset
  - Each  $\mathbf{u}_i$  can capture some “feature”
  - Classify based on  $(\mathbf{u}_1^\top \mathbf{x}, \dots, \mathbf{u}_k^\top \mathbf{x})$ 
    - Rapid recognition
    - Tracking
- **Limitations.**
  - Requires the same size
  - Sensitive to angles
  - Needs “centering”



# Image Compression

- **Goal.** Represent an image using less dimensions
- **Idea.** Do the following:
  - Divide each image in  $12 \times 12$  patches
  - Conduct PCA
  - For each patch, save K digits ( $\mathbf{u}_1^T \mathbf{x}, \dots, \mathbf{u}_k^T \mathbf{x}$ )




144-dimension  
(full)



60-dimension




6-dimension




1-dimension

# Image Compression

- Interestingly, the eigenvectors look similar to cosine transforms (DCT)
  - A version using DCT is called JPEG







Eigenvectors



DCT bases


# Limitations

- Difficult to capture nonlinear dataset
- Does not account for class labels




# Advanced methods


- **Kernel PCA.** Conduct PCA for  $\Phi(\mathbf{x})$ 
  - Requires careful hyperparameter tuning & validation



Spherical Data




No Kernel



Gaussian Kernel ( $\sigma = 20$ )

# Isomap

- Similarly to spectral clustering, build a **graph of points** by connecting each point to  $k$ -nearest neighbors
- Then, find a mapping to a low-dimensional space such that:  
 $\text{distance on graph} \approx \text{distance on embedded space}$





# t-SNE

- Similar to Isomap, but use the neighborhood information

$$p_i(j) = \frac{\exp(-\|\mathbf{x}_i - \mathbf{x}_j\|^2/2\sigma^2)}{\sum_{k \neq i} \exp(-\|\mathbf{x}_i - \mathbf{x}_k\|^2/2\sigma^2)}$$

- Find a low-dimensional embedding such that  $\text{dist}(p_i, p_j) \approx \text{dist}(\mathbf{z}_i, \mathbf{z}_j)$



MNIST embeddings of t-SNE  
(requires computing pairwise  
distances of 60,000 samples)

# Next up

- Decision trees

**</lecture 9>**