

Decision Trees

Motivation

- **Kaggle.** A competition platform for ML and data science
 - People upload data and put bounty to it
 - You solve it

The screenshot shows the Kaggle homepage with a sidebar on the left and a main content area on the right.

Left Sidebar:

- ≡ kaggle
- + Create
- Home
- Competitions
- Datasets
- Models
- <> Code
- Discussions
- Learn
- More

Main Content Area:

Search competitions Filters

All competitions Featured Getting Started Research Community Playground Simulations Analytics

Hotness ▾


Active Competitions

Competition	Bounty	Time Left
Open Problems – Single-Cell Perturbations	\$100,000	2 months to go
Stanford Ribonanza RNA Folding	\$100,000	2 months to go
Optiver - Trading at the Close	\$100,000	2 months to go
CommonLit - Evaluate Student Summaries	\$60,000	4 days to go

Detailed description of the competitions shown:

- Open Problems – Single-Cell Perturbations:** Predict how small molecules change gene...
Featured
431 Teams
\$100,000 2 months to go
- Stanford Ribonanza RNA Folding:** Create a model that predicts the structur...
Research
262 Teams
\$100,000 2 months to go
- Optiver - Trading at the Close:** Predict US stocks closing movements
Featured · Code Competition
1008 Teams
\$100,000 2 months to go
- CommonLit - Evaluate Student Summaries:** Automatically assess summaries written b...
Featured · Code Competition
2044 Teams
\$60,000 4 days to go

Kaggle Survey (2021)



Historical Bits

Historical bits


- Use in modern ML traces back to Morgan & Sonquist (1963)
 - Analyzing survey data on income & savings
 - Data included many demographic subgroups
 - Turned out that the trend was **highly nonlinear**

TABLE 1. SPENDING UNIT INCOME AND THE NUMBER IN THE
UNIT WITHIN VARIOUS SUBGROUPS

Group	Spending unit average (1958) income	Number in unit	Number of cases
Nonwhite, did not finish high school	\$ 2489	3.3	191
Nonwhite, did finish high school	5005	3.4	67
White, retired, did not finish high school	2217	1.7	272
White, retired, did finish high school	4520	1.7	72
White, nonretired farmers, did not finish high school	3950	3.6	87
White nonretired farmers, did finish high school	6750	3.6	24


Historical bits

- Case 1. Multi-collinearity
 - Correlation between income & education, but no interaction




Historical bits

- Case 2. Interaction between features
 - No correlation between income & self-employment




Historical bits

- Case 3. Both



Historical bits

- In each of these examples, having a single linear model doesn't work well
- **Idea.** Take a sequential approach
 - Divide. Partition the data into many subgroups
 - Conquer. Have a simple model for each subgroup (e.g., linear)
- **Example.** High asset?
 - Yes → use **curve 1**
 - No → use **curve 2**







CHART II. Annual Earnings.

Key question


- How do we know if a subgroup needs division?
 - If we know, exactly how do we divide?



Decision Trees

Overview

- Basically a **nested if-else** statement
- A binary tree which recursively partitions and refines the input space
 - **Leaf.** Associated with some **label** \hat{y}
 - If discrete, classification
 - If continuous, regression
 - **Tree.** Associated with some **splitting rule** $g : \mathcal{X} \rightarrow \{0,1\}$



Inference

- Given **x**, recurse down the tree until a leaf is reached
 - Then, output the label of the leaf


```
while(true):  
    if(node == leaf): output label(node)  
    else:  
        if(condition == true): node = right_child(node)  
        else: node = left_child(node)
```

Inference

- When $\mathcal{X} = \mathbb{R}^d$, it is typical to consider only the **axis-aligned splits**

$$g(\mathbf{x}) = \mathbf{1}[x_i \geq t]$$

- Computationally efficient
 - Single index lookup
- Human-interpretable decisions



Training

- Constructing a decision tree requires specifying three elements
 - Prediction rule
 - Stopping rule
 - Splitting rule

until all leaf node is stopped:

visit a leaf node

if(**stopping_rule**(node) = True):


 apply **prediction rule** to label the node
 stop the node

else:

 split the node, using the **splitting rule**

Example: Iris Classification


- For example, consider an iris classification task
 - Input features $\mathcal{X} = \mathbb{R}^2$
 - x_1 : length-width ratio of sepal
 - x_2 : length-width ratio of petal
 - Output labels $\mathcal{Y} = \{1, 2, 3\}$



Example: Iris Classification

- First, construct a single leaf node, using a **prediction rule** that applies to the whole set


$$\hat{y} = 2$$



Example: Iris Classification



- See if the stopping rule is met
 - Very unhomogeneous; continue

$\hat{y} = 2$





Example: Iris Classification

- According to the **splitting rule**, split the leaf to partition the input space for the node





Example: Iris Classification

- According to the **prediction rule**, determine the prediction of new leaf nodes



Example: Iris Classification


- Continue until some **stopping rule** is satisfied for all leaf nodes



Elements


Training: Prediction Rule

- Generating a label for a partitioned set
- Typically very simple
 - Classification. Majority voting
 - Regression. Average, Median, ...



Training: Splitting Rule

- Generating how to partition a set
 - Which axis?
 - Which line?



Training: Splitting Rule

- **Idea.** Minimize some notion of **uncertainty** (a.k.a. impurities) after partitioning the set
- In other words, by dividing some set S into S_1, S_2 , we want to solve:

$$\min_{S_1, S_2 : S_1 \cup S_2 = S, S_1 \cap S_2 = \emptyset} \left(|S_1| \cdot u(S_1) + |S_2| \cdot u(S_2) \right)$$

- Here, $u(\cdot)$ is some measure of uncertainty

Training: Splitting Rule

Example (Binary Classification)

- Suppose that we are given a set S , with $p |S|$ samples labeled as +1
 - Classification Error

$$u(S) = \min\{p, 1 - p\}$$


- Gini Index

$$u(S) = 2p(1 - p)$$

- Entropy

$$u(S) = p \log \frac{1}{p} + (1 - p) \log \frac{1}{1 - p}$$

(G, E are concave upper bounds on C)



Training: Splitting Rule

Example (**Regression**)


- We can simply use variance
 - the minimum mean squared error
 - i.e., the ℓ^2 error of the mean
- Similarly, we can use the minimum mean absolute error, ...

Training: Stopping Rule

- Determining when to stop growing a tree


- Many criteria:

- If splitting does not reduce the uncertainty
- Reaches some pre-specified size of the tree
- Every leaf is “pure”
 - Very prone to overfitting




Pruning

- It is typical to **prune** the tree after growing
 - i.e., remove unnecessary split after training



Pruning

- **Algorithm.**
 - Pick a bottom-level split
 - Remove it
 - If the validation error is improved, leave it pruned
 - Else, restore the subtree
 - Repeat




Pruning

- Note that the iterative algorithm is a “greedy” way to minimize the **total uncertainty**

$$u(\mathcal{T}) := \frac{1}{n} \sum_{\text{leaf } S \in \mathcal{T}} |S| \cdot u(S)$$

- Prone to falling in local minima:
 - Fails on XOR (indifferent to splits)
- **Solution.**
 - Do “random” splits occasionally
 - Then, prune the unnecessary splits



Properties

- **Advantages.**
 - Easy to interpret
 - Fast to execute
- **Limitations.**
 - Difficult to scale up
 - Easy to overfit, if the tree is big

Properties

- Nonparametric
- Based on local regularity
 - Simple locally, complicated globally
- In these senses, similar with nearest neighbors

Forests

Forests

- Scaling up decision trees can be done by growing **multiple trees**




Bagging

- Stands for “Bootstrapped Aggregating”

- **Idea.** Split the data to multiple subsets

- Generate a tree for each subset
- Predictions of the trees are aggregate via
 - Majority voting
 - Averaging



Random Forest


- **Problem.** Bagging leads to highly correlated trees
 - That is, resulting trees look similar to each other
- **Idea.** Decorrelate the trees by using only a **subset of features**
 - To grow each node, randomly select a subset of features and choose the best one among this subset

RANDOMFOREST(\mathcal{D} ; B, m, n)

- 1 **for** $b = 1, \dots, B$
- 2 Draw a bootstrap sample \mathcal{D}_b of size n from \mathcal{D}
- 3 Grow a tree T_b on data \mathcal{D}_b by recursively:
 - 4 Select m variables at random from the d variables
 - 5 Pick the best variable and split point among the m variables
 - 6 Split the node
- 7 **return** tree T_b

Boosting

- **Idea.** Decorrelate by sequentially generating the trees
 - Assign higher weights to samples that other trees got wrong



</lecture 10>