

Optimizing neural nets: SGD & Backpropagation

Recap: Neural networks

- Consider the case of **supervised learning** with neural nets
- We are performing the usual optimization

$$\min_{\theta} \frac{1}{n} \sum_{i=1}^n \ell(y_i, f_{\theta}(\mathbf{x}_i)) =: \min_{\theta} L(\theta)$$

- Predictor is the neural network

$$f_{\theta}(\mathbf{x}) = \mathbf{W}_L \sigma(\mathbf{W}_{L-1} \sigma(\cdots \sigma(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1) \cdots + \mathbf{b}_{L-1}) + \mathbf{b}_L$$

- Parameters are **weights & biases** of each layer


$$\theta = \{(\mathbf{W}_l, \mathbf{b}_l)\}_{l=1}^L$$

Today

- We focus on: How do we solve the **optimization** problem

$$\min_{\theta} L(\theta), \quad f_{\theta}(\mathbf{x}) = \mathbf{W}_L \sigma(\cdots \sigma(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1) \cdots + \mathbf{b}_L)$$

- This is very difficult
 - **Critical point.** Too complicated
 - **Convexity.** Does not hold
- The loss landscape looks like ->




Gradient Descent

- **Solution.** Gradient Descent

- Iteratively update θ in a direction that the loss decreases the fastest

$$\theta^{(t+1)} = \theta^{(t)} - \eta \cdot \nabla_{\theta} L(\theta)$$

Step size (a.k.a., learning rate) Direction of fastest increase



Gradient Descent

- Note that the gradient is the **average of per-sample loss gradients**:

$$\nabla_{\theta} L(\theta) = \frac{1}{n} \sum_{i=1}^n \nabla_{\theta} \ell(y_i, f_{\theta}(\mathbf{x}_i))$$


- **Problem.** Datasets for deep learning involves million–trillion-scale data
 - Examples.
 - ImageNet (Image). 1 million samples
 - Common Crawl (Text). 410 billion tokens
 - Thus, computing gradient of all data at each GD step is expensive

Gradient Descent

- **Solution.** **Stochastic** Gradient Descent (broad)
 - Use gradients of only a few, randomly drawn samples at each step
 - Mini-batch GD. Draw a batch \mathcal{B} of samples and compute



$$\hat{\nabla}_{\theta} L(\theta) = \frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \nabla_{\theta} \ell(y_i, f_{\theta}(\mathbf{x}_i))$$

- SGD (narrow). Mini-batch GD with a single example




Gradient Descent

- Typically, we draw samples without replacement
 - i.e., never use a sample twice unless no sample has been never used



Gradient Descent

- **Epoch.** A set of iterations until every sample has been used once
 - Example. If we use the batch size of 64 for a dataset of size 32,000, we need 500 steps for a single epoch
 - **Batch size** and **learning rate** are key hyperparameters of SGD



Computing per-sample Gradients

Computing Gradients

- The sample-wise loss gradient is a product of
 - (1) the derivative of the loss function, and
 - (2) the gradient w.r.t. the predictor

$$\nabla_{\theta} \left(\ell(y, f_{\theta}(\mathbf{x})) \right) = \frac{\partial \ell(y, z)}{\partial z}(f_{\theta}(\mathbf{x})) \cdot \nabla_{\theta} f_{\theta}(\mathbf{x})$$

loss derivative, evaluated at prediction $f_{\theta}(\mathbf{x})$ | Predictor gradient

- Why?** Recall the chain rule:

$$\frac{\partial}{\partial x} g(f(x)) = g'(f(x)) \cdot f'(x)$$

Computing Gradients

$$\nabla_{\theta} \left(\ell(y, f_{\theta}(\mathbf{x})) \right) = \frac{\partial \ell(y, z)}{\partial z}(f_{\theta}(\mathbf{x})) \cdot \nabla_{\theta} f_{\theta}(\mathbf{x})$$

- The **loss derivative** is typically easy to compute
- **Example.** For squared loss $\ell(y, z) = (y - z)^2$, the loss derivative will be:
$$2(y - f_{\theta}(\mathbf{x}))$$
- Simply do (1) pass the data through the predictor
(2) measure the error
(3) multiply 2

Computing Gradients

$$\nabla_{\theta} \left(\ell(y, f_{\theta}(\mathbf{x})) \right) = \frac{\partial \ell(y, z)}{\partial z}(f_{\theta}(\mathbf{x})) \cdot \nabla_{\theta} f_{\theta}(\mathbf{x})$$

- The **predictor gradient** is much trickier to compute

- The parameter θ is high-dimensional

$$\nabla_{\theta} g(\theta) = \left[\frac{\partial}{\partial \theta_1} g(\theta), \dots, \frac{\partial}{\partial \theta_d} g(\theta) \right]$$

- How do we compute this, for a very complicated function like...?

$$g(\theta) = \mathbf{W}_L \sigma(\cdots \sigma(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1) \cdots + \mathbf{b}_L)$$

Computing Gradients: Numerical Method

$$\nabla_{\theta}g(\theta) = \left[\frac{\partial}{\partial\theta_1}g(\theta), \dots, \frac{\partial}{\partial\theta_d}g(\theta) \right]$$

- One possible way is the **numerical method**

- Note that

$$\frac{\partial}{\partial x}g(x) = \lim_{\epsilon \rightarrow 0} \frac{g(x + \epsilon) - g(x)}{\epsilon}$$

- Make a very small perturbation on the current parameter
 - Do this for the first entry θ_1
 - Do this for the second entry θ_2
 - ...

Computing Gradients: Numerical Method

current W:

[0.34,
-1.11,
0.78,
0.12,
0.55,
2.81,
-3.1,
-1.5,
0.33,...]

loss 1.25347

W + h (first
dim):

[0.34 + 0.0001,
-1.11,
0.78,
0.12,
0.55,
2.81,
-3.1,
-1.5,
0.33,...]

loss 1.25322

gradient dW:

[-2.5,
?,
?,
?,
?,
?,
?,
?,
?,...]

$$(1.25322 - 1.25347)/0.0001 = -2.5$$

$$\frac{df(x)}{dx} = \lim_{h \rightarrow 0} \frac{f(x + h) - f(x)}{h}$$

Computing Gradients: Numerical Method

current W:

[0.34,
-1.11,
0.78,
0.12,
0.55,
2.81,
-3.1,
-1.5,
0.33,...]

loss 1.25347

W + h (second
dim):

[0.34,
-1.11 + 0.0001,
0.78,
0.12,
0.55,
2.81,
-3.1,
-1.5,
0.33,...]
loss 1.25353

gradient dW:

[-2.5,

0.6,

?,

?,

?,

?,

?,

?,

?,...]

$$(1.25353 - 1.25347)/0.0001 = 0.6$$

$$\frac{df(x)}{dx} = \lim_{h \rightarrow 0} \frac{f(x + h) - f(x)}{h}$$

Computing Gradients: Numerical Method

current W:

[0.34,
-1.11,
0.78,
0.12,
0.55,
2.81,
-3.1,
-1.5,
0.33,...]

loss 1.25347

W + h (third
dim):

[0.34,
-1.11,
0.78 + 0.0001,
0.12,
0.55,
2.81,
-3.1,
-1.5,
0.33,...]
loss 1.25347

gradient dW:

[-2.5,
0.6,
0,
?,
?,
?,
?,
?,
?,
?,...]

$$(1.25347 - 1.25347)/0.0001 = 0$$

$$\frac{df(x)}{dx} = \lim_{h \rightarrow 0} \frac{f(x + h) - f(x)}{h}$$

Computing Gradients: Numerical Method

- **Pros.**

- Easy to implement
- Can use for black-box models

- **Cons.**

- Only gives you approximate
 - cannot take the limit $\epsilon \rightarrow 0$, due to the finite precision
- Very slow \leftarrow
 - Requires at least $d + 1$ model inferences

Computing Gradients: Analytic Method

- The most popular method is the **analytic method**

- **Example.** Consider the function

$$g(\theta_1, \theta_2) = \sin(5 \cdot \exp(\theta_1) + 2 \cos(\theta_2))$$

- Then, we know that the gradient will have the formula:

$$\nabla_{\theta_1} g(\theta_1, \theta_2) = 5 \cdot \cos(5 \cdot \exp(\theta_1) + 2 \cdot \cos(\theta_2)) \cdot \exp(\theta_1)$$

$$\nabla_{\theta_2} g(\theta_1, \theta_2) = -2 \cdot \cos(5 \cdot \exp(\theta_1) + 2 \cdot \cos(\theta_2)) \cdot \sin(\theta_2)$$

- We can simply evaluate these functions

Computing Gradients: Analytic Method


- **Pros.**
 - Exact
- **Cons.**
 - Requires deriving the gradient formula for all parameters
 - Still needs computing the gradients for each parameters
- Luckily, for neural nets, the cons become easy to solve
 - Derivation can be automatized
 - Computing the gradients can be grouped and simplified

Backpropagation

Analytic Form of Gradients

- **Question.** How do we derive an analytic form of $\nabla_{\theta} f_{\theta}(\mathbf{x})$, for...?

$$f_{\theta}(\mathbf{x}) = \mathbf{W}_L \sigma(\mathbf{W}_{L-1} \sigma(\cdots \sigma(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1) \cdots + \mathbf{b}_{L-1}) + \mathbf{b}_L$$



Analytic Form of Gradients

- **Idea.** View this as a composition of elementary operations

$$f_{\theta}(\mathbf{x}) = f_{\mathbf{b}_L} \circ f_{\mathbf{W}_L} \circ f_{\sigma_L} \circ \dots \circ f_{\mathbf{W}_1}(\mathbf{x})$$

- $f_{\mathbf{W}_i}(\mathbf{x}) = \mathbf{W}_i \mathbf{x}$
- $f_{\mathbf{b}_i}(\mathbf{x}) = \mathbf{x} + \mathbf{b}_i$
- $f_{\sigma}(\mathbf{x}) = \sigma(\mathbf{x})$
- Then:
 - Derivatives of each elementary op can be hard-coded
 - Use **chain rule** to combine these

Example

- Consider a function

$$g(x, y, z) = (x + y) \cdot z$$


- This is a composition of two elementary operations

$$g(x, y, z) = g_2(g_1(x, y), z)$$

- Addition:

$$g_1(a, b) = a + b$$


- Multiplication: $g_2(a, b) = ab$



Example

- Each elementary operation has an easy-to-write gradient


- $\frac{\partial g_1}{\partial a} = 1, \frac{\partial g_1}{\partial b} = 1$
- $\frac{\partial g_2}{\partial a} = b, \frac{\partial g_1}{\partial b} = a$



Example

- Each elementary operation has an easy-to-write gradient

- $\frac{\partial g_1}{\partial a} = 1, \frac{\partial g_1}{\partial b} = 1$
- $\frac{\partial g_2}{\partial a} = b, \frac{\partial g_1}{\partial b} = a$



- Chain rule tells you that:


$$\frac{\partial g}{\partial x}(x, y, z) = \left[\frac{\partial g_2}{\partial a}(g_1(x, y), z) \right] \cdot \left[\frac{\partial g_1}{\partial a}(x, y) \right]$$

$= z$

Example

- Each elementary operation has an easy-to-write gradient

- $\frac{\partial g_1}{\partial a} = 1, \frac{\partial g_1}{\partial b} = 1$
- $\frac{\partial g_2}{\partial a} = b, \frac{\partial g_1}{\partial b} = a$



- Chain rule tells you that:

$$\frac{\partial g}{\partial x}(x, y, z) = \frac{\partial g_2}{\partial a}(g_1(x, y), z) \cdot \frac{\partial g_1}{\partial a}(x, y)$$


$$\frac{\partial g}{\partial y}(x, y, z) = \frac{\partial g_2}{\partial a}(g_1(x, y), z) \cdot \frac{\partial g_1}{\partial b}(x, y)$$

$= z$

Example

- Each elementary operation has an easy-to-write gradient

- $\frac{\partial g_1}{\partial a} = 1, \frac{\partial g_1}{\partial b} = 1$
- $\frac{\partial g_2}{\partial a} = b, \frac{\partial g_1}{\partial b} = a$



- Chain rule tells you that:

$$\frac{\partial g}{\partial x}(x, y, z) = \frac{\partial g_2}{\partial a}(g_1(x, y), z) \cdot \frac{\partial g_1}{\partial a}(x, y)$$

$$\frac{\partial g}{\partial y}(x, y, z) = \frac{\partial g_2}{\partial a}(g_1(x, y), z) \cdot \frac{\partial g_1}{\partial b}(x, y)$$


$$\frac{\partial g}{\partial z}(x, y, z) = \begin{cases} \frac{\partial g_2}{\partial b}(g_1(x, y), z) \\ = g_1(x, y) \end{cases}$$

Example

$$\frac{\partial g}{\partial x}(x, y, z) = \frac{\partial g_2}{\partial a}(g_1(x, y), z) \cdot \frac{\partial g_1}{\partial a}(x, y)$$

$$\frac{\partial g}{\partial y}(x, y, z) = \frac{\partial g_2}{\partial a}(g_1(x, y), z) \cdot \frac{\partial g_1}{\partial b}(x, y) \quad \frac{\partial g}{\partial z}(x, y, z) = \frac{\partial g_2}{\partial b}(g_1(x, y), z)$$

- **Observation 1.** Computing gradients involves intermediate states of the composite function
 - **Idea.** Compute all intermediate states and store them. Later, we can combine these intermediate values




Example

$$\frac{\partial g}{\partial x}(x, y, z) = \frac{\partial g_2}{\partial a}(g_1(x, y), z) \cdot \frac{\partial g_1}{\partial a}(x, y)$$


$$\frac{\partial g}{\partial y}(x, y, z) = \frac{\partial g_2}{\partial a}(g_1(x, y), z) \cdot \frac{\partial g_1}{\partial b}(x, y) \quad \frac{\partial g}{\partial z}(x, y, z) = \frac{\partial g_2}{\partial b}(g_1(x, y), z)$$

- **Observation 2.** The computed gradients themselves can be reused
 - In particular, the gradient of the later block is used for computing the gradients of earlier-block parameters




Example: Backpropagation

- Inspired by these, we can think of a three-step algorithm
- **1. Forward Pass.** Compute the output, storing all intermediate states in the memory
 - From input to output




Example: Backpropagation

- Inspired by these, we can think of a three-step algorithm
- **1. Forward Pass.** Compute the output, storing all intermediate states in the memory
 - From input to output




Example: Backpropagation

- Inspired by these, we can think of a three-step algorithm
- **1. Forward Pass.** Compute the output, storing all intermediate states in the memory
 - From input to output




Example: Backpropagation

- **2. Backward Pass.** Compute the gradient using stored states
 - From output to input




Example: Backpropagation

- **2. Backward Pass.** Compute the gradient using stored states
 - From output to input



Example: Backpropagation

- **2. Backward Pass.** Compute the gradient using stored states
 - From output to input




Example: Backpropagation

- 3. GD. Update the parameters

$$x \leftarrow x - \eta \cdot 2, \quad y \leftarrow y - \eta \cdot 2, \quad z \leftarrow z - \eta \cdot 8$$


- Then, repeat 1–3 over and over...



Another example

- Consider a function


$$f_w(\mathbf{x}) = \frac{1}{1 + \exp(- (w_0x_0 + w_1x_1 + w_2))}$$



Another example

- Consider a function


$$f_w(\mathbf{x}) = \frac{1}{1 + \exp(- (w_0x_0 + w_1x_1 + w_2))}$$



Another example

- Consider a function


$$f_w(\mathbf{x}) = \frac{1}{1 + \exp(- (w_0x_0 + w_1x_1 + w_2))}$$



Another example

- Consider a function


$$f_w(\mathbf{x}) = \frac{1}{1 + \exp(- (w_0x_0 + w_1x_1 + w_2))}$$



Another example

- Consider a function


$$f_w(\mathbf{x}) = \frac{1}{1 + \exp(- (w_0x_0 + w_1x_1 + w_2))}$$



Another example

- Consider a function


$$f_w(\mathbf{x}) = \frac{1}{1 + \exp(- (w_0x_0 + w_1x_1 + w_2))}$$



Another example


- Consider a function

$$f_w(\mathbf{x}) = \frac{1}{1 + \exp(- (w_0x_0 + w_1x_1 + w_2))}$$




Computational Graph

- For simple neural networks, the computation graph will look like:



Computational Graph

- For larger models, the computation graph will be like:



Computational Graph

- Fortunately, deep learning frameworks will automatically construct the computational graph for you
 - PyTorch
 - TensorFlow

Remarks

- **Computation.** Backpropagation requires **a lot of memory!**
 - Additional memory needed is typically twice the model size (keep the gradients & intermediate states)
 - Sometimes, we discard the intermediate states (activations) and rematerialize them whenever needed
 - Gradients of some activation functions are cheaper to compute/store
 - e.g., ReLU

Next up

- More about optimization
 - Advanced optimizers
 - Training strategies
 - Network initialization

</lecture 12>