

Reinforcement Learning - 1

Recap

- **So far.**
 - Supervised Learning
 - Unsupervised Learning
 - Generative Modeling
- **Today.**
 - Reinforcement Learning

Reinforcement learning

- Different from other ML paradigms in several aspects:
 - No supervisor – only a reward
 - Delayed feedback – not instantaneous
 - Time matters – Sequential, non-i.i.d. data
 - Agent's action affect the subsequent data it receives

Overview: A Maze Example

Problem

- Suppose that we want to **solve maze**
- **Task.** Given any maze, train a model that can find a path: Start -> End

$$f(\text{maze}) = \hat{\text{path}} = (\text{loc}_1, \text{loc}_2, \dots)$$

- **Goal.**

(1) Path should **solve the maze**:


$$\hat{\text{path}}[0] = \text{start}, \hat{\text{path}}[-1] = \text{end}$$

(2) Path should be **valid**

$$\hat{\text{path}} \cup \text{wall} = \emptyset$$

(3) Make it **shortest** as possible


$$\min \mathbb{E}[\text{len}(\hat{\text{path}})]$$



Supervised Learning

- How do we solve this problem?
- **Naïve.** Try a **supervised learning**
 - Collect many samples of mazes
 $\text{maze}_1, \text{maze}_2, \dots, \text{maze}_n$
 - Label the samples (e.g., hire maze experts)
 $(\text{maze}_1, \text{path}_1), \dots, (\text{maze}_n, \text{path}_n)$
 - Train a predictor using some distance measure (e.g., MSE loss)

$$\min_f \frac{1}{n} \sum_{i=1}^n \text{diff}(\text{path}_i, f(\text{maze}_i)) + \lambda \cdot \text{len}(f(\text{maze}_i))$$



Limitations


- This approach suffers from many limitations

1. Annotation. Costly to collect

- Expensive to hire expert maze solvers
- Some mazes are too difficult
- Multiple solutions can be true

2. Complexity. Label space is too complicated



- Requires a lot of data



Limitations

3. Generalization. Often generalizes very poorly

- Unseen data types and styles
 - Circular maze
- Incomplete observables
 - Maze in the dark



Idea

- Formulate the problem as a **sequential** decision-making
 - Greatly simplifies the prediction task
- Given your current **state** (i.e., position), predict **action** (i.e., direction)


$$g(\text{loc}_i; \text{maze}) = \hat{\text{dir}}_i \in \{\text{up}, \text{down}, \text{left}, \text{right}\}$$

- The next state is determined by both current state and action

$$\text{loc}_{i+1} = \text{loc}_i + \text{dir}_i$$

- Repeat until you arrive (or reach max runtime)

$$\text{loc}_t = \text{end} \quad \text{or} \quad t \geq T$$



Idea

- **Objective.** Maximize the **reward** (i.e., success rate)

$$\mathbb{E} \left[\mathbf{1}[g^t(\text{start}; \text{maze}_i) = \text{end}, \text{for some } t] \right]$$

- **Data collection.** *Deploy the model* on sample mazes to collect many **episode** (i.e., path) unrolled by the model

$\text{maze}_i, (\text{loc}_1, \text{loc}_2, \dots, \text{loc}_T)$


- If succeeded, use it as a successful sample
- If failed, use it as a wrong sample
- Mostly wrong samples at first, but will get better as training goes on

Idea

- **Note.** No human supervision at all!
 - Unlike image/text cases, the success depends on objective criterion
 - We have access to an **environment**, which gives us reward
 - If we have, we can also use human-labeled samples
 - Guide for early training
- **Question.** How do we encourage shorter path?
 - Select some discount factor $\gamma \in (0,1]$ and try to maximize
$$\mathbb{E} \left[\gamma^t \cdot \mathbf{1}[g^t(\text{start}; \text{maze}_i) = \text{end, for some } t] \right]$$

Applications

- Robotic control, especially when “human supervision” does not work
 $\text{movement}_i = g(\text{image}_i, \text{instruction}_i, \text{robot state}_i)$



Applications

- Stock trading

$$\text{trading decision}_i = g(\text{market}_i, \text{portfolio}_i)$$



Applications



- Drug discovery
 - Next material to try can be generated by adding/subtracting atoms
 $\text{material to try}_i = g(\text{materials tried}_i)$



Applications


- Macro-level placement plans for chips

$$(\text{element}_i, \text{location}_i) = g(\text{current plan}_i)$$




Key issues

- Often, we have an **exploration-exploitation** tradeoff
 - If we find a successful predictor –e.g., right-hand policy
 - Then we can either
 - **Explore.** Try a new policy, potentially finding a better one
 - **Exploit.** Stick to the policy, to maximize the success rate



Key issues

- Rewards are **sparse** in many cases
 - Difficult to know if your intermediate actions have been correct
 - Especially problematic for long-horizon tasks
 - Can fall into myopic actions




In what follows

- Formalisms
 - Markov Decision Process
 - Bellman Equation
- Algorithms
 - As time permits...

Formalism

Reinforcement learning

- Learning a behavior strategy (a **policy**), which maximizes the long term sum of rewards (**delayed reward**) by a direct interaction (**trial-and-error**) with an unknown and uncertain environments



Interaction protocol

- **Agent.** Generates action, given the current state (plus reward)


$$a_i = f_{\theta}(s_i)$$

- **Environment.** Generates next state, given the action and current state

$$s_{i+1} = g(s_i, a_i)$$

- **Critic.** Generates reward


$$r_i = h(s_i, s_{i+1}, a_i)$$



```
for t = 1, ..., n do
    The agent perceives state st
    The agent performs action at
    The environment evolves to st+1
    The agent receives reward rt
end for
```

Markov chain

- A key property of this framework is that the future state s_{i+1} depends only on (s_i, a_i)
 - It does not depend on $(s_{i-1}, a_{i-1}), (s_{i-2}, a_{i-2}), \dots$



- All information about the past is contained inside s_i

Markov chain

- Let's make this argument a little more formal
 - Mathematical base for how to model “states”

Definition (**Markov Chain**).

Suppose that we have three random variables X, Y, Z . We say that they form a Markov chain $X - Y - Z$, whenever

$$P(Z | Y, X) = P(Z | Y)$$

That is, Z and X are conditionally independent, given Y

Markov chain

Example (**Weather**).


Suppose that in Pohang, it is only “sunny” ☀ or “rainy” ☔

- If it was sunny yesterday, then it will stay sunny with 90% chance
- If it was rainy yesterday, then it will stay rainy with 50% chance

Then, letting

- X : Yesterday's weather
- Y : Today's weather
- Z : Tomorrow's weather

We have $X - Y - Z$ a Markov chain.



Markov chain


- For this weather model, let $p_1^{(i)}$ be the probability of being sunny on i -th day and $p_2^{(i)}$ be the probability of being rainy.
 - Then, we know that

State
Transition
Matrix

$$\begin{bmatrix} 0.9 & 0.5 \\ 0.1 & 0.5 \end{bmatrix} \begin{bmatrix} p_1^{(i)} \\ p_2^{(i)} \end{bmatrix} = \begin{bmatrix} p_1^{(i+1)} \\ p_2^{(i+1)} \end{bmatrix}$$

- If there is some p_1, p_2 such that

$$\begin{bmatrix} 0.9 & 0.5 \\ 0.1 & 0.5 \end{bmatrix} \begin{bmatrix} p_1 \\ p_2 \end{bmatrix} = \begin{bmatrix} p_1 \\ p_2 \end{bmatrix}$$



then the weather distribution is **stationary** – stays same everyday

Markov decision process

- The evolution of environment can be formalized as follows:

Definition (Markov Decision Process; MDP).

A Markov decision process is defined as a tuple $M = (S, A, p, r)$, where

- t : **time** clock
- S : **state space**
- A : **action space**
- $p(y | x, a)$: **transition probability**, with

$$p(y | x, a) = \mathbf{P}(s_{t+1} = y | s_t = x, a_t = a)$$

- $r(x, y, a)$: **reward** of transition (x, y, a)

Policy

- The behavior of agents can be characterized by policy


Definition (Policy).

A decision rule π_t can be:

- Deterministic: $\pi_t : S \rightarrow A$
- Stochastic: $\pi_t : S \rightarrow \Delta(A)$

A policy can be:

- Non-stationary: $\pi = (\pi_0, \pi_1, \dots)$
- Stationary: $\pi = (\pi, \pi, \dots)$



MDP + Stationary Policy

- **Remark.** Suppose that ...
 - Our state evolves as an MDP
 - Our policy is stationary

Then, we have a **Markov chain** of state S and transition probability

$$p(y | x) = p(y | x, \pi(x))$$

- This framework is powerful enough, capturing many sequential decision-making problems.

Interaction protocol

- The environment may differ by:
- **Controllability.** Fully (e.g., chess)
Partially (e.g., portfolio optimization)
- **Uncertainty.** Deterministic (e.g., chess)
Stochastic (e.g., dice-involving)
- **Reactive.** Adversarial (e.g., chess)
Fixed (tetris)
- **Observability.** Full (e.g., chess)
Partial (e.g., starcraft)
- **Availability.** Known (e.g., chess)
Unknown (e.g., robotics)

Next class

- Markov reward process
- Bellman equation
- Basic algorithms

</lecture 22>