


Reinforcement Learning - 2

Recap

- **Last class.**
 - Hand-wavy introduction to **reinforcement learning**
 - How it differs from supervised learning
 - Formalisms
 - Markov chain
 - Policy
- **Today.**
 - Markov Process / MRP / MDP
 - Bellman's equation

Recap: RL Framework

- Learning to solve **sequential decision-making**, without supervision
 - Can break down complicated problems into easier sub-problems
 - Can adapt to any change in environment
 - Easier to generalize
- **Goal.** Select actions to maximize the total expected future reward



```
for  $t = 1, \dots, n$  do
    The agent perceives state  $s_t$ 
    The agent performs action  $a_t$ 
    The environment evolves to  $s_{t+1}$ 
    The agent receives reward  $r_t$ 
end for
```

Markov Process

- Let t be the time clock
- A **Markov process** is a tuple (S, p)
 - S : state space
 - p : state transition probability
- We make a **Markov assumption** on p
 - The next state s_{i+1} is determined solely by the current state s_i ,
$$p(s_{t+1} | s_1, \dots, s_t) = p(s_{t+1} | s_t)$$
 - That is, “state” is rich enough to contain all relevant information
- Note. No rewards, no actions


Markov Process

- If the number of states is finite, the transition dynamics can be expressed via **state transition matrix**
 - If states are $S = \{z_1, \dots, z_N\}$, then we'll have

$$P = \begin{bmatrix} p(z_1 | z_1) & P(z_2 | z_1) & \cdots & p(z_N | z_1) \\ p(z_1 | z_2) & P(z_2 | z_2) & \cdots & p(z_N | z_2) \\ \cdots & \cdots & \cdots & \cdots \\ p(z_1 | z_N) & P(z_2 | z_N) & \cdots & p(z_N | z_N) \end{bmatrix}$$

- **Example.** Weather transition model

$$\begin{bmatrix} p_{\text{sunny}}^{(t+1)} \\ p_{\text{rainy}}^{(t+1)} \end{bmatrix} = \begin{bmatrix} 0.9 & 0.5 \\ 0.1 & 0.5 \end{bmatrix} \begin{bmatrix} p_{\text{sunny}}^{(t)} \\ p_{\text{rainy}}^{(t)} \end{bmatrix}$$



Markov Reward Process

- A **Markov reward process (MRP)** is a tuple (S, p, r)
 - a Markov process (S, p)
 - a **reward** r
- The reward is earned by transitioning from one state to another
 - Deterministic: $r_t = r(s_t, s_{t+1})$
 - Stochastic: $r_t \sim r(\cdot | s_t, s_{t+1})$
- For simplicity, assume that it depends on the current state only
 - We can define the **reward function**

$$R(s) = \mathbb{E}[r_t | s_t = s]$$

Markov Reward Process

- We are interested in the **expected accumulated rewards** (or “return”)

$$G_t = r_t + \gamma r_{t+1} + \cdots + \gamma^{H-1} r_{t+H-1}$$

- γ : discount factor
- H : “horizon,” i.e., the number of steps in each episode
 - can be infinite
- The **state value function** is the **expected return**, starting at state s

$$\begin{aligned} V(s) &= \mathbb{E}[G_t | s_t = s] \\ &= \mathbb{E}[G_t | r_t + \gamma r_{t+1} + \cdots + \gamma^{H-1} r_{t+H-1} | s_t = s] \end{aligned}$$

Markov Reward Process

- If we have an infinite horizon, then we have:

$$V(s) = R(s) + \gamma \sum_{s' \in S} p(s' | s)V(s)$$

- Value = immediate reward + discounted sum of future rewards
- Note. This is called the **Bellman equation** (for MRPs)
- For a finite-state MRP, we have a neat matrix form

$$\begin{bmatrix} V(s_1) \\ \dots \\ V(s_N) \end{bmatrix} = \begin{bmatrix} R(s_1) \\ \dots \\ R(s_N) \end{bmatrix} + \gamma \begin{bmatrix} p(s_1 | s_1) & \dots & p(s_N | s_1) \\ & \dots & \\ p(s_1 | s_N) & \dots & p(s_N | s_N) \end{bmatrix} \begin{bmatrix} V(s_1) \\ \dots \\ V(s_N) \end{bmatrix}$$

- More simply, we can write

$$V = R + \gamma PV$$

Markov Reward Process

$$V = R + \gamma PV$$

- Solving the matrix equation, we have

$$V = (I - \gamma P)^{-1}R$$

- Fortunately, the matrix $(I - \gamma P)$ is invertible
- Computation.**

- Direct inverse: Requires $O(N^3)$

- Iterative solution: Requires $O(N^2)$ per step

- Initialize $V_{(0)} = 0$

- Repeat $V_{(k+1)} = R + \gamma PV_{(k)}$



Converges, theoretically
(Fixed point theorem)

Markov Decision Process

- A **Markov decision process (MRP)** is a tuple (S, A, p, r)

- an **action space** A

- a Markov reward process (S, p, r) , with
 - Action-dependent transition model

$$p(s_{t+1} \mid s_t, a_t)$$

- Action-dependent reward

$$r_t \sim r(\cdot \mid s_t, s_{t+1}, a_t)$$

- Again, for simplicity, assume independence w.r.t. s_{t+1}

$$R(s, a) = \mathbb{E}[r_t \mid s_t = s, a_t = a]$$

MDP Policies

- Agent's behaviors are specified by **policies**
 - Deterministic: $a_t = \pi(s_t)$
 - Stochastic: $a_t \sim \pi(\cdot | s_t)$
- If we fix the policy π , then MDP becomes a MRP
 - The return function and the state transitions are given as

$$R^\pi(s) = \sum_{a \in A} \pi(a | s) R(s, a)$$

$$p^\pi(s' | s) = \sum_{a \in A} \pi(a | s) p(s' | s, a)$$

MDP Policies

- Now, we can compute the value function $V(s)$ for this MRP
 - Similarly to MRP, we have a Bellman equation

$$V^\pi(s) = \sum_{a \in A} \pi(a | s) \left(R(a, s) + \gamma \sum_{s' \in S} p(s' | s, a) V^\pi(s') \right)$$

- Again, the solution can be found by an iterative method:
 - Initialize: $V_{(0)}^\pi(s) = 0$
 - Iterate:

$$V_{(k+1)}^\pi(s) = \sum_{a \in A} \pi(a | s) \left(R(s, a) + \gamma \sum_{s' \in S} p(s' | s, a) V_{(k)}^\pi(s') \right)$$

Optimal Policy for MDP

- Given the initial state s , the expected return of a policy π is: $V^\pi(s)$
 - We are interested in the **optimal policy**:
$$\pi^*(s) = \arg \max_{\pi} V^\pi(s)$$
- Mathematically, we can show that:
 - Exists a solution
 - The solution is unique
- For an infinite-horizon problem, the solution is:
 - Deterministic
 - Stationary

Finding the optimal policy

$$\pi^*(s) = \arg \max_{\pi} V^\pi(s)$$

- **Question.** How do we solve the optimization?
- **Naïve.** Search all deterministic policies

$$\pi : S \rightarrow A$$

- Computationally very difficult
 - The search space has the size $|A|^{|S|}$

Finding the optimal policy

$$\pi^*(s) = \arg \max_{\pi} V^\pi(s)$$

- **Policy iteration.** Usually more efficient
 - Initialize $\pi_0(s)$ randomly for all states s
 - Until convergence, do:
 - Evaluate π_i to get the function V^{π_i}
 - Update π_{i+1} to be an **improved version**
- Recall: Gradient descent, K-means, E-M

State-Action Value Q

- Define the **state-action value** of a policy

$$Q^\pi(s, a) = R(s, a) + \gamma \sum_{s' \in S} p(s' | s, a) V^\pi(s')$$

- The expected return from:

- taking an action a at state s
- then continuing with the policy π

- Satisfies the equality

$$V^\pi(s) = \sum_{a \in A} \pi(a | s) Q^\pi(a, s)$$

← Fix



Policy iteration

- Similar to the expectation-maximization, we do:

- Compute the Q function for the current policy π_i

$$Q^{\pi_i}(s, a) = R(s, a) + \gamma \sum_{s' \in S} p(s' | s, a) V^{\pi_i}(s')$$

- Compute the new policy π_{i+1} for all $s \in S$

$$\pi_{i+1}(s) = \arg \max_{a \in A} Q^{\pi_i}(s, a), \quad \forall s \in S$$

- **Theorem (w/o proof).** We have $V^{\pi_{i+1}}(s) \geq V^{\pi_i}(s)$

- Somewhat surprising, as the optimization of π_{i+1} was based on the assumption that we'll use π_i from the next step.

Value iteration

- Policy iteration operates as:
 - computes the infinite horizon value of a policy
 - somewhat heavy – computing value requires iterations
 - improves that policy
- Value iteration is another technique:
 - Compute the best next state, when $H = 1$
 - Compute the best next state, when $H = 2$
 - ...

Value iteration

- More formally, the value iteration is:

- Initialize $V_{(0)}(s) = 0$ for all states s
- Iterate:
 - Update the value function

$$V_{(i+1)}(s) = \max_a \left(R(s, a) + \gamma \sum_{s' \in S} p(s' | s, a) V_{(i)}(s') \right)$$

- Update the policy function

$$\pi_{(i+1)}(s) = \arg \max_a \left(R(s, a) + \gamma \sum_{s' \in S} p(s' | s, a) V_{(i)}(s') \right)$$

Value iteration vs. Policy iteration

- **Value iteration**
 - Lighter
 - Better for large state-/action spaces
 - Can compute finite-horizon policies
- **Policy iteration**
 - Requires less iteration, usually
 - Guaranteed and stable convergence
 - Especially when γ is very large

Next class

- Q-learning
- Policy Gradient
- Wrap-up

</lecture 24>