

Sparsity – 1

EECE695D: Efficient ML Systems

Spring 2025

Agenda

- **Last Class**
 - Matmuls
 - Computation vs. Memory
- **W2 & W3**
 - Reducing computation & memory at the matmul level
- **Today**
 - Sparsity & Pruning

Basic idea

Goal

- We want to reduce the computational cost of **matrix multiplication**
 - Well-trained linear model with $d_{\text{in}} = d_{\text{out}} = 3$ and the dataset size $N = 3$

$$\mathbf{WX} = \begin{bmatrix} w_1 & w_2 & w_3 \\ w_4 & w_5 & w_6 \\ w_7 & w_8 & w_9 \end{bmatrix} \begin{bmatrix} x_1 & x_2 & x_3 \\ x_4 & x_5 & x_6 \\ x_7 & x_8 & x_9 \end{bmatrix}$$

- Compute. $2d_{\text{in}}d_{\text{out}}N = 54 \text{ FLOPs}$
- Memory I/O. $3 \times 3 \text{ FP32 weights} = 36 \text{ Bytes}$
(loading weights)

Sparsity

- Remove **less important** entries of the weight, thus skipping associated ops
 - Suppose that we “prune out” 4 entries, to get a 5-sparse matrix

$$\mathbf{W}_{\text{pruned}}\mathbf{X} = \begin{bmatrix} w_1 & w_2 & 0 \\ 0 & w_5 & 0 \\ w_7 & 0 & w_9 \end{bmatrix} \begin{bmatrix} x_1 & x_2 & x_3 \\ x_4 & x_5 & x_6 \\ x_7 & x_8 & x_9 \end{bmatrix}$$

- Fancy put, we take a Hadamard product with some mask matrix

$$\mathbf{W}_{\text{pruned}} = \mathbf{M} \odot \mathbf{W}, \quad \mathbf{M} = \begin{bmatrix} 1 & 1 & 0 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{bmatrix}$$

Quiz. The matrix $\mathbf{W}_{\text{pruned}}$ has ...

(a) 44.4% Sparsity

(b) 55.5% Sparsity

$$\begin{bmatrix} w_1 & w_2 & 0 \\ 0 & w_5 & 0 \\ w_7 & 0 & w_9 \end{bmatrix}$$

Advantages


$$\begin{bmatrix} w_1 & w_2 & 0 \\ 0 & w_5 & 0 \\ w_7 & 0 & w_9 \end{bmatrix} \begin{bmatrix} x_1 & x_2 & x_3 \\ x_4 & x_5 & x_6 \\ x_7 & x_8 & x_9 \end{bmatrix}$$

- Compute & memory decreases proportionally to the sparsity
 - Compute. $(1 - \text{sparsity}) \times (\text{dense FLOPs}) = 30 \text{ FLOPs}$
 - Memory I/O. $(1 - \text{sparsity}) \times (\text{dense I/O}) = 20 \text{ Bytes}$

Note. There are certain overheads, as we will see in the next class

Pruning neural networks

- We apply pruning at a model level
 - Layer 1 pruned to xx% sparsity
 - Layer 2 pruned to yy% sparsity
 - (...)
- ⇒ Model achieves zz% global sparsity



- Typically, one can achieve 20%—80% global sparsity without accuracy drop

Algorithm

Problem formulation

- Minimize the **training risk** of the pruned model, given the **sparsity constraint**

$$\text{minimize}_{\mathbf{w}_{\text{pruned}}} \hat{L}(\mathbf{w}_{\text{pruned}}) \quad \text{subject to} \quad \|\mathbf{w}_{\text{pruned}}\|_0 \leq \tau$$

- \mathbf{w} : all neural net weights, vectorized
- $\hat{L}(\cdot)$: training risk
- $\|\cdot\|_0$: ℓ_0 norm (i.e., the number of nonzero entries)
- τ : sparsity constraint

Note. We are using a global sparsity constraint, just for simplicity.

Problem formulation

- Alternatively, view it as a **joint optimization** of weights and mask

$$\text{minimize}_{\mathbf{m}, \mathbf{w}} \quad \hat{L}(\mathbf{m} \odot \mathbf{w})$$

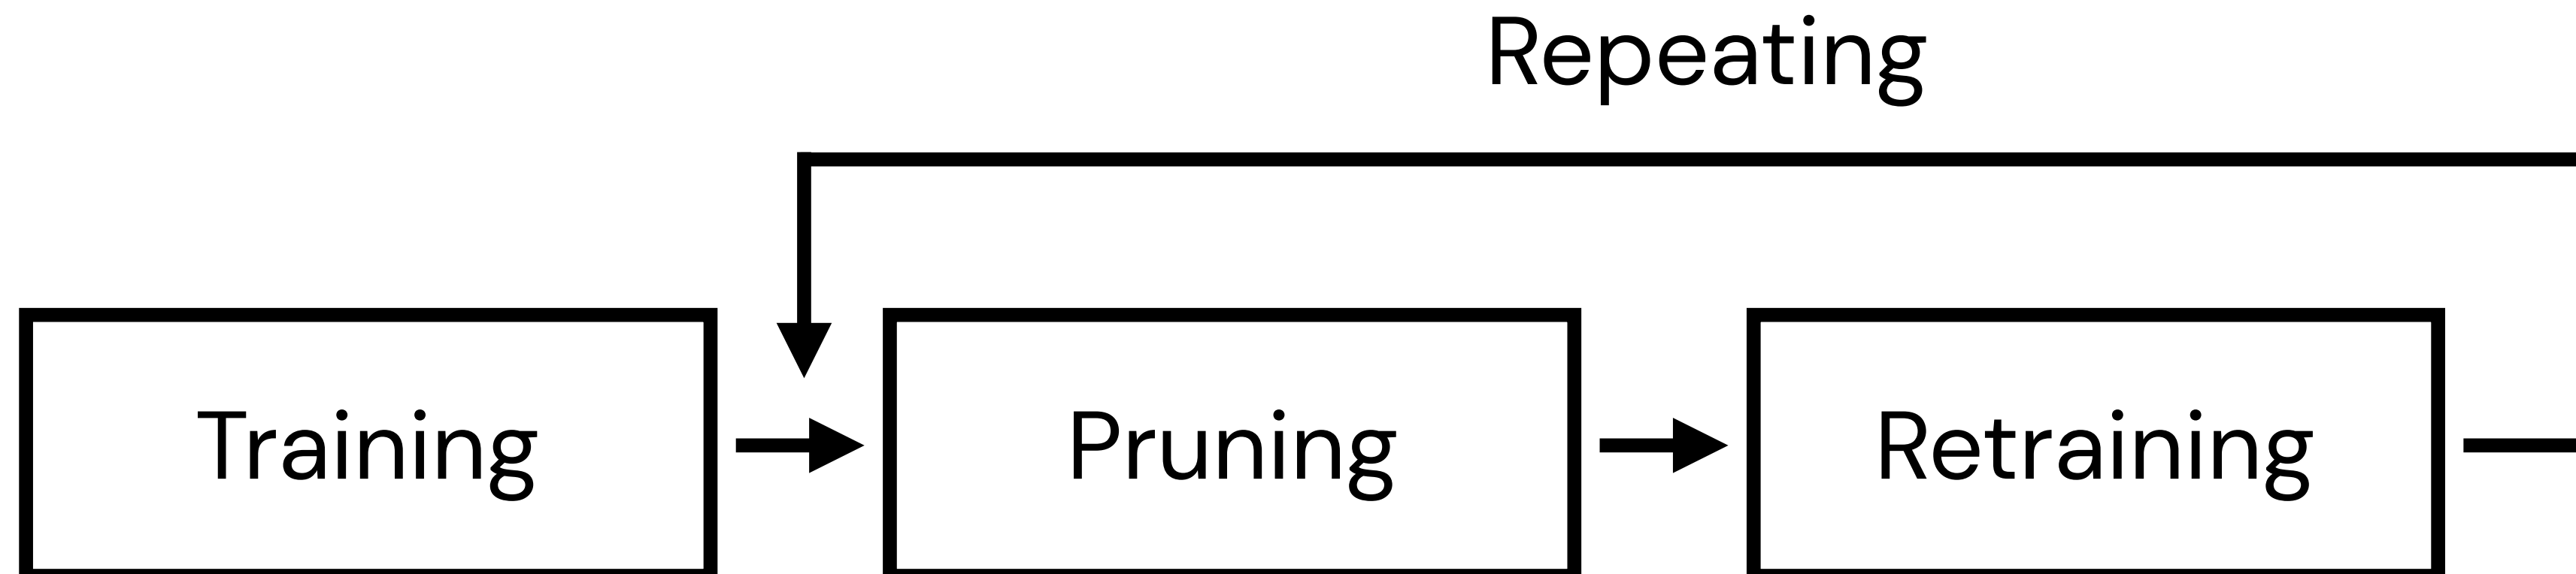
$$\text{subject to} \quad \|\mathbf{m}\|_0 \leq \tau, \quad m_{ij} \in \{0, 1\}$$

- By doing so, we have decomposed this into two **subproblems**
 - Optimizing \mathbf{w} : Unconstrained, continuous optimization
 - Optimizing \mathbf{m} : Constrained, discrete optimization

Algorithm

- Typical pruning algorithms solve this via **alternating optimization**:

1. Training. Train the model for some steps (optimize w)
2. Pruning. Remove some weights, using some criterion (fix w , optimize m)
3. Retraining. Retrain the model for some steps, to recover from damage. (fix m , optimize w)
4. Repeating. Repeat steps 2–3 for some iterations



Algorithm

- Two key elements:
 - **Saliency**. How can we identify less important weight?
 - **Schedule**. When do we introduce the sparsity?
 - Often reflects operational constraints, e.g., training cost

1. Training. Train the model for **some steps**

2. Pruning. Remove **some weights**, using **some criterion**

3. Retraining. Retrain the model for **some steps**, to recover from damage.

4. Repeating. Repeat steps 2–3 for **some iterations**

Saliency

Saliency

- At the pruning phase, we are solving the **mask optimization**:

$$\text{minimize}_{\mathbf{m}} \quad \hat{L}(\mathbf{m} \odot \mathbf{w}) \quad \text{subject to} \quad \|\mathbf{m}\|_0 \leq \tau, \quad m_i \in \{0,1\}$$

- This is **NP-hard**, and thus we typically use heuristics...
 - Hessian
 - Gradient
 - Magnitude

Hessian-based pruning

- **Idea.** Express the training risk using the **Taylor approximation**
 - Suppose that pruning changes the weight $\mathbf{w} \rightarrow \mathbf{w} + \mathbf{u}$.
 - Example. Removing i -th weight makes $\mathbf{u} = -w_i \mathbf{e}_i$
 - Then, we can write:

$$\hat{L}(\mathbf{w} + \mathbf{u}) \approx \hat{L}(\mathbf{w}) + \mathbf{g}^\top \mathbf{u} + \frac{\mathbf{u}^\top \mathbf{H} \mathbf{u}}{2}$$

- \mathbf{g} : First-order derivative (gradient), evaluated at \mathbf{w}
- \mathbf{H} : Second-order derivative (Hessian), evaluated at \mathbf{w}

Hessian-based pruning

- Now, assume that the weight \mathbf{w} is well-trained.
 - Then, the gradient is near-zero, making:

$$\hat{L}(\mathbf{w} + \mathbf{u}) \approx \hat{L}(\mathbf{w}) + \frac{\mathbf{u}^\top \mathbf{H} \mathbf{u}}{2}$$

- As the first term on RHS is independent of mask, the mask optimization can be approximated by:

$$\min_{\mathbf{u}} \frac{\mathbf{u}^\top \mathbf{H} \mathbf{u}}{2}$$

with appropriate constraints on \mathbf{u} .

Optimal Brain Damage

- LeCun et al. (1989) simplifies this as follows:
 - Suppose that we remove **only one weight**
 - Then, removing i-th layer makes $\mathbf{u} = -w_i \mathbf{e}_i$ and thus

$$\frac{\mathbf{u}^T \mathbf{H} \mathbf{u}}{2} = \frac{|w_i|^2 \mathbf{H}_{ii}}{2}$$

- Simply compute **this score** for all i, and remove bottom-k weights
 - Requires some calibration data to compute Hessian
- **Question (hw).** How do we compute Hessian diagonals for neural nets?

Optimal Brain Damage

- **Problem.** Hessians have $(\text{\#weight})^2$ entries
 - 1B-scale model will have 10^{18} entries for Hessian = 4 exabytes
- Fortunately, OBD only need **Hessian diagonals**, with (\#weight) entries
 - Can be computed in a similar way to backpropagation (Homework. Derive the formula)

Optimal Brain Surgeon

- Hassibi & Stork (1992) considers a slightly involved version:
 - Suppose that we remove only one weight, but can also **update other weights to compensate** for the removed weight.
 - Then, we are solving:

$$\min_i \left\{ \min_{\mathbf{u}_i} \left\{ \frac{\mathbf{u}_i^\top \mathbf{H} \mathbf{u}_i}{2} \right\} \quad \text{subject to} \quad \mathbf{e}_i^\top \mathbf{u}_i + w_i = 0 \right\}$$

- The Lagrangian form is:

$$L_i = \frac{\mathbf{u}_i^\top \mathbf{H} \mathbf{u}_i}{2} + \lambda(e_i^\top \mathbf{u}_i - w_i)$$

Optimal Brain Surgeon

- For fixed i , the solution and the Lagrangian is:

$$\mathbf{u}_i = - \frac{w_i}{[\mathbf{H}^{-1}]_{ii}} \mathbf{H}^{-1} \cdot \mathbf{e}_i$$

$$L_i = \frac{w_i^2}{2[\mathbf{H}^{-1}]_{ii}}$$

- We can select one weights with the smallest Lagrangian, make corresponding updates, and repeat...
- **Problem.** This requires computing the inverse Hessian!
 - How can we do this, without requiring extremely large matrix inverse?

Computing the inverse Hessian

- Suppose that we use the squared loss for the risk

$$\hat{L}(\mathbf{w}) = \frac{1}{2N} \sum_{i=1}^N (y_i - f(x_i; \mathbf{w}))^2$$

- Then, the loss gradient can be written as:

$$\mathbf{g} = \frac{1}{N} \sum_{i=1}^N (f(x_i; \mathbf{w}) - y_i) \frac{\partial f}{\partial \mathbf{w}}(x_i; \mathbf{w})$$

Samplewise
Error

Samplewise
Gradient

Computing the inverse Hessian

- The Hessian is:

$$\mathbf{H} = \frac{1}{N} \sum_{i=1}^N \frac{\partial f}{\partial \mathbf{w}}(x_i; \mathbf{w}) \frac{\partial f^\top}{\partial \mathbf{w}}(x_i; \mathbf{w}) - \frac{1}{N} \sum_{i=1}^N (f(x_i; \mathbf{w}) - y_i) \frac{\partial^2 f}{\partial \mathbf{w}^2}(x_i; \mathbf{w})$$

- If our model is good enough, we can approximate:

$$\begin{aligned} \mathbf{H} &\approx \frac{1}{N} \sum_{i=1}^N \frac{\partial f}{\partial \mathbf{w}}(x_i; \mathbf{w}) \frac{\partial f^\top}{\partial \mathbf{w}}(x_i; \mathbf{w}) \\ &\triangleq \frac{1}{N} \sum_{i=1}^N \mathbf{q}_i \mathbf{q}_i^\top \end{aligned}$$

Samplewise
Gradient

Computing the inverse Hessian

- This gives us a recursive formula for computing the Hessian

$$\mathbf{H}_m = \mathbf{H}_{m-1} + \frac{1}{N} \mathbf{q}_m \mathbf{q}_m^\top$$

- Combine this with the matrix inversion formula:

$$(\mathbf{A} + \mathbf{BCD})^{-1} = \mathbf{A}^{-1} - \mathbf{A}^{-1} \mathbf{B} (\mathbf{C}^{-1} + \mathbf{DA}^{-1} \mathbf{B})^{-1} \mathbf{DA}^{-1}$$

- Then, we get a recursive formula for computing inverse Hessian

$$\mathbf{H}_m^{-1} = \mathbf{H}_{m-1}^{-1} - \frac{\mathbf{H}_{m-1}^{-1} \mathbf{q}_m \mathbf{q}_m^\top \mathbf{H}_{m-1}^{-1}}{N + \mathbf{q}_m^\top \mathbf{H}_{m-1}^{-1} \mathbf{q}_m}$$

Other techniques

- Still, Hessian is too large to compute & hold on RAM for large models
- **Solution.**
 - Compute Hessian layer-by-layer (Dong et al., 2017; Layerwise OBS)
 - <https://arxiv.org/abs/1705.07565>
 - Recent works on LLM develop more involved techniques:
 - Will be discussed in future lectures; <https://arxiv.org/abs/2301.00774>

Gradient-based pruning

- In many cases, we cannot simply assume that $\text{gradient} = 0$
 - Pruning at initialization
 - e.g., SNIP (Lee et al., 2019)
 - Pruning pre-trained models before fine-tuning
 - e.g., Movement Pruning (Sanh et al., 2020)
 - Pruning underfitting models
 - e.g., large language models

Gradient-based pruning

- **Idea.** Use the first-order approximation:

$$\hat{L}(\mathbf{w} + \mathbf{u}) \approx \hat{L}(\mathbf{w}) + \mathbf{g}^\top \mathbf{u}$$

- Choose i weights with the smallest values of the gradient score:

$$-w_i \mathbf{g}_i$$

- In fact, taking an absolute value is a good idea (**why?**)


$$|w_i \mathbf{g}_i|$$

Magnitude-based pruning

- Suppose that we cannot compute Hessian
 - Too much memory & computation needed
 - No calibration data
- **Idea.** Blindly assume that the Hessian diagonal \mathbf{H}_{ii} is identical for all weights.
 - i.e., use the saliency score w_i^2
 - i.e., remove weights with bottom-k weight magnitudes $|w_i|$

Magnitude-based pruning

- **Problem.** Prone to layer collapse, on global pruning
 - Suppose that we have two layer MLP, with:




Magnitude-based pruning

- **Solution.**
 - Layerwise heuristics (Gale et al., 2019)
 - Score scaling (Lee et al., 2021)

To understand the limits of the magnitude pruning heuristic, we modify our ResNet-50 training setup to leave the first convolutional layer fully dense, and only prune the final fully-connected layer to 80% sparsity. This heuristic is reasonable for ResNet-50, as the first layer makes up a small fraction of the total parameters in the model and the final layer makes up only .03% of the total FLOPs. While tuning

Magnitude-based pruning


- With gradual pruning & good HP, magnitude-based pruning is good enough.
 - With limited retraining, Hessian-based methods are better



Schedule

Pruning schedules


- There are many different sparsity schedules, with different purposes.
- **Gradual Pruning.**
 - Best in terms of the accuracy vs. inference cost tradeoff
 - Requires lengthy training (2x – 10x of the original training)
 - Needs joint tuning of learning rate and sparsity schedules



Pruning schedules

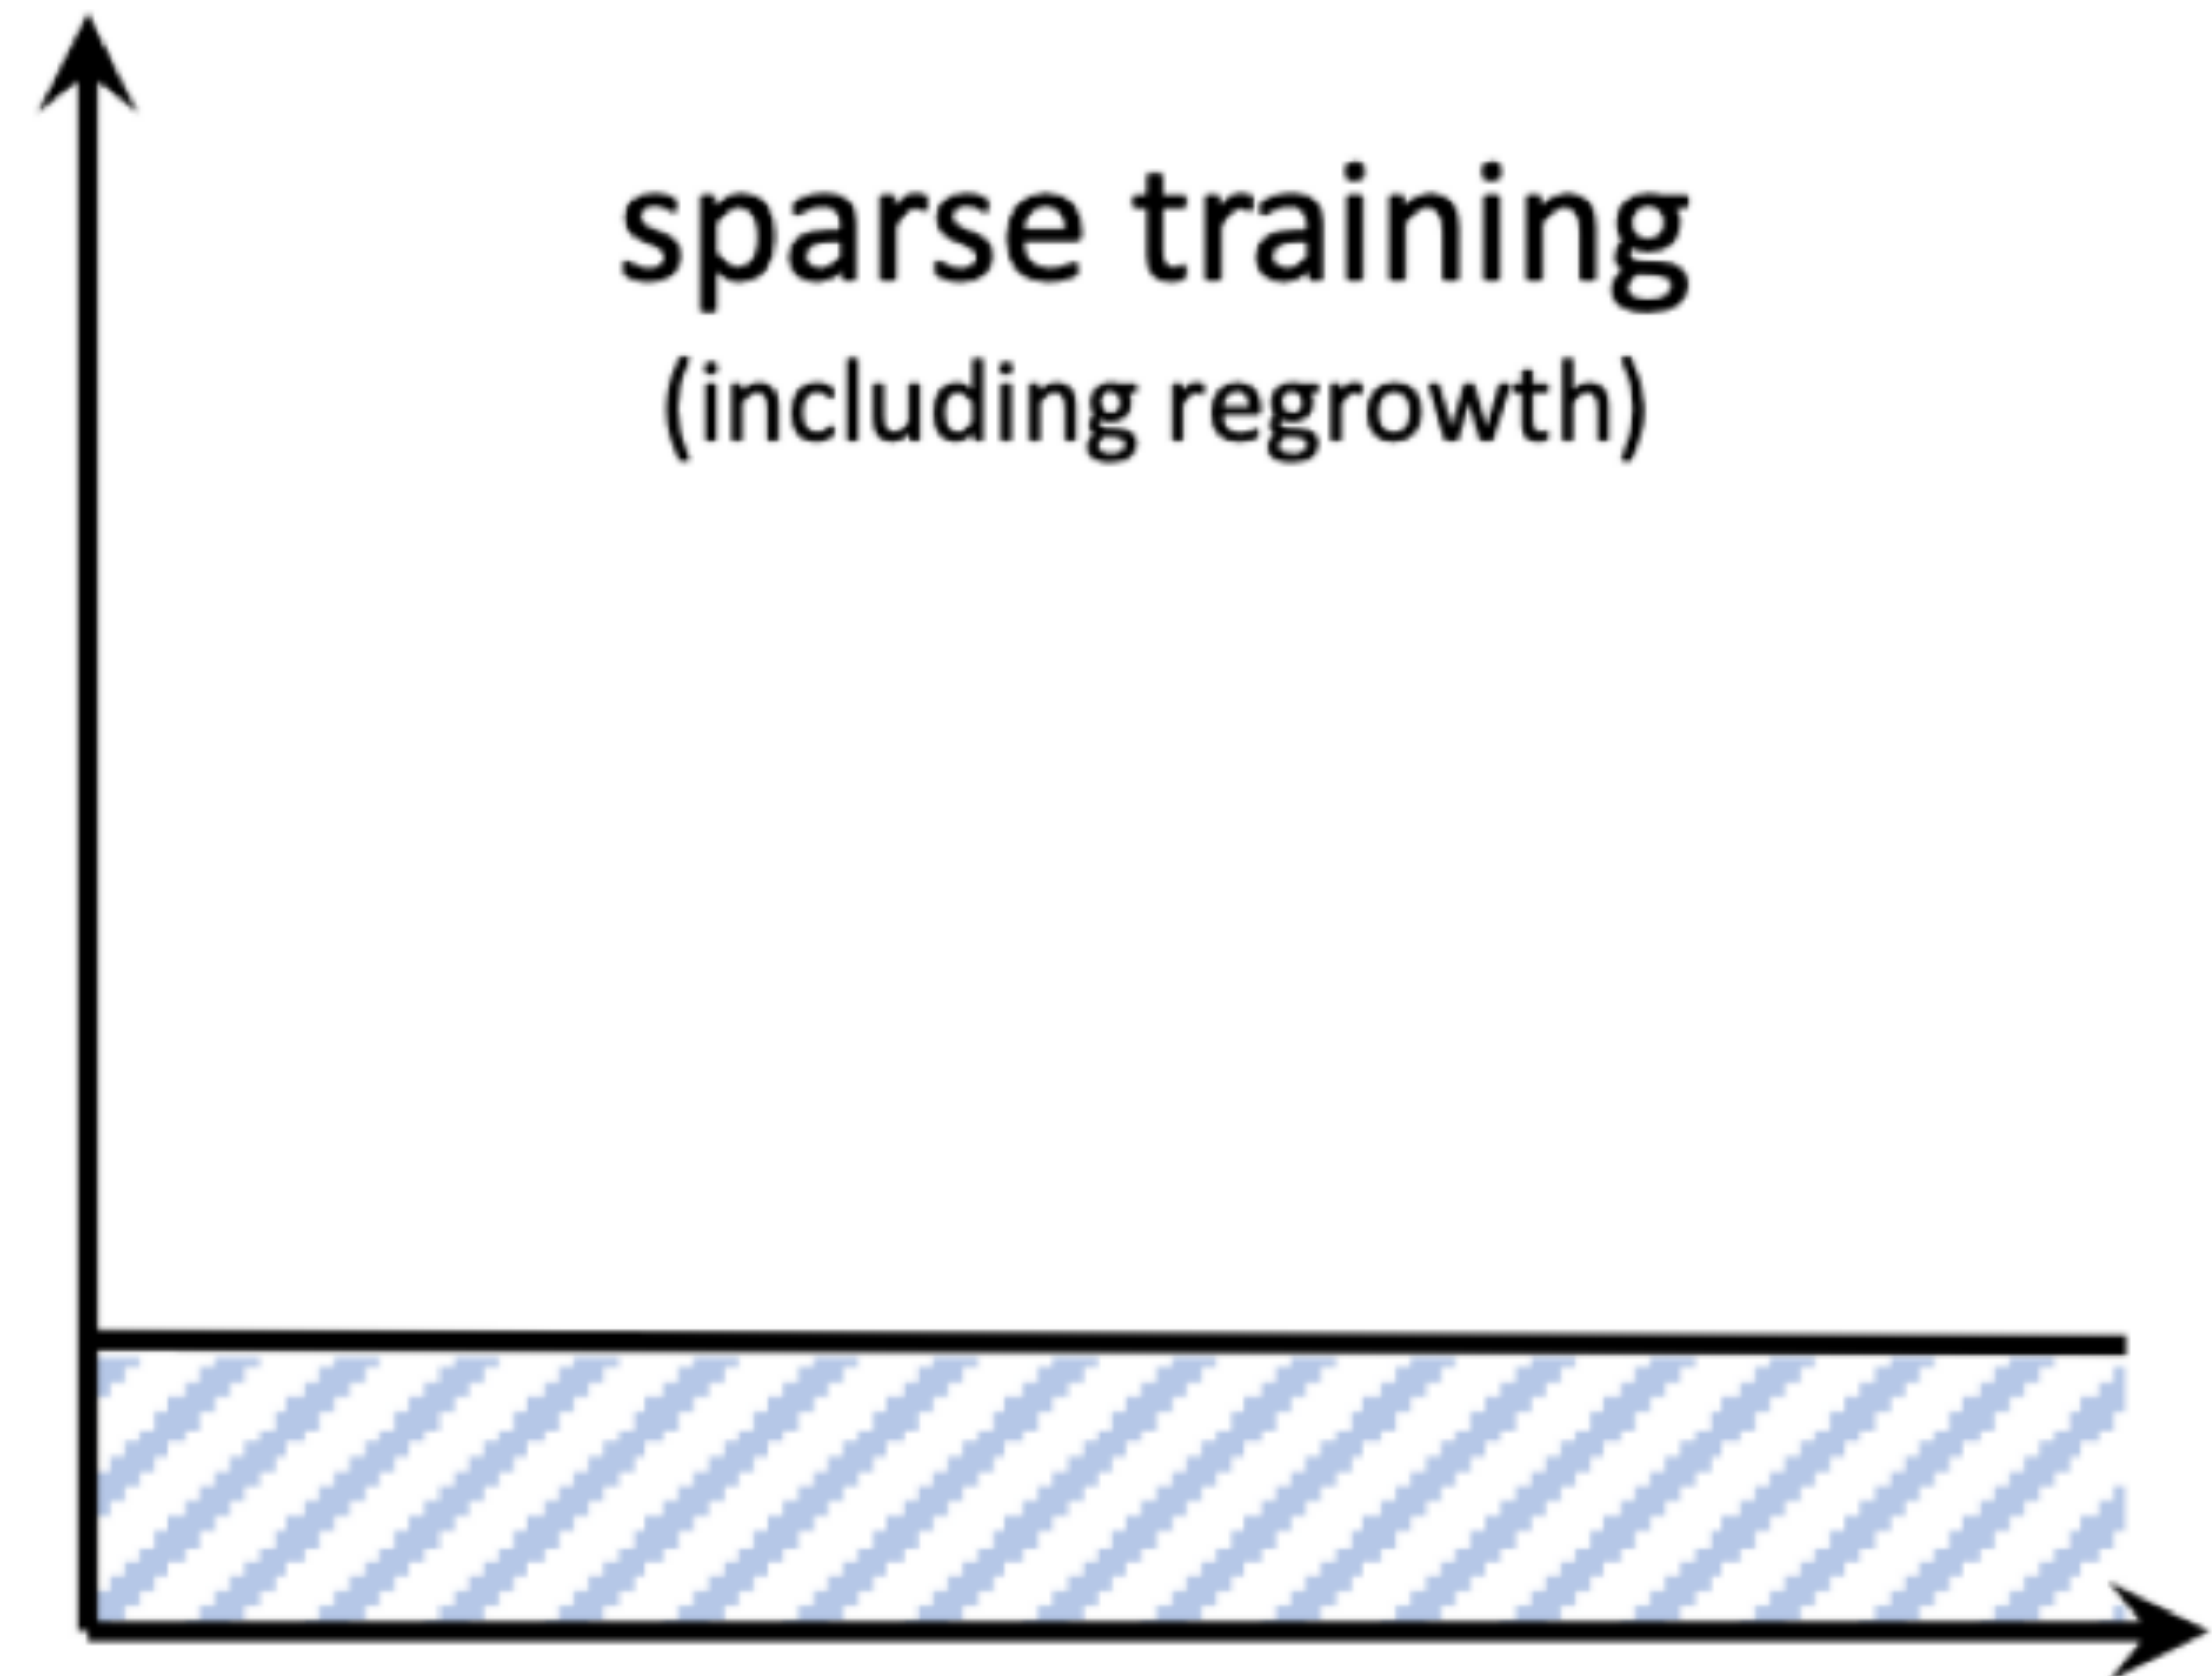
- e.g., cubic schedule (Zhu & Gupta, 2017; Google default)

$$s_t = s_f + (s_i - s_f) \left(1 - \frac{t - t_0}{n\Delta t}\right)^3 \quad \text{for } t \in \{t_0, t_0 + \Delta t, \dots, t_0 + n\Delta t\}$$



Pruning schedules


- **Sparse Training (Pruning at init.)**
 - Less GPU memory & training cost
 - Can train larger (sparse) models, theoretically
 - Requires lengthy training
 - Very unstable performance
 - Usually requires small sparsity or regrowth



Pruning schedules

- **One-shot Pruning (or Post-Training Sparsity)**

- Little or no retraining.
 - Suitable for LLM-scale models
- Bad performance, usually.
- Can exploit pretrained checkpoints
 - End-user friendly



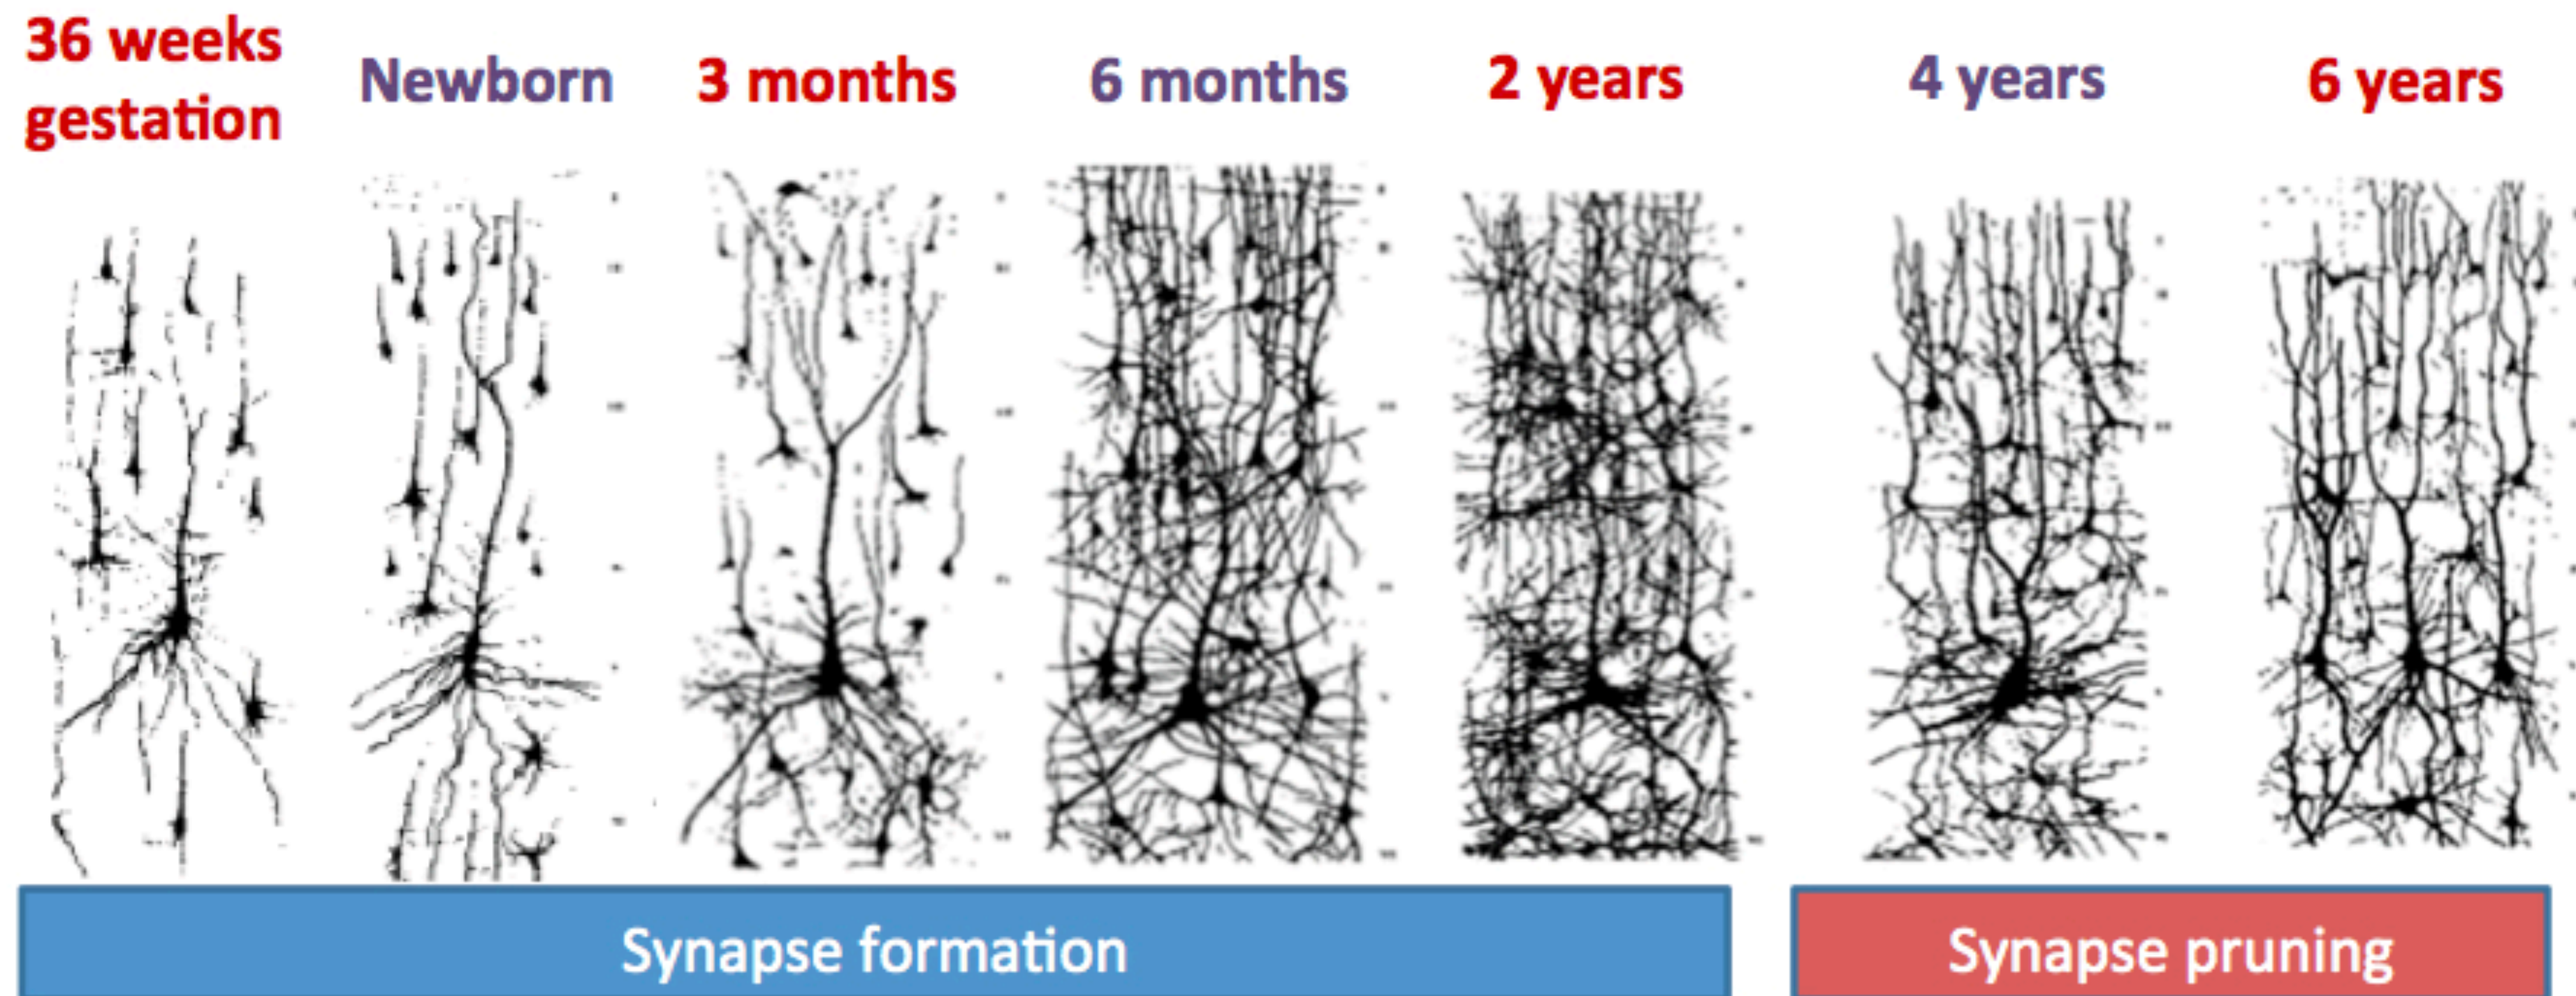
Why do sparse nets work?

Why should it work?

- Question. Why do we expect sparse models to work as well as dense models?
 - Answer. No concrete justification 😓
 - Nevertheless, there are some motivations...

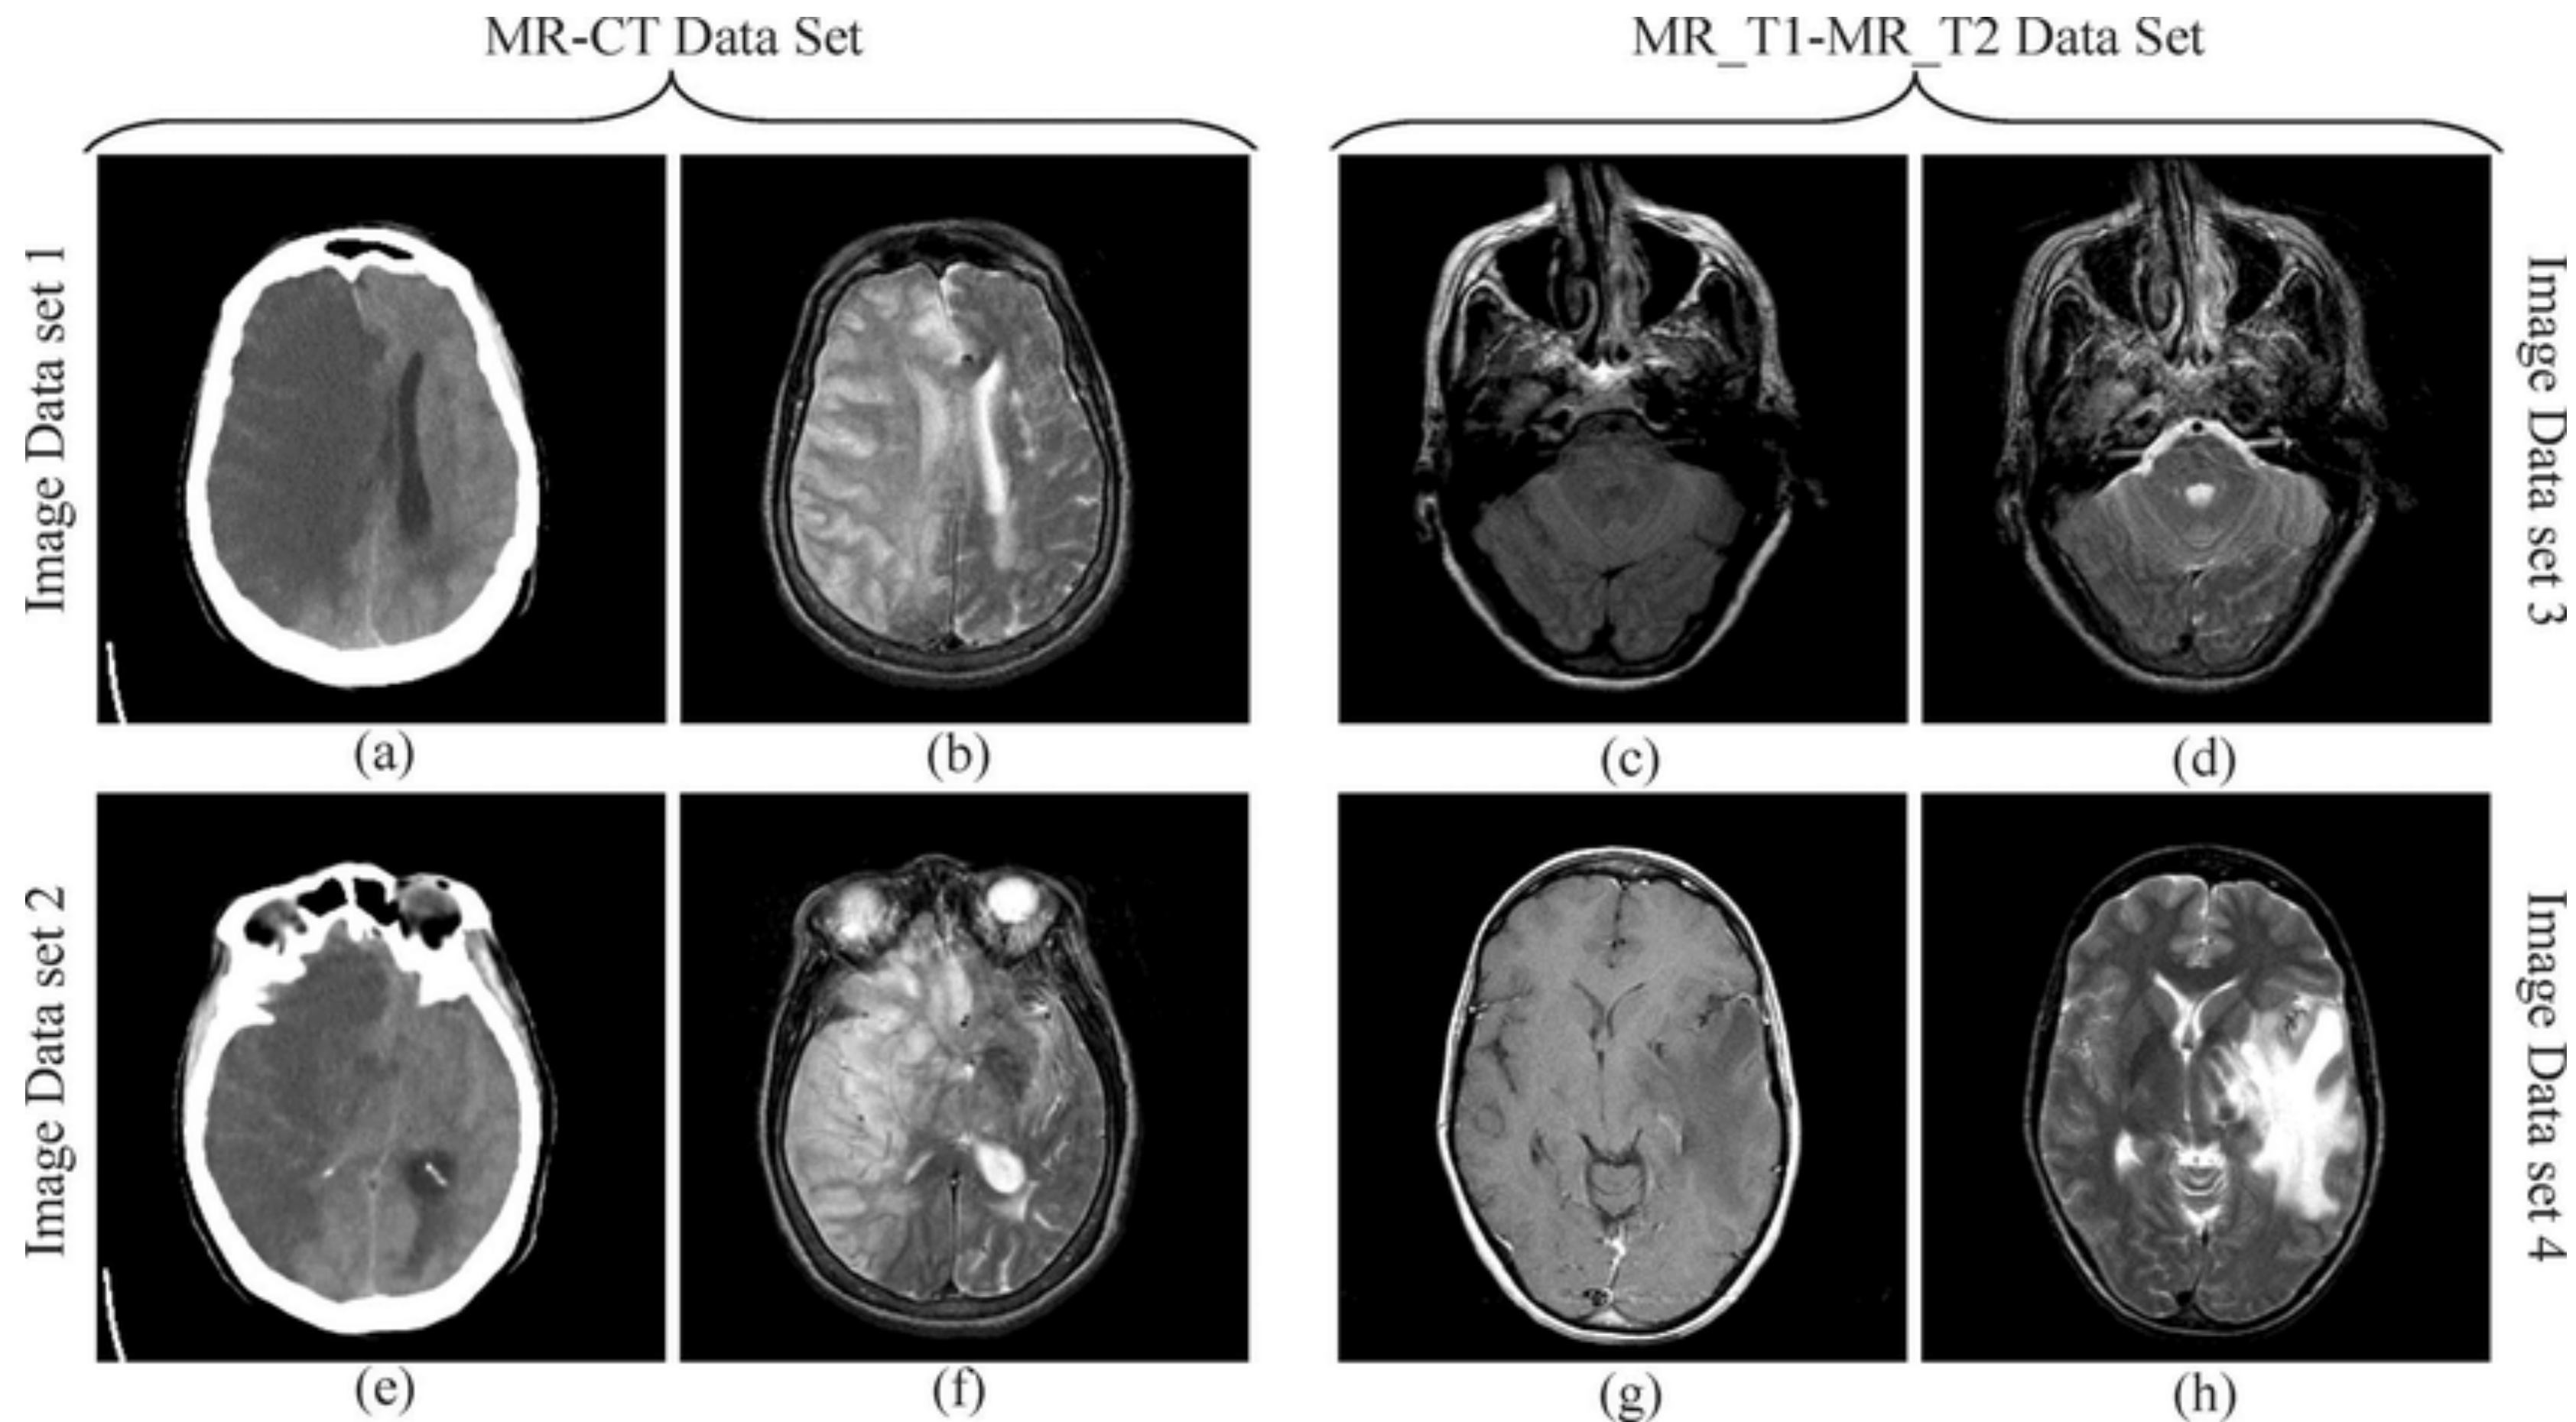
Why should it work?

- **Biological motivation.** Human brain also does some sort of pruning.



Why should it work?

- **Natural sparsity.** Many natural data or relationships are actually sparse
 - e.g., simply irrelevant input features




Why should it work?

- **Theoretical guarantees.** We use much more parameters than what is theoretically sufficient.
 - We need only $\tilde{O}(\sqrt{N})$ weights to achieve zero training loss on N samples.

Theorem 1.1 (informal statement). *Let $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N) \in \mathbb{R}^d \times \{1, \dots, C\}$ be a set of N labeled samples of a constant dimension d , with $\|\mathbf{x}_i\| \leq r$ for every i and $\|\mathbf{x}_i - \mathbf{x}_j\| \geq \delta$ for every $i \neq j$. Then, there exists a ReLU neural network $F : \mathbb{R}^d \rightarrow \mathbb{R}$ with width 12, depth $\tilde{O}(\sqrt{N})$, and $\tilde{O}(\sqrt{N})$ parameters, such that $F(\mathbf{x}_i) = y_i$ for every $i \in [N]$, where the notation $\tilde{O}(\cdot)$ hides logarithmic factors in N, C, r, δ^{-1} .*

Why should it work?

- **Generalization (depracated)**. In the past, it was believed that less parameters will lead to better generalization, by avoiding overfitting.
 - This no longer seems to be a valid logic, and is empirically not true.



Remarks

Next class

- **Today.** Algorithmic aspects of pruning
 - Which entries to prune?
 - How to recover from the damage?
 - When to prune?
- **Next Class.** System-level complications
 - How can we process sparse matrices?
 - What structures can we impose?

Further Readings

- Lottery ticket hypothesis
 - <https://arxiv.org/abs/1803.03635>
- What is the state of neural network pruning?
 - <https://arxiv.org/abs/2003.03033>

That's it for today 🙌