

# Parallelism - 2

EECE695D: Efficient ML Systems

Spring 2025


# Recap

- **Last class.**
  - Data parallelism
  - Model parallelism
    - Pipeline, Tensor, Expert
- **Today.** Advanced topics
  - Sequence parallelism
  - ZeRO, Gradient compression
  - Automated parallelism

# Sequence parallelism


# Motivation

- Training a transformer-based generative model
  - Want to generate high-dimensional data with an **extremely long context**
  - Example. High-resolution video generation
    - Spatio-temporal tokens as an input
- **Problem.** Not holdable on one device, even for a small batch



# Basic idea

- **Solution.** Each GPU processes a **fraction of input tokens**
  - FFN. Easy, because tokens are handled separately anyways
  - MHSA. Requires additional communication



# The case of MHSA

- **Goal.** Compute the output of each token:

$$\mathbf{o}_i = \sum_{i=1}^L \mathbf{s}_i \mathbf{v}_i$$


- $\mathbf{s}_i$  are the attention scores:

$$\mathbf{s}_i = \text{SoftMax} \left( \frac{1}{\sqrt{d}} [\mathbf{q}_i^\top \mathbf{k}_1, \mathbf{q}_i^\top \mathbf{k}_2, \dots, \mathbf{q}_i^\top \mathbf{k}_L] \right)$$

- $\mathbf{q}_i, \mathbf{k}_i, \mathbf{v}_i$  are query/key/values.
- **Problem.** Tokens are distributed among devices (k&v, in particular)


# Ring Self-Attention

- **Idea.** Transmit the key and value embeddings of the sequence
- Step 1. Compute and transmit keys
  - Each node can start computing the  $q_i^T k_j$ , as soon as they receive any fraction of the key embeddings
  - After the full ring, can compute the softmax to get attention scores



# Ring Self-Attention

- Step 2. Compute and transmit values
- Step 3. Now everybody has the full KV, and can compute the full output



# Further readings


- Combined framework with other notions of parallelism
  - Megatron-SP (NVIDIA)
    - Combines with tensor parallelism
    - <https://arxiv.org/abs/2205.05198>
  - DeepSpeed-Ulysses (Microsoft)
    - <https://arxiv.org/abs/2309.14509>

**ZeRO**

# Motivation

- If we use optimizers like AdamW, we need to keep various **optimizer states**
- **Example.** Optimizing a model with  $M$  parameter with Adam, in FP16


- Param.  $2M$  bytes
- Grad.  $2M$  bytes
- Variance.  $4M$  bytes
- Momentum.  $4M$  bytes
- FP32 Params.  $4M$  bytes



⇒ High redundancy in GPUs, when we do **data-parallel**


# Idea

- Partition the states and gradients on many GPUs
- **ZeRO-1.** The optimizer states are distributed (~4x memory reduction)
  - Gradients for each GPU are partitioned and sent to corresponding GPUs
  - Updated parameters are sent to all GPUs




# Idea

- **ZeRO-2.** Gradients are partitioned as well ( $\sim 8x$  memory reduction)
  - On **GPU i**, the gradients for **layer j** is:
    - Kept, if the GPU i is responsible for layer j
    - Discarded, otherwise,  
after computing the gradient for layer  $j-1$   
and transmitting to the responsible GPU



# Idea

- ZeRO-3. Even parameters are partitioned
  - Significant communication load; use when extremely memory-poor



M is the number of parameters, N is the number of devices.

	Optimizer States (12M)	Gradients (2M)	Model Weights (2M)	Memory Cost	Communication Cost
Data Parallelism	Replicated	Replicated	Replicated	$16M$	all-reduce(2M)
ZeRO Stage 1	Partitioned	Replicated	Replicated	$4M + \frac{12M}{N}$	all-reduce(2M)
ZeRO Stage 2	Partitioned	Partitioned	Replicated	$2M + \frac{14M}{N}$	all-reduce(2M)
ZeRO Stage 3	Partitioned	Partitioned	Partitioned	$\frac{16M}{N}$	1.5 all-reduce(2M)


# Further materials

- Cool explanatory video
  - <https://www.microsoft.com/en-us/research/blog/zero-deepspeed-new-system-optimizations-enable-training-models-with-over-100-billion-parameters/>
- Other advances
  - Memory checkpointing, offloading, and so on
  - <https://arxiv.org/abs/1910.02054>

# Gradient compression


# Motivation

- Recall that in DP, the key bottleneck is the **communication bandwidth**
  - Transmitting gradients & model updates




# Basic idea

- Compress the gradients using model compression techniques
- **Remark.** No longer need to take “inference efficiency” into account (e.g., no stringent need for linear quantization)
- Instead, encoding / decoding cost may be an issue



# Sparsity

- Select only **top-K gradients** (i.e., magnitude pruning)
  - What is not transmitted (“residuals”) are stored, for the next communication round




# Sparsity: Nitty-gritty details

- **Momentum.** Update the momentums based on the pruned gradient, not the original ones
- **Gradient Clipping.** Clip the gradients before adding the residuals
- **Warm-up.** Warm up both step size and sparsity

Task	Baseline	Deep Gradient Compression
<b>ResNet-50 On ImageNet</b>	Top-1 Accuracy	75.96% <b>76.15% (+0.19%)</b>
	Top-5 Accuracy	92.91% <b>92.97% (+0.06%)</b>
	Gradient Compression Ratio	1 × <b>277 ×</b>
<b>5-Layer GRU On LibriSpeech</b>	Word Error Rate (WER)	9.45% <b>9.06% (-0.39%)</b>
	Word Error Rate (WER)	27.07% <b>27.04% (-0.03%)</b>
	Gradient Compression Ratio	1 × <b>608 ×</b>
<b>2-Layer LSTM Language Model On Penn Treebank</b>	Perplexity	72.30 <b>72.24 (-0.06)</b>
	Gradient Compression Ratio	1 × <b>462 ×</b>


# Quantization

- In **1-bit SGD**, the gradients are quantized to binary values
  - Allocate column-wise scaling factors
  - Accumulate quantization errors




# All-reducing compressed gradients

- **Problem.** Suppose that we use all-reduce to aggregate gradient signals
  - Sparsity. No longer sparse
  - Quantization. No longer low-bit
  - Repeated pruning/quantization leads to much noise / order-dependency



# PowerSGD

- Apply low-rank approximation to gradients
    - Free of the order-dependency issue



---

**Algorithm 1** Rank- $r$  POWERSGD compression

```

1: The update vector  $\Delta_w$  is treated as a list of tensors corresponding to individual model parameters.  

   Vector-shaped parameters (biases) are aggregated uncompressed. Other parameters are reshaped  

   into matrices. The functions below operate on such matrices independently. For each matrix  

    $M \in \mathbb{R}^{n \times m}$ , a corresponding  $Q \in \mathbb{R}^{m \times r}$  is initialized from an i.i.d. standard normal distribution.  

2: function COMPRESS+AGGREGATE(update matrix  $M \in \mathbb{R}^{n \times m}$ , previous  $Q \in \mathbb{R}^{m \times r}$ )  

3:    $P \leftarrow MQ$   

4:    $P \leftarrow \text{ALL REDUCE MEAN}(P)$                                  $\triangleright$  Now,  $P = \frac{1}{W}(M_1 + \dots + M_W)Q$   

5:    $\hat{P} \leftarrow \text{ORTHOGONALIZE}(P)$                                  $\triangleright$  Orthonormal columns  

6:    $Q \leftarrow M^\top \hat{P}$   

7:    $Q \leftarrow \text{ALL REDUCE MEAN}(Q)$                                  $\triangleright$  Now,  $Q = \frac{1}{W}(M_1 + \dots + M_W)^\top \hat{P}$   

8:   return the compressed representation  $(\hat{P}, Q)$ .  

9: end function  

10: function DECOMPRESS( $\hat{P} \in \mathbb{R}^{n \times r}$ ,  $Q \in \mathbb{R}^{m \times r}$ )  

11:   return  $\hat{P}Q^\top$   


12: end function

```

# Automating model parallelism


# Inter-op vs Intra-op

- Roughly, there are two ways to distribute operations:
- **Inter-op.** Assign different operators to different devices  
(e.g., Pipeline parallel)
  - Good. Less communication
  - Bad. Much idle time




# Inter-op vs Intra-op

- **Intra-op.** Assign different regions of one operator to different devices (e.g., tensor parallel, data parallel)
  - Good. Devices stay busy all the time
  - Bad. Much communication
    - Replication & all-reduce



# Motivation

- **Question.** Which parallelism should I adopt, for my own model & cluster?




# Formulation

- Abstractly put, we want to solve:

$$\min_{\text{strategy}} \text{Cost}(\text{model}, \text{cluster}; \text{strategy})$$

- strategy is any possible combination of inter-op & intra-op parallelism




# Approaches

- There are quite many approaches:
  - **MCMC.** FlexFlow (2018)
  - **RL.** ColocRL (2017)
  - (...)
- A popular approach is called **Alpa**
  - Hierarchical optimization-based method


# Alpa

- Prioritize performing:
  - **Inter-op.** Between nodes (as it requires less comm)
  - **Intra-op.** Between devices, inside a node (as it requires more comm)




# Alpa

- The search space thus becomes smaller and structured




# Alpa

- Roughly, the search is done by a two-stage iterative optimization
  - **Inter-op.** Determine the group of ops to be done in a node
  - **Intra-op.** How to conduct tensor/data parallel inside a node




# Inter-op


- Given a computational graph,



- Determine the partition of the graph



or




or

...


# Inter-op

- Then, assign the nodes for each partition, via dynamic programming
  - **Required.** For this to be accurate, need a good latency estimate of each partition on the nodes



# Intra-op


- In each intra-op pass, we solve an optimization problem
  - Assignment problem (discrete decisions) with linear costs
  - a **mixed integer-linear programming!**




# Remarks


# Can we parallelize to infinity?

- Suppose that we can use infinite amount of GPUs
- **Question.** Can we make the batch size infinity, and finish training in seconds?
  - Answer. Unfortunately, no. We lose generalizability







(a) Simple CNN on MNIST




(b) Simple CNN on Fashion MNIST




(c) ResNet-8 on CIFAR-10




(d) ResNet-50 on ImageNet




(e) ResNet-50 on Open Images




(f) Transformer on LM1B



(g) Transformer on Common Crawl




(h) VGG-11 on ImageNet



(i) LSTM on LM1B

# Why?

- No complete answer, but some speculations...
  - Large batch → Small SGD noise → Trapped in local minima (narrow valley)



That's it for today

