

Matryoshka Quantization

Google DeepMind

2025. 04. 21

Efficient ML

Seung-taek Woo, Chiwoong Lee, Byeongho Yu

Preliminary

Quantization?

FP16

1.3	-0.5	-1.2
-0.4	0.2	1.1
0.1	-0.1	0.7

Preliminary

Quantization?

FP16			INT8		
1.3	-0.5	-1.2	127	-49	-117
-0.4	0.2	1.1	-39	20	108
0.1	-0.1	0.7	10	-10	68

* 97.7

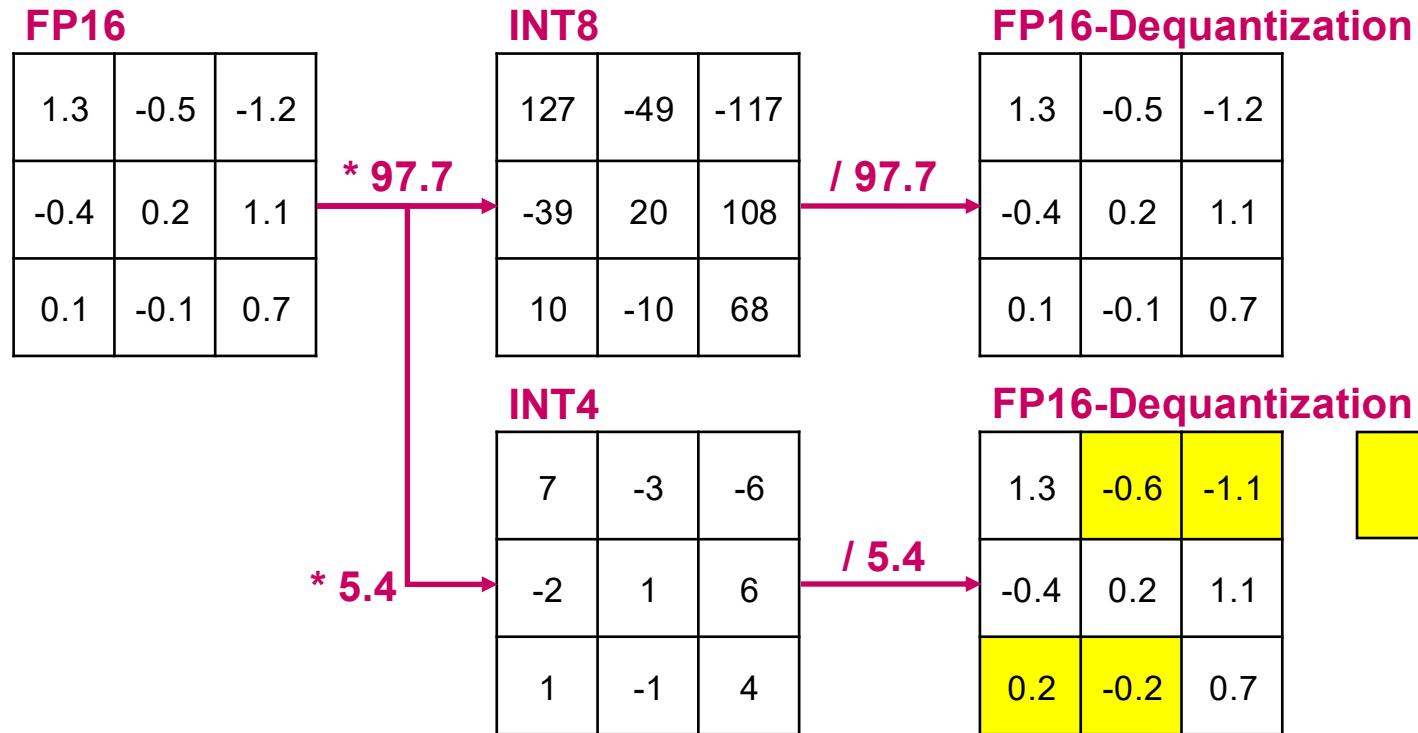
Preliminary

Quantization?



Preliminary

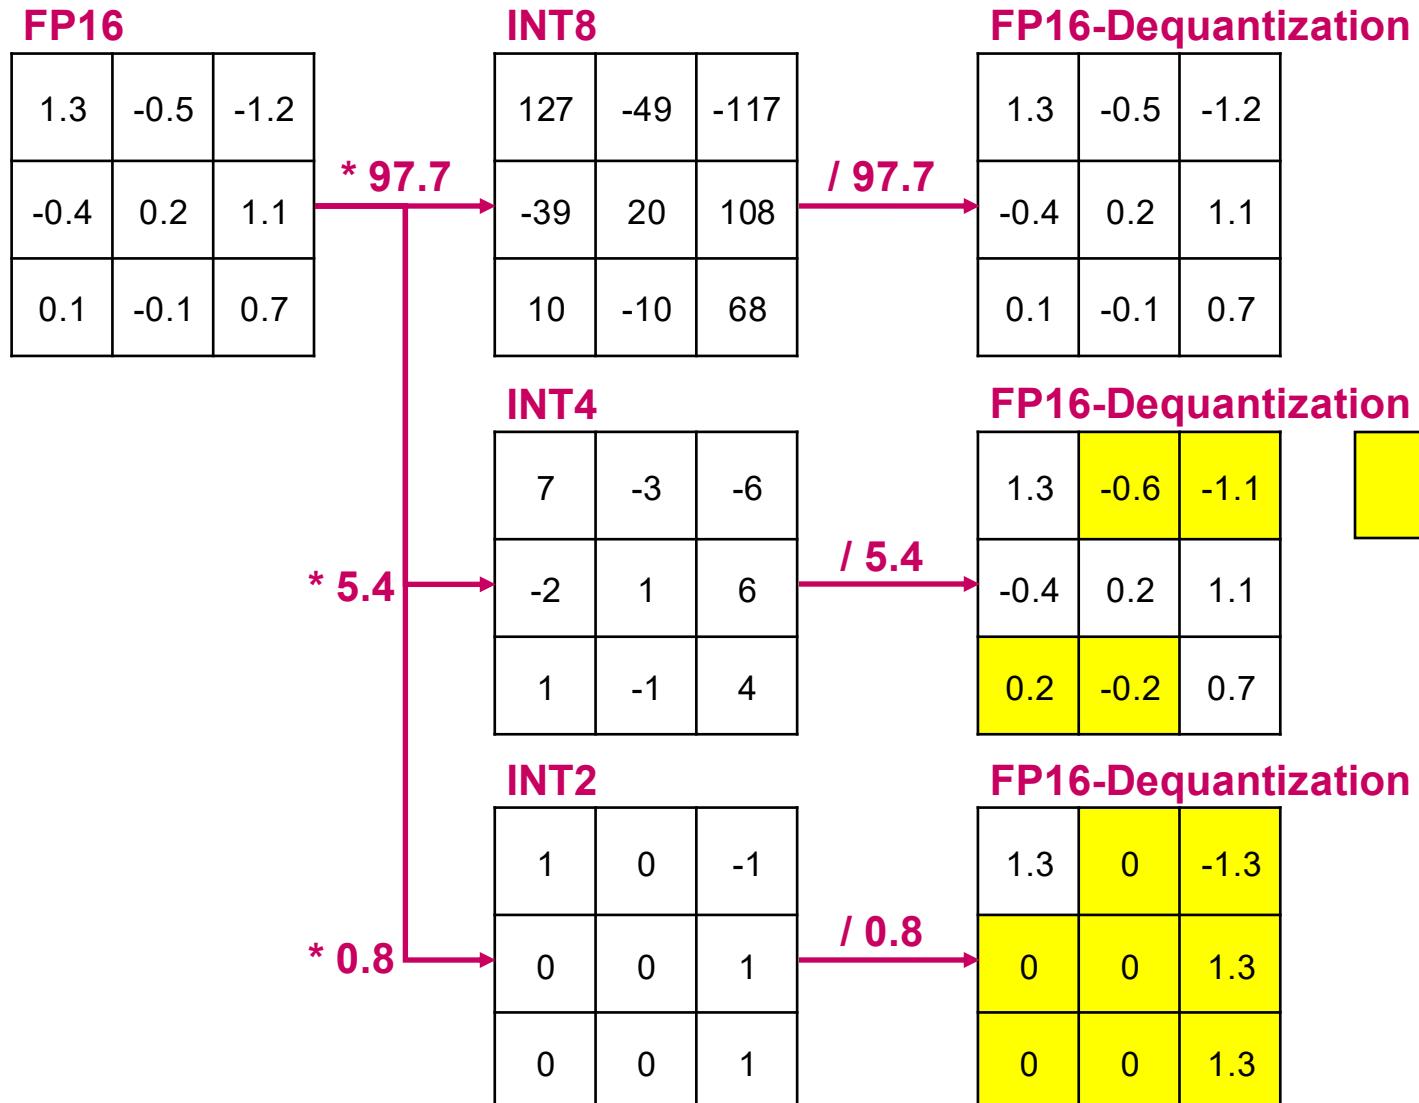
Quantization?



: Quantization Error

Preliminary

Quantization?



: Quantization Error

Introduction

LLM Service



ChatGPT



Claude

Gemini



Copilot

Foundation LLM



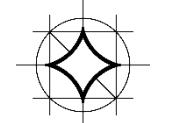
Llama



deepseek



Qwen



Gemma

Introduction

LLM Service



ChatGPT



100B



~1T



Copilot

Foundation LLM



Llama



1B~



700B



Gemma

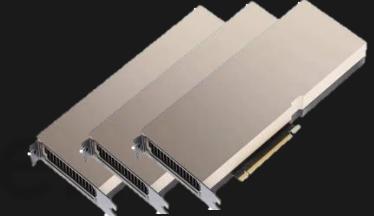
Introduction

LLM Service

100B(FP16) \approx 200GB

Foundation LLM

Three A100-80Gs are needed for
inference only.



Introduction

Problem: Need many GPUs.

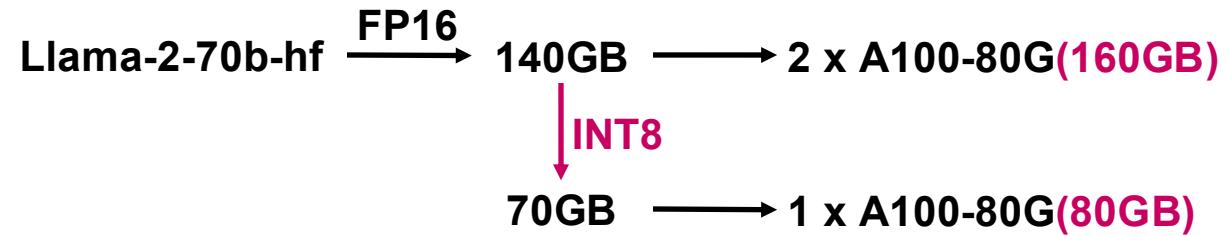
Quantization is the solution!

Llama-2-70b-hf $\xrightarrow{\text{FP16}}$ 140GB \longrightarrow 2 x A100-80G(160GB)

Introduction

Problem: Need many GPUs.

Quantization is the solution!



Introduction

Problem: Need many GPUs.

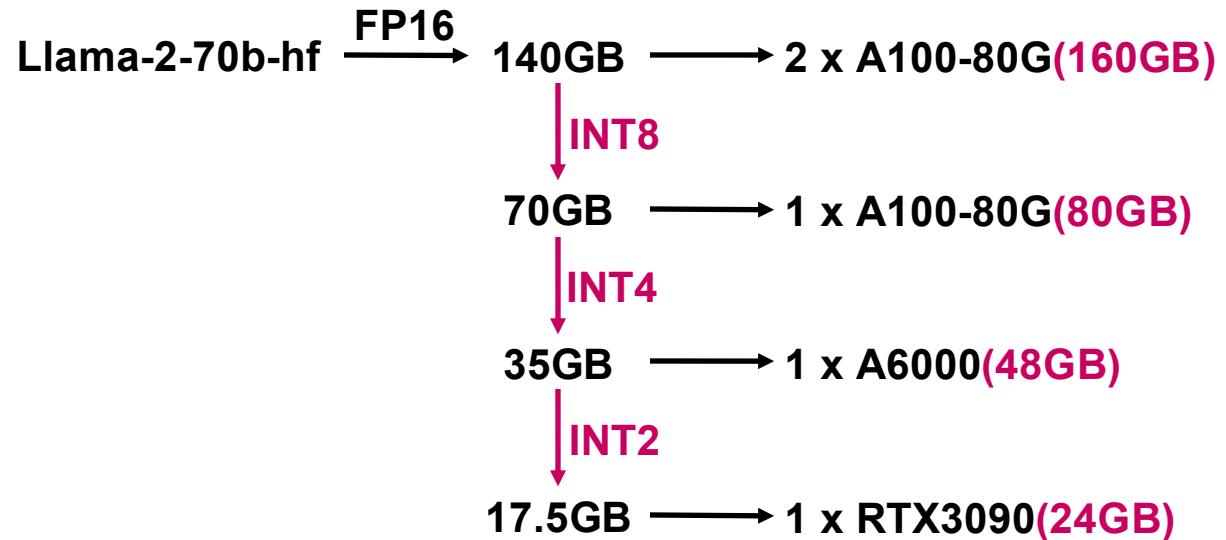
Quantization is the solution!



Introduction

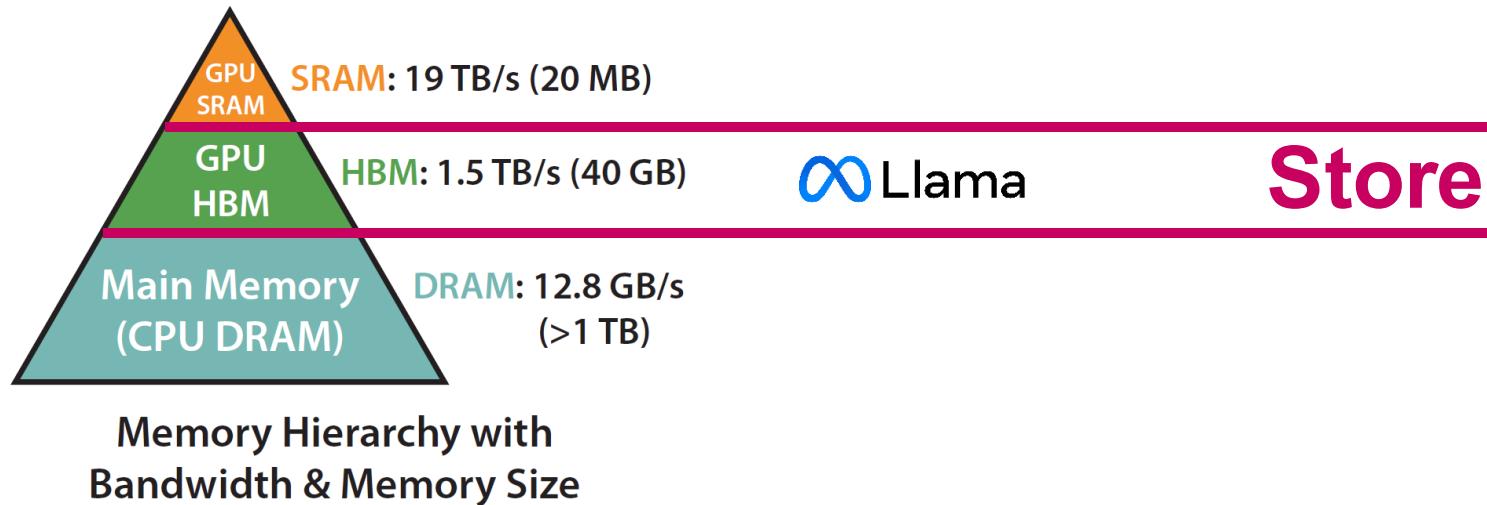
Problem: Need many GPUs.

Quantization is the solution!



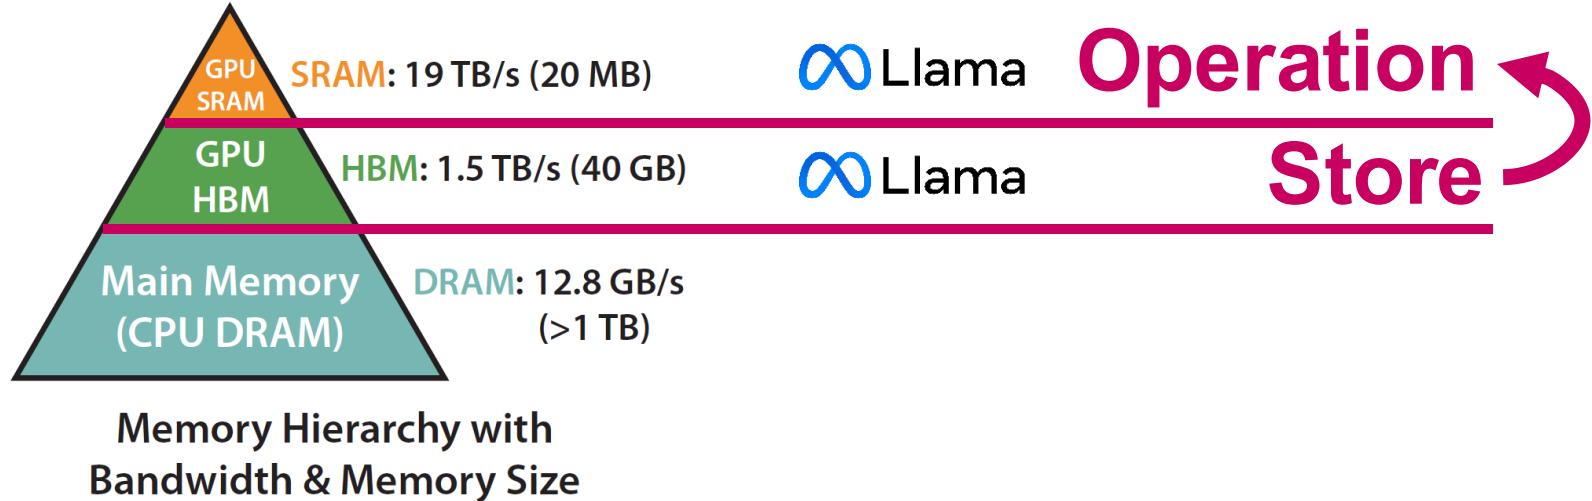
Introduction

GPU Memory Hierarchy

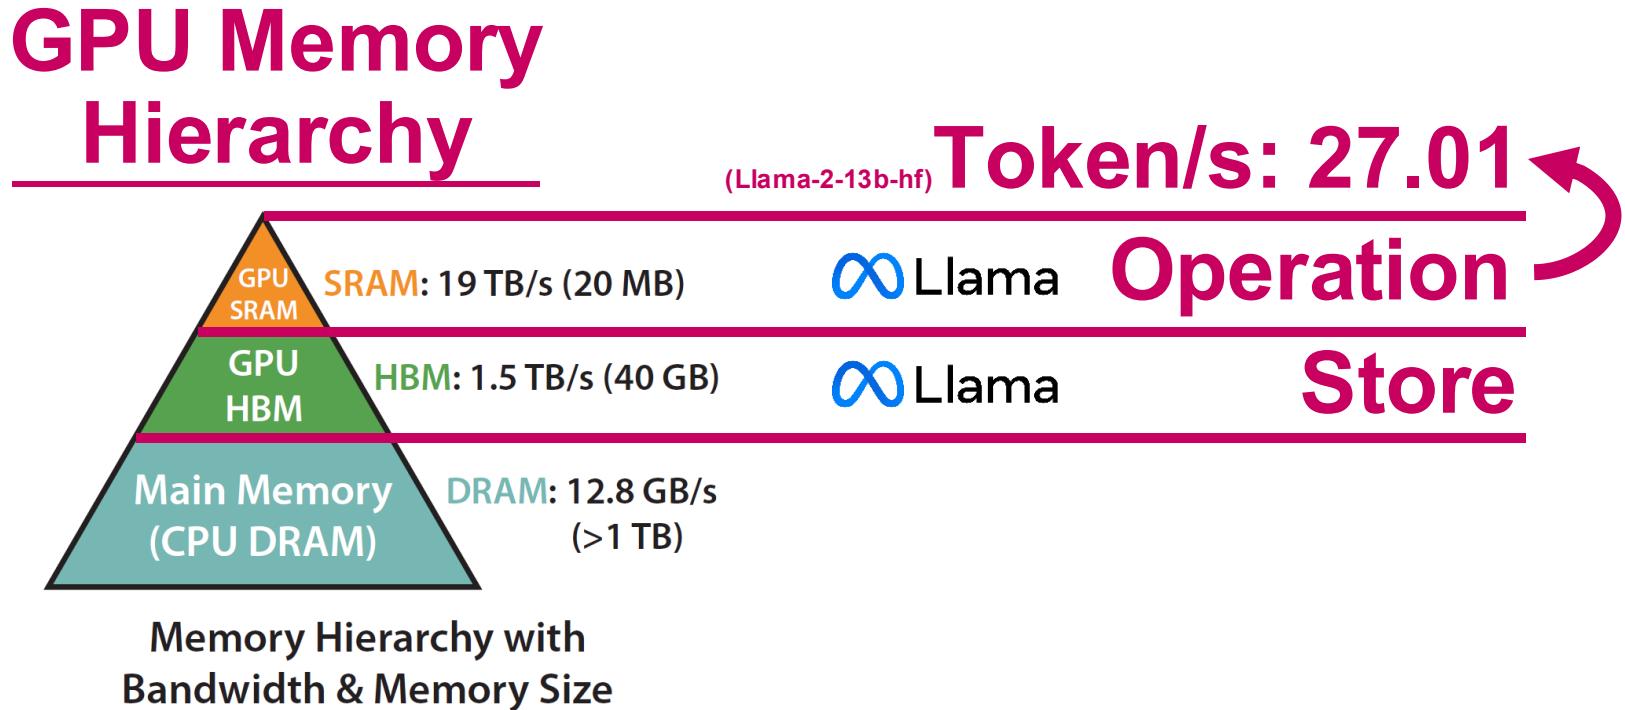


Introduction

GPU Memory Hierarchy



Introduction



Introduction

GPU Memory

13B(FP16) \approx 26GB

Hierarchy

Token/s: 27.01

Decoding latency is dominated
by Memory Bound.

Memory Hierarchy with
Bandwidth & Memory Size

Introduction

Problem: Low Speed.

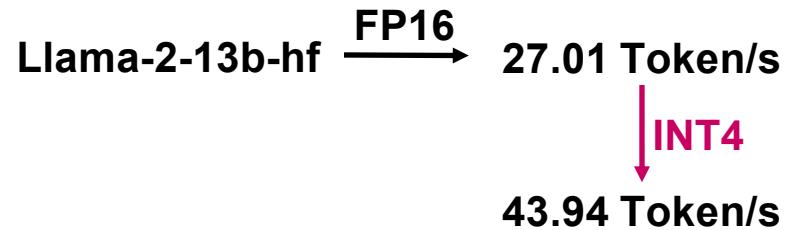
Quantization is the solution!

Llama-2-13b-hf $\xrightarrow{\text{FP16}}$ 27.01 Token/s

Introduction

Problem: Low Speed.

Quantization is the solution!



Introduction

Problem: Low Speed.

Quantization is the solution!



Introduction

However, current quantization methods^[1, 2, 3]...

FP16

1.3	-0.5	-1.2
-0.4	0.2	1.1
0.1	-0.1	0.7

Need to optimize independently to target precision.

[1] Lee, Changhun, et al. "Owq: Outlier-aware weight quantization for efficient fine-tuning and inference of large language models." *Proceedings of the AAAI Conference on Artificial Intelligence*. Vol. 38. No. 12. 2024.

[2] Lin, Ji, et al. "Awq: Activation-aware weight quantization for on-device lilm compression and acceleration." *Proceedings of Machine Learning and Systems* 6 (2024): 87-100.

[3] Frantar, Elias, et al. "Gptq: Accurate post-training quantization for generative pre-trained transformers." *arXiv preprint arXiv:2210.17323* (2022).

Introduction

However, current quantization methods^[1, 2, 3]...



FP16		
1.3	-0.5	-1.2
-0.4	0.2	1.1
0.1	-0.1	0.7

INT8		
127	-49	-117
-39	20	108
10	-10	68

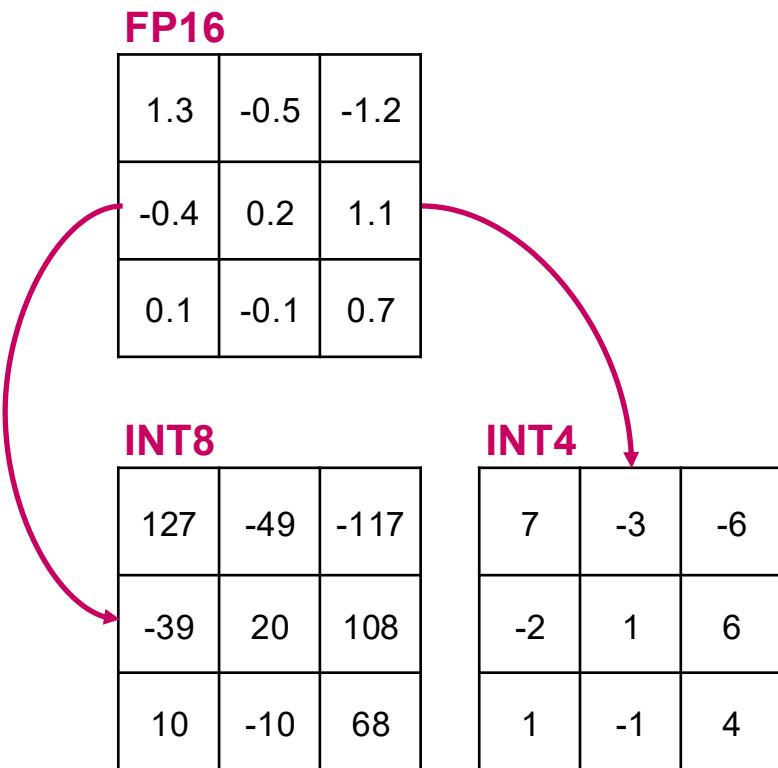
[1] Lee, Changhun, et al. "Owq: Outlier-aware weight quantization for efficient fine-tuning and inference of large language models." *Proceedings of the AAAI Conference on Artificial Intelligence*. Vol. 38. No. 12. 2024.

[2] Lin, Ji, et al. "Awq: Activation-aware weight quantization for on-device llm compression and acceleration." *Proceedings of Machine Learning and Systems* 6 (2024): 87-100.

[3] Frantar, Elias, et al. "Gptq: Accurate post-training quantization for generative pre-trained transformers." *arXiv preprint arXiv:2210.17323* (2022).

Introduction

However, current quantization methods^[1, 2, 3]...



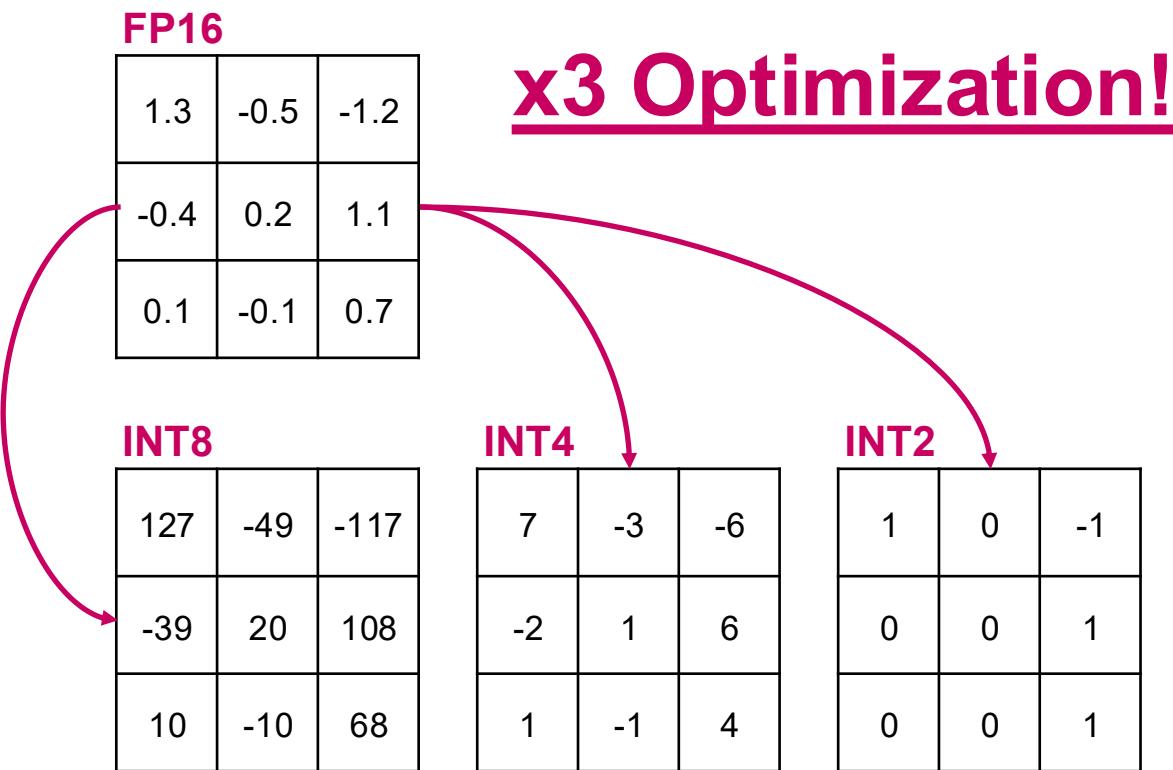
[1] Lee, Changhun, et al. "Owq: Outlier-aware weight quantization for efficient fine-tuning and inference of large language models." *Proceedings of the AAAI Conference on Artificial Intelligence*. Vol. 38. No. 12. 2024.

[2] Lin, Ji, et al. "Awq: Activation-aware weight quantization for on-device llm compression and acceleration." *Proceedings of Machine Learning and Systems* 6 (2024): 87-100.

[3] Frantar, Elias, et al. "Gptq: Accurate post-training quantization for generative pre-trained transformers." *arXiv preprint arXiv:2210.17323* (2022).

Introduction

However, current quantization methods^[1, 2, 3]...



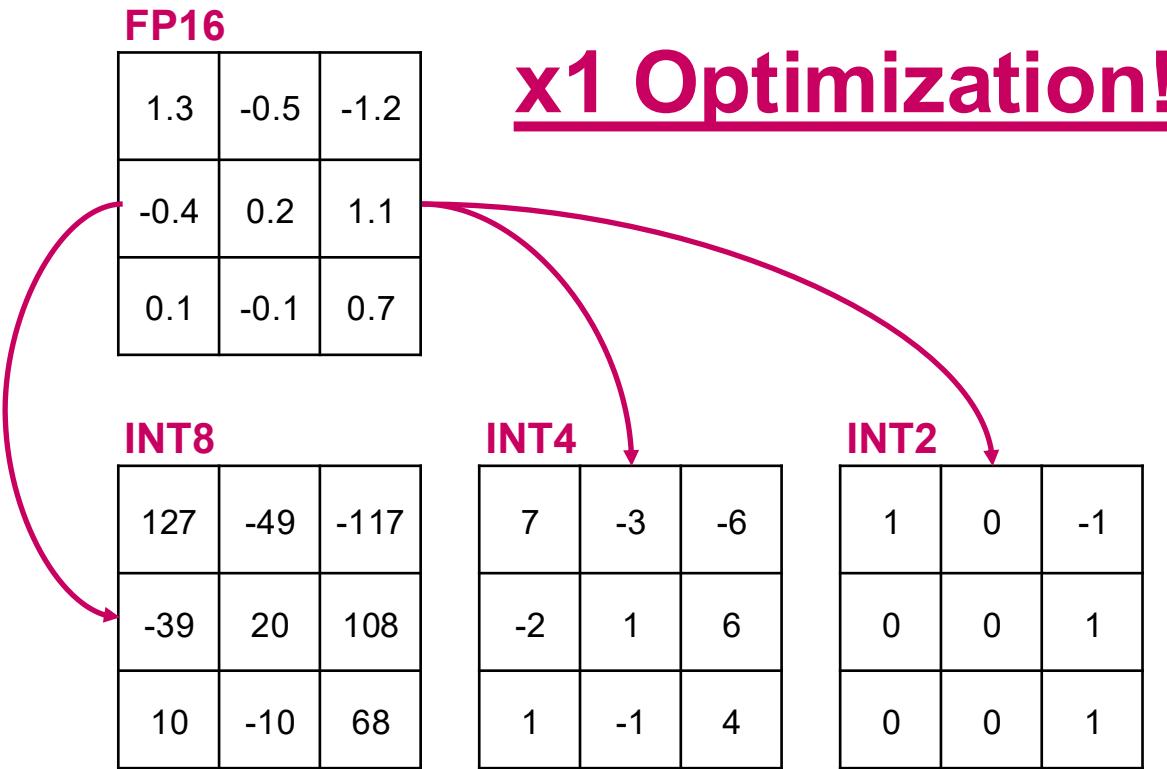
[1] Lee, Changhun, et al. "Owq: Outlier-aware weight quantization for efficient fine-tuning and inference of large language models." *Proceedings of the AAAI Conference on Artificial Intelligence*. Vol. 38. No. 12. 2024.

[2] Lin, Ji, et al. "Awq: Activation-aware weight quantization for on-device llm compression and acceleration." *Proceedings of Machine Learning and Systems* 6 (2024): 87-100.

[3] Frantar, Elias, et al. "Gptq: Accurate post-training quantization for generative pre-trained transformers." *arXiv preprint arXiv:2210.17323* (2022).

Introduction

Question: Can I extract multiple low-precision from a single optimization?



Introduction

Matryoshka



Introduction

Matryoshka



INT8

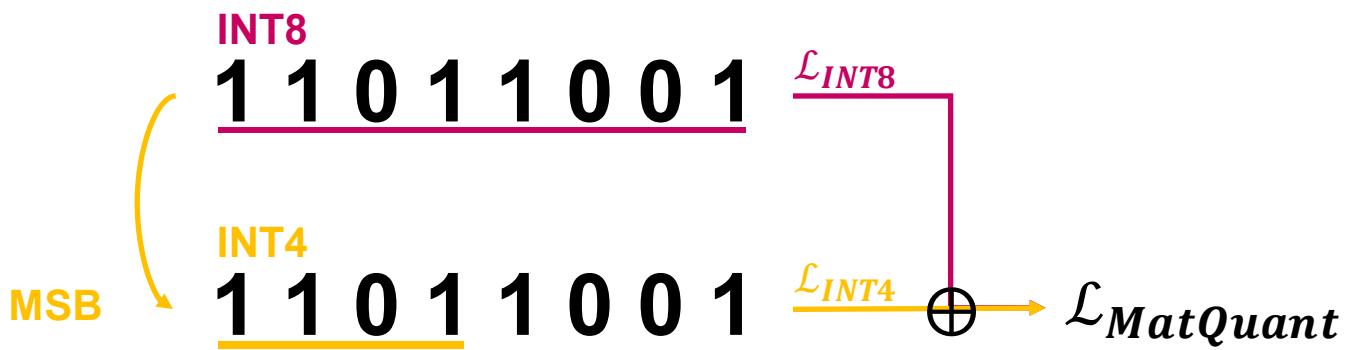
1 1 0 1 1 0 0 1

\mathcal{L}_{INT8}

$\mathcal{L}_{MatQuant}$

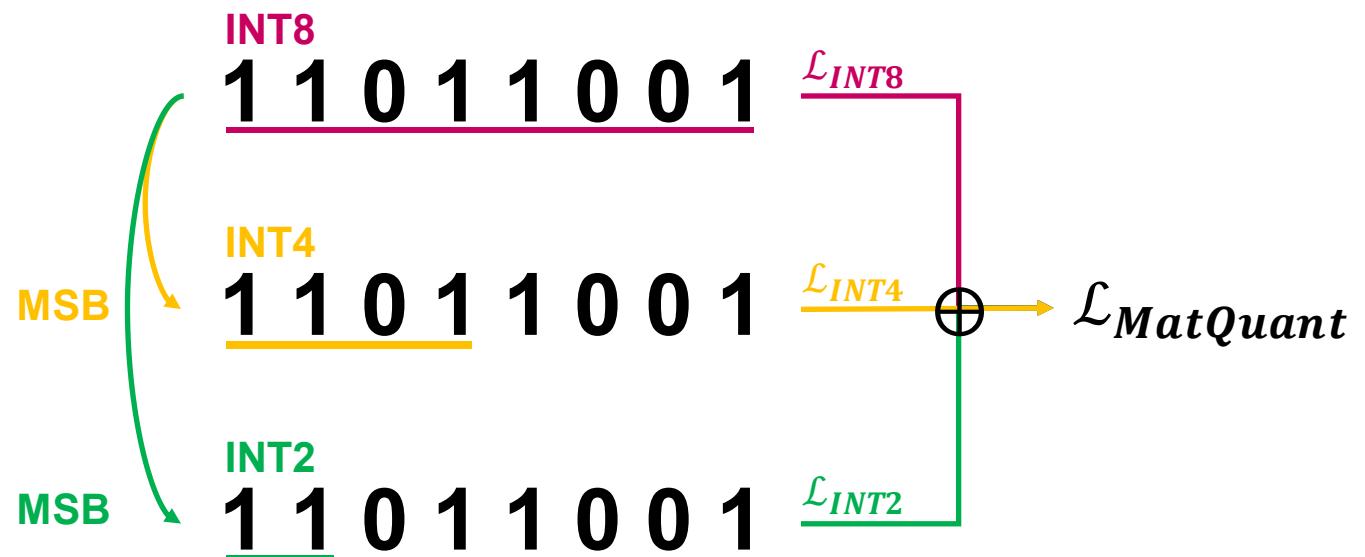
Introduction

Matryoshka



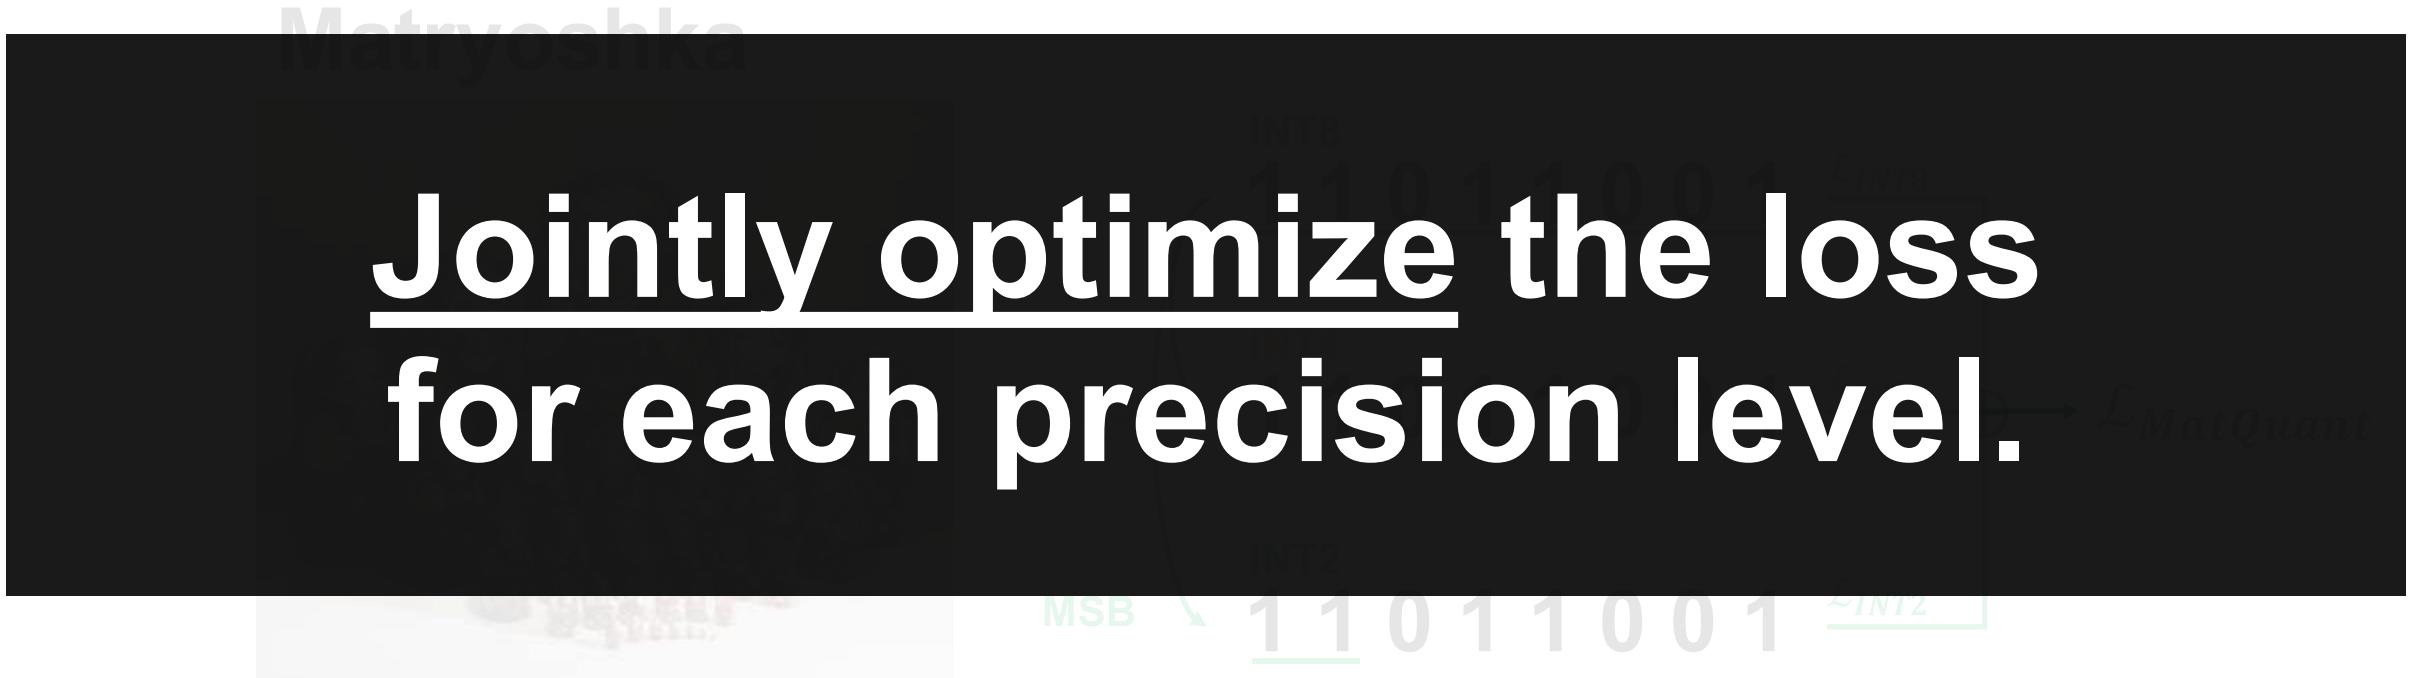
Introduction

Matryoshka



Introduction

Jointly optimize the loss
for each precision level.



Preliminaries

Quantization Aware Training (QAT)

- Quantization Aware Training (QAT) learns a c-bit quantized model by minimizing end-to-end cross-entropy loss via gradient descent.
- It uses quantized weights during the forward pass and applies a **Straight-Through Estimator (STE)** to backpropagate gradients through the non-differentiable quantization operation.

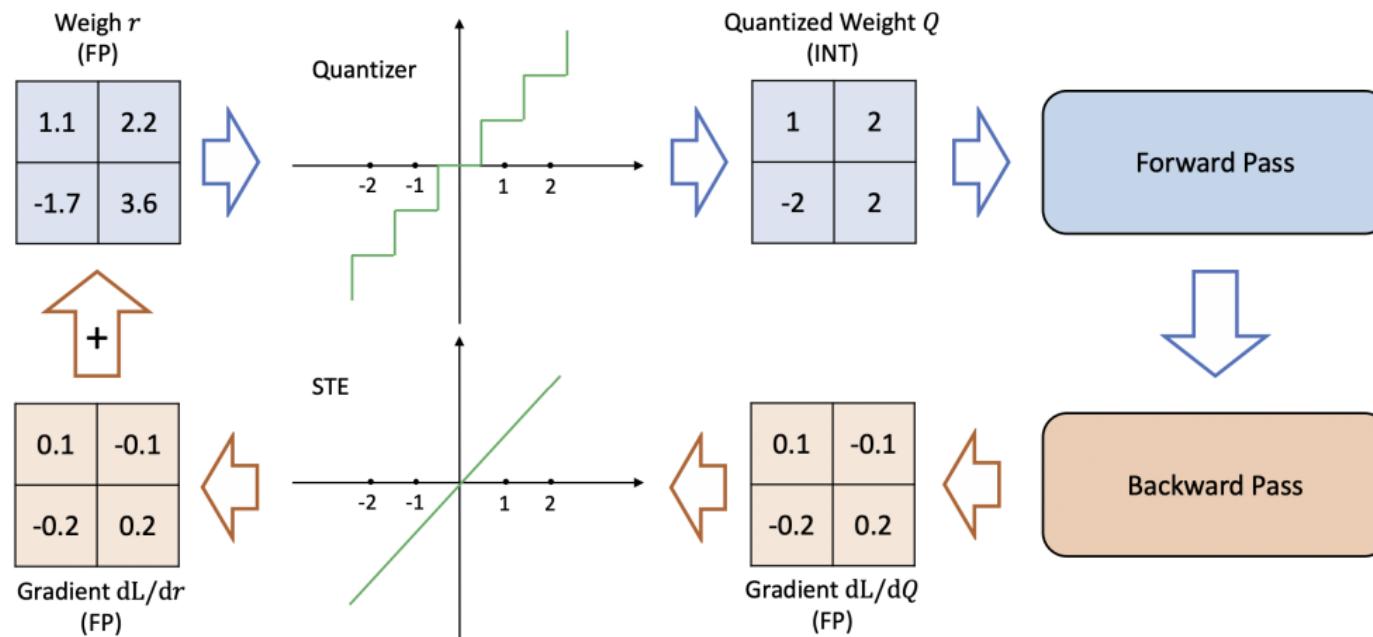


Figure 5: Illustration of Quantization-Aware Training procedure, including the use of Straight Through Estimator (STE).

Preliminaries

Quantization Aware Training (QAT)

- Quantization Aware Training (QAT) minimizes end-to-end cross-entropy loss via gradient descent.
- It uses quantized weights during the non-differentiable quantization step.
$$\alpha = \frac{\max(w) - \min(w)}{2^c - 1}, \quad z = -\frac{\min(w)}{\alpha}$$
$$Q_{MM}(w, c) = \text{clamp}\left(\left\lfloor \frac{w}{\alpha} + z \right\rfloor, 0, 2^c - 1\right)$$

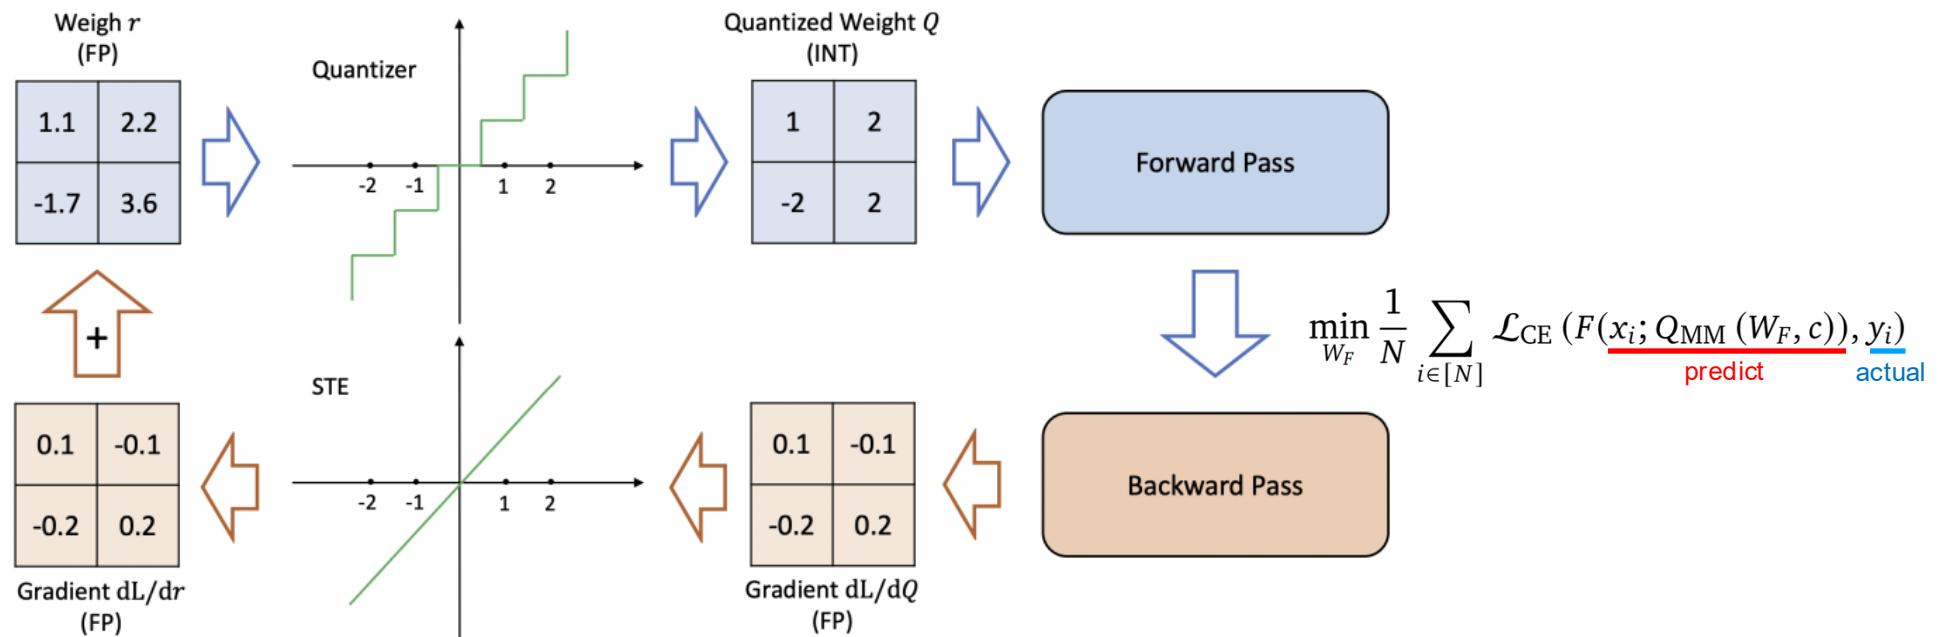


Figure 5: Illustration of Quantization-Aware Training procedure, including the use of Straight Through Estimator (STE).

Preliminaries

Quantization Aware Training (QAT)

- Quantization Aware Training (QAT) minimizes end-to-end cross-entropy loss via gradient descent.
 - It uses quantized weights during the non-differentiable quantization through Estimator (STE) to backpropagate gradients through STE to backpropagate gradients through
- $$Q_{\text{MM}}(w, c) = \text{clamp}\left(\left\lfloor \frac{w}{\alpha} + z \right\rfloor, 0, 2^c - 1\right)$$
- $$\alpha = \frac{\max(w) - \min(w)}{2^c - 1}, \quad z = -\frac{\min(w)}{\alpha}$$

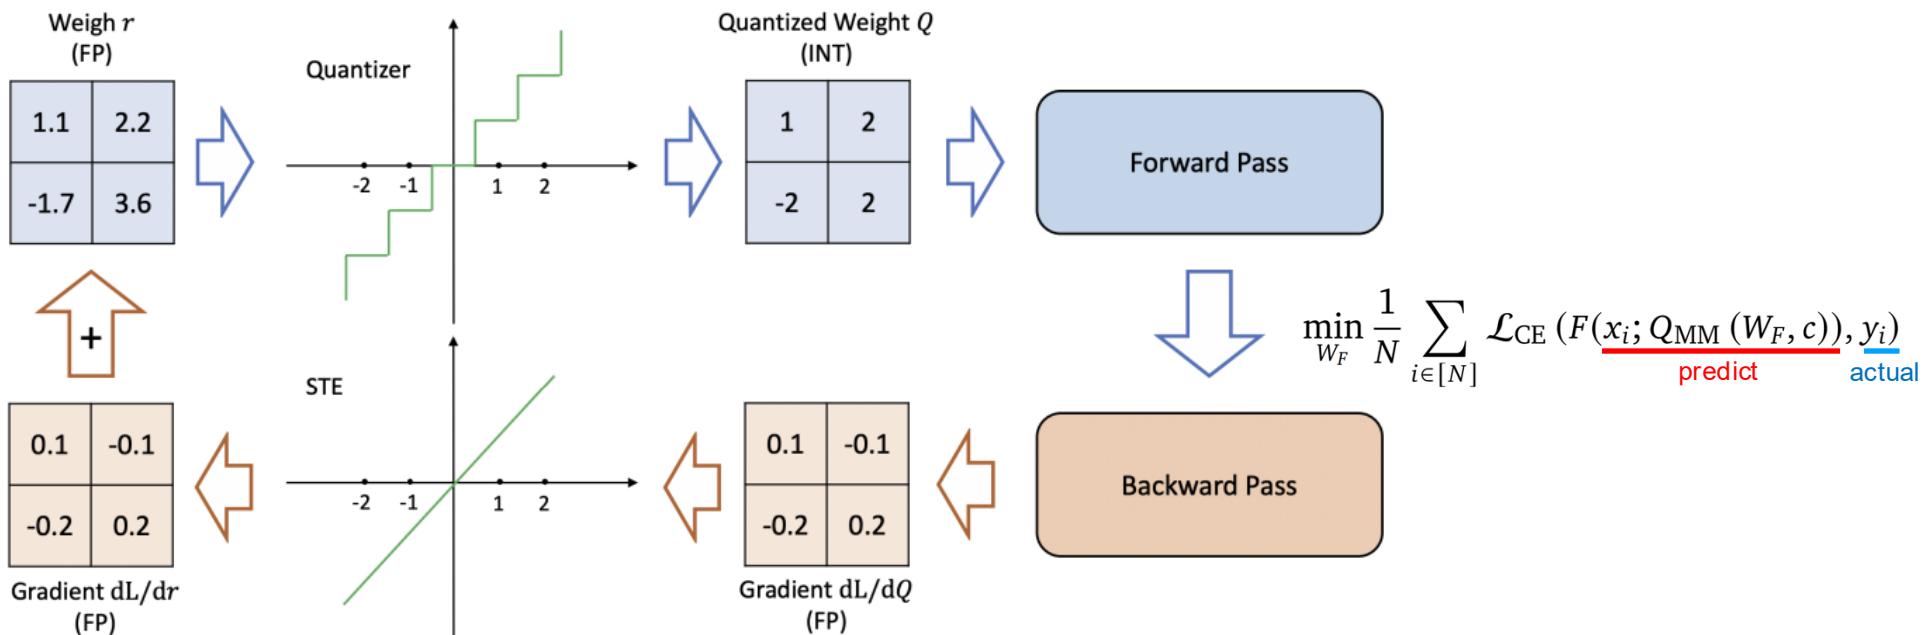


Figure 5: Illustration of Quantization-Aware Training procedure, including the use of Straight Through Estimator (STE).

Preliminaries

Quantization Aware Training (QAT)

- Quantization Aware Training (QAT) minimizes end-to-end cross-entropy loss via gradient descent.
 - It uses quantized weights during the non-differentiable quantization through Estimator (STE) to backpropagate gradients through Estimator (STE) to backpropagate gradients through
- $$Q_{\text{MM}}(w, c) = \text{clamp}\left(\left[\frac{w}{\alpha} + z\right], 0, 2^c - 1\right)$$
- $$\alpha = \frac{\max(w) - \min(w)}{2^c - 1}, \quad z = -\frac{\min(w)}{\alpha}$$

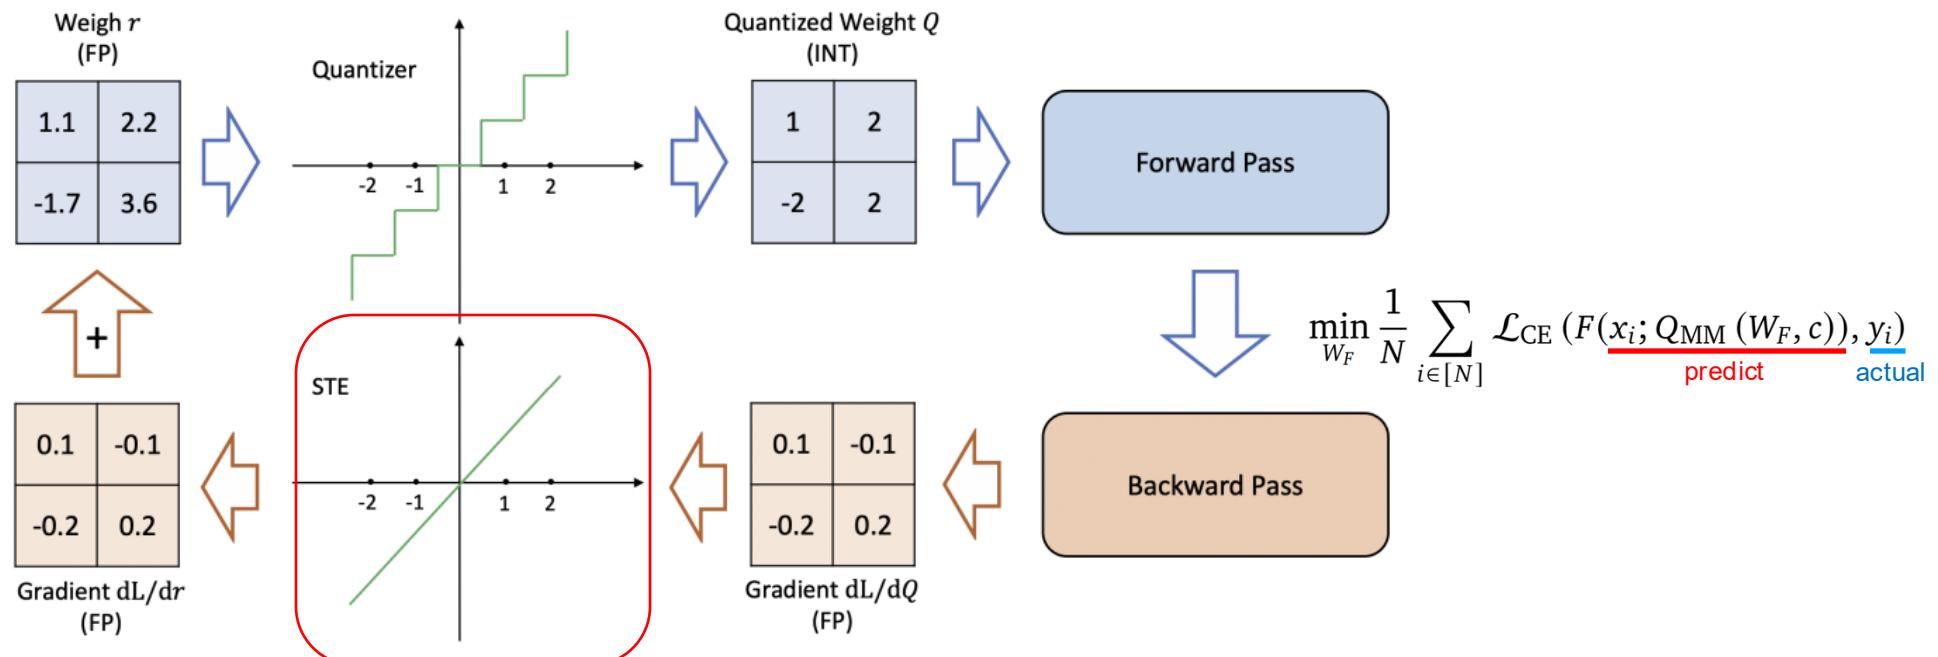


Figure 5: Illustration of Quantization Aware Training (QAT) procedure, including the use of Straight Through Estimator (STE).

$$\frac{d\mathcal{L}}{dr} = \frac{d\mathcal{L}}{dQ} \cdot \frac{dQ}{dr} \approx \frac{d\mathcal{L}}{dQ}$$

Preliminaries

OmniQuant (ICLR2024, Spotlight)

- Unlike QAT, OmniQuant does not update the model parameters.
- Instead, it learns additional scaling and shifting parameters through gradient descent over layer-wise L2 error reconstruction.

$$Q_{\text{MM}}(w, c) = \text{clamp}\left(\left\lfloor \frac{w}{\alpha} + z \right\rfloor, 0, 2^c - 1\right)$$
$$\alpha = \frac{\max(w) - \min(w)}{2^c - 1}, \quad z = -\frac{\min(w)}{\alpha}$$

$$Q_{\text{Omni}}(w, c) = \text{clamp}\left(\left\lfloor \frac{w}{\alpha} + z \right\rfloor, 0, 2^c - 1\right)$$
$$\alpha = \frac{\gamma \cdot \max(w) - \beta \cdot \min(w)}{2^c - 1}, \quad z = -\frac{\beta \cdot \min(w)}{\alpha}$$

QAT

OmniQuant

Preliminaries

OmniQuant (ICLR2024, Spotlight)

- Unlike QAT, OmniQuant does not update the model parameters.
- Instead, it learns additional scaling and shifting parameters through gradient descent over layer-wise L2 error reconstruction.

$$Q_{\text{MM}}(w, c) = \text{clamp}\left(\left\lfloor \frac{w}{\alpha} + z \right\rfloor, 0, 2^c - 1\right)$$
$$\alpha = \frac{\max(w) - \min(w)}{2^c - 1}, \quad z = -\frac{\min(w)}{\alpha}$$

$$XW+b \rightarrow X \cdot Q_{\text{MM}}(W) + b$$

QAT

$$Q_{\text{Omni}}(w, c) = \text{clamp}\left(\left\lfloor \frac{w}{\alpha} + z \right\rfloor, 0, 2^c - 1\right)$$
$$\alpha = \frac{\gamma \cdot \max(w) - \beta \cdot \min(w)}{2^c - 1}, \quad z = -\frac{\beta \cdot \min(w)}{\alpha}$$

$$XW+b \rightarrow ((X - \underbrace{\delta}_{\text{Shifting Factor}}) \oslash \underbrace{s}_{\text{Smoothing Factor}}) \cdot Q_{\text{Omni}}(W \oslash s) + b + \underbrace{\delta \cdot W}_{\text{Shifting Factor}}$$

$$X \in \mathbb{R}^{n \times d}$$
$$W \in \mathbb{R}^{d \times d_0}$$

$$b \in \mathbb{R}^{d_0}$$

$$\delta \in \mathbb{R}^d$$

$$s \in \mathbb{R}^d$$

OmniQuant

Preliminaries

OmniQuant (ICLR2024, Spotlight)

- Unlike QAT, OmniQuant does not update the model parameters.
- Instead, it learns additional scaling and shifting parameters through gradient descent over layer-wise L2 error reconstruction.

$$Q_{\text{MM}}(w, c) = \text{clamp}\left(\left\lfloor \frac{w}{\alpha} + z \right\rfloor, 0, 2^c - 1\right)$$
$$\alpha = \frac{\max(w) - \min(w)}{2^c - 1}, \quad z = -\frac{\min(w)}{\alpha}$$

$$XW+b \rightarrow X \cdot Q_{\text{MM}}(W) + b$$

Cross Entropy Loss

$$\min_{W_F} \frac{1}{N} \sum_{i \in [N]} \mathcal{L}_{\text{CE}}(F(x_i; Q_{\text{MM}}(W_F, c)), y_i)$$

QAT

$$Q_{\text{Omni}}(w, c) = \text{clamp}\left(\left\lfloor \frac{w}{\alpha} + z \right\rfloor, 0, 2^c - 1\right)$$
$$\alpha = \frac{\gamma \cdot \max(w) - \beta \cdot \min(w)}{2^c - 1}, \quad z = -\frac{\beta \cdot \min(w)}{\alpha}$$

$$XW+b \rightarrow ((X - \underbrace{\delta}_{\text{Shifting Factor}}) \oslash \underbrace{s}_{\text{Smoothing Factor}}) \cdot Q_{\text{Omni}}(W \oslash s) + b + \underbrace{\delta \cdot W}_{\text{Shifting Factor}}$$

Layer-Wise L2 Error

$$\min_{\gamma, \beta, \delta, s} \|F_l(W_F^l, X_l) - F_l(Q_{\text{Omni}}(W_F^l), X_l)\|_2^2$$

OmniQuant

$$X \in \mathbb{R}^{n \times d}$$
$$W \in \mathbb{R}^{d \times d_o}$$

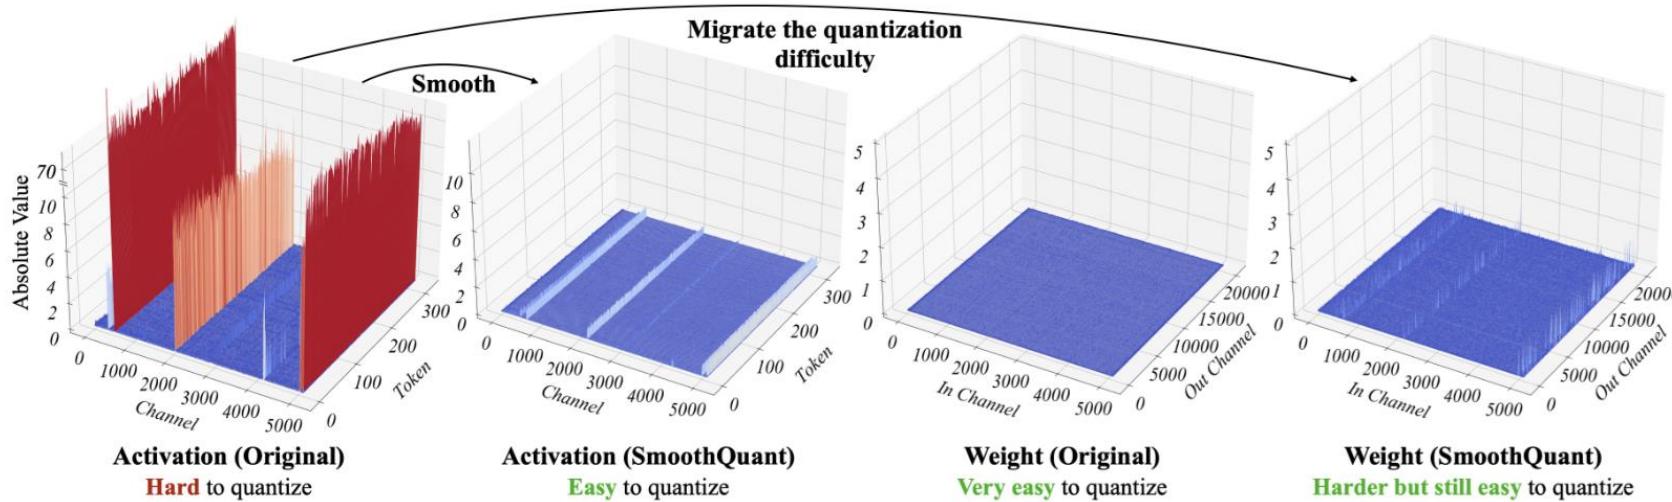
$$b \in \mathbb{R}^{d_o}$$
$$\delta \in \mathbb{R}^d$$
$$s \in \mathbb{R}^d$$

Preliminaries

Smoothing Factor ; s

$$XW+b \rightarrow ((X - \delta) \oslash s) \cdot Q_{\text{Omni}}(W \oslash s) + b + \delta \cdot W$$

- The smoothing factor redistributes the quantization difficulty caused by activation outliers to the weights.
- The smoothing factor enables a mathematically equivalent transformation. $Y = (X \text{diag}(s)^{-1}) \cdot (\text{diag}(s)W) = \hat{X}\hat{W}$



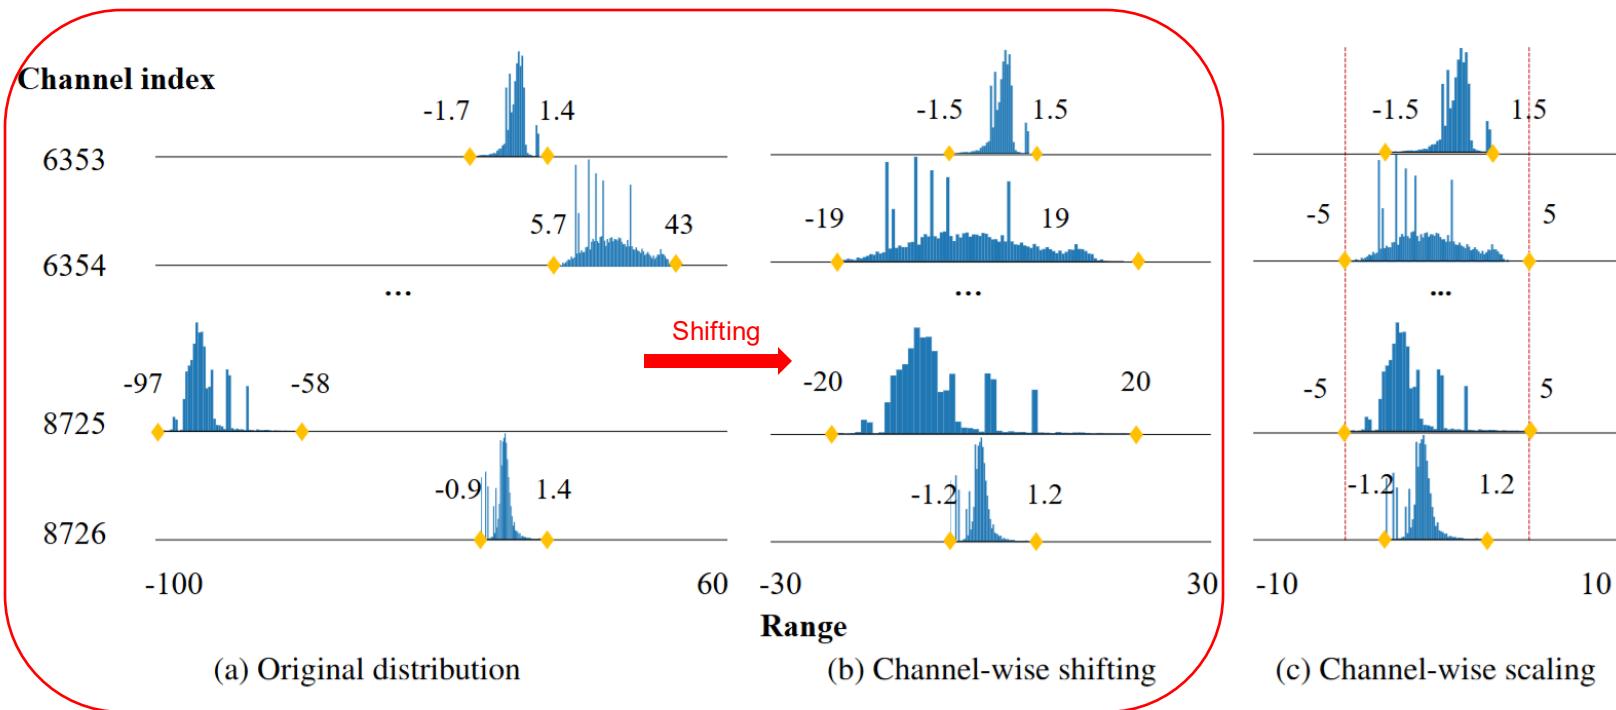
SmoothQuant: Accurate and Efficient Post-Training Quantization for Large Language Models

Preliminaries

Shifting Factor ; δ

$$XW+b \rightarrow ((X - \delta) \oslash s) \cdot Q_{\text{Omni}}(W \odot s) + b + \delta \cdot W$$

- The shifting factor aligns channel centers to remove asymmetric outliers, making the distribution easier to quantize.
- The shifting factor enables a mathematically equivalent transformation.



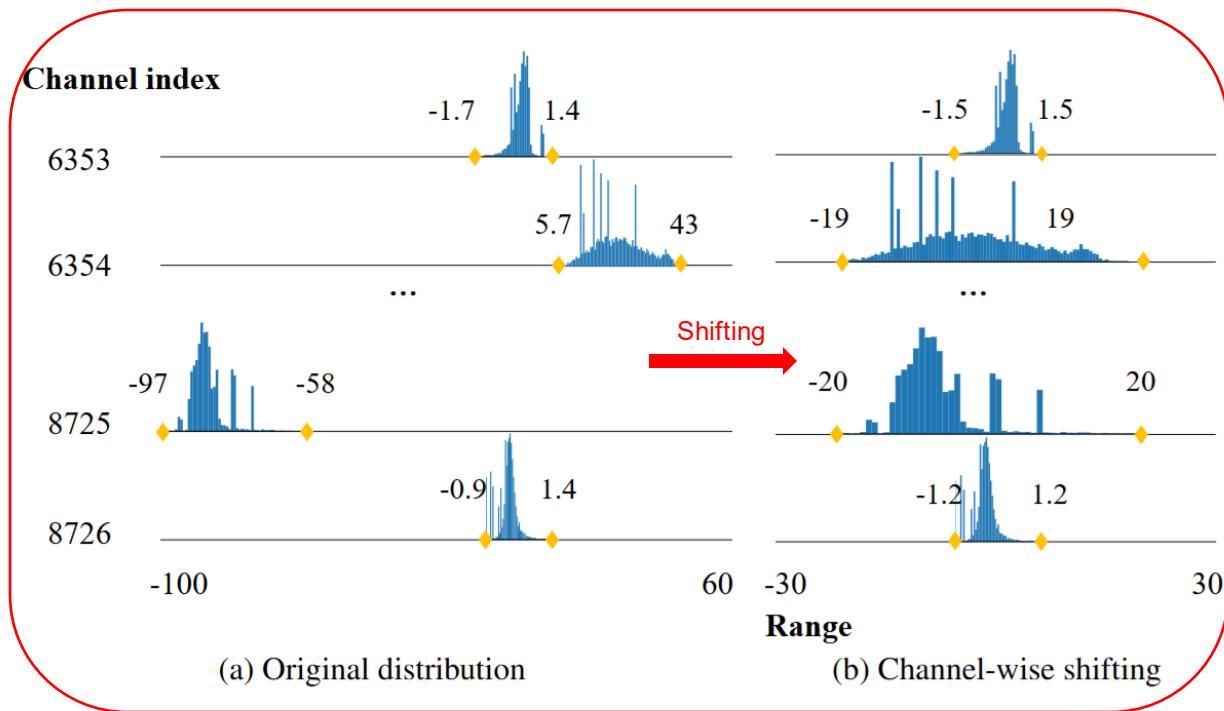
Outlier Suppression+ : Accurate quantization of large language models by equivalent and optimal shifting and scaling

Preliminaries

Shifting Factor ; δ

$$XW+b \rightarrow ((X - \delta) \oslash s) \cdot Q_{\text{Omni}}(W \odot s) + b + \delta \cdot W$$

- The shifting factor aligns channel centers to remove asymmetric outliers, making the distribution **easier to quantize**.
- The shifting factor enables a **mathematically equivalent** transformation.



$$\alpha = \frac{\max(w) - \min(w)}{2^c - 1}$$

Assuming $c = 8$ (bit)

(Before shifting)

$$\alpha = \frac{43 - (-97)}{255} = 0.549$$

(After shifting)

$$\alpha = \frac{20 - (-20)}{255} = 0.157$$

Outlier Suppression+ : Accurate quantization of large language models by equivalent and optimal shifting and scaling

Method

MatQuant

- If we want to extract a **r-bit** model from a **c-bit** model ($0 < r < c$), we can just **slice out** the r most significant bits (MSBs) – using a right shift, followed by a left shift of the same order.

$$q^c = Q(w, c) = \text{clamp}\left(\left\lfloor \frac{w}{\alpha} + z \right\rfloor, 0, 2^c - 1\right)$$

$$S(q^c, r) = \text{clamp}\left(\left\lfloor \frac{q^c}{2^{c-r}} \right\rfloor, 0, 2^r - 1\right) * 2^{c-r}$$

- Example
 - $c=8, r=4$ (8bit \rightarrow 4bit)
 - $q_8=234$

$$\underbrace{\left\lfloor \frac{234}{16} \right\rfloor}_{= \lfloor 14.625 \rfloor = 14} \xrightarrow{\text{clamp}} 14 \xrightarrow{\times 16} 224$$

1 1 1 0 1 0 1 0 \rightarrow 1 1 1 0 \rightarrow 1 1 1 0 0 0 0 0
INT8 INT4

Method

MatQuant

- If we want to extract a **r -bit** model from a **c -bit** model ($0 < r < c$), we can just **slice out** the r most significant bits (MSBs) – using a right shift, followed by a left shift of the same order.

$$q^c = Q(w, c) = \text{clamp}\left(\left\lfloor \frac{w}{\alpha} + z \right\rfloor, 0, 2^c - 1\right)$$

$$S(q^c, r) = \text{clamp}\left(\left\lfloor \frac{q^c}{2^{c-r}} \right\rfloor, 0, 2^r - 1\right) * 2^{c-r}$$

- MatQuant's overall objective (Weight Quantization on FFN)

$$\min_P \frac{1}{N} \sum_{i \in [N]} \sum_{r \in R} \lambda_r \cdot \mathcal{L}(F(S(Q(\theta, c), r), x'_i), y'_i)$$

$$R = \{8, 4, 2\}$$

λ_r = Loss reweighing factor for bit-width r

Experiment Setting

MatQuant working with two popular **learning based quantization methods**:

1. **OmniQuant**
2. **QAT**

Models & Target Bit precisions

- Gemma-2 2B, 9B / Mistral 7B models.
- Default target quantization precisions : **int8, int4, int2**
 - + the interpolative nature of MatQuant through evaluations on **int6** and **int3**

Training

OmniQuant

- 128 examples with a sequence length of 2048 from the **C4 dataset** train using a batch size of 4
- train for a total of 10M tokens for all models except the int2 baseline,
where we train the model for 20M tokens

QAT

- sample a fixed set of 100M tokens from the **C4 dataset** ,
and train all our models using a batch size of 16 and a sequence length of 8192 for a single epoch

Evaluation Datasets

Calculating Perplexity with C4's test set

Downstream evaluations with zero-shot accuracy

- ARC-c, ARC-e.
- BoolQ
- HellaSwag
- PIQA
- Winogrande

Q. What is PPL ?

A. Perplexity (PPL) is a metric that measures **how well a language model predicts a sequence**. **lower PPL values indicate better performance.**

MatQuant with OmniQuant

Data type	Method	Gemma-2 2B		Gemma-2 9B		Mistral 7B	
		OmniQuant	Task Avg.	log pplx.	Task Avg.	log pplx.	Task Avg.
bfloat16		68.21	2.551	74.38	2.418	73.99	2.110
int8	Baseline	68.25	2.552	74.59	2.418	73.77	2.110
	MatQuant	68.02	2.570	74.05	2.438	73.65	2.125
int4	Sliced int8	62.87	2.730	72.26	2.480	38.51	4.681
	Baseline	67.03	2.598	74.33	2.451	73.62	2.136
	MatQuant	66.58	2.618	73.83	2.491	73.06	2.153
int2	Sliced int8	39.78	17.030	38.11	15.226	37.29	11.579
	Baseline	51.33	3.835	60.24	3.292	59.74	3.931
	MatQuant	52.37	3.800	63.35	3.187	62.75	3.153

MatQuant with OmniQuant

Data type	Method	Gemma-2 2B		Gemma-2 9B		Mistral 7B	
		OmniQuant	Task Avg.	log pplx.	Task Avg.	log pplx.	Task Avg.
bfloat16		68.21	2.551	74.38	2.418	73.99	2.110
int8	Baseline	68.25	2.552	74.59	2.418	73.77	2.110
	MatQuant	68.02	2.570	74.05	2.438	73.65	2.125
Baseline (OmniQuant) is better, but MatQuant shows comparable performance							81
int4	Baseline	67.03	2.598	74.33	2.451	73.62	2.136
	MatQuant	66.58	2.618	73.83	2.491	73.06	2.153
int2	Sliced int8	39.78	17.030	38.11	15.226	37.29	11.579
	Baseline	51.33	3.835	60.24	3.292	59.74	3.931
	MatQuant	52.37	3.800	63.35	3.187	62.75	3.153

MatQuant with OmniQuant

Data type	Method	Gemma-2 2B		Gemma-2 9B		Mistral 7B	
		OmniQuant	Task Avg.	log pplx.	Task Avg.	log pplx.	Task Avg.
bfloat16			68.21	2.551	74.38	2.418	73.99
int8	Baseline	68.25	2.552	74.59	2.418	73.77	2.110
	MatQuant	68.02	2.570	74.05	2.438	73.65	2.125
int4	Sliced int8	62.87	2.730	72.26	2.480	38.51	4.681
	Baseline	67.03	2.598	74.33	2.451	73.62	2.136
	MatQuant	66.58	2.618	73.83	2.491	73.06	2.153
int2	Sliced int8	39.78	17.030	38.11	15.226	37.29	11.579
	Baseline	51.33	3.835	60.24	3.292	59.74	3.931
	MatQuant	52.37	3.800	63.35	3.187	62.75	3.153

In int2, MatQuant shows more accurate performance

MatQuant with OmniQuant

Data type	Method	Gemma-2 2B		Gemma-2 9B		Mistral 7B	
		OmniQuant	Task Avg.	log pplx.	Task Avg.	log pplx.	Task Avg.
bfloat16			68.21	2.551	74.38	2.418	73.99
int8	Baseline	68.25	2.552	74.59	2.418	73.77	2.110
		Naïve bit slicing shows significant drop in accuracy					
int4	Sliced int8	62.87	2.730	72.26	2.480	38.51	4.681
	Baseline	67.03	2.598	74.33	2.451	73.62	2.136
	MatQuant	66.58	2.618	73.83	2.491	73.06	2.153
int2	Sliced int8	39.78	17.030	38.11	15.226	37.29	11.579
	Baseline	51.33	3.835	60.24	3.292	59.74	3.931
	MatQuant	52.37	3.800	63.35	3.187	62.75	3.153

MatQuant with OmniQuant

Sliced Interpolation.

- Beyond the target quantization granularities (int8, int4, and int2),
MatQuant allows for bit-width interpolation to bit-widths not optimized during training

Data type	Method	Gemma-2 2B		Gemma-2 9B		Mistral 7B	
		OmniQuant	Task Avg.	log pplx.	Task Avg.	log pplx.	Task Avg.
int6	Sliced int8	67.72	2.497	74.64	2.353	73.00	2.071
	Baseline	68.06	2.554	74.23	2.420	74.10	2.112
	MatQuant	67.52	2.574	73.92	2.440	73.63	2.127
int3	Sliced int8	41.35	6.024	54.18	3.977	39.21	10.792
	Baseline	64.37	2.727	73.23	2.549	71.68	2.211
	MatQuant	64.47	2.618	72.87	2.607	71.16	2.238

MatQuant with QAT

Data type	Method	Gemma-2 2B		Gemma-2 9B		Mistral 7B	
		QAT	Task Avg.	log pplx.	Task Avg.	log pplx.	Task Avg.
bfloat16			68.21	2.551	74.38	2.418	73.99
int8	Baseline		67.82	2.458	74.17	2.29	73.48
	MatQuant		67.44	2.449	74.52	2.262	72.58
int4	Sliced int8		67.13	2.483	73.36	2.276	71.76
	Baseline		67.03	2.512	73.26	2.324	72.13
	MatQuant		66.59	2.499	73.24	2.429	71.99
int2	Sliced int8		39.27	10.217	40.40	7.259	37.41
	Baseline		47.74	3.433	56.02	2.923	54.95
	MatQuant		52.20	3.055	62.29	2.265	61.97

MatQuant with QAT

Data type	Method	Gemma-2 2B		Gemma-2 9B		Mistral 7B	
		QAT	Task Avg.	log pplx.	Task Avg.	log pplx.	Task Avg.
bfloat16			68.21	2.551	74.38	2.418	73.99
int8	Baseline	67.82	2.458	74.17	2.29	73.48	2.084
	MatQuant	67.44	2.449	74.52	2.262	72.58	2.104
Baseline (QAT) is better, but MatQuant shows comparable performance							
int4	Baseline	67.03	2.512	73.26	2.324	72.13	2.105
	MatQuant	66.59	2.499	73.24	2.429	71.99	2.148
int2	Sliced int8	39.27	10.217	40.40	7.259	37.41	9.573
	Baseline	47.74	3.433	56.02	2.923	54.95	2.699
	MatQuant	52.20	3.055	62.29	2.265	61.97	2.524

MatQuant with QAT

Data type	Method	Gemma-2 2B		Gemma-2 9B		Mistral 7B	
		QAT	Task Avg.	log pplx.	Task Avg.	log pplx.	Task Avg.
bfloat16			68.21	2.551	74.38	2.418	73.99
int8	Baseline		67.82	2.458	74.17	2.29	73.48
	MatQuant		67.44	2.449	74.52	2.262	72.58
int4	Sliced int8		67.13	2.483	73.36	2.276	71.76
	Baseline		67.03	2.512	73.26	2.324	72.13
	MatQuant		66.59	2.499	73.24	2.429	71.99
int2	Sliced int8		39.27	10.217	40.40	7.259	37.41
	Baseline		47.74	3.433	56.02	2.923	54.95
	MatQuant		52.20	3.055	62.29	2.265	61.97

In int2, MatQuant shows more accurate performance

MatQuant with QAT

Data type	Method	Gemma-2 2B		Gemma-2 9B		Mistral 7B	
		QAT	Task Avg.	log pplx.	Task Avg.	log pplx.	Task Avg.
bfloat16			68.21	2.551	74.38	2.418	73.99
int8	Baseline		67.82	2.458	74.17	2.29	73.48
		Naïve bit slicing shows significant drop in accuracy					
int4	Sliced int8		67.13	2.483	73.36	2.276	71.76
	Baseline		67.03	2.512	73.26	2.324	72.13
	MatQuant		66.59	2.499	73.24	2.429	71.99
int2	Sliced int8		39.27	10.217	40.40	7.259	37.41
	Baseline		47.74	3.433	56.02	2.923	54.95
	MatQuant		52.20	3.055	62.29	2.265	61.97

MatQuant with QAT

Sliced Interpolation.

- Models trained using MatQuant with QAT exhibit strong interpolative performance similar to that of MatQuant with OmniQuant.

Data type	Method	Gemma-2 2B		Gemma-2 9B		Mistral 7B	
		QAT	Task Avg.	log pplx.	Task Avg.	log pplx.	Task Avg.
int6	Sliced int8	67.72	2.497	74.64	2.353	73.00	2.071
	Baseline	68.06	2.554	74.23	2.420	74.10	2.112
	MatQuant	67.52	2.574	73.92	2.440	73.63	2.127
int3	Sliced int8	41.35	6.024	54.18	3.977	39.21	10.792
	Baseline	64.37	2.727	73.23	2.549	71.68	2.211
	MatQuant	64.47	2.618	72.87	2.607	71.16	2.238

Comparison OmniQuant vs QAT

- While OmniQuant only trains the auxiliary parameters needed for quantization,
QAT also updates the weight parameters.

Data type	Method	Gemma-2 2B		Data type	Method	Gemma-2 2B	
		OmniQuant	Task Avg.			QAT	Task Avg.
bfloat16		68.21	2.551	bfloat16		68.21	2.551
int8	Baseline	68.25	2.552	int8	Baseline	67.82	2.458
	MatQuant	68.02	2.570		MatQuant	67.44	2.449
int4	Sliced int8	62.87	2.730	int4	Sliced int8	67.13	2.483
	Baseline	67.03	2.598		Baseline	67.03	2.512
	MatQuant	66.58	2.618		MatQuant	66.59	2.499
int2	Sliced int8	39.78	17.030	int2	Sliced int8	39.27	10.217
	Baseline	51.33	3.835		Baseline	47.74	3.433
	MatQuant	52.37	3.800		MatQuant	52.20	3.055
int6	Sliced int8	67.72	2.497	int6	Sliced int8	67.53	2.401
	Baseline	68.06	2.554		Baseline	67.75	2.460
	MatQuant	67.52	2.574		MatQuant	67.33	2.453
int3	Sliced int8	41.35	6.024	int3	Sliced int8	59.56	2.882
	Baseline	64.37	2.727		Baseline	61.75	2.678
	MatQuant	64.47	2.618		MatQuant	60.76	2.734

Comparison OmniQuant vs QAT

- While OmniQuant only trains the auxiliary parameters needed for quantization,
QAT also updates the weight parameters.

Data type	Method	Gemma-2 2B		Data type	Method	Gemma-2 2B	
	OmniQuant	Task Avg.	log pplx.		QAT	Task Avg.	log pplx.
bfloat16		68.21	2.551	bfloat16		68.21	2.551
int8	Baseline	68.25	2.552	int8	Baseline	67.82	2.458
	MatQuant	68.02	2.570		MatQuant	67.44	2.449
int4	Sliced int8	62.87	2.730	int4	Sliced int8	67.13	2.483
	Baseline	67.03	2.598		Baseline	67.03	2.512
	MatQuant	66.58	2.618		MatQuant	66.59	2.499
int2	Sliced int8	39.78	17.030	int2	Sliced int8	39.27	10.217
	Baseline	51.33	3.835		Baseline	47.74	3.433
	MatQuant	52.37	3.800		MatQuant	52.20	3.055
int6	Sliced int8	67.72	2.497	int6	Sliced int8	67.53	2.401
	Baseline	68.06	2.554		Baseline	67.75	2.460
	MatQuant	67.52	2.574		MatQuant	67.33	2.453
int3	Sliced int8	41.35	6.024	int3	Sliced int8	59.56	2.882
	Baseline	64.37	2.727		Baseline	61.75	2.678
	MatQuant	64.47	2.618		MatQuant	60.76	2.734

QAT exhibits lower ppl than OmniQuant

Comparison OmniQuant vs QAT

- While OmniQuant only trains the auxiliary parameters needed for quantization,
QAT also updates the weight parameters.

Data type	Method	Gemma-2 2B		Data type	Method	Gemma-2 2B	
		OmniQuant	Task Avg.		QAT	Task Avg.	log pplx.
OmniQuant exhibits higher Task Accuracy than QAT	bfloat16		68.21	bfloat16		68.21	2.551
		Baseline	68.25		Baseline	67.82	2.458
	int8	MatQuant	68.02	int8	MatQuant	67.44	2.449
		Sliced int8	62.87		Sliced int8	67.13	2.483
	int4	Baseline	67.03	int4	Baseline	67.03	2.512
		MatQuant	66.58		MatQuant	66.59	2.499
	int2	Sliced int8	39.78	int2	Sliced int8	39.27	10.217
		Baseline	51.33		Baseline	47.74	3.433
		MatQuant	52.37		MatQuant	52.20	3.055
	int6	Sliced int8	67.72	int6	Sliced int8	67.53	2.401
		Baseline	68.06		Baseline	67.75	2.460
		MatQuant	67.52		MatQuant	67.33	2.453
	int3	Sliced int8	41.35	int3	Sliced int8	59.56	2.882
		Baseline	64.37		Baseline	61.75	2.678
		MatQuant	64.47		MatQuant	60.76	2.734

Comparison OmniQuant vs QAT

- While OmniQuant only trains the auxiliary parameters needed for quantization,
QAT also updates the weight parameters.

Data type	Method	Gemma-2 2B		Data type	Method	Gemma-2 2B	
	OmniQuant	Task Avg.	log pplx.		QAT	Task Avg.	log pplx.
OmniQuant exhibits higher Task Accuracy than QAT	bfloat16	68.21	2.551	bfloat16	68.21	2.551	
		68.25	2.552		Baseline	67.82	2.458
	int8	68.02	2.570	int8	MatQuant	67.44	2.449
		62.87	2.730		Sliced int8	67.13	2.483
	int4	67.03	2.598	int4	Baseline	67.03	2.512
		66.58	2.618		MatQuant	66.59	2.499
QAT exhibits lower ppl than OmniQuant	int2	39.78	17.030	int2	Sliced int8	39.27	10.217
		51.33	3.835		Baseline	47.74	3.433
		52.37	3.800		MatQuant	52.20	3.055
	int6	67.72	2.497	int6	Sliced int8	67.53	2.401
		68.06	2.554		Baseline	67.75	2.460
		67.52	2.574		MatQuant	67.33	2.453
	int3	41.35	6.024	int3	Sliced int8	59.56	2.882
		64.37	2.727		Baseline	61.75	2.678
		64.47	2.618		MatQuant	60.76	2.734

Comparison OmniQuant vs QAT

- While OmniQuant only trains the auxiliary parameters needed for quantization,
QAT also updates the weight parameters.

OmniQuant exhibits
higher Task Accuracy
than QAT

QAT exhibits
lower ppl
than OmniQuant

Data type	Method	Gemma-2 2B		Data type	Method	Gemma-2 2B	
	OmniQuant	Task Avg.	log pplx.		QAT	Task Avg.	log pplx.
bfloat16		68.21	2.551	bfloat16		68.21	2.551
int8				int8		QAT → overfitting to the C4 subset	
	Sliced int8	62.87	2.730		Sliced int8	67.13	2.483
	Baseline	67.03	2.598	int4	Baseline	67.03	2.512
	MatQuant	66.58	2.618		MatQuant	66.59	2.499
	Sliced int8	39.78	17.030		Sliced int8	39.27	10.217
	Baseline	51.33	3.835	int2	Baseline	47.74	3.433
	MatQuant	52.37	3.800		MatQuant	52.20	3.055
	Sliced int8	67.72	2.497		Sliced int8	67.53	2.401
	Baseline	68.06	2.554	int6	Baseline	67.75	2.460
	MatQuant	67.52	2.574		MatQuant	67.33	2.453
	Sliced int8	41.35	6.024		Sliced int8	59.56	2.882
	Baseline	64.37	2.727	int3	Baseline	61.75	2.678
	MatQuant	64.47	2.618		MatQuant	60.76	2.734

Comparison OmniQuant vs QAT

- While OmniQuant only trains the auxiliary parameters needed for quantization,
QAT also updates the weight parameters.

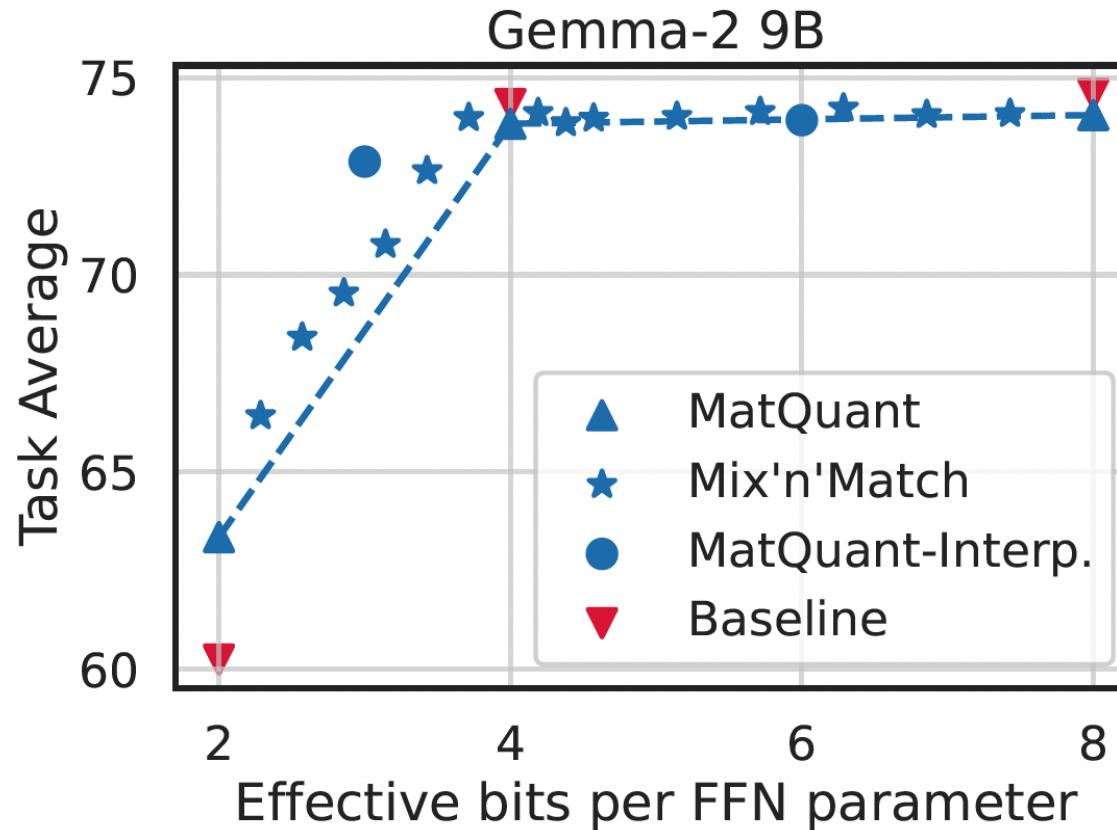
OmniQuant exhibits
higher Task Accuracy
than QAT

Data type	Method	Gemma-2 2B		Data type	Method	Gemma-2 2B	
	OmniQuant	Task Avg.	log pplx.		QAT	Task Avg.	log pplx.
bfloat16		68.21	2.551	bfloat16		68.21	2.551
int8	M	QAT → overfitting to the C4 subset				2.458	
	S	1. the need for high-quality data for QAT				2.449	
	S	2. Users are better off using resource-friendly methods like OmniQuant.				2.483	
	MatQuant	52.37	3.800		MatQuant	52.20	3.055
int6	Sliced int8	67.72	2.497		Sliced int8	67.53	2.401
	Baseline	68.06	2.554	int6	Baseline	67.75	2.460
	MatQuant	67.52	2.574		MatQuant	67.33	2.453
int3	Sliced int8	41.35	6.024		Sliced int8	59.56	2.882
	Baseline	64.37	2.727	int3	Baseline	61.75	2.678
	MatQuant	64.47	2.618		MatQuant	60.76	2.734

QAT exhibits
lower ppl
than OmniQuant

Additional: Layerwise Mix'n'Match

- Mix'n'Match provides a mechanism to **obtain a combinatorial number of strong models** by using **layerwise different quantization granularities**, from the target bit-widths – i.e., int8, int4, and int2 across layers



Ablation studies: Weightings (λ_r) for MatQuant

$$\min_P \frac{1}{N} \sum_{i \in [N]} \sum_{r \in R} \lambda_r \cdot \mathcal{L}(F(S(Q(\theta, c), r), x'_i), y'_i)$$

Loss coefficient for each target bits (8,4,2 bits)

< overall objective of MatQuant >

Data type	Weightings	Gemma-2 2B	Gemma-2 9B	Mistral 7B	
		8	4	2	Task Avg.
int8	(0.1, 0.1, 1)	68.02		74.05	73.27
	(0.2, 0.2, 1)	67.91		73.91	73.44
	(0.3, 0.3, 1)	68.01		73.88	73.56
	(0.4, 0.4, 1)	67.95		73.84	73.65
int4	(0.1, 0.1, 1)	66.58		73.83	72.76
	(0.2, 0.2, 1)	67.47		73.8	73.16
	(0.3, 0.3, 1)	66.97		73.25	73.47
	(0.4, 0.4, 1)	67.48		74.32	73.66
int2	(0.1, 0.1, 1)	52.37		63.35	63.25
	(0.2, 0.2, 1)	51.88		64.04	63.99
	(0.3, 0.3, 1)	51.05		64.1	63.6
	(0.4, 0.4, 1)	51.69		61.98	62.75

Ablation studies: Weightings (λ_r) for MatQuant

$$\min_P \frac{1}{N} \sum_{i \in [N]} \sum_{r \in R} \lambda_r \cdot \mathcal{L}(F(S(Q(\theta, c), r), x'_i), y'_i)$$

Loss coefficient for each target bits (8,4,2 bits)

< overall objective of MatQuant >

Data type	Weightings	Gemma-2 2B	Gemma-2 9B	Mistral 7B	
		8	4	2	Task Avg.
int8	(0.1, 0.1, 1)	68.02	74.05	73.27	
	(0.2, 0.2, 1)	67.91	73.91	73.44	
	(0.3, 0.3, 1)	68.01	73.88	73.56	
	(0.4, 0.4, 1)	67.95	73.84	73.65	
→ Higher accuracy in int8/int4	(0.1, 0.1, 1)	66.58	73.83	72.76	
	(0.2, 0.2, 1)	67.47	73.8	73.16	
	(0.3, 0.3, 1)	66.97	73.25	73.47	
	(0.4, 0.4, 1)	67.48	74.32	73.66	
→ Lower accuracy in int2	(0.1, 0.1, 1)	52.37	63.35	63.25	
	(0.2, 0.2, 1)	51.88	64.04	63.99	
	(0.3, 0.3, 1)	51.05	64.1	63.6	
	(0.4, 0.4, 1)	51.69	61.98	62.75	

Low coefficient for 8bit/4bit

→ Higher accuracy in int8/int4

→ Lower accuracy in int2

Ablation studies: Weightings (λ_r) for MatQuant

$$\min_P \frac{1}{N} \sum_{i \in [N]} \sum_{r \in R} \boxed{\lambda_r} \cdot \mathcal{L} (F(S(Q(\theta, c), r), x'_i), y'_i)$$

Loss coefficient for each target bits (8,4,2 bits)

< overall objective of MatQuant >

Data type	Weightings	Gemma-2 2B	Gemma-2 9B	Mistral 7B		
		8	4	2	Task Avg.	
int8	(0.1, 0.1, 1)	68.02		74.05	73.27	High coefficient for 8bit/4bit
	(0.2, 0.2, 1)	67.91		73.91	73.44	
	(0.3, 0.3, 1)	68.01		73.88	73.56	
	(0.4, 0.4, 1)	67.95		73.84	73.65	
int4	(0.1, 0.1, 1)	66.58		73.83	72.76	→ Higher accuracy in int2
	(0.2, 0.2, 1)	67.47		73.8	73.16	→ Lower accuracy in int8/int4
	(0.3, 0.3, 1)	66.97		73.25	73.47	
	(0.4, 0.4, 1)	67.48		74.32	73.66	
int2	(0.1, 0.1, 1)	52.37		63.35	63.25	
	(0.2, 0.2, 1)	51.88		64.04	63.99	
	(0.3, 0.3, 1)	51.05		64.1	63.6	
	(0.4, 0.4, 1)	51.69		61.98	62.75	

Ablation studies: Single Precision (S.P.) MatQuant

- Eliminate other target bits loss (8bit & 4bit),
except for 2bit loss

$$\min_P \frac{1}{N} \sum_{i \in [N]} \sum_{r \in R} \lambda_r \cdot \mathcal{L}(F(S(Q(\theta, c), r), x'_i), y'_i)$$

Loss a for each target bits (8,4,2 bits)

< overall objective of MatQuant>

λ_r : r is a target bit, $\lambda_8, \lambda_4 : 0$, $\lambda_2 : 1$

int2		Gemma-2 2B		Gemma-2 9B		Mistral 7B	
Method	Task Avg.	log pplx.	Task Avg.	log pplx.	Task Avg.	log pplx.	
OmniQuant	51.33	3.835	60.24	3.292	59.74	3.931	
S.P. MatQuant	53.42	3.631	64.02	3.171	63.58	2.976	
MatQuant	52.37	3.800	63.35	3.187	62.75	3.153	
QAT	47.74	3.433	56.02	2.923	54.95	2.699	
S.P. MatQuant	52.08	3.054	62.66	2.656	61.48	2.509	
MatQuant	52.20	3.055	62.29	2.660	61.97	2.524	

Ablation studies: Co-distillation for MatQuant

- Outputs from a higher-precision model → used for lower-precision nested model training.
either in a standalone fashion or alongside the ground truth target (weighted equally).

Data type	Config.	Gemma-2 9B	OmniQuant		QAT	
		Task Avg.	log pplx.	Task Avg.	log pplx.	
int8	[8, 4, 2]	74.05	2.438	74.52	2.262	
	[8, 4, 8 → 2]	72.76	2.473	74.75	2.242	
	[8, 4, 2, 8 → 2]	73.99	2.435	74.87	2.240	
	[8, 4, 2, 8 → 4; 2]	73.85	2.437	74.81	2.240	
int4	[8, 4, 2]	73.83	2.491	73.24	2.295	
	[8, 4, 8 → 2]	72.65	2.519	73.76	2.279	
	[8, 4, 2, 8 → 2]	73.63	2.486	73.77	2.276	
	[8, 4, 2, 8 → 4; 2]	73.55	2.478	73.93	2.277	
int2	[8, 4, 2]	63.35	3.187	62.29	2.660	
	[8, 4, 8 → 2]	62.64	3.289	62.31	2.670	
	[8, 4, 2, 8 → 2]	62.91	3.138	62.70	2.673	
	[8, 4, 2, 8 → 4; 2]	64.32	3.227	62.60	2.670	

Ablation studies: Co-distillation for MatQuant

- Outputs from a higher-precision model → used for lower-precision nested model training.
either in a standalone fashion or alongside the ground truth target (weighted equally).

Data type	Gemma-2 9B	OmniQuant		QAT	
	Config.	Task Avg.	log pplx.	Task Avg.	log pplx.
8→4: use int8 outputs for int4 training	[8, 4, 2]	74.05	2.438	74.52	2.262
8→2 : use int8 outputs for int2 training	[8, 4, 8 → 2]	72.76	2.473	74.75	2.242
	[8, 4, 2, 8 → 2]	73.99	2.435	74.87	2.240
	[8, 4, 2, 8 → 4; 2]	73.85	2.437	74.81	2.240
int4	[8, 4, 2]	73.83	2.491	73.24	2.295
	[8, 4, 8 → 2]	72.65	2.519	73.76	2.279
	[8, 4, 2, 8 → 2]	73.63	2.486	73.77	2.276
	[8, 4, 2, 8 → 4; 2]	73.55	2.478	73.93	2.277
int2	[8, 4, 2]	63.35	3.187	62.29	2.660
	[8, 4, 8 → 2]	62.64	3.289	62.31	2.670
	[8, 4, 2, 8 → 2]	62.91	3.138	62.70	2.673
	[8, 4, 2, 8 → 4; 2]	64.32	3.227	62.60	2.670

Ablation studies: Co-distillation for MatQuant

- Outputs from a higher-precision model → used for lower-precision nested model training.
either in a standalone fashion or alongside the ground truth target (weighted equally).

Data type	Config.	Gemma-2 9B	OmniQuant		QAT	
		Task Avg.	log pplx.	Task Avg.	log pplx.	
int8	[8, 4, 2]	74.05	2.438	74.52	2.262	
	[8, 4, 8 → 2]	72.76	2.473	74.75	2.242	
	[8, 4, 2, 8 → 2]	73.99	2.435	74.87	2.240	
	[8, 4, 2, 8 → 4; 2]	73.85	2.437	74.81	2.240	
int4	[8, 4, 2]	73.83	2.491	73.24	2.295	
	[8, 4, 8 → 2]	72.65	2.519	73.76	2.279	
	[8, 4, 2, 8 → 2]	73.63	2.486	73.77	2.276	
	[8, 4, 2, 8 → 4; 2]	73.55	2.478	73.93	2.277	
int2	[8, 4, 2]	63.35	3.187	62.29	2.660	
	[8, 4, 8 → 2]	62.64	3.289	62.31	2.670	
	[8, 4, 2, 8 → 2]	62.91	3.138	62.70	2.673	
	[8, 4, 2, 8 → 4; 2]	64.32	3.227	62.60	2.670	

Ablation studies: FFN + ATTN Weight Quantization

- Using QAT, apply MatQuant to **FFN**, and **also ATTN**

Data type	Method	Gemma-2 9B		Mistral 7B		
		QAT	Task Avg.	log pplx.	Task Avg.	log pplx.
bfloat16			74.38	2.418	73.99	2.110
int8	Baseline	74.61	2.353	73.73	2.091	
	MatQuant	74.85	2.333	73.88	2.182	
int4	Sliced int8	73.15	2.362	71.46	2.290	
	Baseline	72.98	2.40	71.87	2.132	
	MatQuant	74.01	2.396	71.44	2.441	
int2	Sliced int8	38.97	23.467	35.06	10.640	
	Baseline	-	-	-	-	
	S.P. MatQuant	45.69	3.780	35.35	7.761	
	MatQuant	44.19	3.826	38.36	10.971	
int6	Sliced int8	74.49	2.290	73.61	2.104	
	Baseline	74.65	2.357	73.72	2.093	
	MatQuant	74.57	2.340	74.04	2.161	
int3	Sliced int8	64.19	2.895	39.01	6.018	
	Baseline	-	-	-	-	
	S.P. MatQuant	67.68	2.520	67.59	2.335	
	MatQuant	63.63	2.937	40.55	4.776	

Ablation studies: FFN + ATTN Weight Quantization

- Using QAT, apply MatQuant to **FFN**, and **also ATTN**

< FFN MatQaunt >

Data type	Method	Gemma-2 9B		Mistral 7B	
		QAT	Task Avg.	log pplx.	Task Avg.
bfloat16		74.38	2.418	73.99	2.110
int8	Baseline	74.17	2.29	73.48	2.084
	MatQuant	74.52	2.262	72.58	2.104
int4	Sliced int8	73.36	2.276	71.76	2.18
	Baseline	73.26	2.324	72.13	2.105
int2	MatQuant	73.24	2.429	71.99	2.148
	Sliced int8	40.40	7.259	37.41	9.573
int6	Baseline	56.02	2.923	54.95	2.699
	MatQuant	62.29	2.265	61.97	2.524
int3	Sliced int8	74.15	2.232	73.35	2.097
	Baseline	74.31	2.293	72.71	2.077
int3	MatQuant	74.30	2.265	72.59	2.106
	Sliced int8	68.70	2.512	64.33	2.493
int3	Baseline	69.9	2.43	68.82	2.197
	MatQuant	70.41	2.429	67.16	2.324

< ATTN + FFN MatQaunt >

Data type	Method	Gemma-2 9B		Mistral 7B		
		QAT	Task Avg.	log pplx.	Task Avg.	log pplx.
bfloat16			74.38	2.418	73.99	2.110
int8	Baseline	74.61	2.353	73.73	2.091	
	MatQuant	74.85	2.333	73.88	2.182	
int4	Sliced int8	73.15	2.362	71.46	2.290	
	Baseline	72.98	2.40	71.87	2.132	
int2	MatQuant	74.01	2.396	71.44	2.441	
	Sliced int8	38.97	23.467	35.06	10.640	
int6	Baseline	-	-	-	-	
	S.P. MatQuant	45.69	3.780	35.35	7.761	
int3	MatQuant	44.19	3.826	38.36	10.971	
	Sliced int8	74.49	2.290	73.61	2.104	
int6	Baseline	74.65	2.357	73.72	2.093	
	MatQuant	74.57	2.340	74.04	2.161	
int3	Sliced int8	64.19	2.895	39.01	6.018	
	Baseline	-	-	-	-	
int3	S.P. MatQuant	67.68	2.520	67.59	2.335	
	MatQuant	63.63	2.937	40.55	4.776	

Ablation studies: FFN + ATTN Weight Quantization

- Using QAT, apply MatQuant to **FFN**, and **also ATTN**

< FFN MatQaunt >

Data type	Method	Gemma-2 9B		Mistral 7B	
		QAT	Task Avg.	log pplx.	Task Avg.
bfloat16		74.38	2.418	73.99	2.110
int8	Baseline	74.17	2.29	73.48	2.084
	MatQuant	74.52	2.262	72.58	2.104
int4	Sliced int8	73.36	2.276	71.76	2.18
	Baseline	73.26	2.324	72.13	2.105
int2	MatQuant	73.24	2.429	71.99	2.148
	Sliced int8	40.40	7.259	37.41	9.573
int6	Baseline	56.02	2.923	54.95	2.699
	MatQuant	62.29	2.265	61.97	2.524
int3	Sliced int8	74.15	2.232	73.35	2.097
	Baseline	74.31	2.293	72.71	2.077
int3	MatQuant	74.30	2.265	72.59	2.106
	Sliced int8	68.70	2.512	64.33	2.493
int3	Baseline	69.9	2.43	68.82	2.197
	MatQuant	70.41	2.429	67.16	2.324

< ATTN + FFN MatQaunt >

Data type	Method	Gemma-2 9B		Mistral 7B		
		QAT	Task Avg.	log pplx.	Task Avg.	log pplx.
bfloat16			74.38	2.418	73.99	2.110
int8	Baseline	74.61	2.353	73.73	2.091	
	MatQuant	74.85	2.333	73.88	2.182	
int4	Sliced int8	73.15	2.362	71.46	2.290	
	Baseline	72.98	2.40	71.87	2.132	
int2	MatQuant	74.01	2.396	71.44	2.441	
	Sliced int8	38.97	23.467	35.06	10.640	
int6	Baseline	-	-	-	-	
	S.P. MatQuant	45.69	3.780	35.35	7.761	
int3	MatQuant	44.19	3.826	38.36	10.971	
	Sliced int8	74.49	2.290	73.61	2.104	
int6	Baseline	74.65	2.357	73.72	2.093	
	MatQuant	74.57	2.340	74.04	2.161	
int3	Sliced int8	64.19	2.895	39.01	6.018	
	Baseline	-	-	-	-	
int3	S.P. MatQuant	67.68	2.520	67.59	2.335	
	MatQuant	63.63	2.937	40.55	4.776	

Very Poor Performance when Quantize ATTN & FFN Both!!

Additional Consideration

Deployment Considerations

- Current hardware accelerators → native support for int8 and int4 quantized models.
- Custom-implemented CUDA kernels → can support int2 and int3

Additional Consideration

Deployment Considerations

- Current hardware accelerators → native support for int8 and int4 quantized models.
- Custom-implemented CUDA kernels → can support int2 and int3
- To apply it to various settings as above, you need to be prepared for all kinds of configurations.

Additional Consideration

Deployment Considerations

- Current hardware accelerators → native support for int8 and int4 quantized models.
- Custom-implemented CUDA kernels → can support int2 and int3
- To apply it to various settings as above, you need to be prepared for all kinds of configurations.

MatQuant can be a simple solution for deployment!

- MatQuant can generate a large number of models at inference time.
- Depending on the serving environment,
we can choose between Mix'n'Match models and homogeneous sliced models.

Additional Consideration

Extension to Floating Point

- Extending MatQuant to floating-point representations, such as FP8 and FP4, presents significant challenges. → Why?

Additional Consideration

Extension to Floating Point

- Extending MatQuant to floating-point representations, such as FP8 and FP4, presents significant challenges. → Why?

For example,

- slicing the first two or 4bits from 8bits is easy
- However, this would not be the case when slicing two exponent bits from FP8.

Additional Consideration

Extension to Floating Point

- Extending MatQuant to floating-point representations, such as FP8 and FP4, presents significant challenges. → Why?

For example,

- slicing the first two or 4bits from 8bits is easy
- However, this would not be the case when slicing two exponent bits from FP8.
→ needs further research !

Summary for MatQuant

Strength

1. Eliminates the need to perform quantization optimization multiple times for different bit precisions, as it can be handled with a single optimization.

Summary for MatQuant

Strength

1. Eliminates the need to perform quantization optimization multiple times for different bit precisions, as it can be handled with a single optimization.
2. For various deployment environments, the required bit precision can be allocated at the inference. In other words, specific optimization for each environment is not necessary.

Summary for MatQuant

Strength

1. Eliminates the need to perform quantization optimization multiple times for different bit precisions, as it can be handled with a single optimization.
2. For various deployment environments, the required bit precision can be allocated at the inference. In other words, specific optimization for each environment is not necessary.
3. Even with int8 and int4, it shows performance comparable to the baseline, and in particular, it demonstrates clear performance improvements over the baseline at int2.

Summary for MatQuant

Weakness

1. Poor Performance

- Most recent quantized models are deployed with 8-bit or 4-bit precision.
→ because the performance degradation with 2-bit quant is too severe to justify the memory savings.
- However, MatQuant shows little to no performance improvement at int8 or int4, raising concerns about its practicality in real-world deployment scenarios.

Data type	Method	Gemma-2 2B		Gemma-2 9B		Mistral 7B	
		OmniQuant	Task Avg. log pplx.				
bfloat16		68.21	2.551	74.38	2.418	73.99	2.110
int8	Baseline	68.25	2.552	74.59	2.418	73.77	2.110
	MatQuant	68.02	2.570	74.05	2.438	73.65	2.125
int4	Sliced int8	62.87	2.730	72.26	2.480	38.51	4.681
	Baseline	67.03	2.598	74.33	2.451	73.62	2.136
	MatQuant	66.58	2.618	73.83	2.491	73.06	2.153
int2	Sliced int8	39.78	17.030	38.11	15.226	37.29	11.579
	Baseline	51.33	3.835	60.24	3.292	59.74	3.931
	MatQuant	52.37	3.800	63.35	3.187	62.75	3.153

< MatQuant with

Summary for MatQuant

Weakness

2. no justification for poor performance in ATTN/FFN Quant

- The paper merely states that applying QAT to both the attention and FFN modules leads to instability at extremely low bit settings.
- However, it does not provide any justification or further explanation for this observation.

< FFN MatQaunt >						
Data type	Method	Gemma-2 9B		Mistral 7B		
	QAT	Task Avg.	log pplx.	Task Avg.	log pplx.	
bfloating16	QAT	74.38	2.418	73.99	2.110	
	MatQuant	74.52	2.262	72.58	2.104	
int8	Baseline	74.17	2.29	73.48	2.084	
	MatQuant	74.52	2.262	72.58	2.104	
int4	Sliced int8	73.36	2.276	71.76	2.18	
	Baseline	73.26	2.324	72.13	2.105	
int2	MatQuant	73.24	2.429	71.99	2.148	
	Sliced int8	40.40	7.259	37.41	9.573	
int6	Baseline	56.02	2.923	54.95	2.699	
	MatQuant	62.29	2.265	61.97	2.524	
int3	Sliced int8	74.15	2.232	73.35	2.097	
	Baseline	74.31	2.293	72.71	2.077	
int6	MatQuant	74.30	2.265	72.59	2.106	
	Sliced int8	68.70	2.512	64.33	2.493	
int3	Baseline	69.9	2.43	68.82	2.197	
	MatQuant	70.41	2.429	67.16	2.324	

< ATTN + FFN MatQaunt >						
Data type	Method	Gemma-2 9B		Mistral 7B		
	QAT	Task Avg.	log pplx.	Task Avg.	log pplx.	
bfloating16	QAT	74.38	2.418	73.99	2.110	
	MatQuant	74.85	2.333	73.88	2.182	
int8	Baseline	74.61	2.353	73.73	2.091	
	MatQuant	74.85	2.333	73.88	2.182	
int4	Sliced int8	73.15	2.362	71.46	2.290	
	Baseline	72.98	2.40	71.87	2.132	
int2	MatQuant	74.01	2.396	71.44	2.441	
	Sliced int8	38.97	23.467	35.06	10.640	
int6	Baseline	-	-	-	-	
	S.P. MatQuant	45.69	3.780	35.35	7.761	
int3	MatQuant	44.19	3.826	38.36	10.971	
	Sliced int8	74.49	2.290	73.61	2.104	
int6	Baseline	74.65	2.357	73.72	2.093	
	MatQuant	74.57	2.340	74.04	2.161	
int3	Sliced int8	64.19	2.895	39.01	6.018	
	Baseline	-	-	-	-	
int6	S.P. MatQuant	67.68	2.520	67.59	2.335	
	MatQuant	63.63	2.937	40.55	4.776	

Very Poor Performance when Quantize ATTN & FFN Both!!

Thank you.

Appendix

