

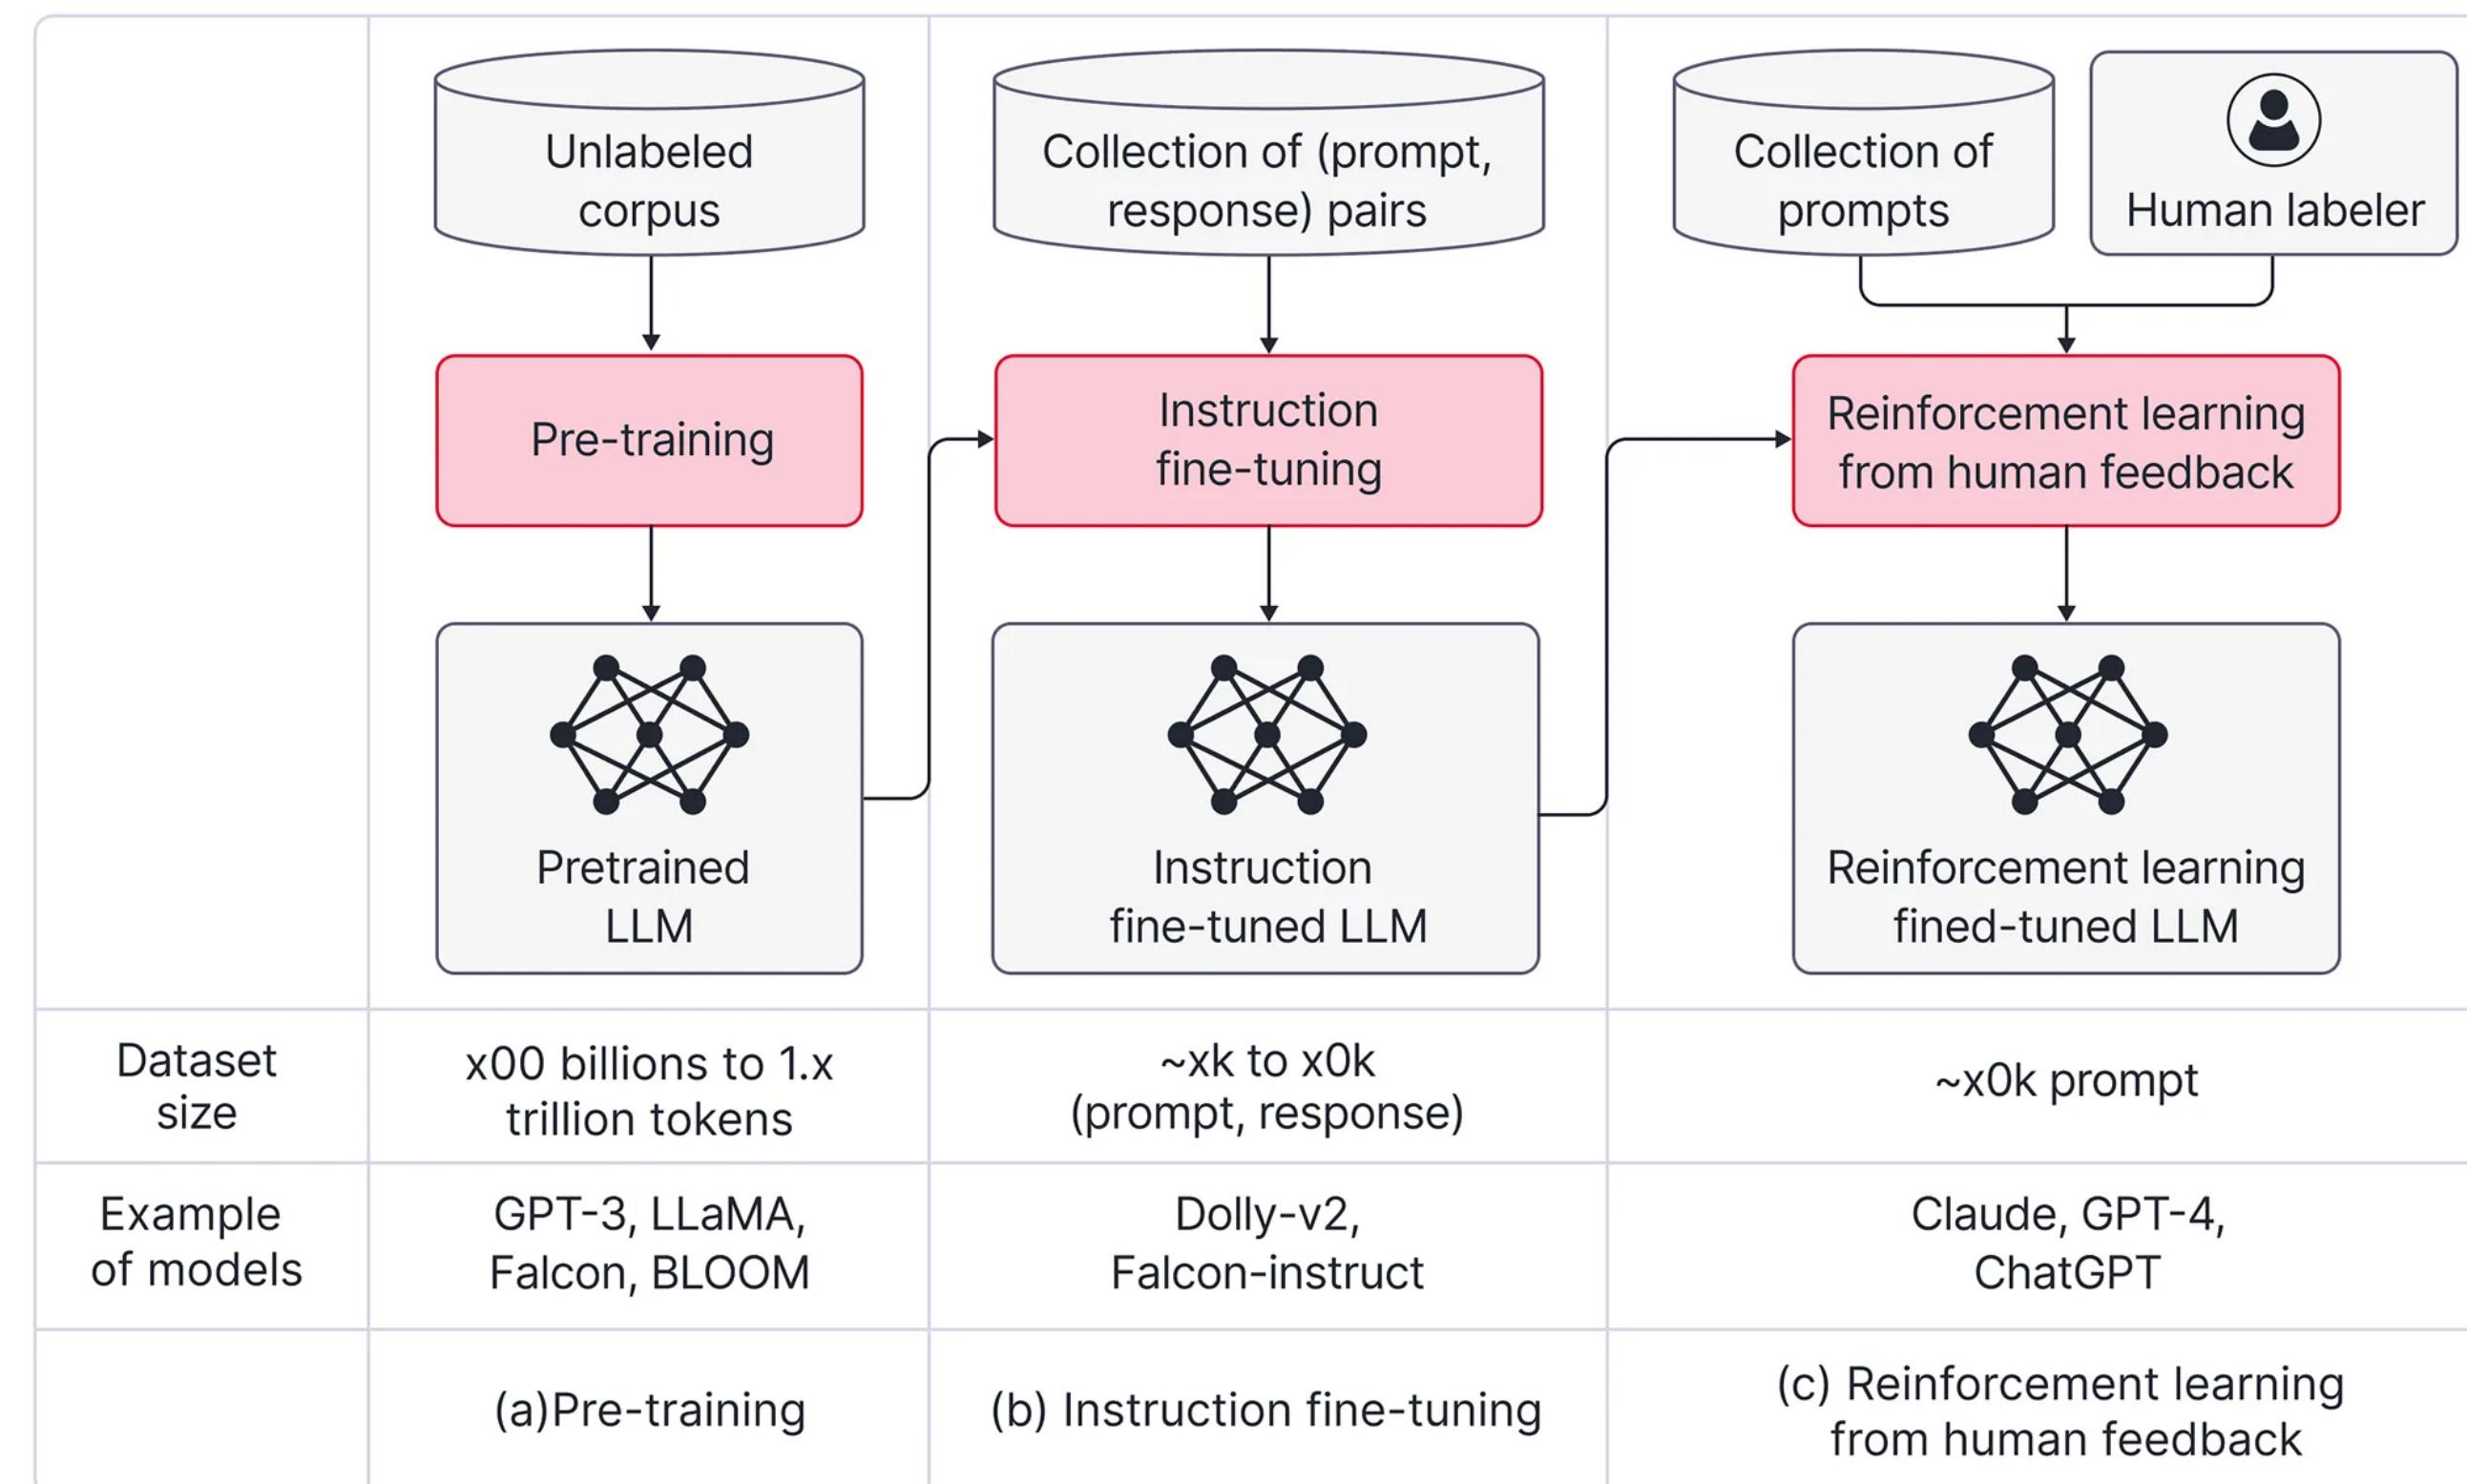
LLM compression

EECE695D: Efficient ML Systems

Spring 2025

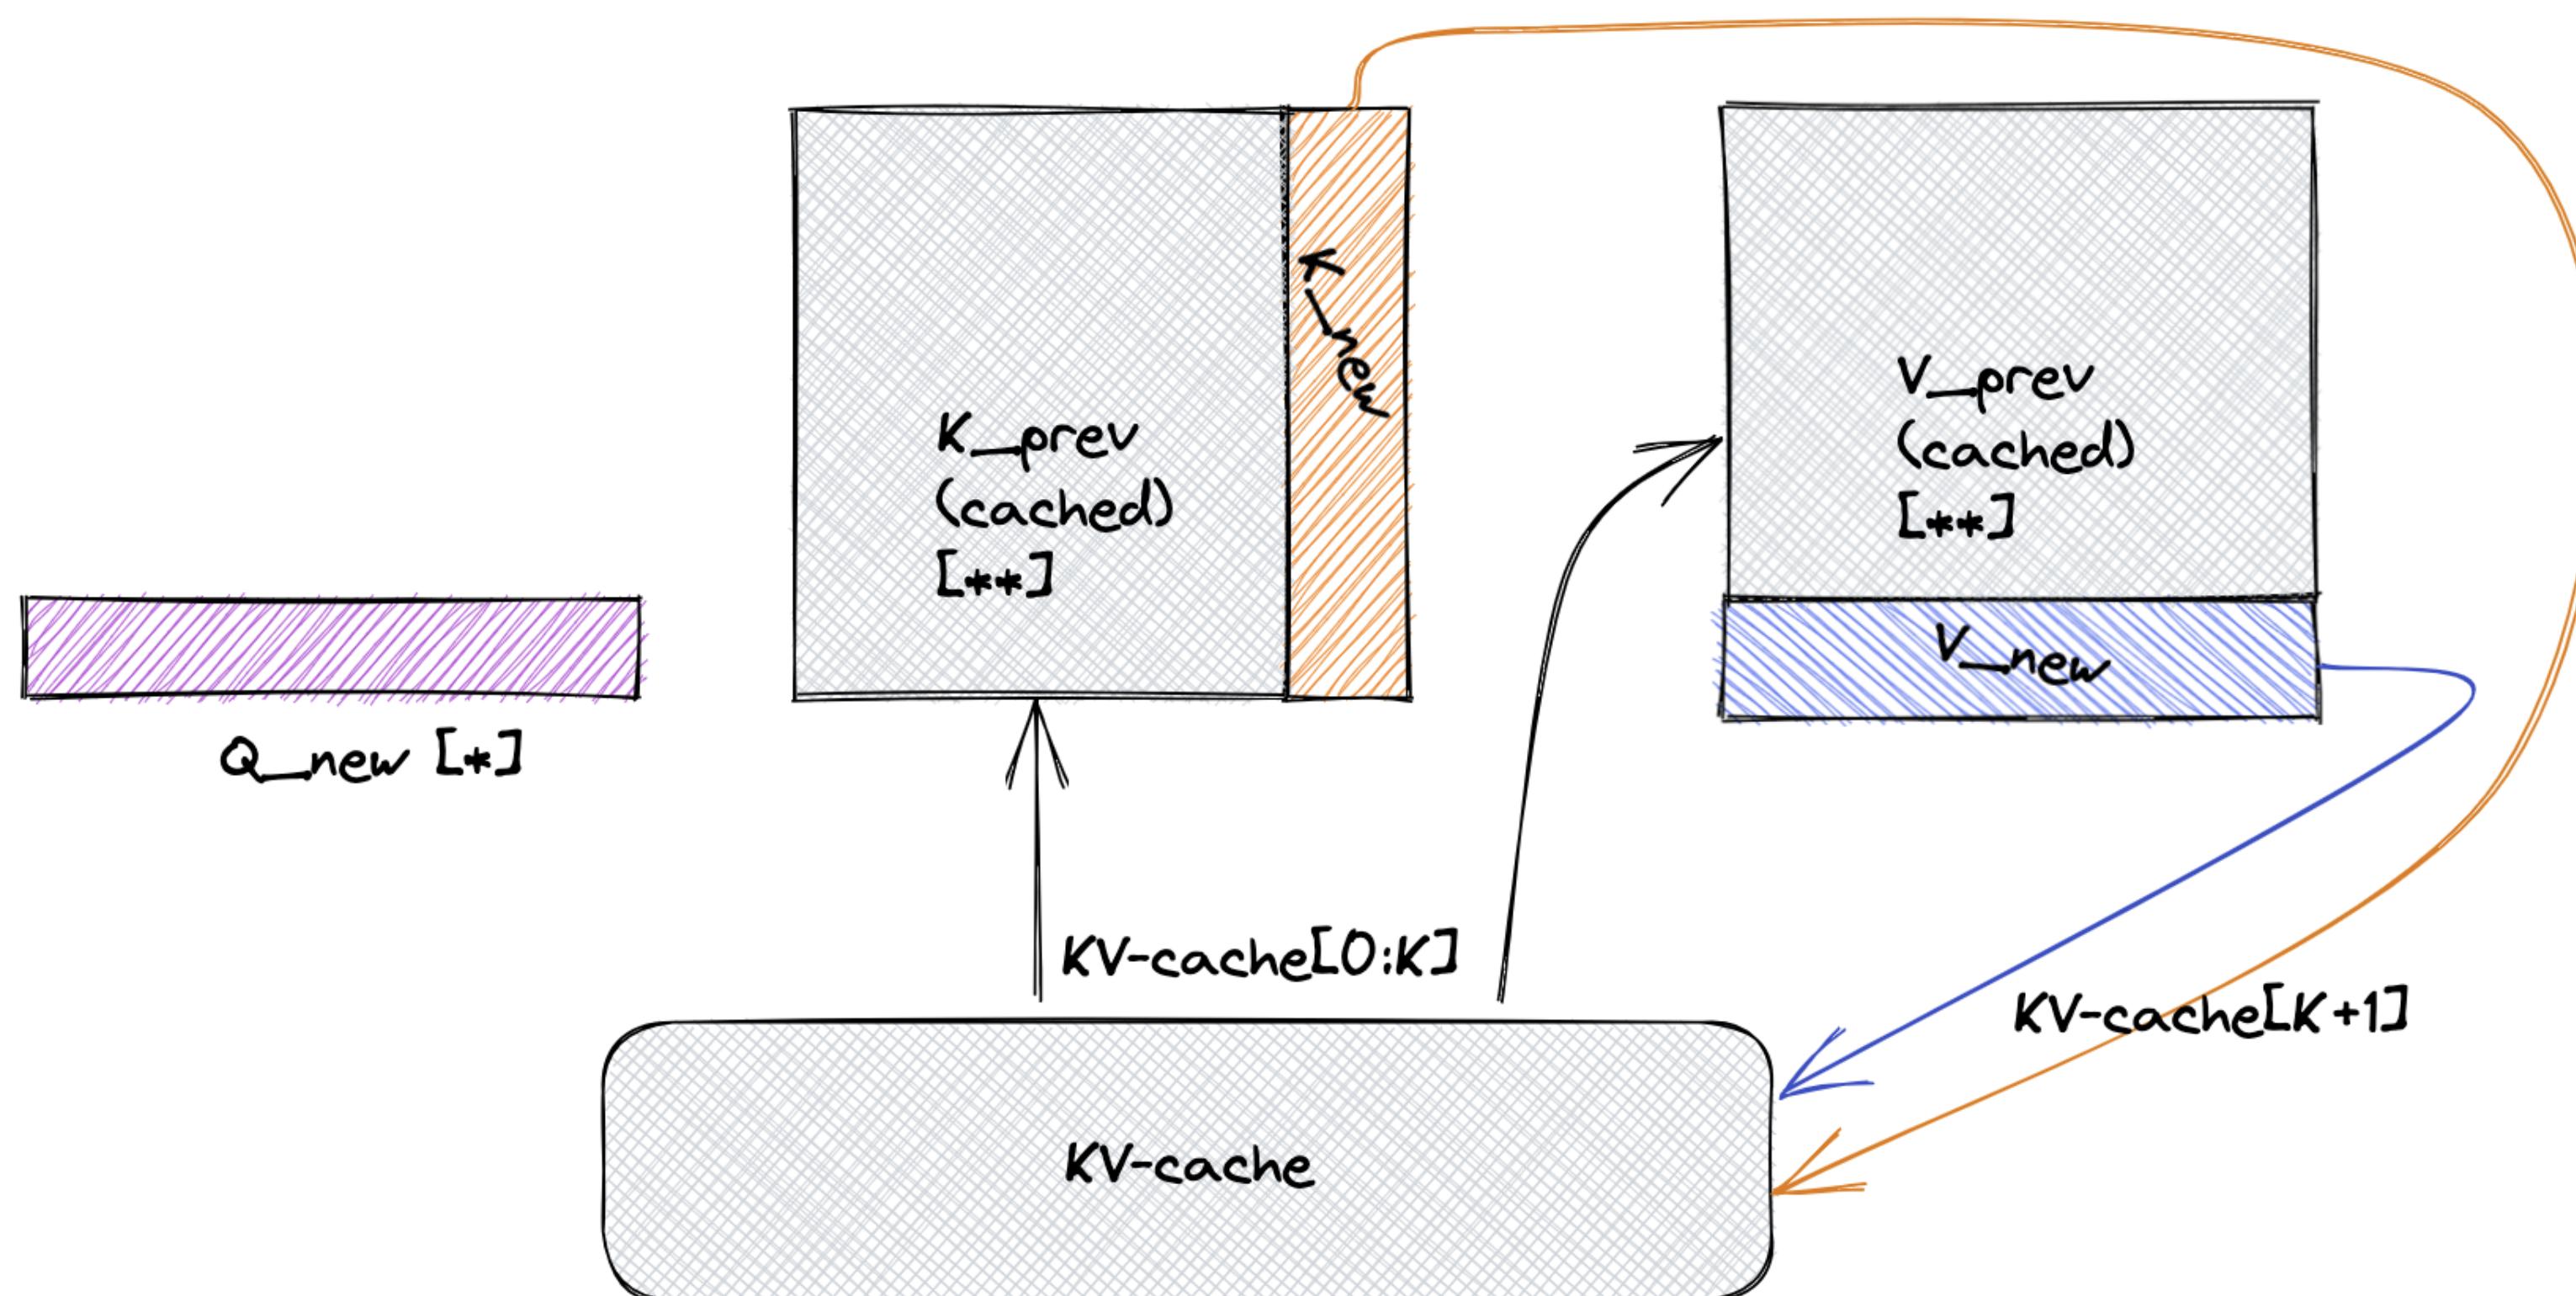
Compressing LLMs

- **Post-training compression** has been the mainstream
 - Retraining cost is too large, including alignment
 - 2025. Shifting toward methods involving retraining (e.g., Gemma QAT)



Compressing LLMs

- Also, more efforts on resolving **memory bottleneck**
 - e.g., weight quantization > weight & activation quantization
 - 2025. Activation quantization as well, especially the KV cache



Popular ideas

- Thus, much emphasis on finding a good **approximation** of original model

$$\min_{\hat{\theta}:\text{compressed}} \|f(x; w) - f(x; \hat{w})\|^2$$

- Two mainstream approaches:
 - **Hessian-based**
“Minimize the compression error very carefully, using Hessians”
 - **Outlier-driven**
“Identify outliers, and use these to keep compression error small”

Hessian-based Approach

Basic idea

- Already discussed Hessian-based loss approximation in sparsity
 - **Optimal Brain Damage**
 - Approximate the loss of removing a weight as:
$$f(x; w + \delta) - f(x; w) \approx \delta^\top \mathbf{H} \delta$$
 - **Optimal Brain Surgeon**
 - Prune the weight with minimal score: $w_i^2 / 2[\mathbf{H}^{-1}]_{ii}$
 - Update other weights by $-w_i \mathbf{H}^{-1} \mathbf{e}_i / [\mathbf{H}^{-1}]_{ii}$
 - **Idea.** Perform OBS for LLMs

Basic idea

- **Challenge.** Computing Hessian inverse for LLMs, multiple times
 - Very heavy: trillion x trillion matrix
 - **Idea.** Approximate by the **layerwise** subproblem

$$\min_{\hat{\mathbf{W}}} \|\mathbf{W}\mathbf{X} - \hat{\mathbf{W}}\mathbf{X}\|_2^2$$

- Further approximate it by the **rowwise** subproblem

$$\min_{\hat{\mathbf{w}}} \|\mathbf{w}_{i,:}^\top \mathbf{X} - \hat{\mathbf{w}}^\top \mathbf{X}\|_2^2$$

- Then the Hessian becomes: $\mathbf{H} = 2\mathbf{X}\mathbf{X}^\top \in \mathbb{R}^{d_{\text{col}} \times d_{\text{col}}}$
(same for all rows)

Basic idea

- Again, computing the Hessian inverse can be done recursively by using the matrix inversion formula:

$$\mathbf{H}_m^{-1} = \mathbf{H}_{m-1}^{-1} - \frac{\mathbf{H}_{m-1}^{-1} \mathbf{x}_m \mathbf{x}_m^\top \mathbf{H}_{m-1}^{-1}}{N + \mathbf{x}_m \mathbf{H}_{m-1}^{-1} \mathbf{x}_m}$$

- **Problem.** Need to compute Hessian after removing each weight
 - LLMs are very large, so a lot of repetitions!

Basic idea

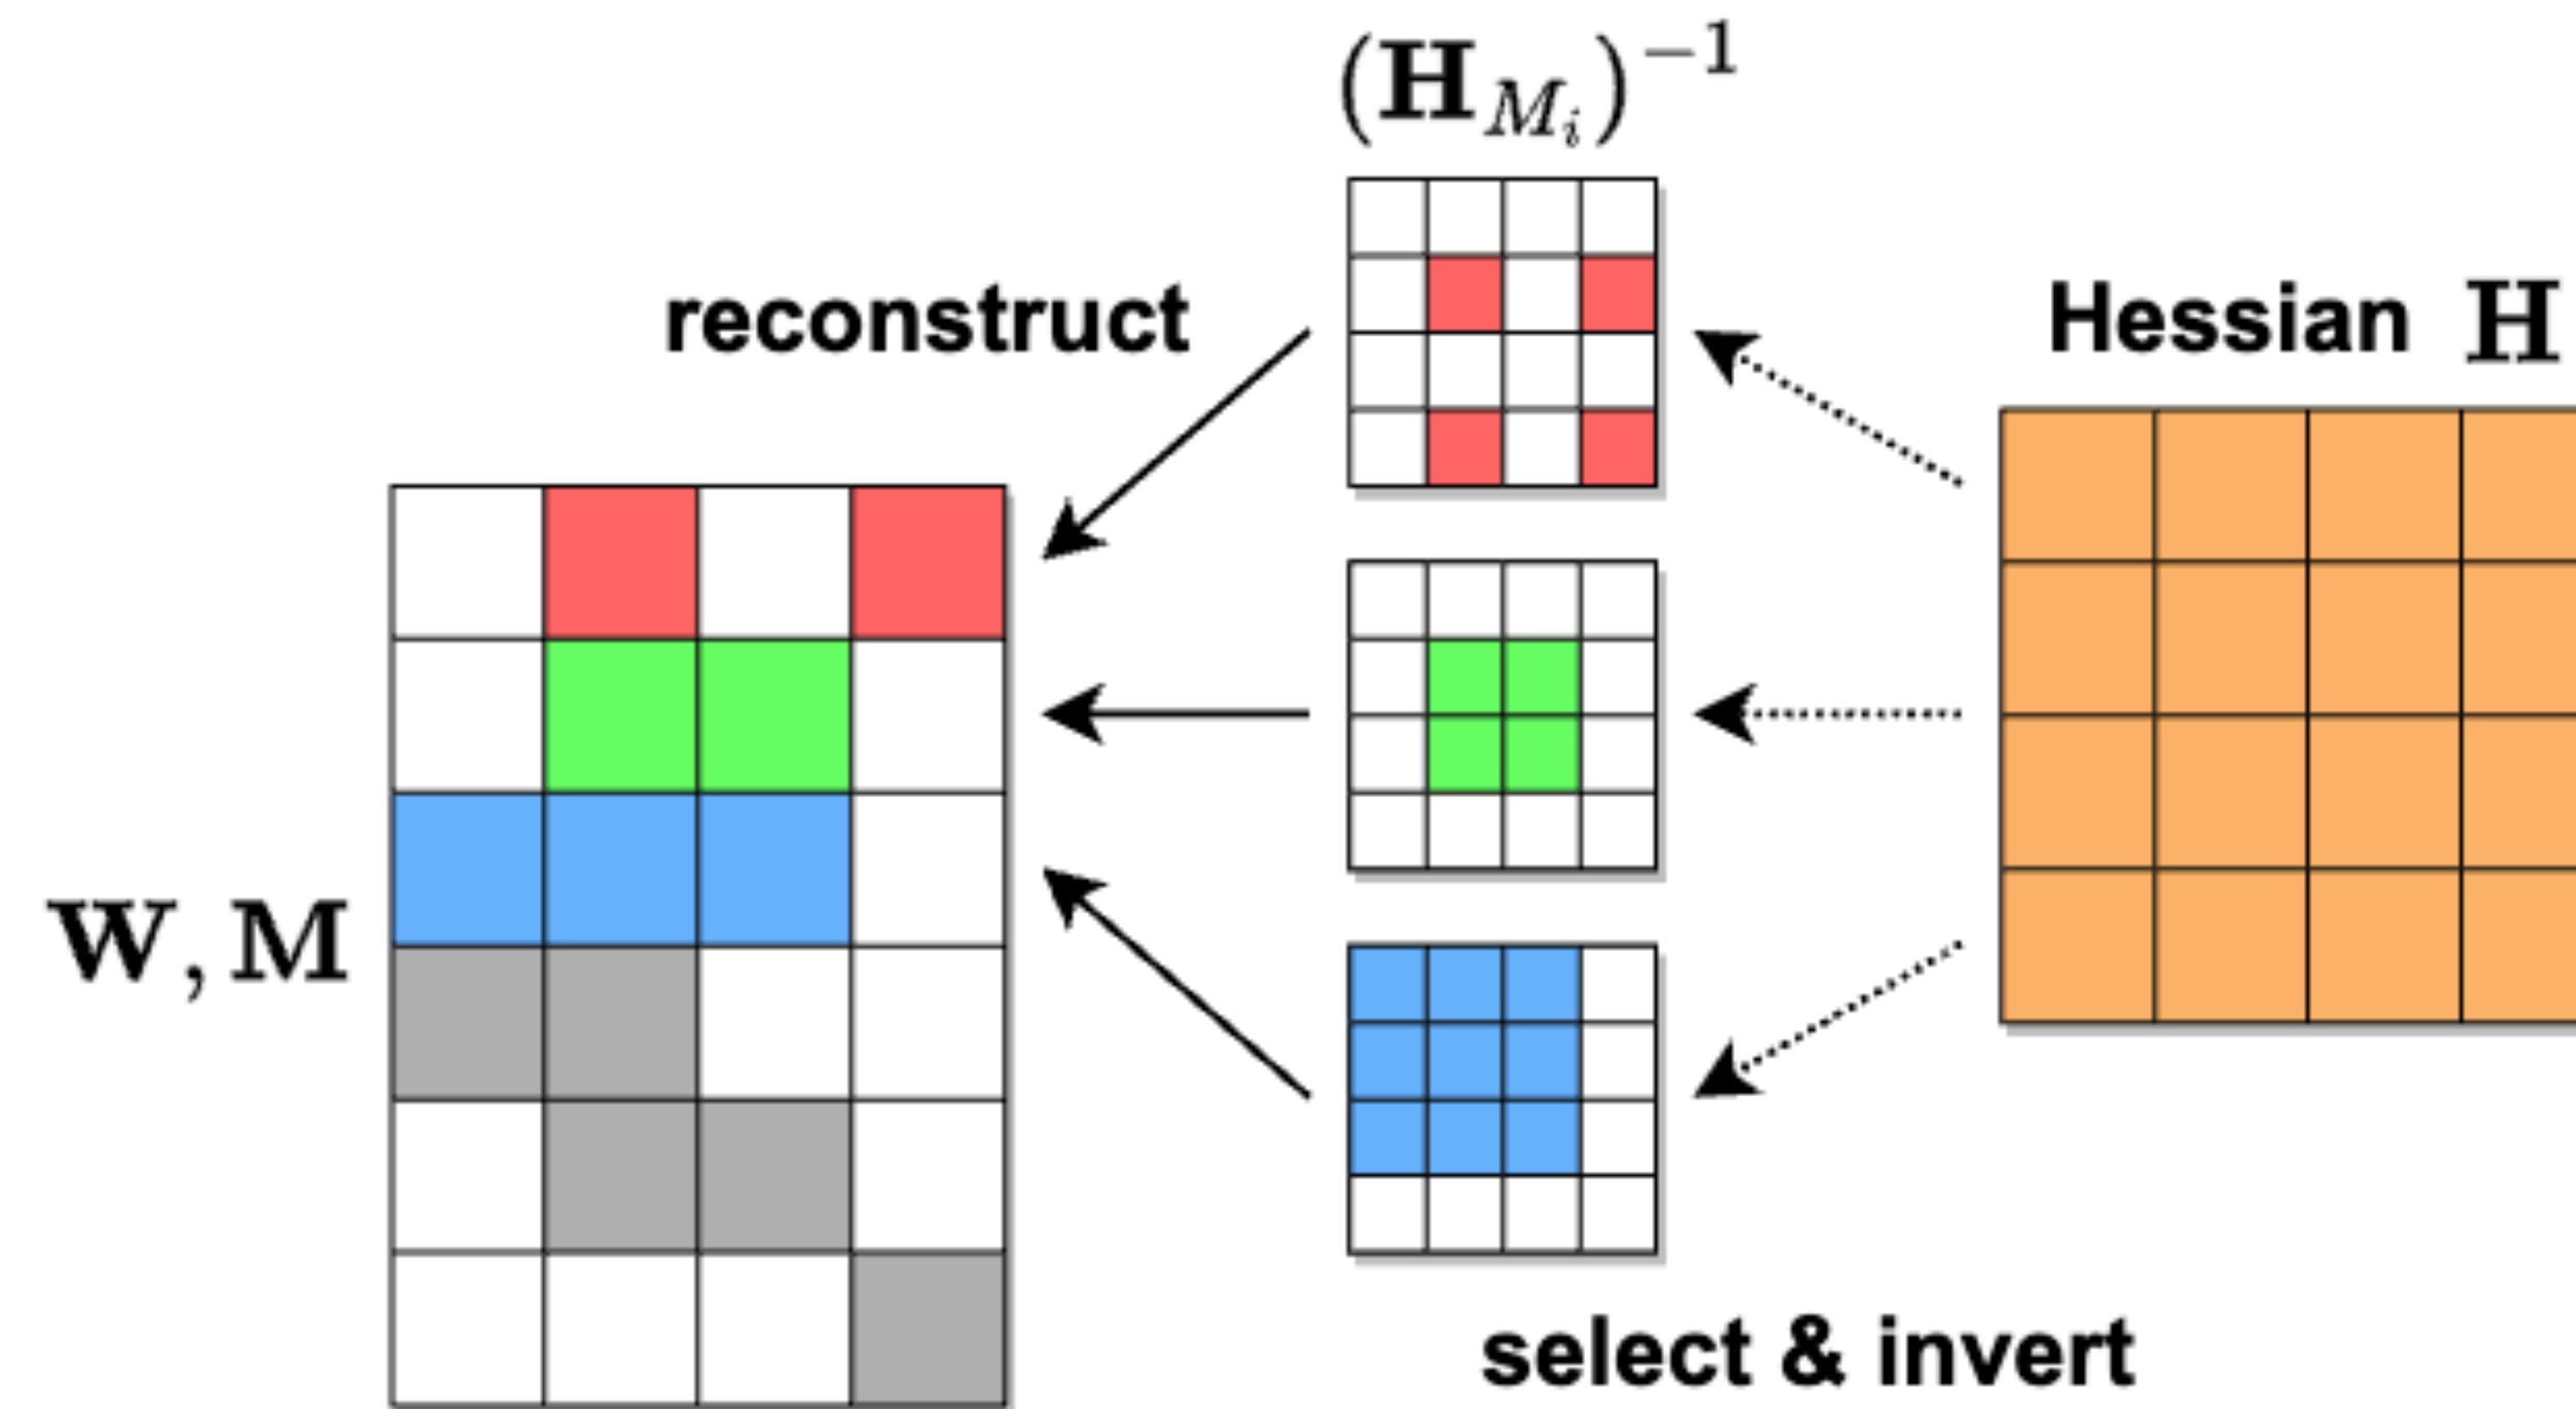
- Fortunately, we have the following lemma:
- **Lemma.** Given an invertible $d_{\text{col}} \times d_{\text{col}}$ matrix \mathbf{H} and its inverse \mathbf{H}^{-1} , let \mathbf{H}_{-i} denote the Hessian with row and column i removed. Then, we have:

$$\mathbf{H}_{-i}^{-1} = \left(\mathbf{H}^{-1} - \frac{1}{[\mathbf{H}^{-1}]_{ii}} \mathbf{H}_{:,i}^{-1} \mathbf{H}_{i,:}^{-1} \right)_{-i}$$

- Gaussian elimination of row/col i in \mathbf{H}^{-1} , then remove them completely
- Complexity of $\Theta(d_{\text{col}}^2)$
- Allows parallel processing of rows

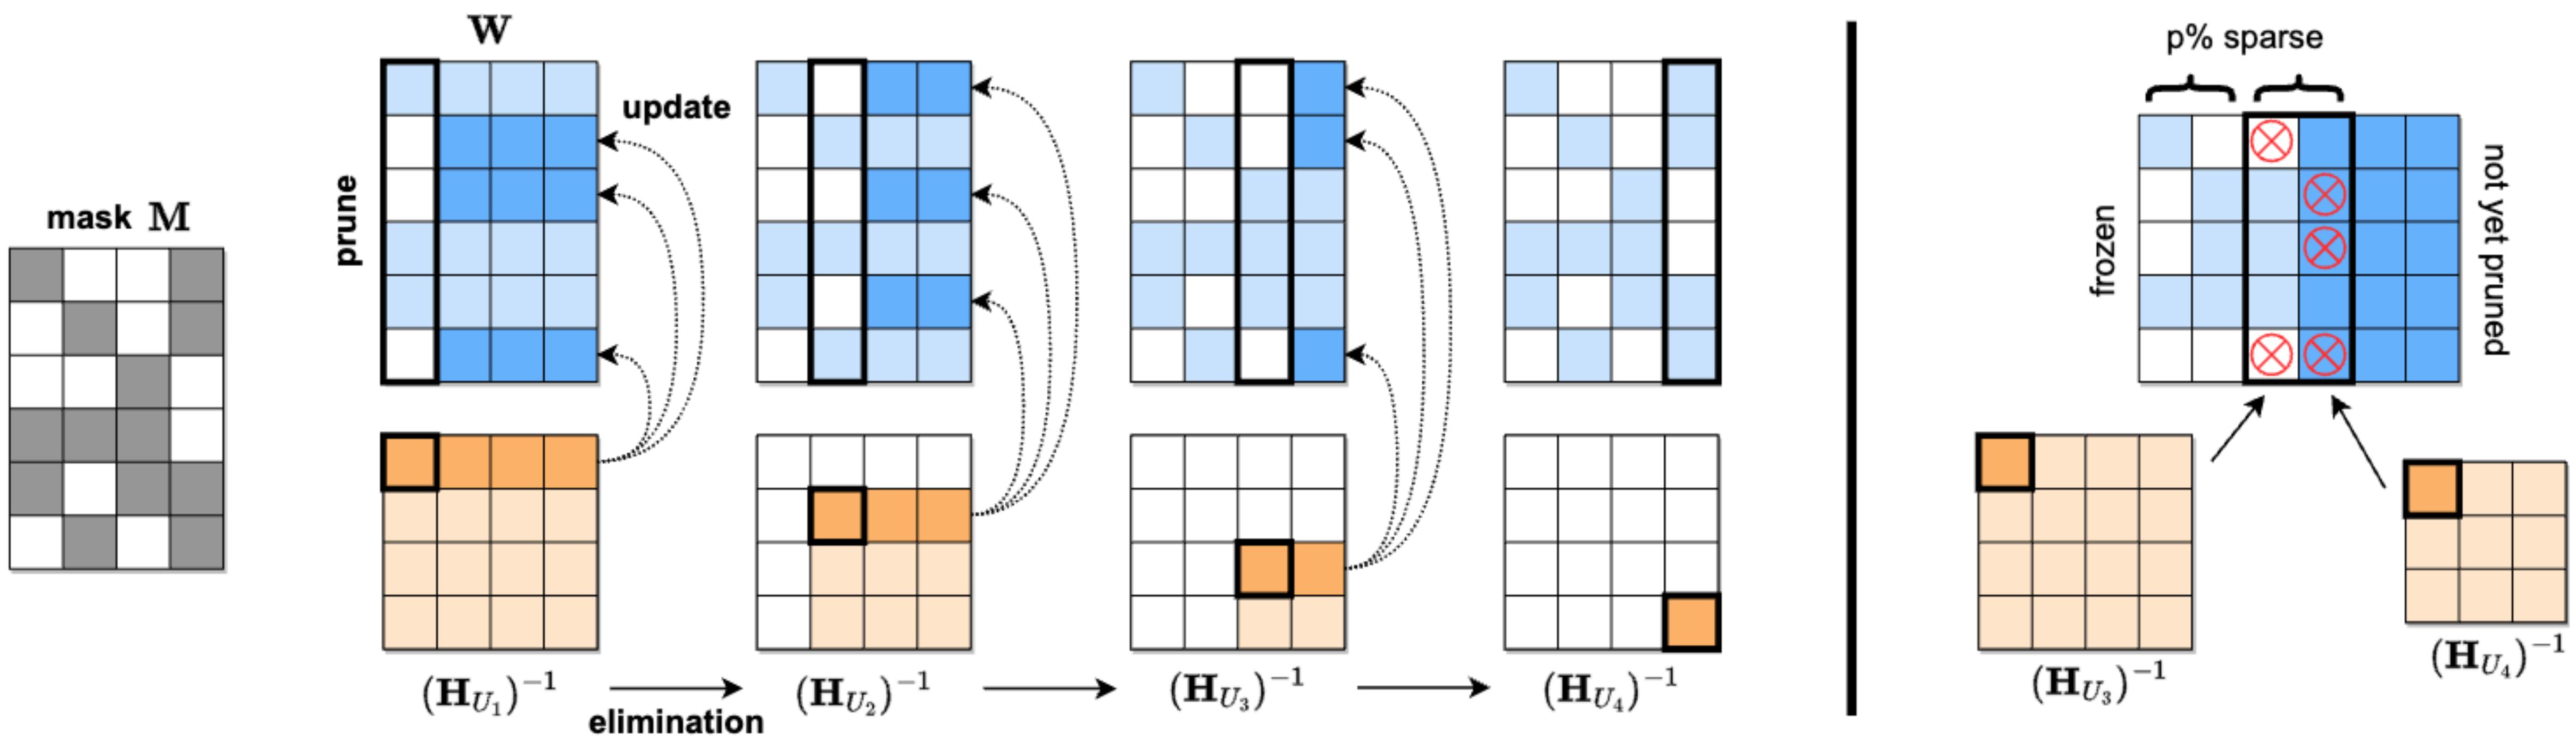
SparseGPT

- Problem. Keeping track of **row-wise Hessian inverse** is memory-heavy



SparseGPT

- **Idea.** Iterate over columns, removing certain fraction at a time and freezing the surviving weights



Quantization

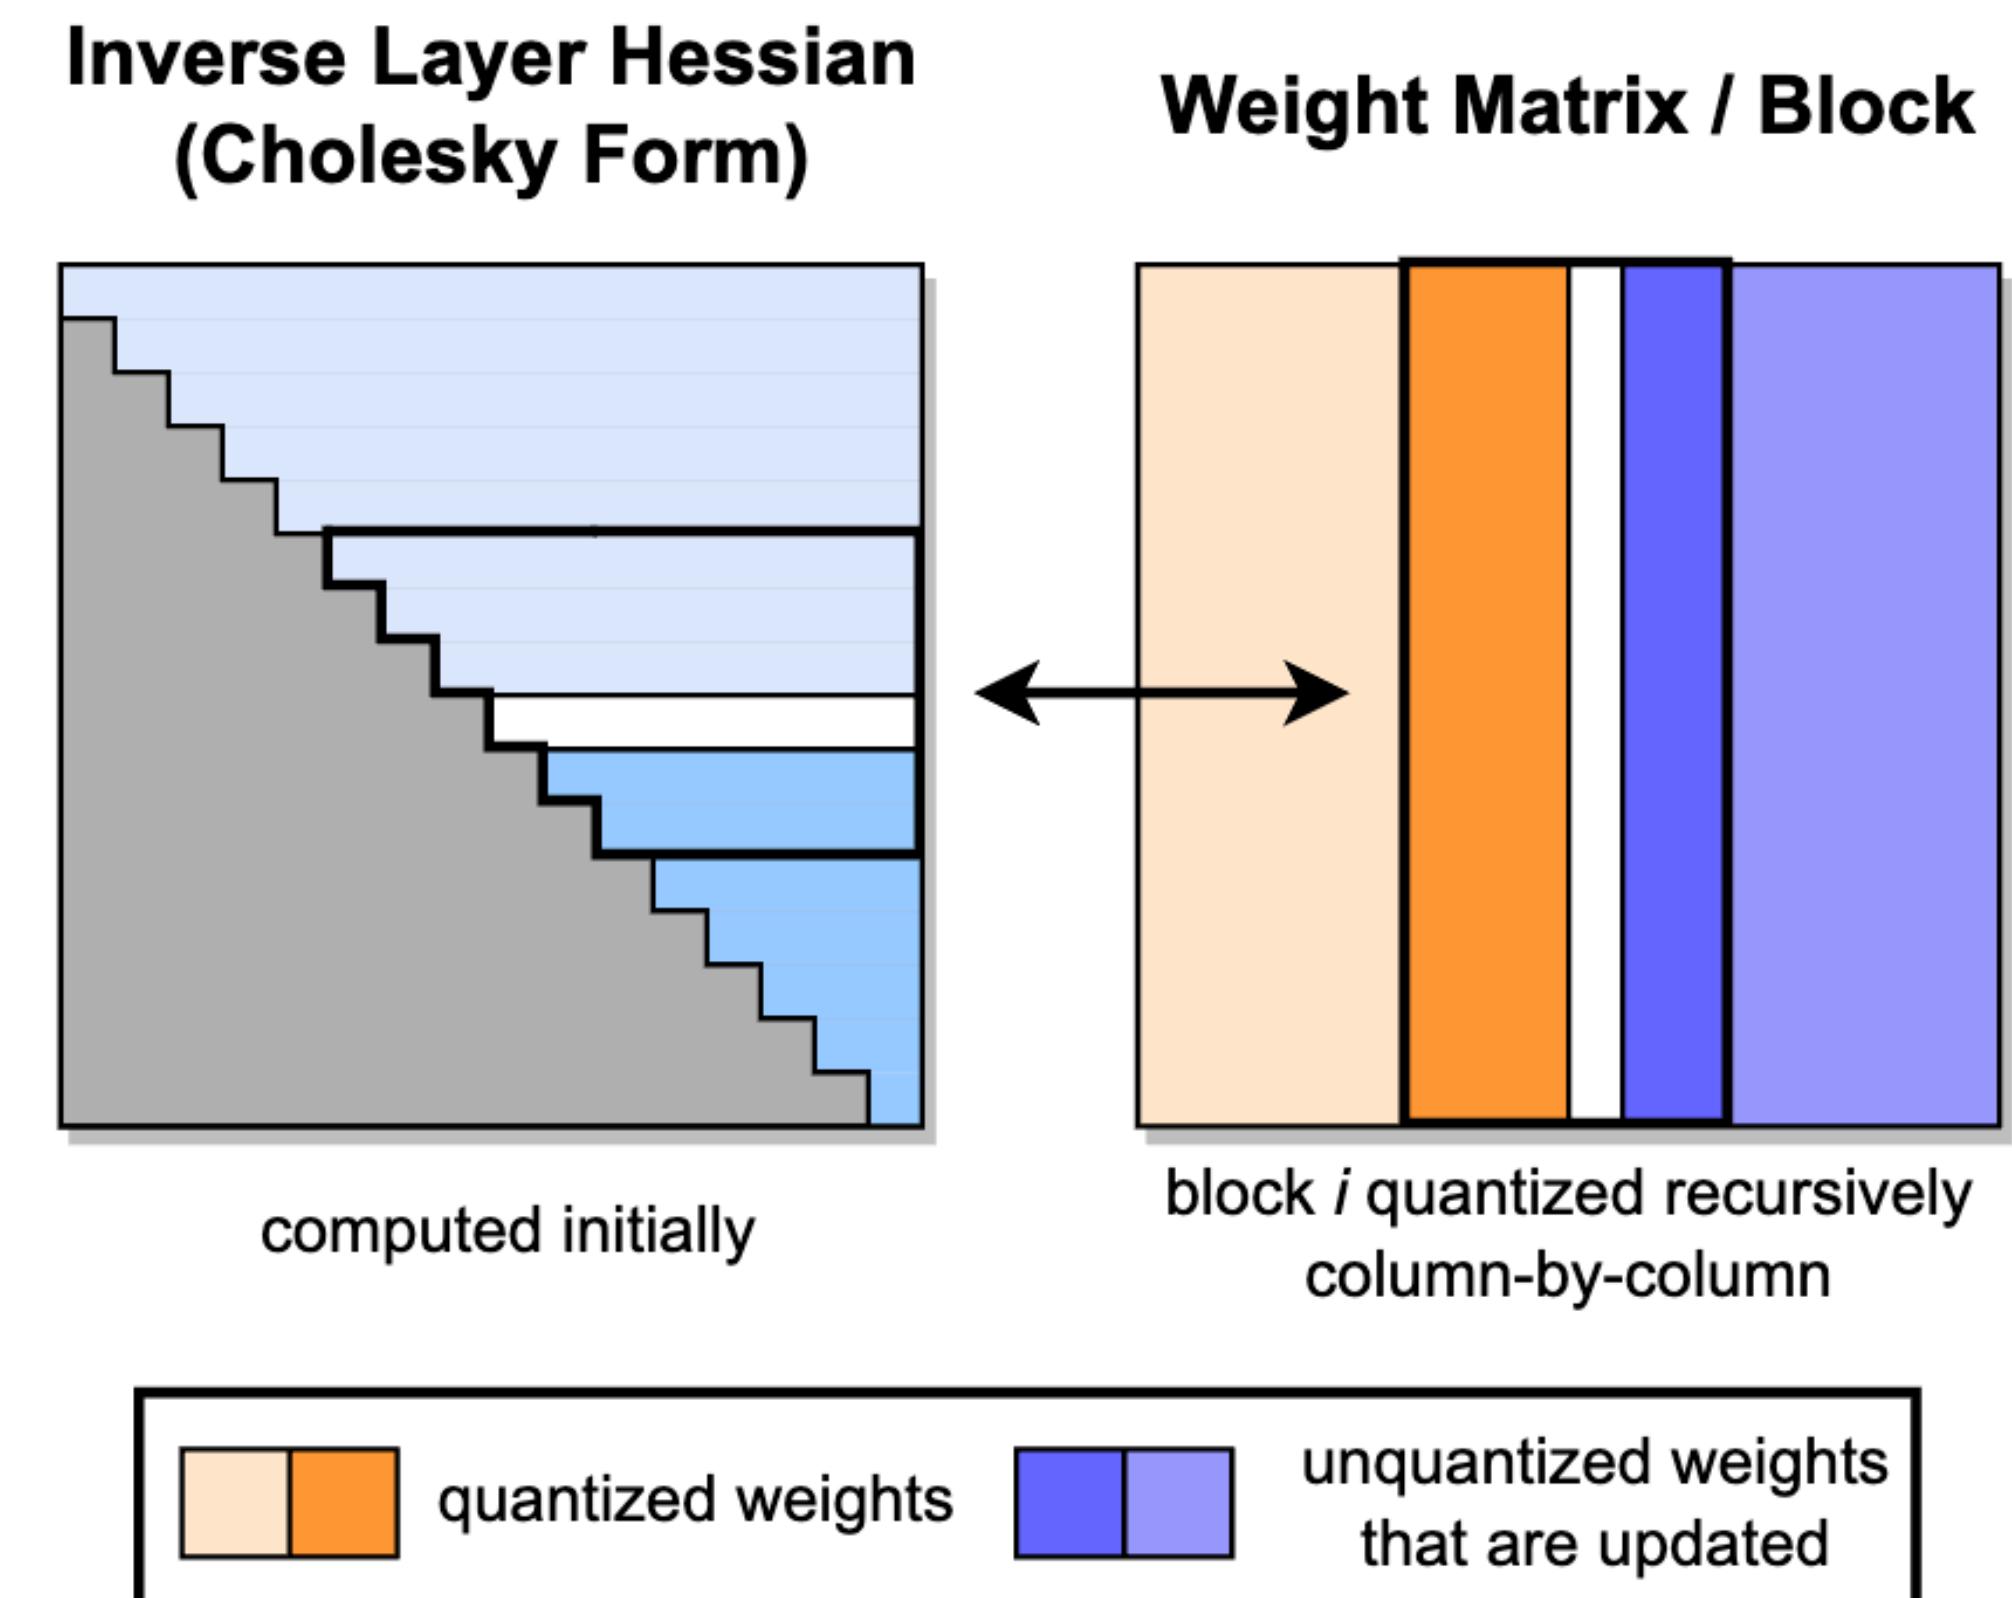
- Similarly, the **weight quantization** can be done using Hessians
 - For simplicity, assume that we have a fixed grid, and put weights on grid one-by-one.
 - Similar derivation gives that the weight updates needed to compensate for quantizing w_q is:

$$\delta = \frac{\text{quant}(w_q) - w_q}{[\mathbf{H}^{-1}]_{qq}} \mathbf{H}^{-1} e_q$$

- Weights quantized later are likelier to change more
- **Heuristic.** Quantize the outliers, as soon as they appear

GPTQ

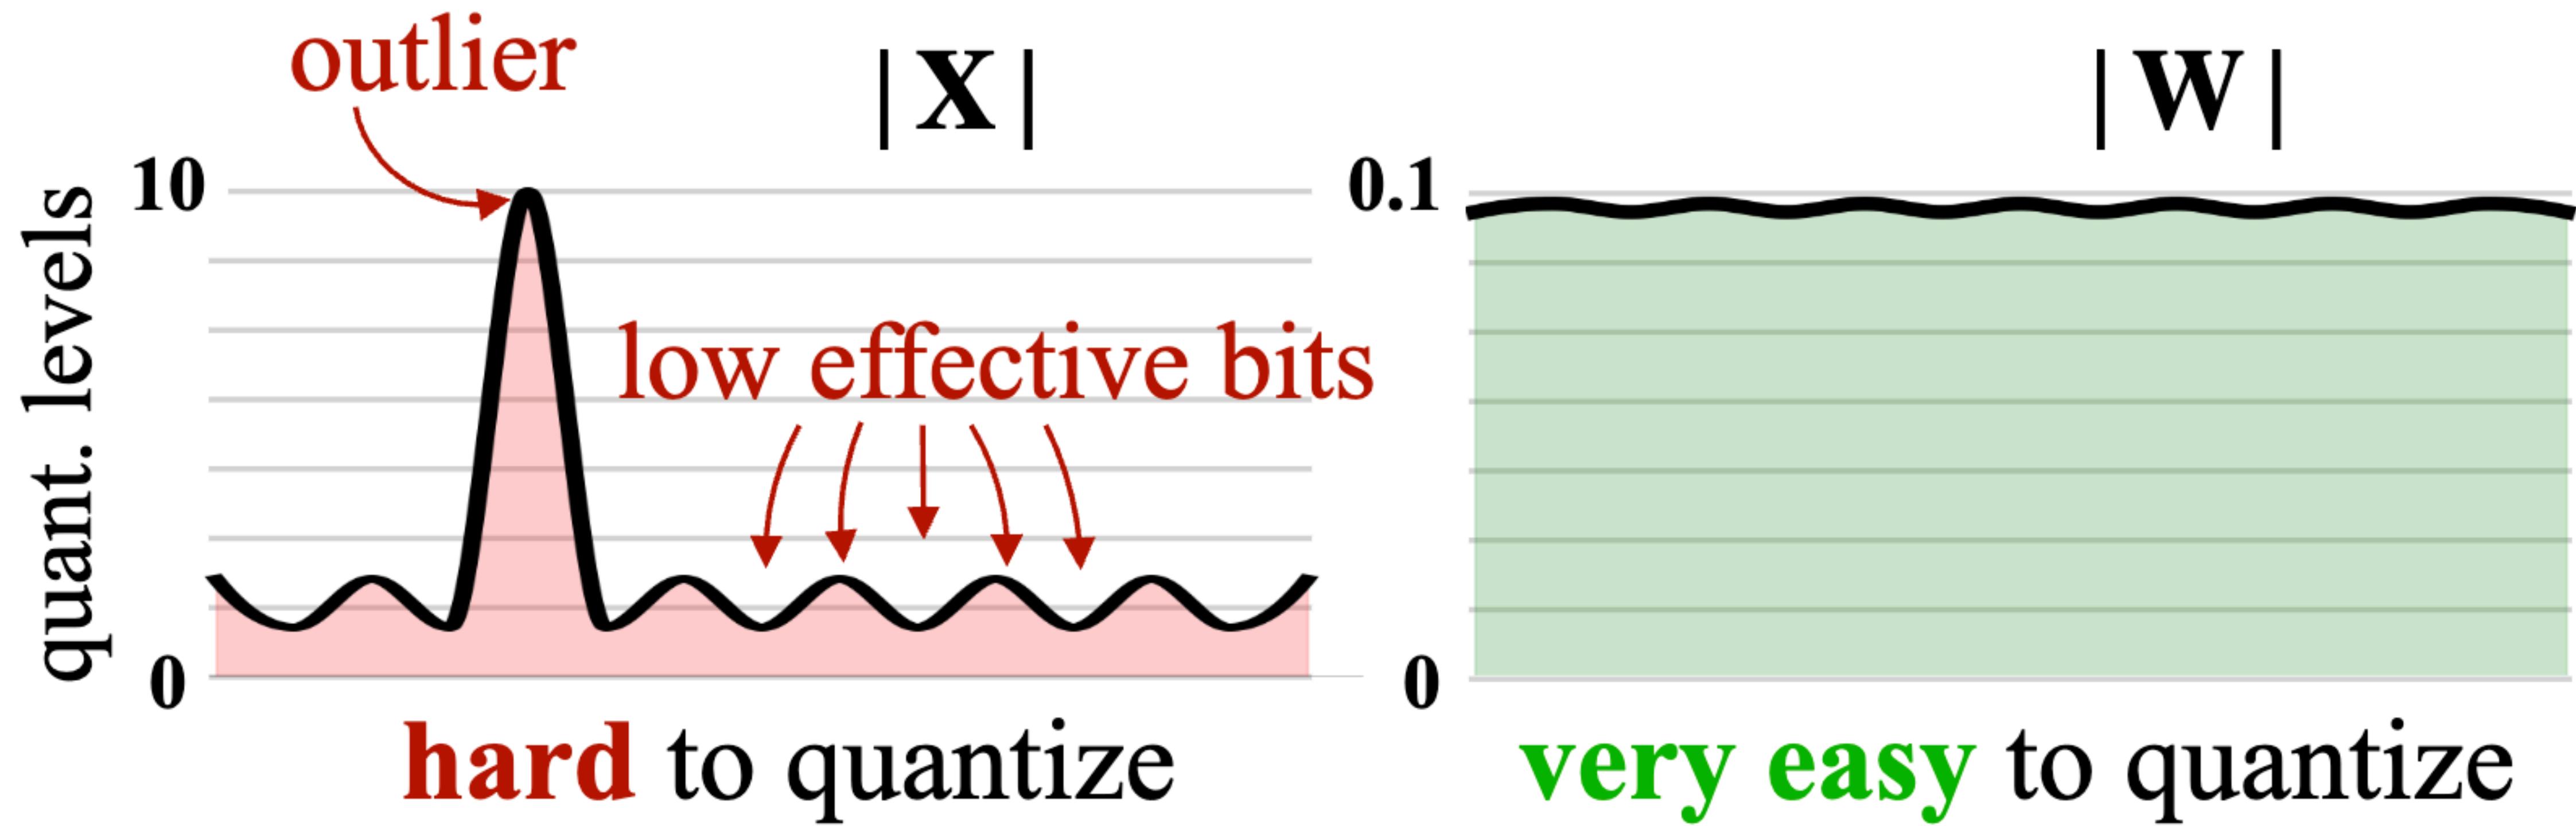
- **Order.** Later, people discovered that the order does not matter much
 - Simple quantized in an order, in parallel, save time
- **More tricks** (see the paper)
 - Lazy batching: Update later columns a bit slowly to prevent frequent memory access
 - Cholesky reformulation: Improved numerical stability, avoiding the accumulation of errors from repeated updates



Outlier-driven Approach

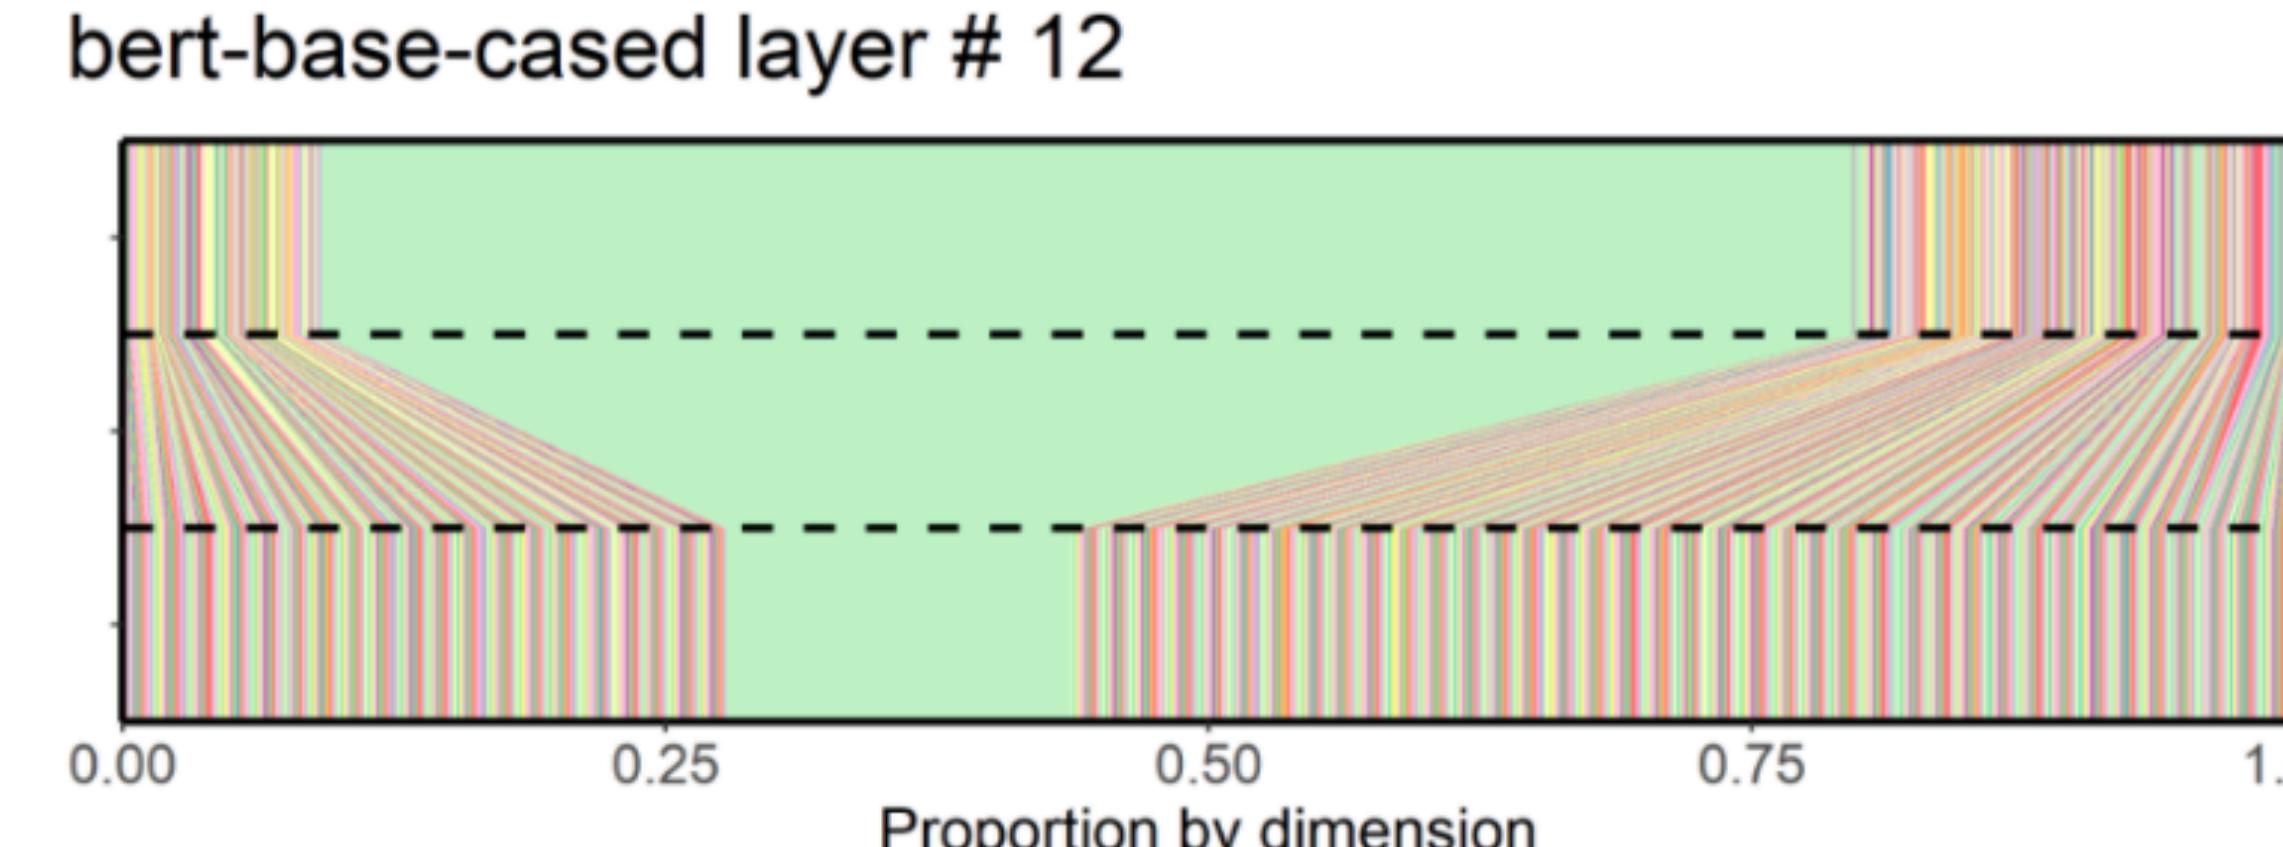
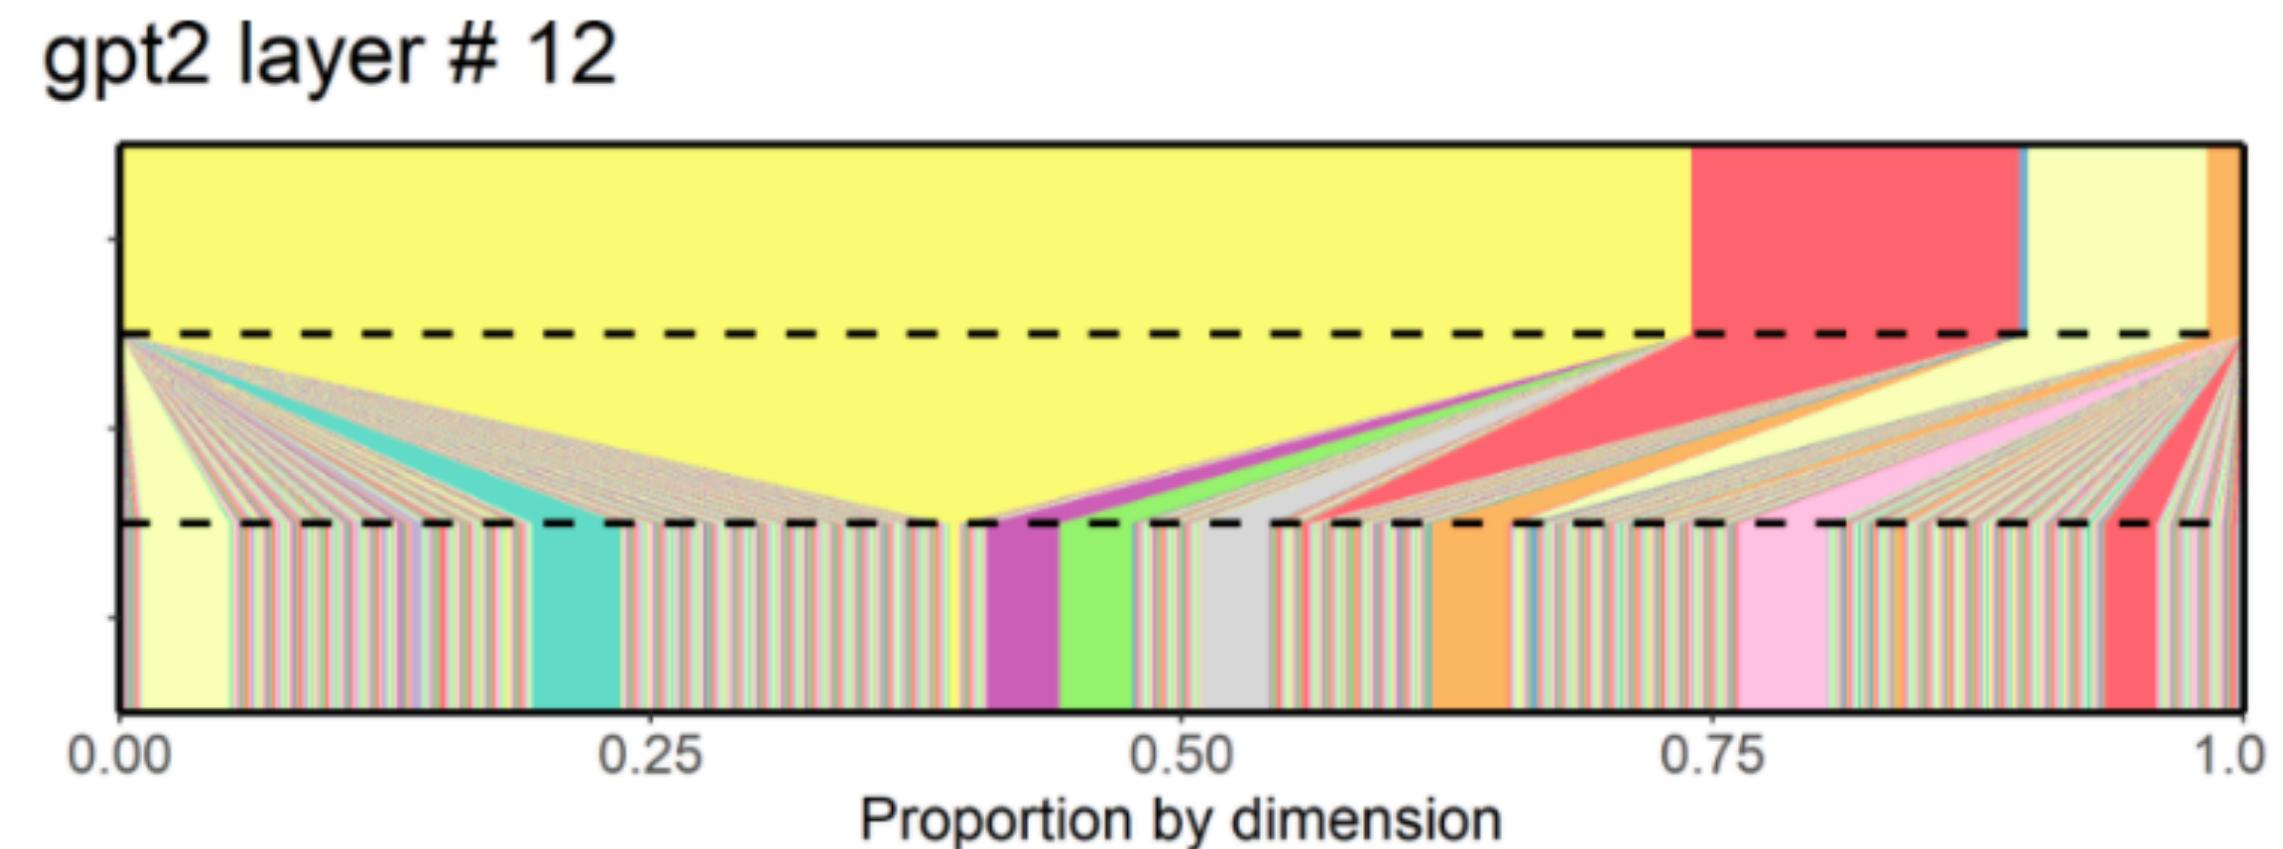
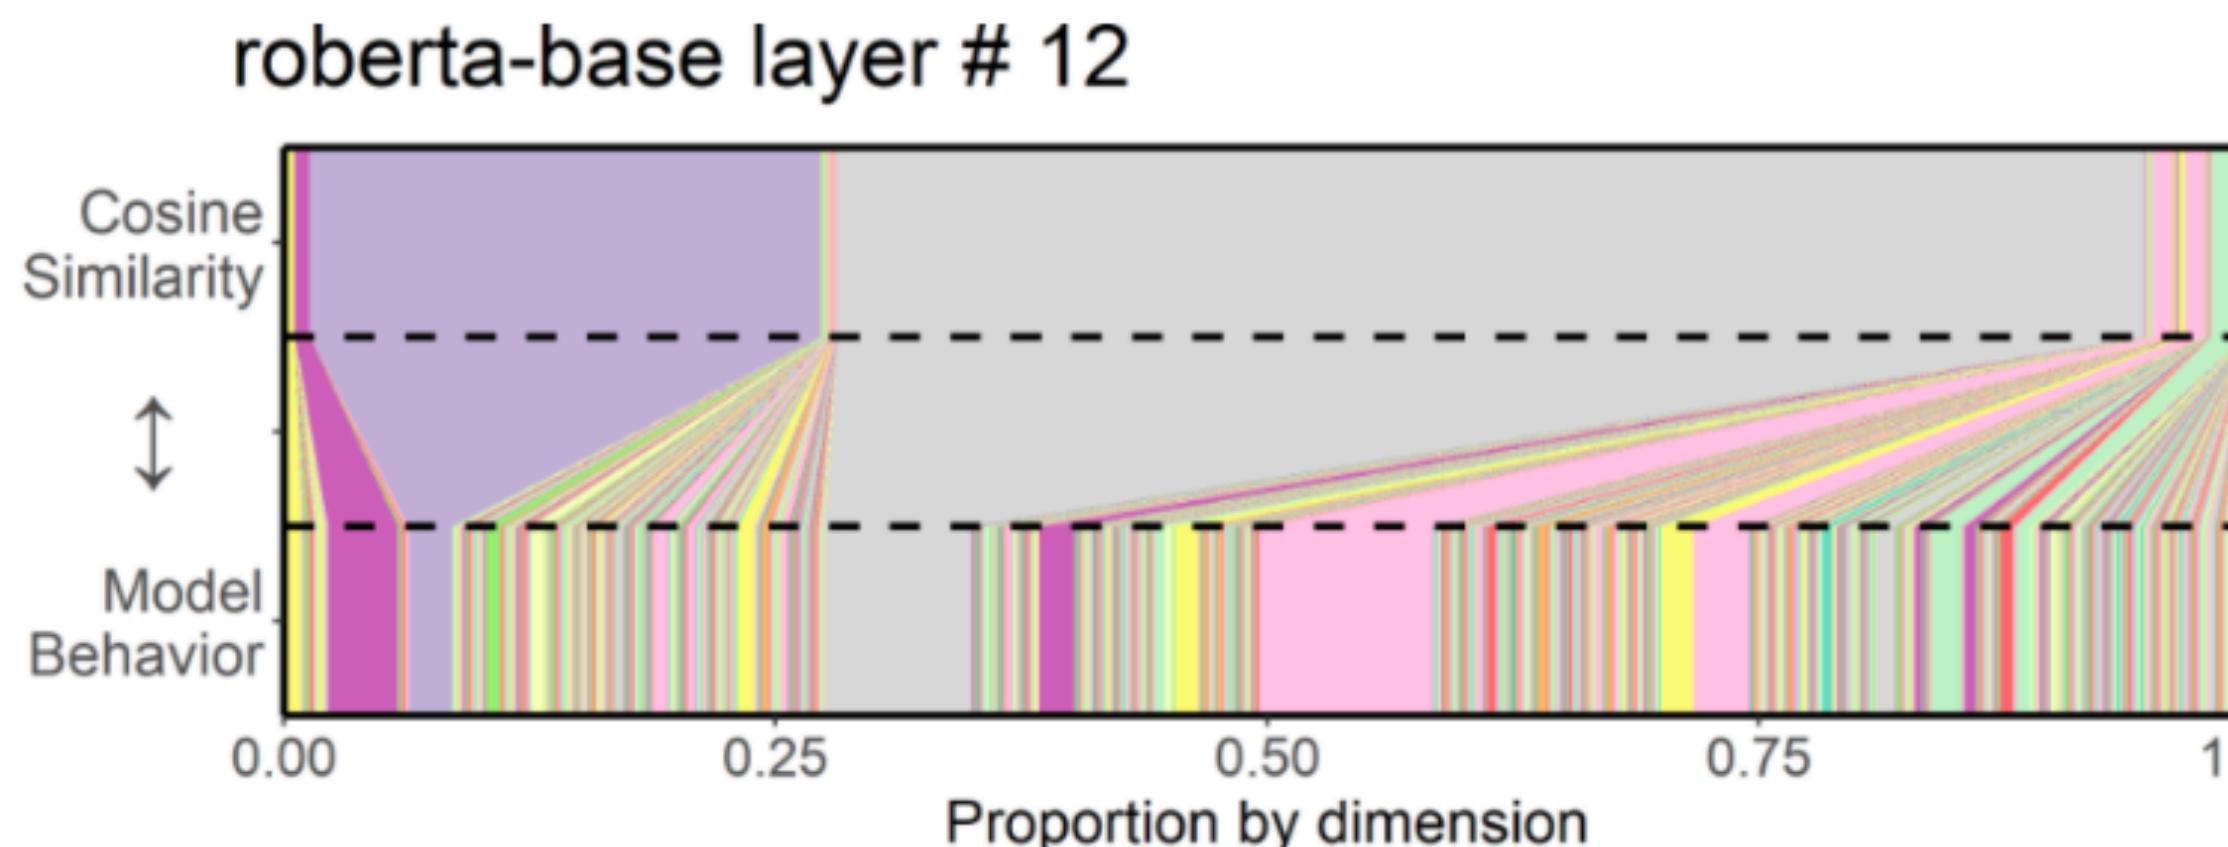
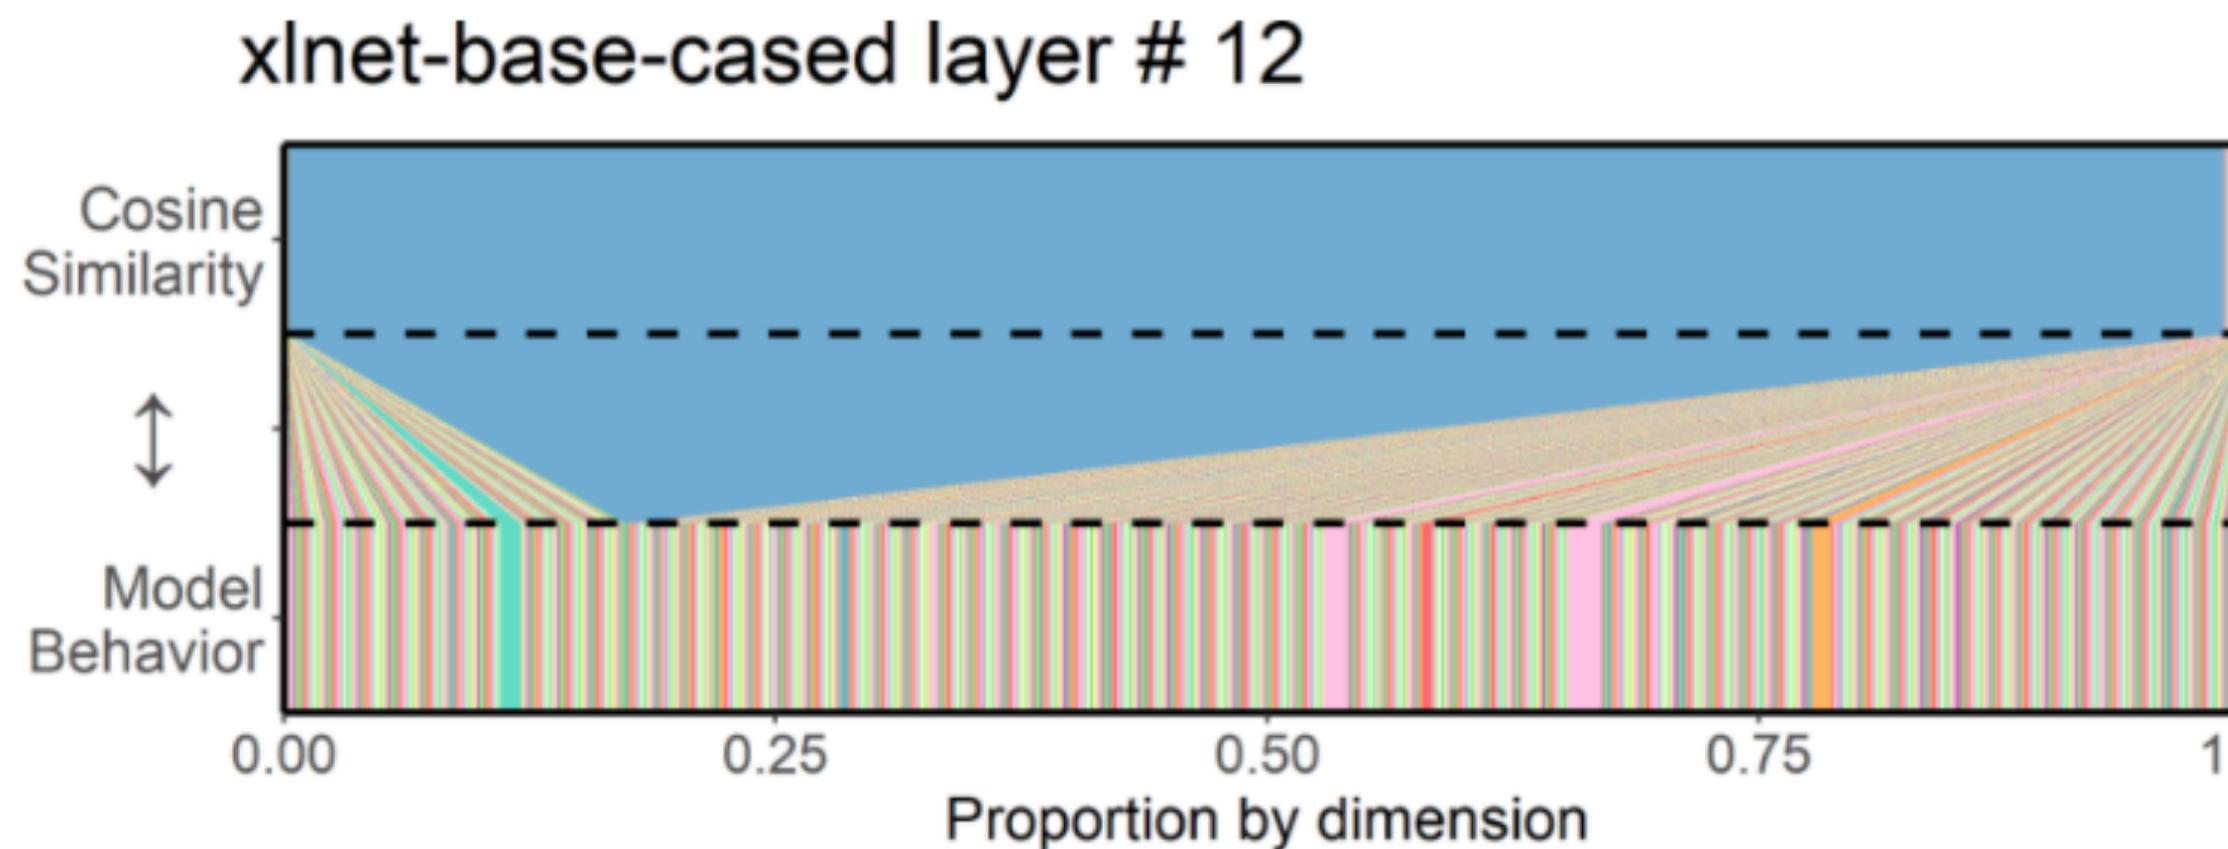
Outliers?

- Transformer activations of very large magnitude
 - Happens in a small number of **channels** / tokens / layers



Timkey & van Schijndel (EMNLP 2021)

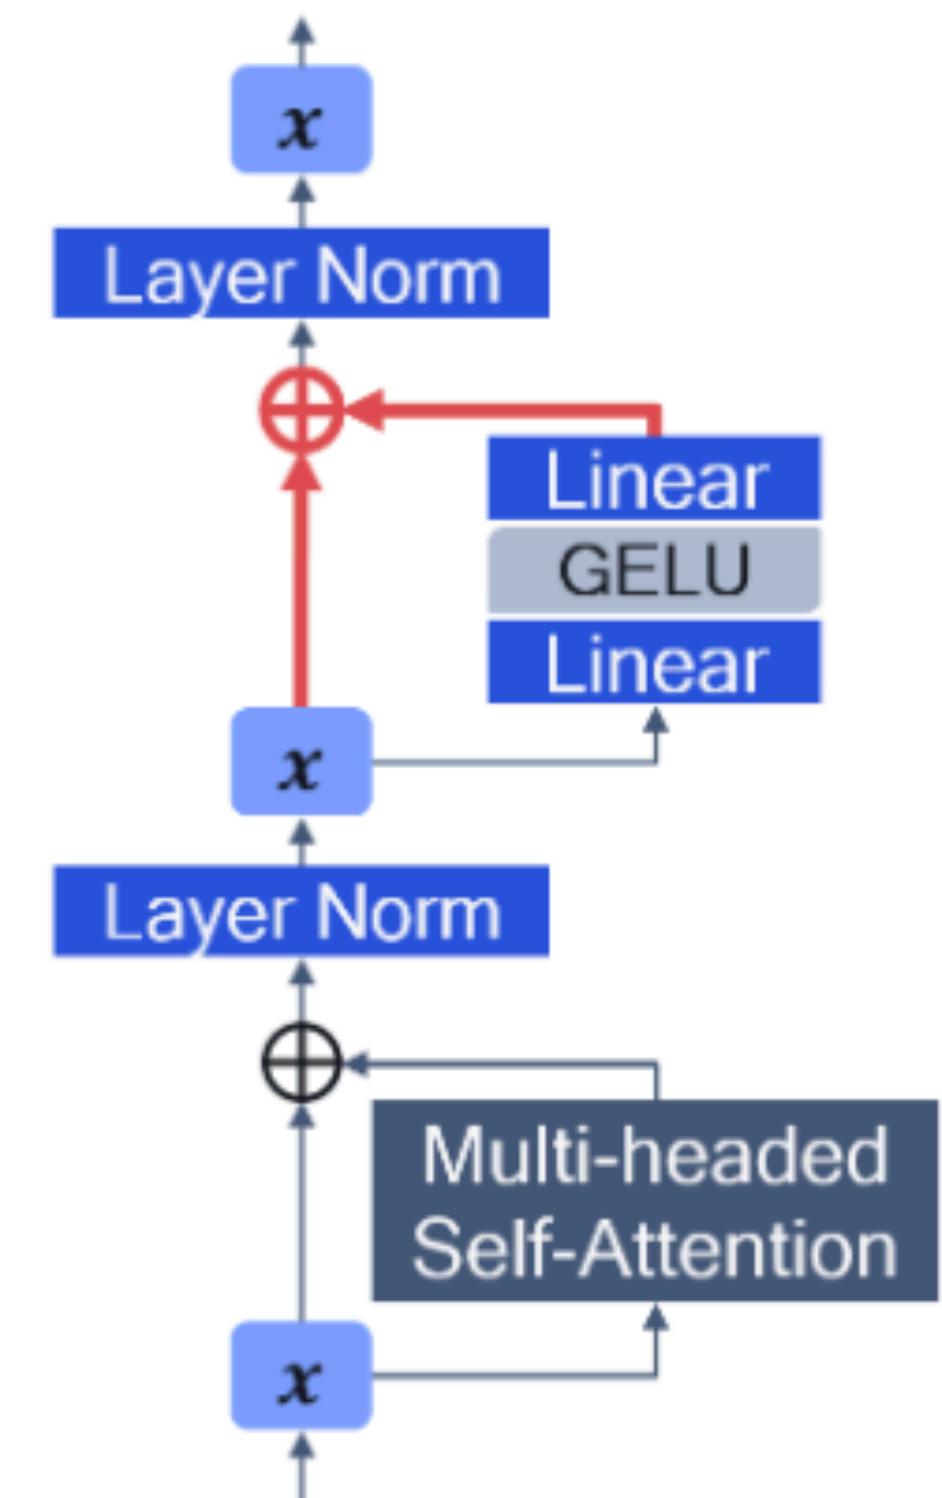
- Observes that 1-3 channels in later layers are outliers
 - Dominates the cosine similarity computation



Bondarenko et al. (EMNLP 2021)

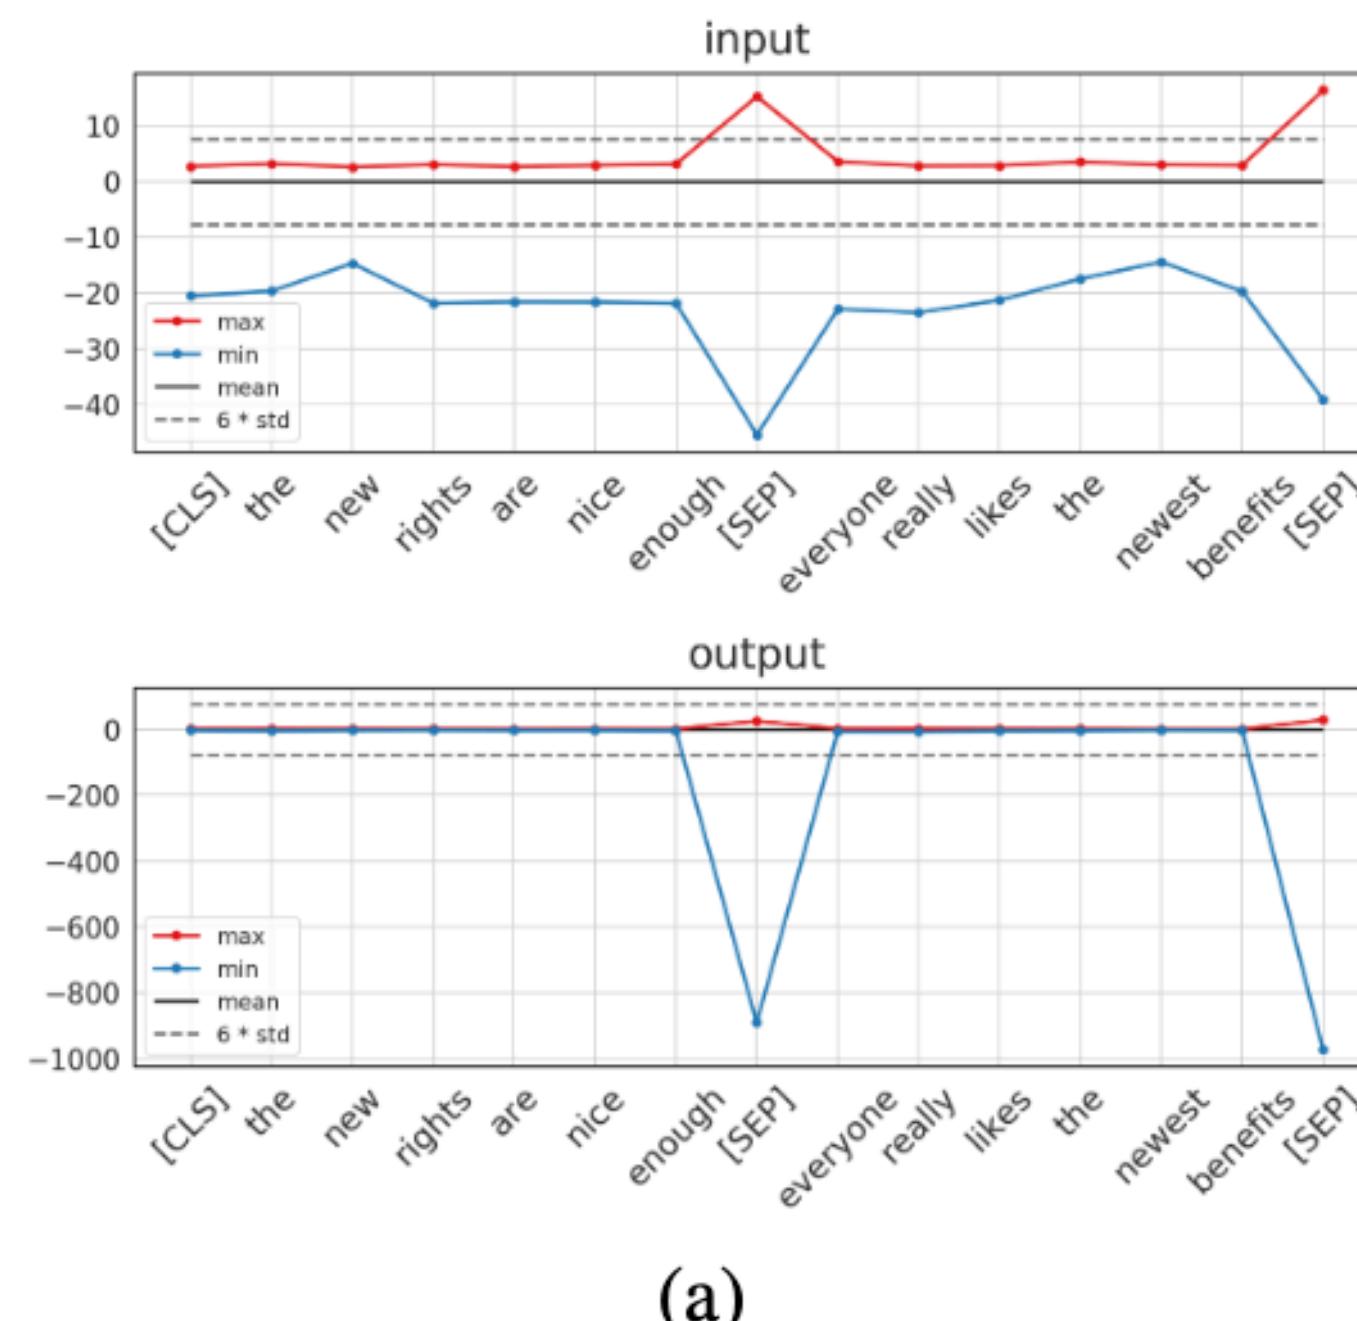
- Outliers are the residual sum after FFN of layer 10 & 11
 - Large accuracy drop of W32A8 / W8A8 performance on BERT

Quantized activations	STS-B	MNLI	QNLI	RTE
none (FP32 model)	89.09	84.91	91.58	70.40
all	62.64	42.67	50.74	48.74
all, except softmax input	70.92	42.54	51.84	48.74
all, except sum of embeddings	67.57	46.82	51.22	51.26
all, except self-attention output	70.47	46.57	50.98	50.90
all, except softmax output	72.83	50.35	50.23	49.46
all, except residual connections after FFN	81.57	82.56	89.73	67.15
same as above, but for layers 10, 11 only	79.40	81.24	88.03	63.90

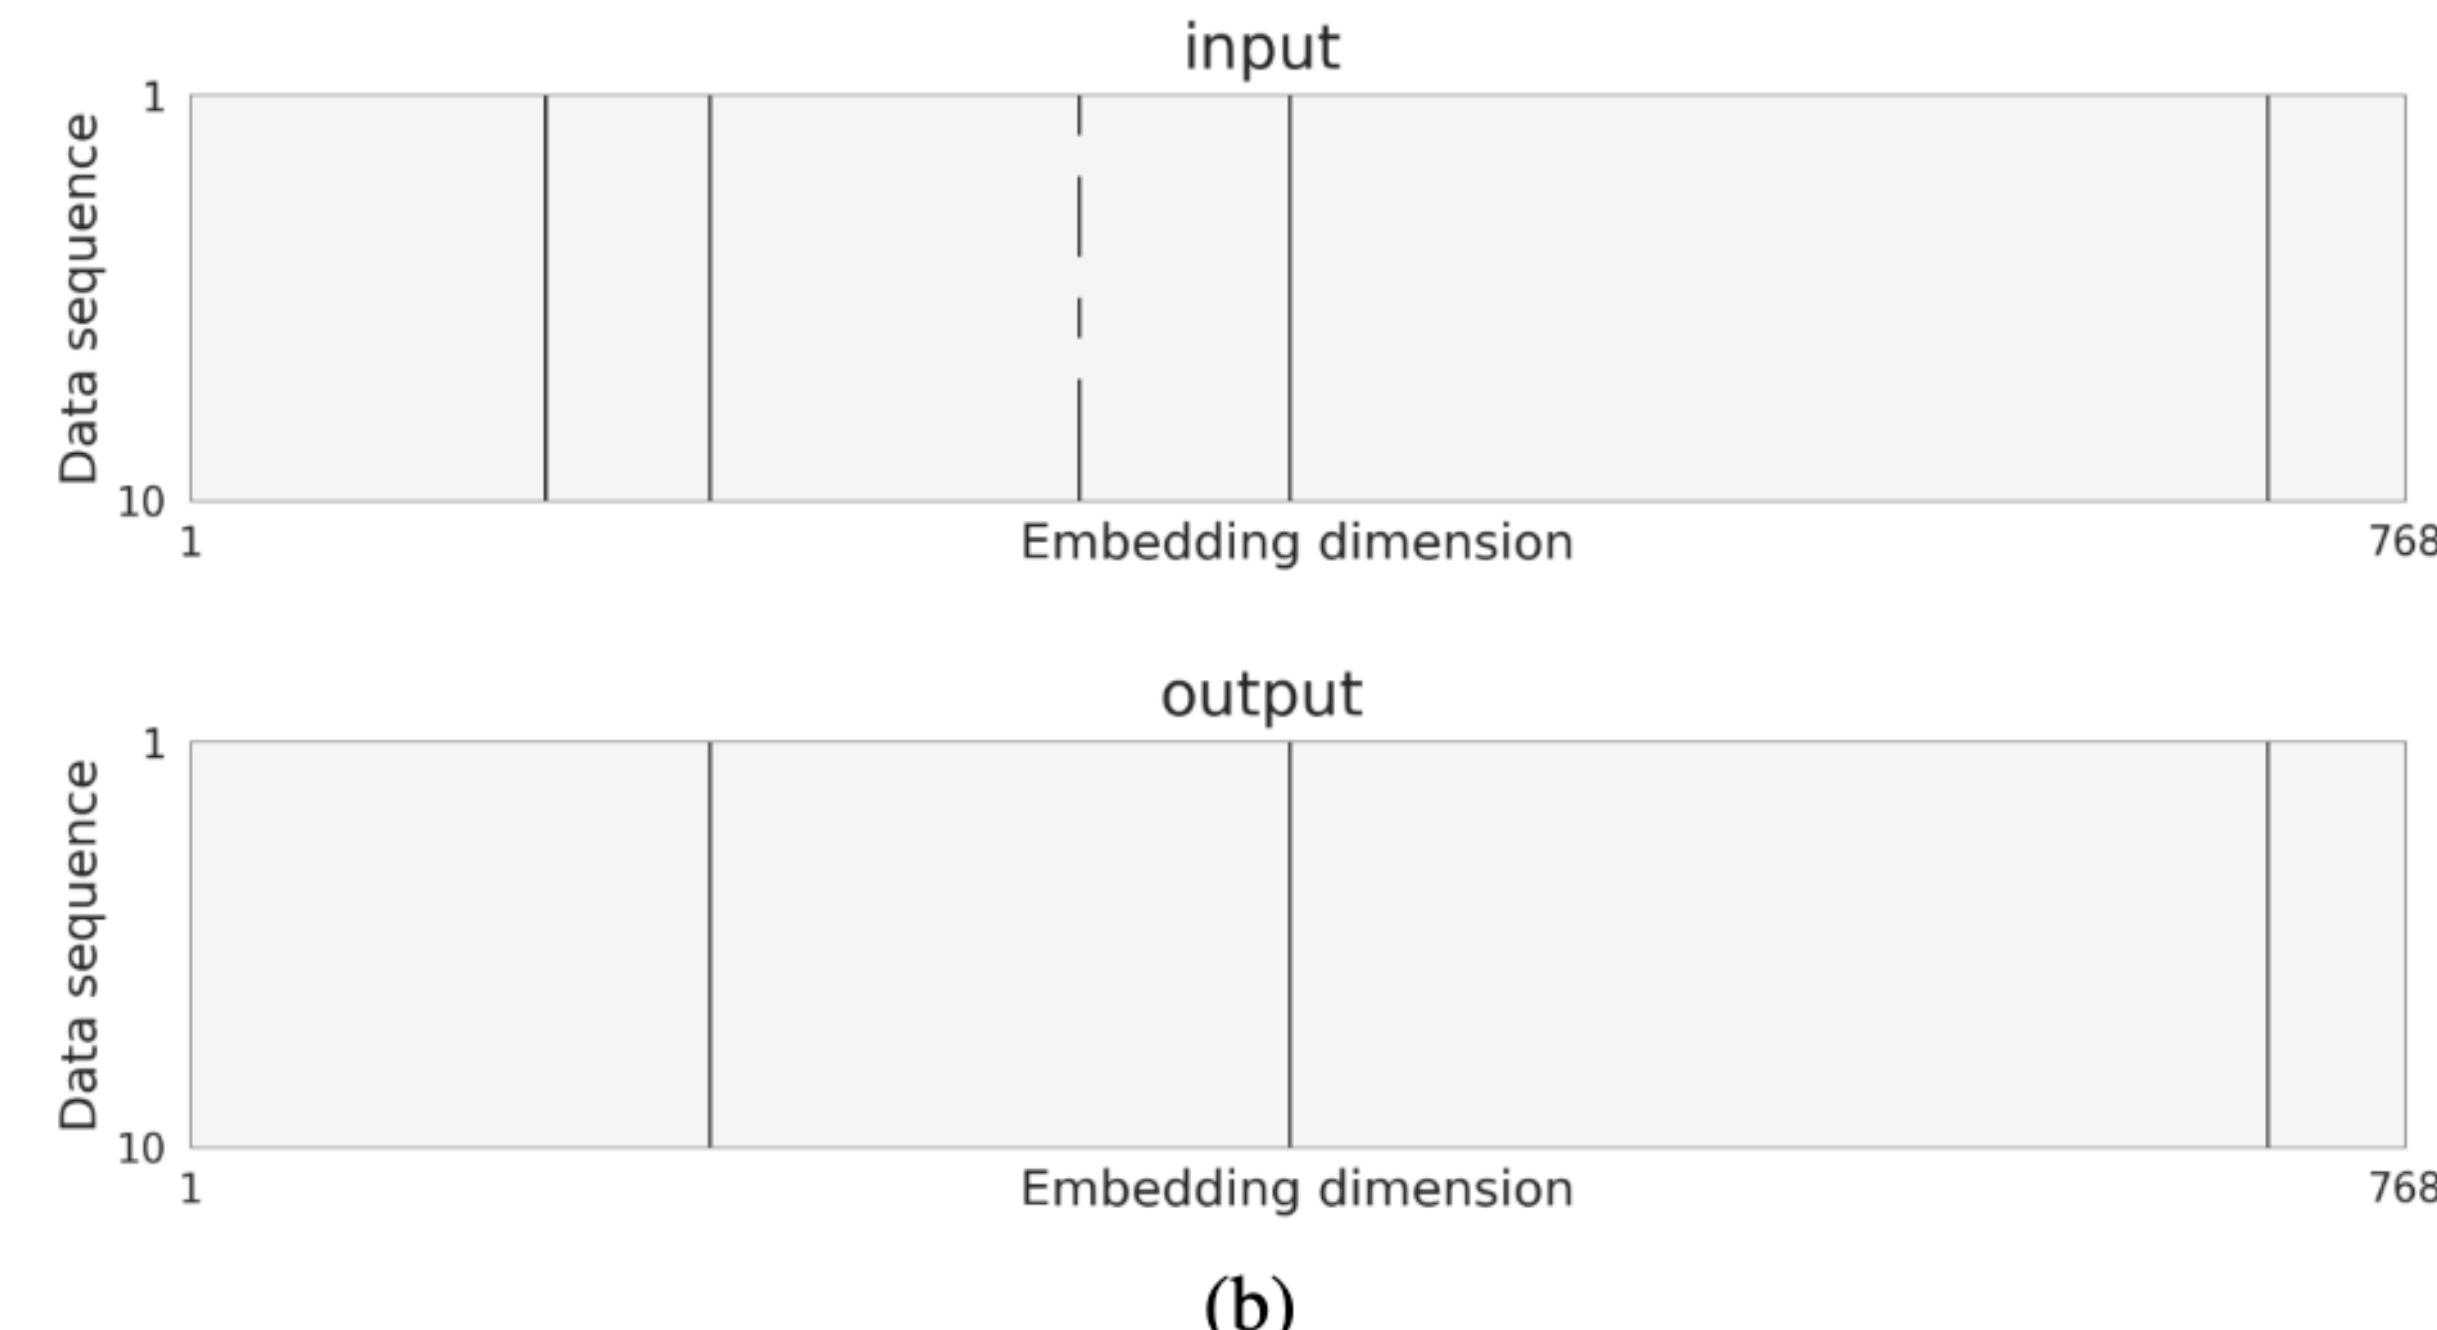


Bondarenko et al. (EMNLP 2021)

- Happens for [SEP] token, and small number of channels
 - Outliers have small variance, though



(a)

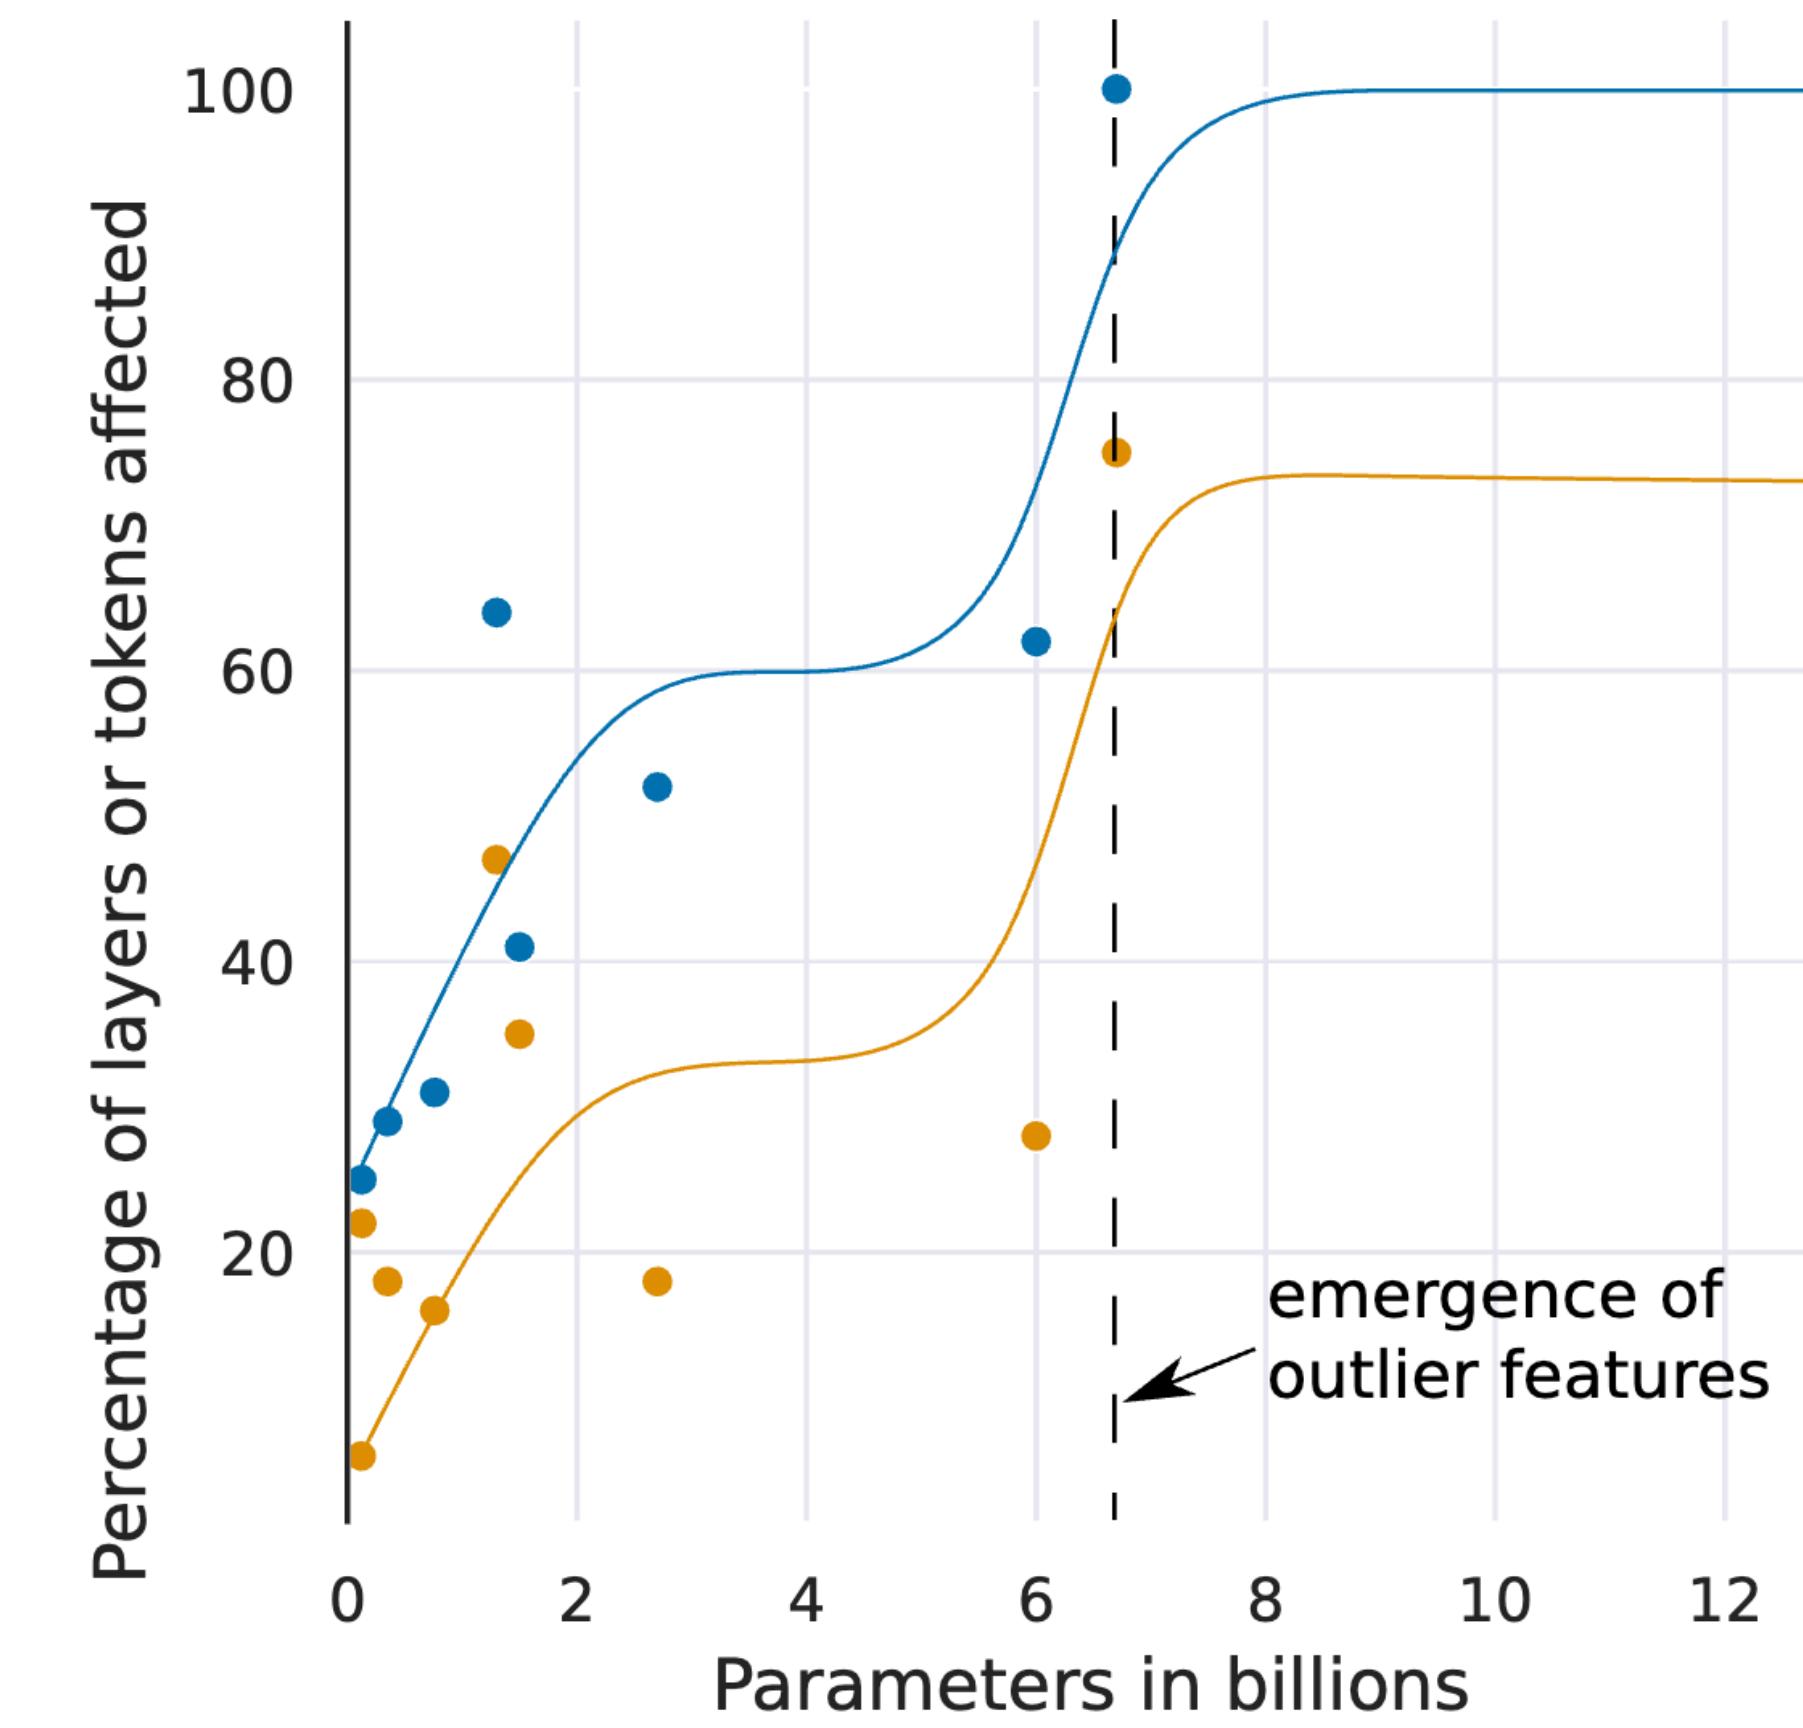
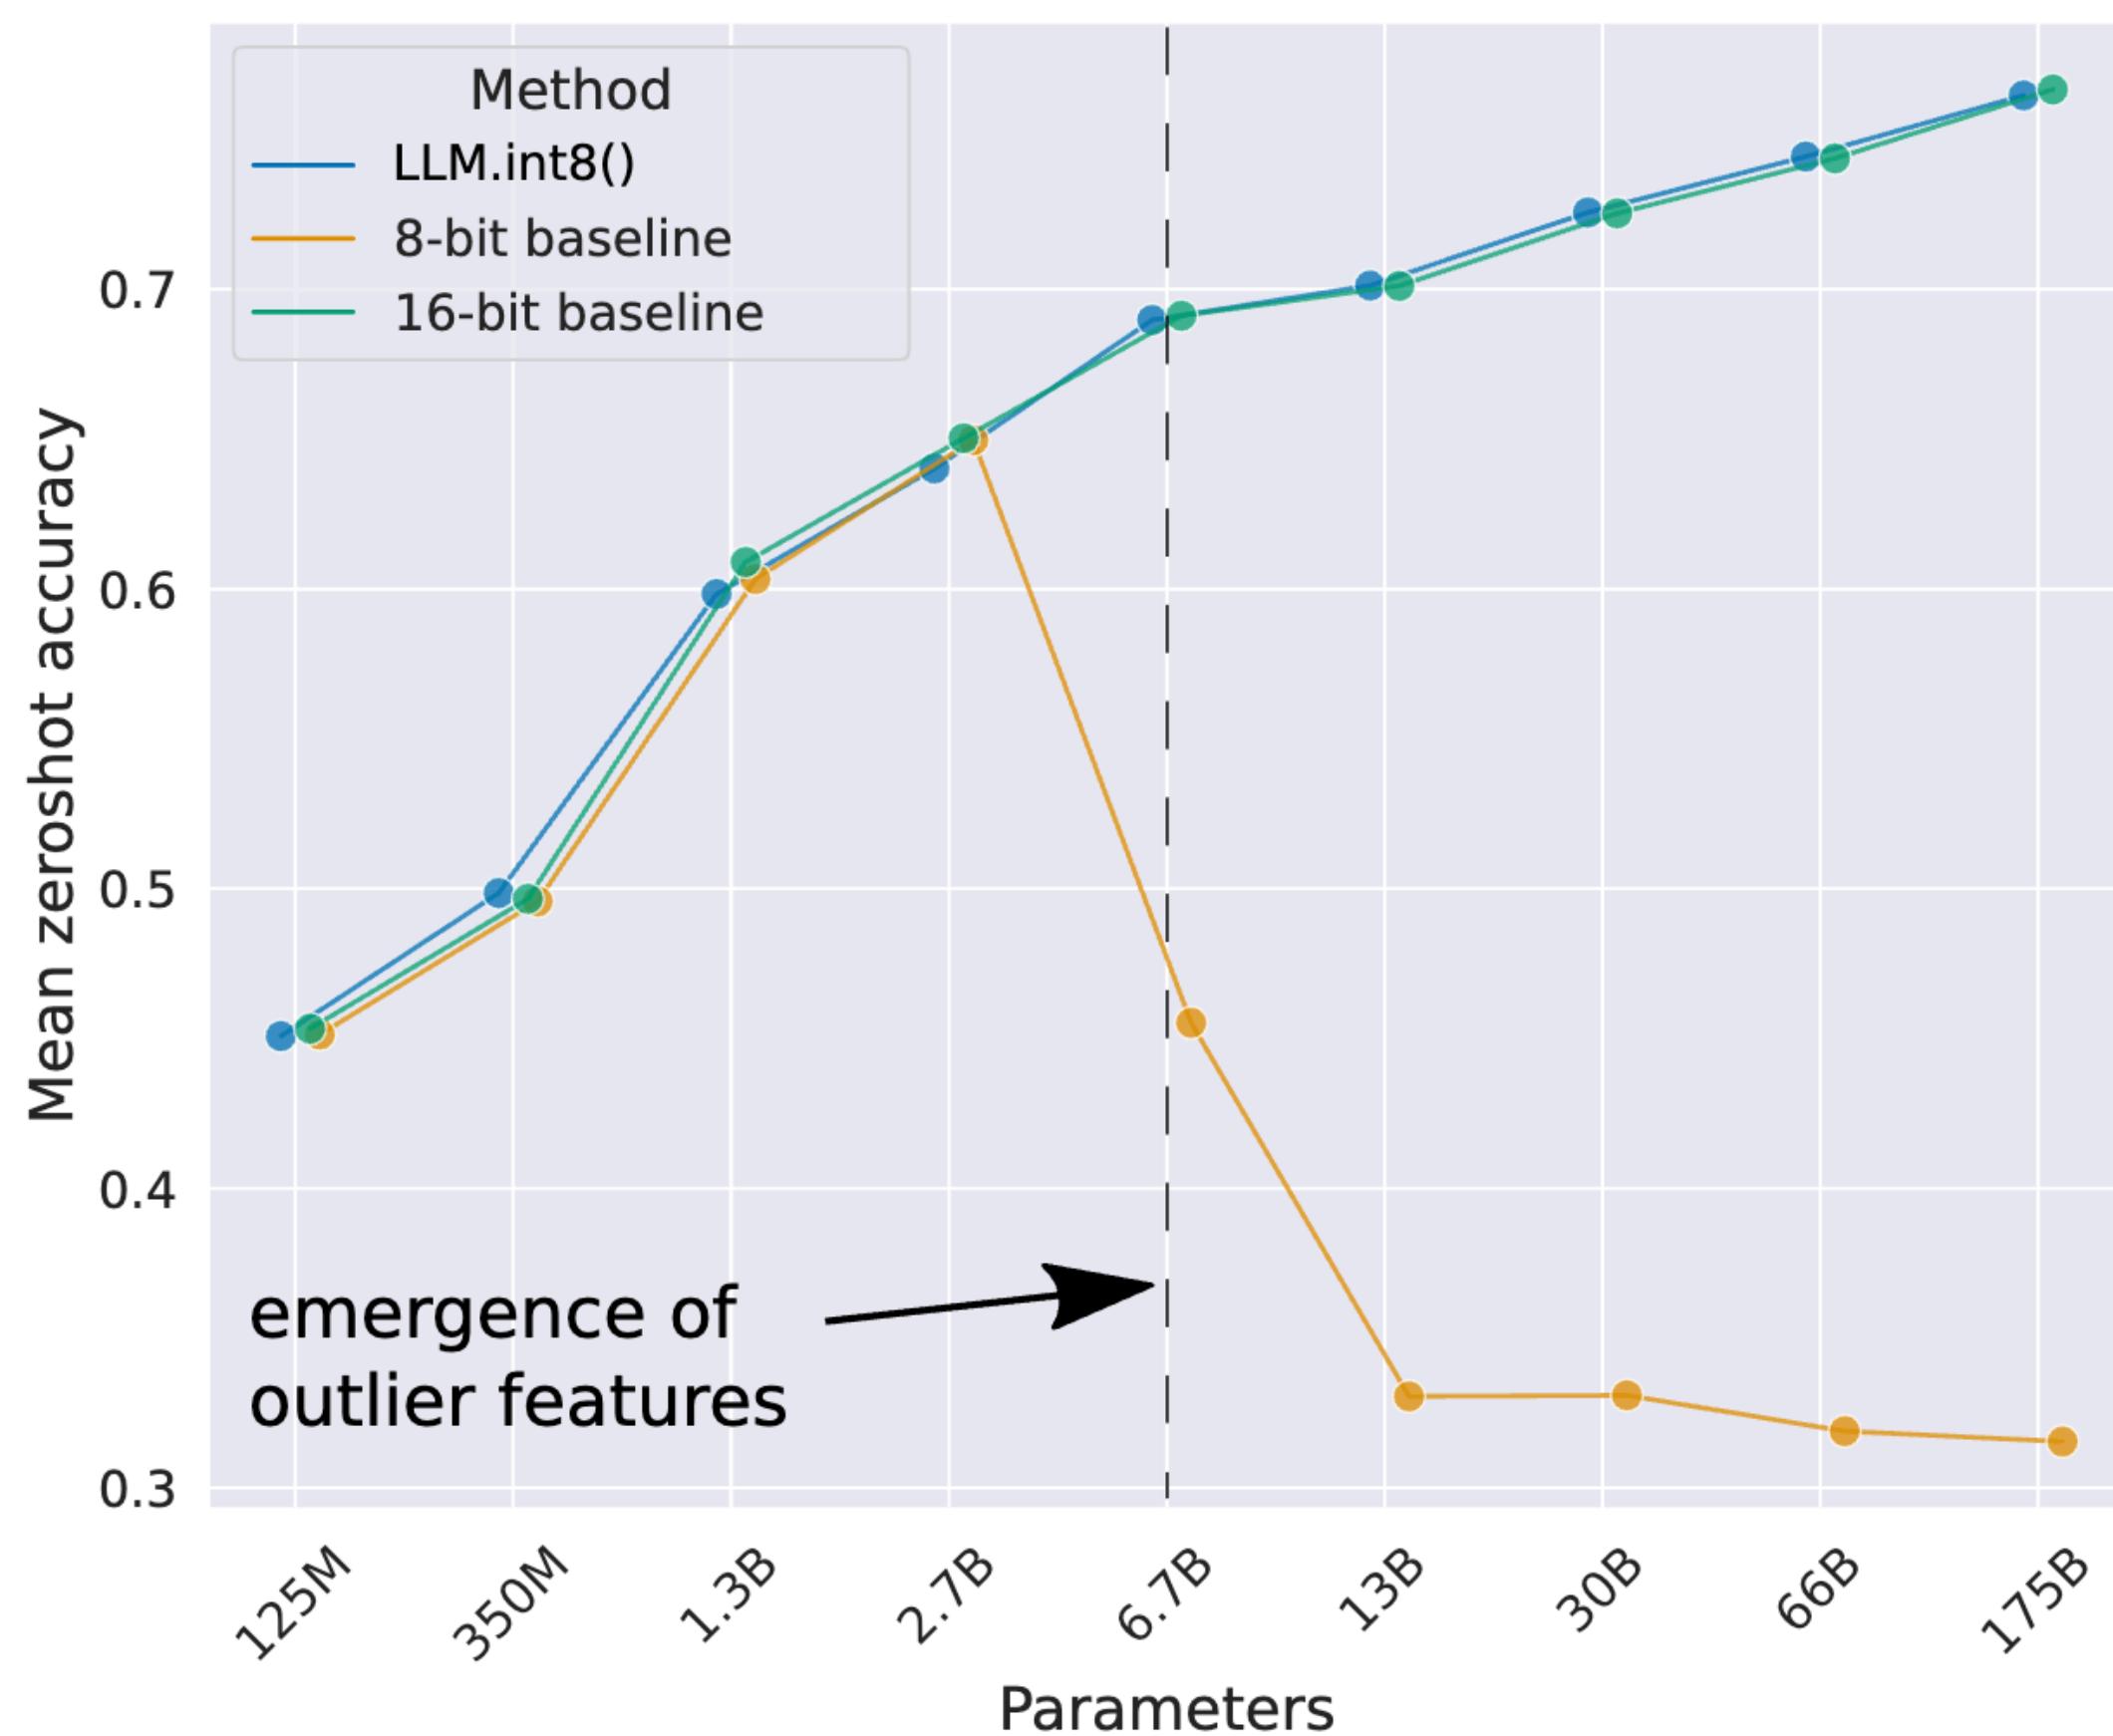


(b)

Figure 2: Full-precision FFN input (top row) and output (bottom row) in 11th layer of BERT. (a) Per-token ranges for first data sequence in the MNLI development set. (b) Visualization of outliers across embedding dimension for the first ten data sequences in the MNLI development set. Dark grey color indicates values that exceed six standard deviations from the mean of the activation tensor.

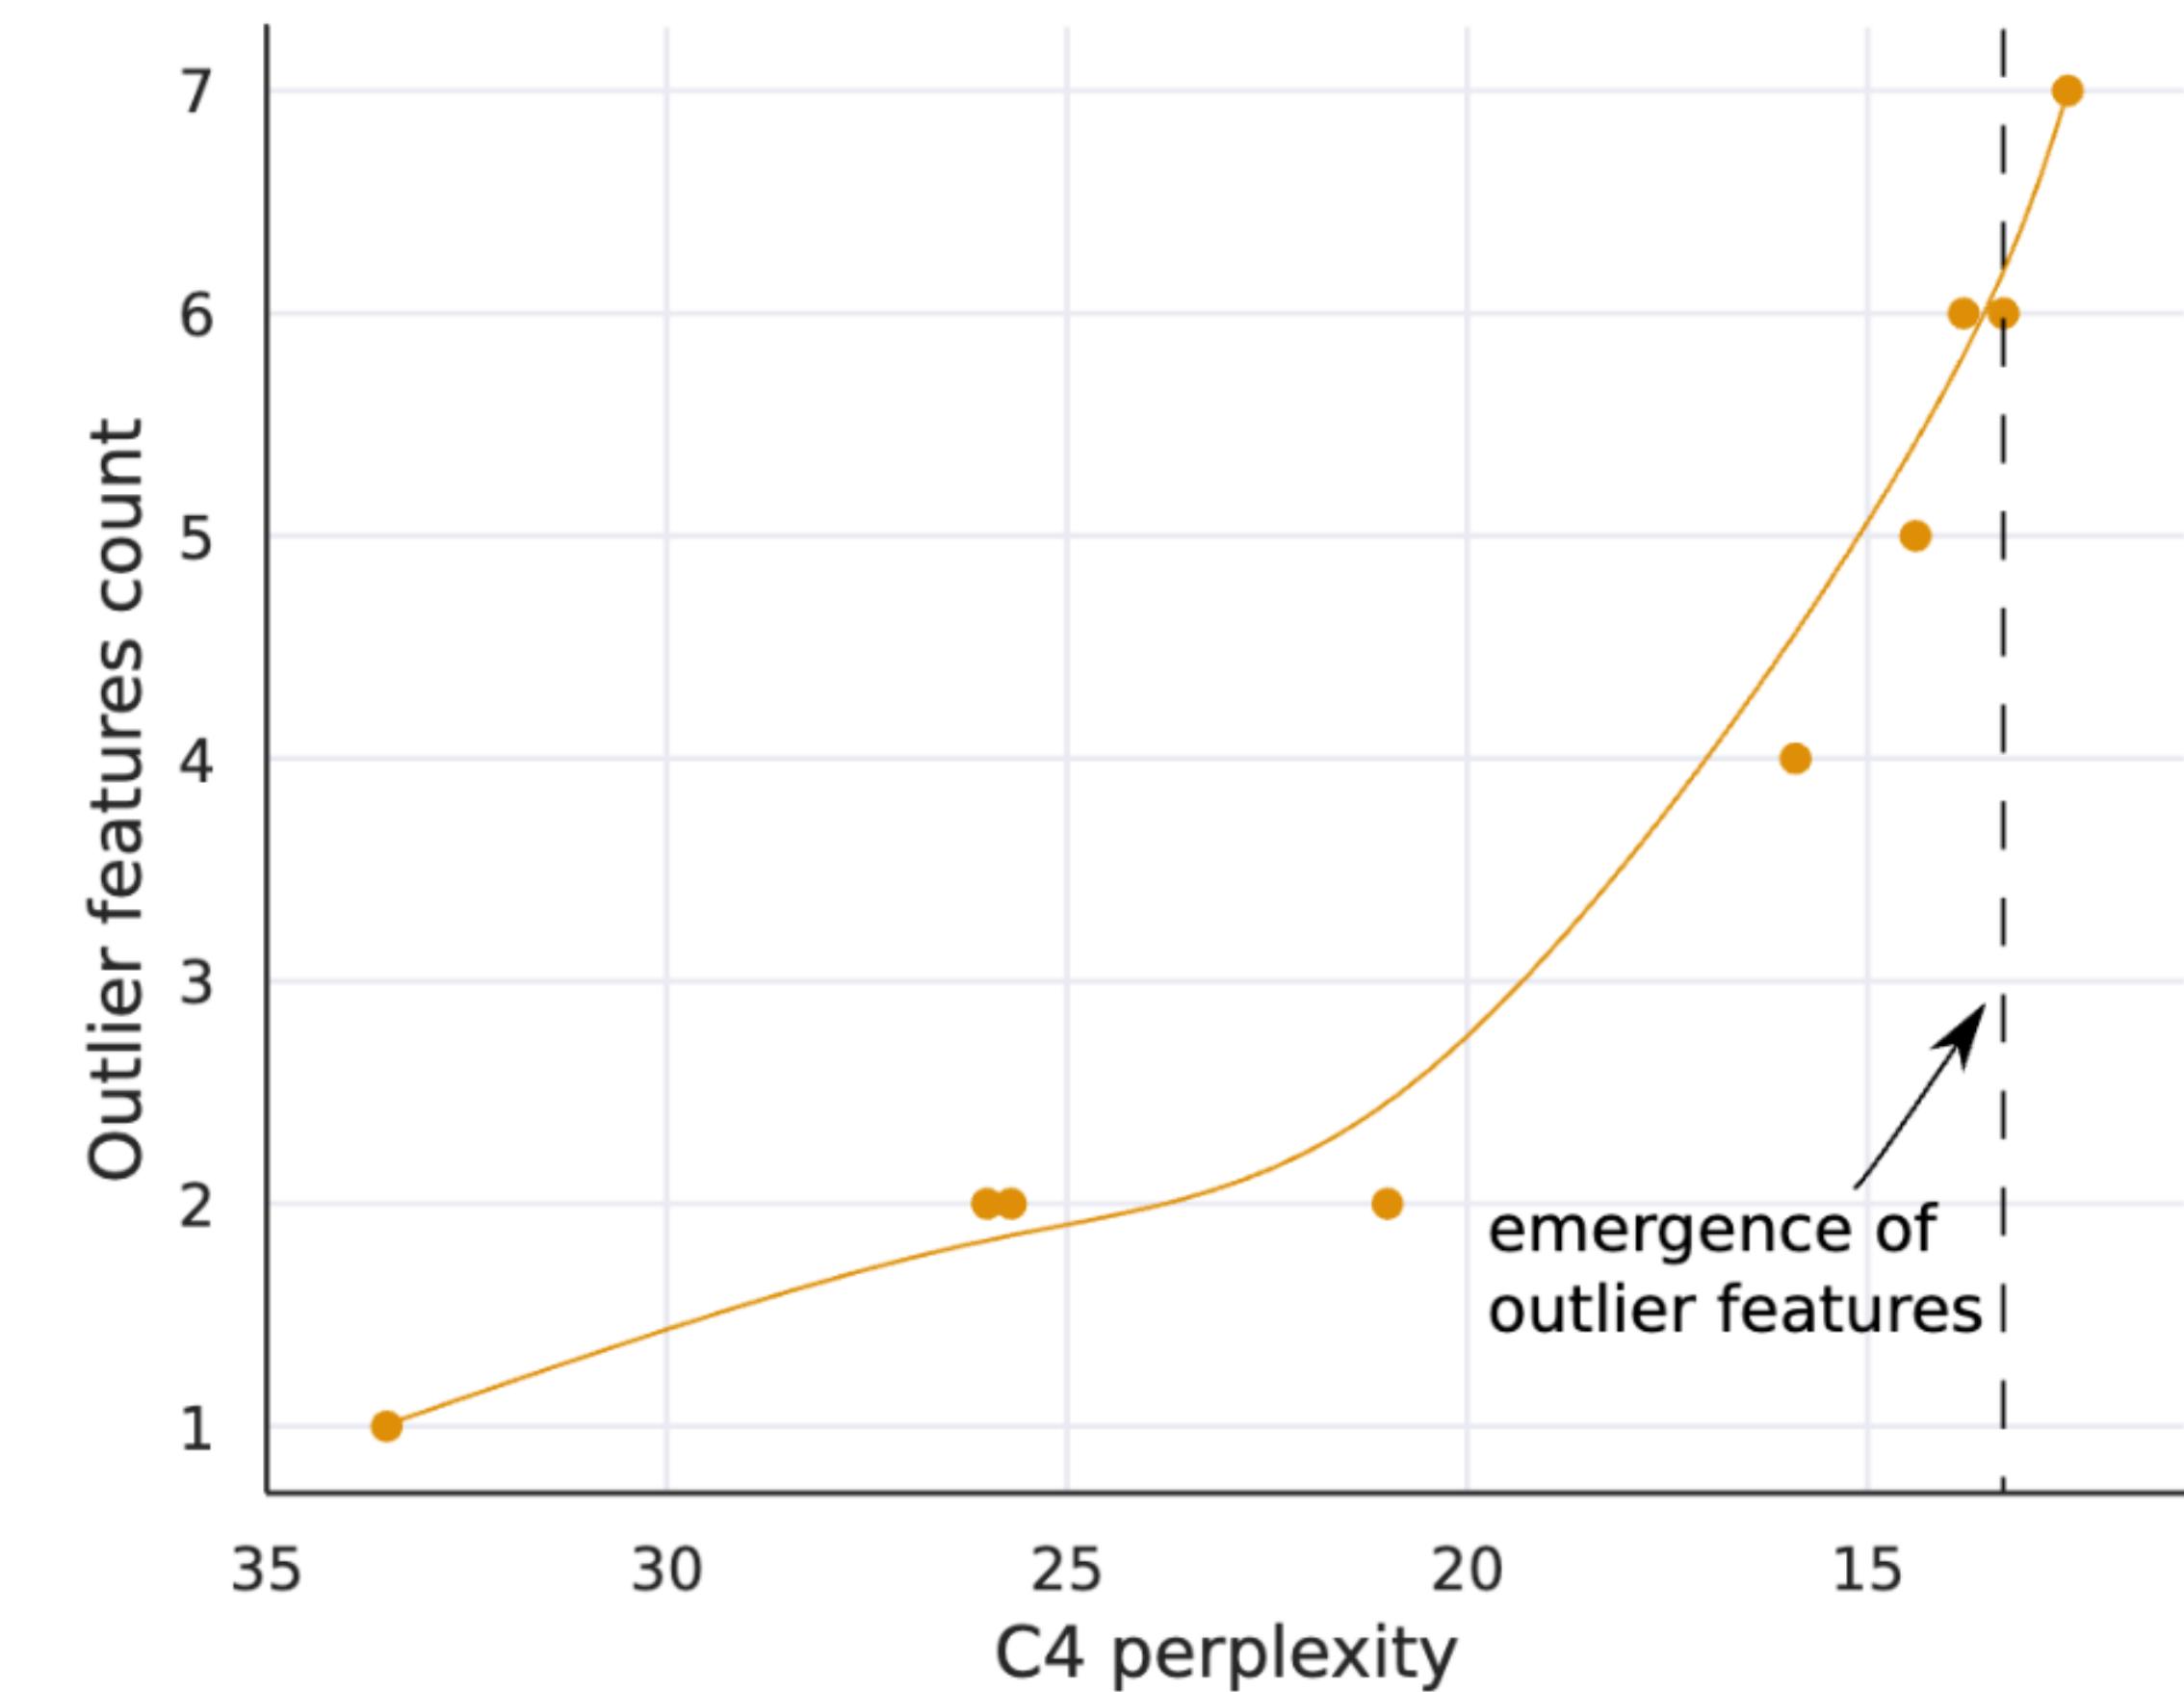
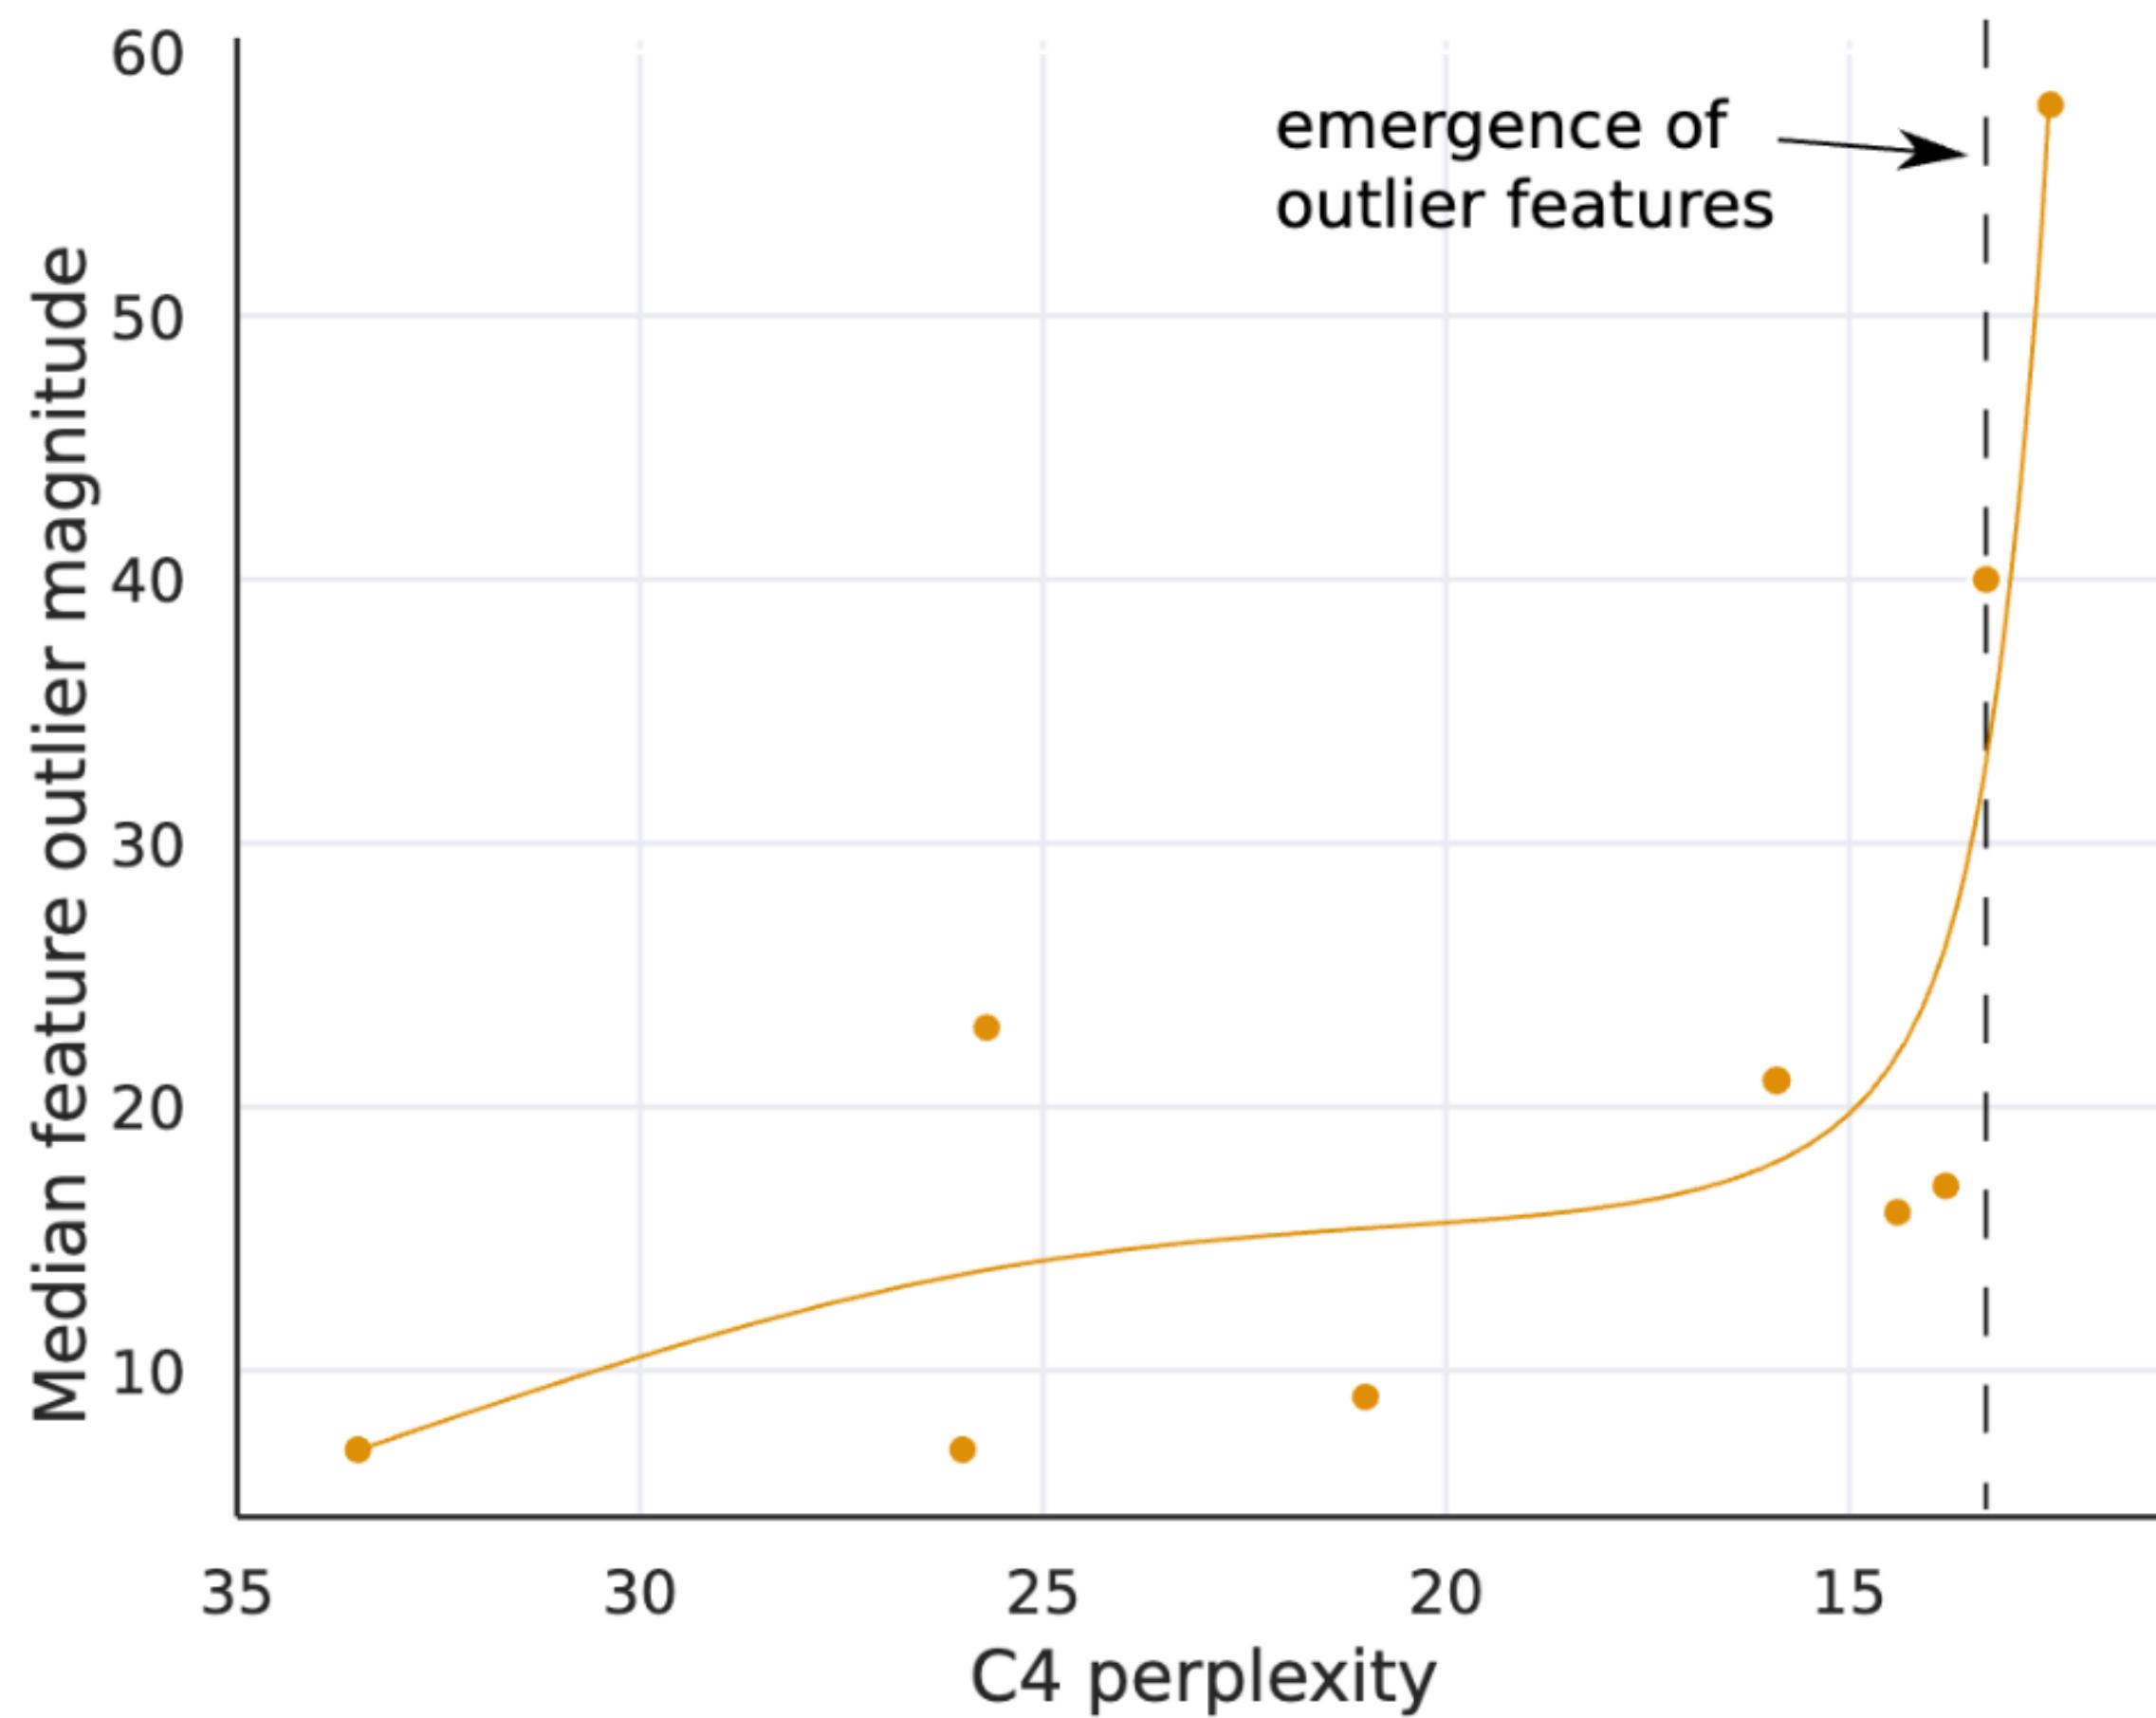
Dettmers et al. (NeurIPS 2022)

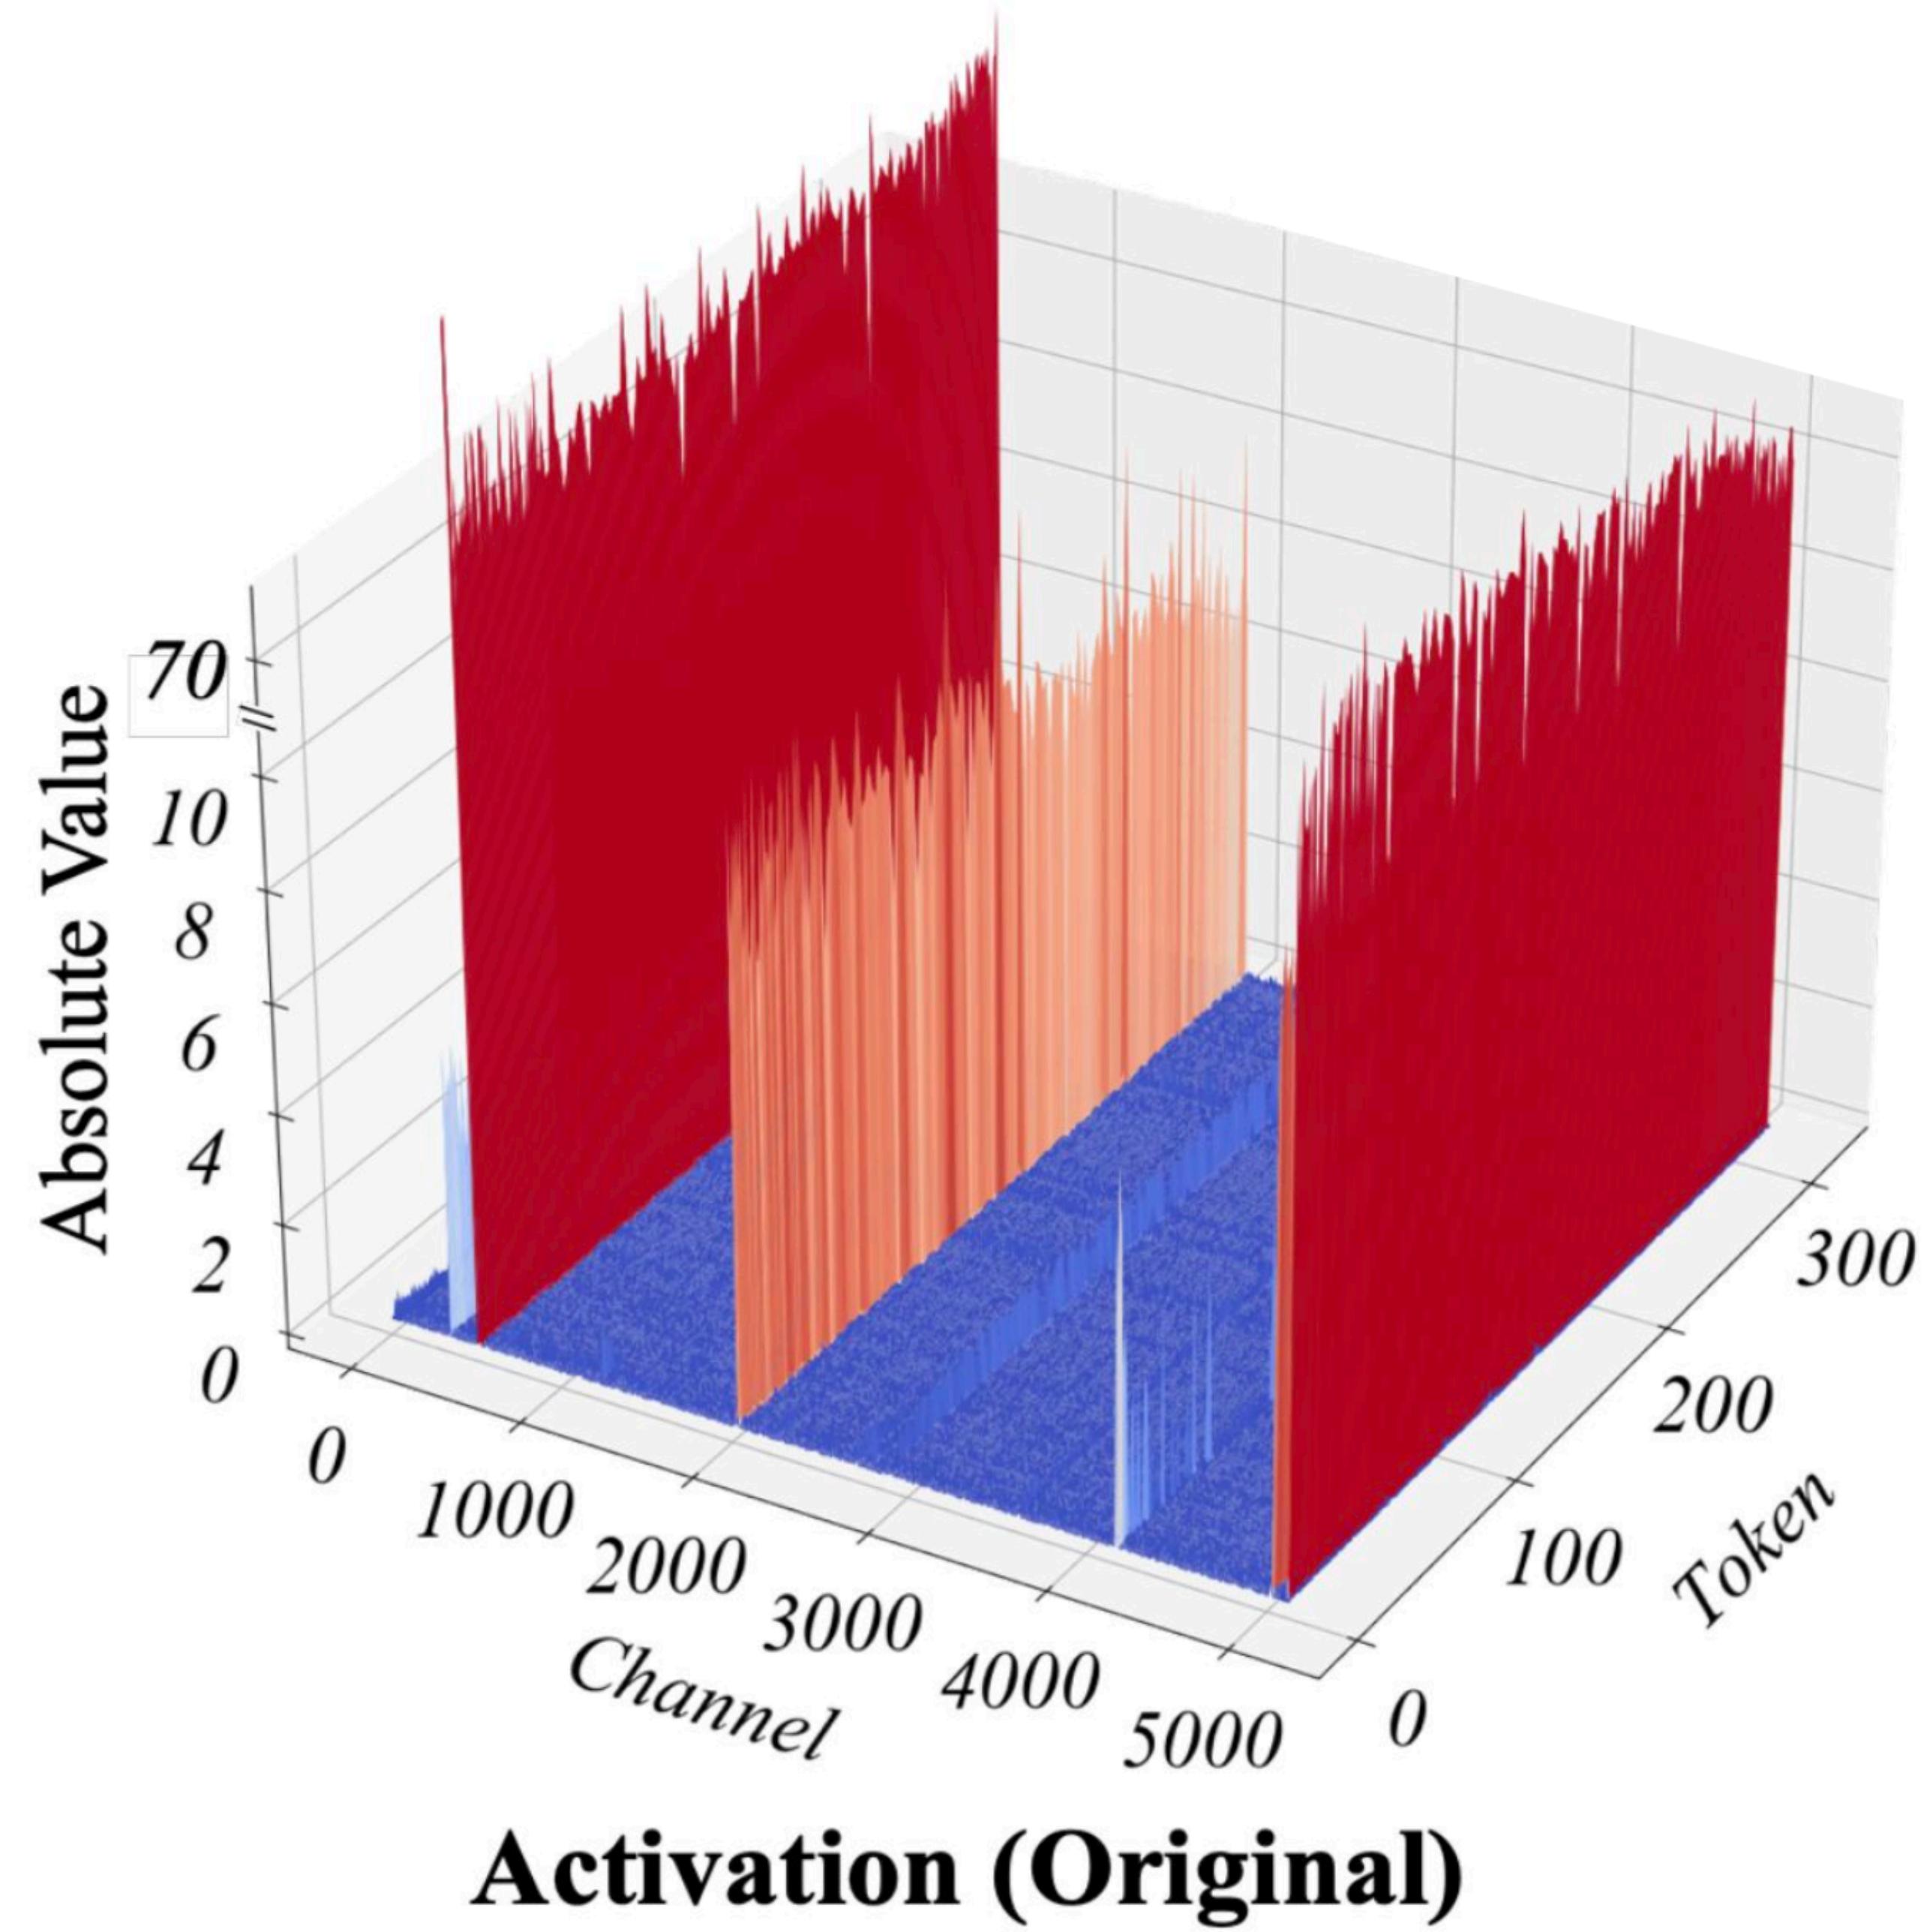
- More significant in **>6.7B models** (attention projection & 1st FFN output)
 - At 6.7B scale, outliers occur in all layers and 75% of all tokens



Dettmers et al. (NeurIPS 2022)

- Larger model => Larger outlier magnitude, greater number of outlier features



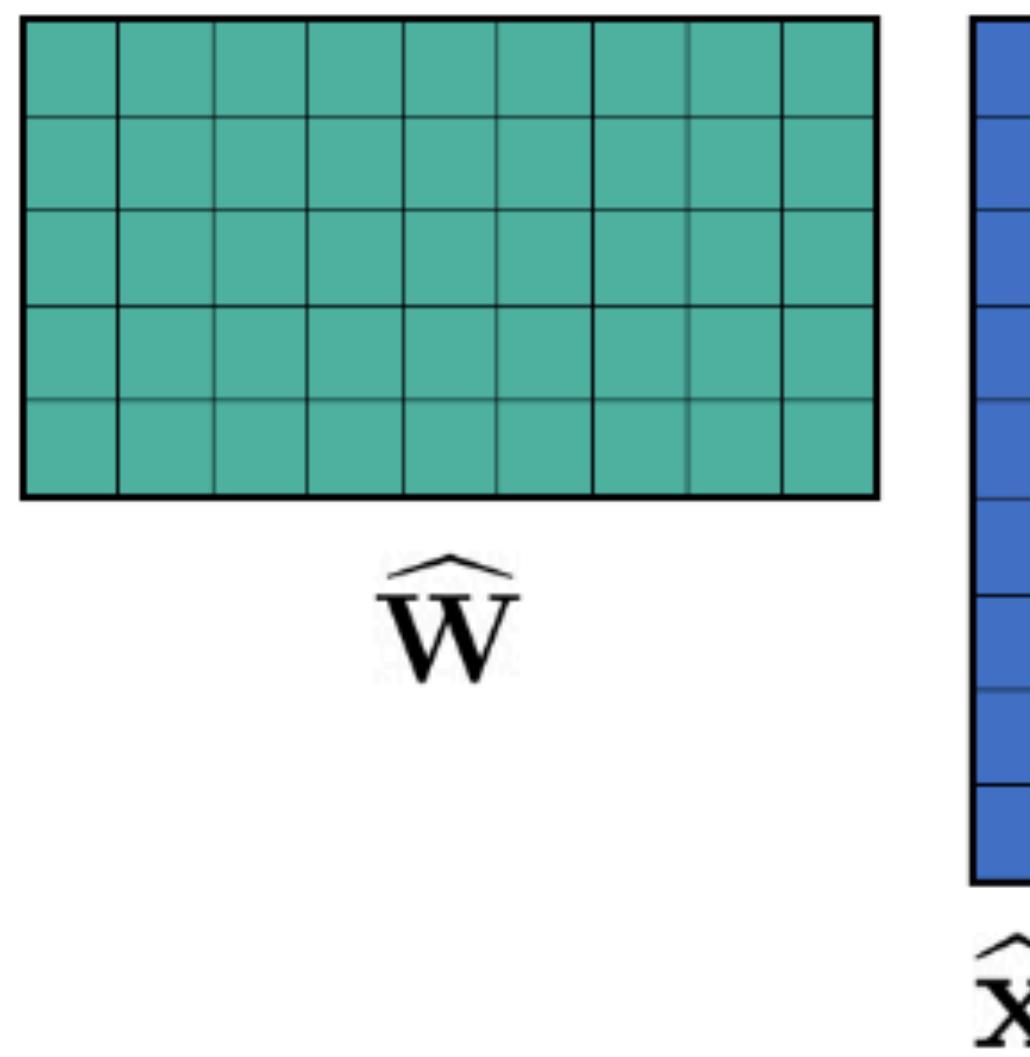


Basic idea

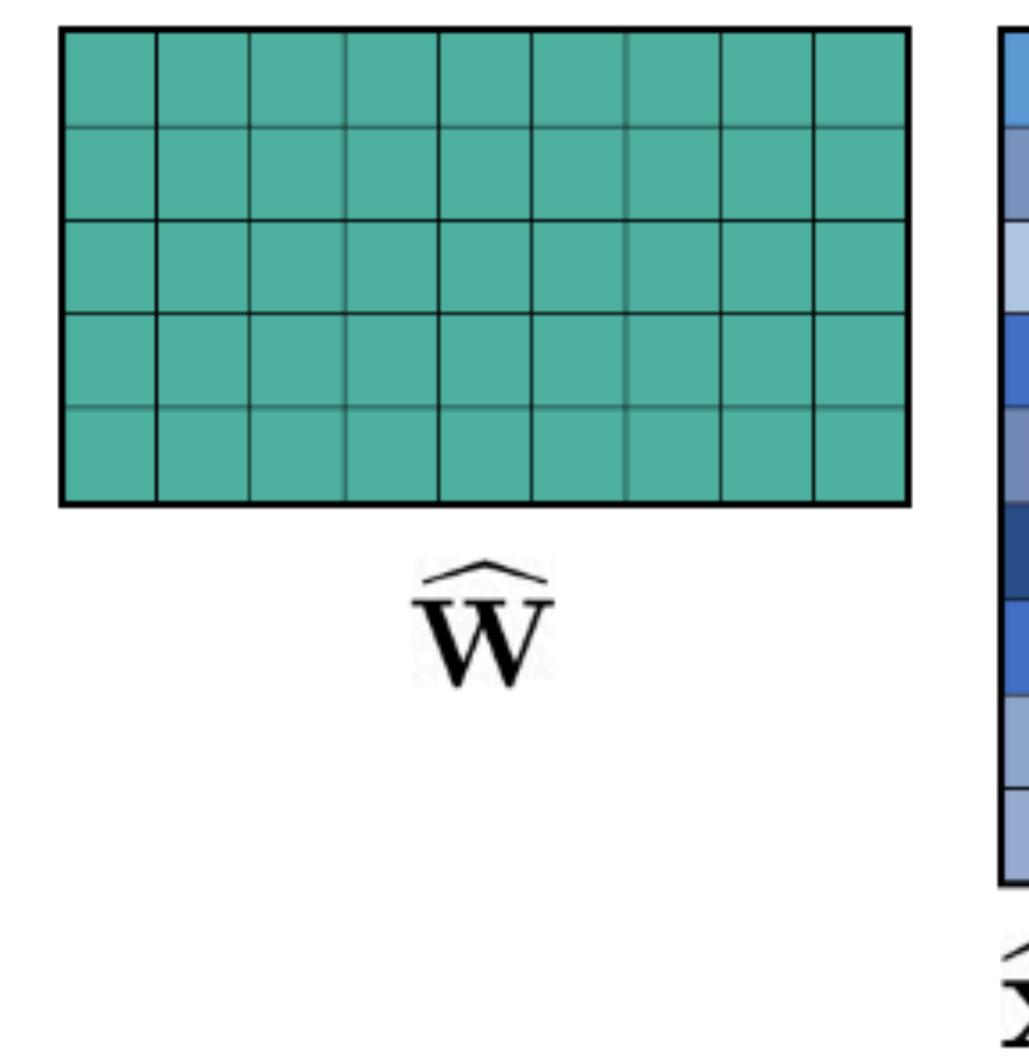
- **Idea.** If certain channels are likelier to have larger input:
 - pruning / quantizing the weights connected to the channel is likelier to hurt the model accuracy more
$$\|WX - \hat{W}X\|^2$$
 - selecting quantization range can be problematic: all non-outliers are likely to be quantized to zero
- **Question.** If we know certain channels are outlier-prone, what can we do?

Activation Quantization

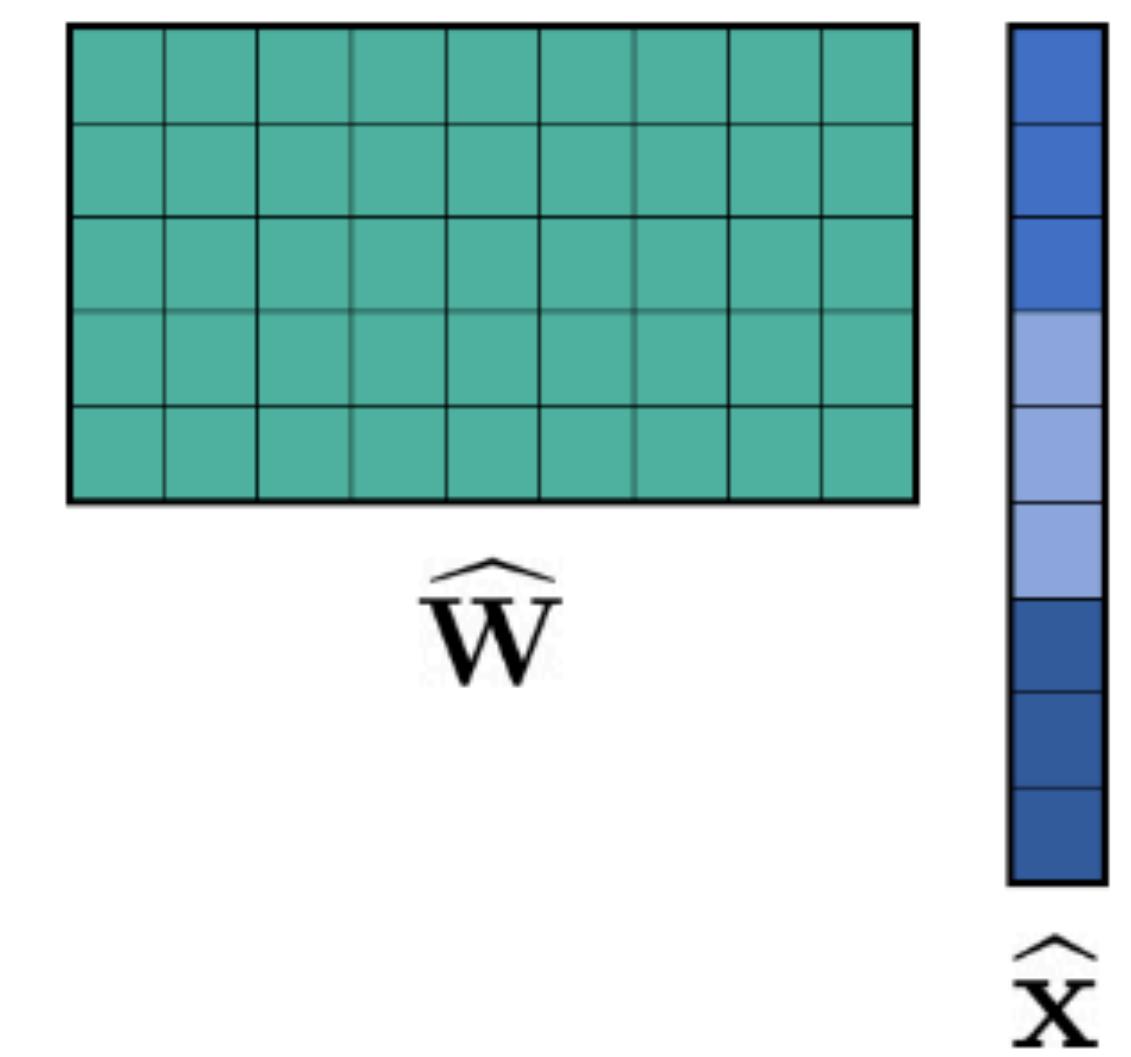
- (1) Divide-and-Quantize
 - Assign **different quantization range** to different groups of channels



(a) Per-tensor



(b) Per-embedding

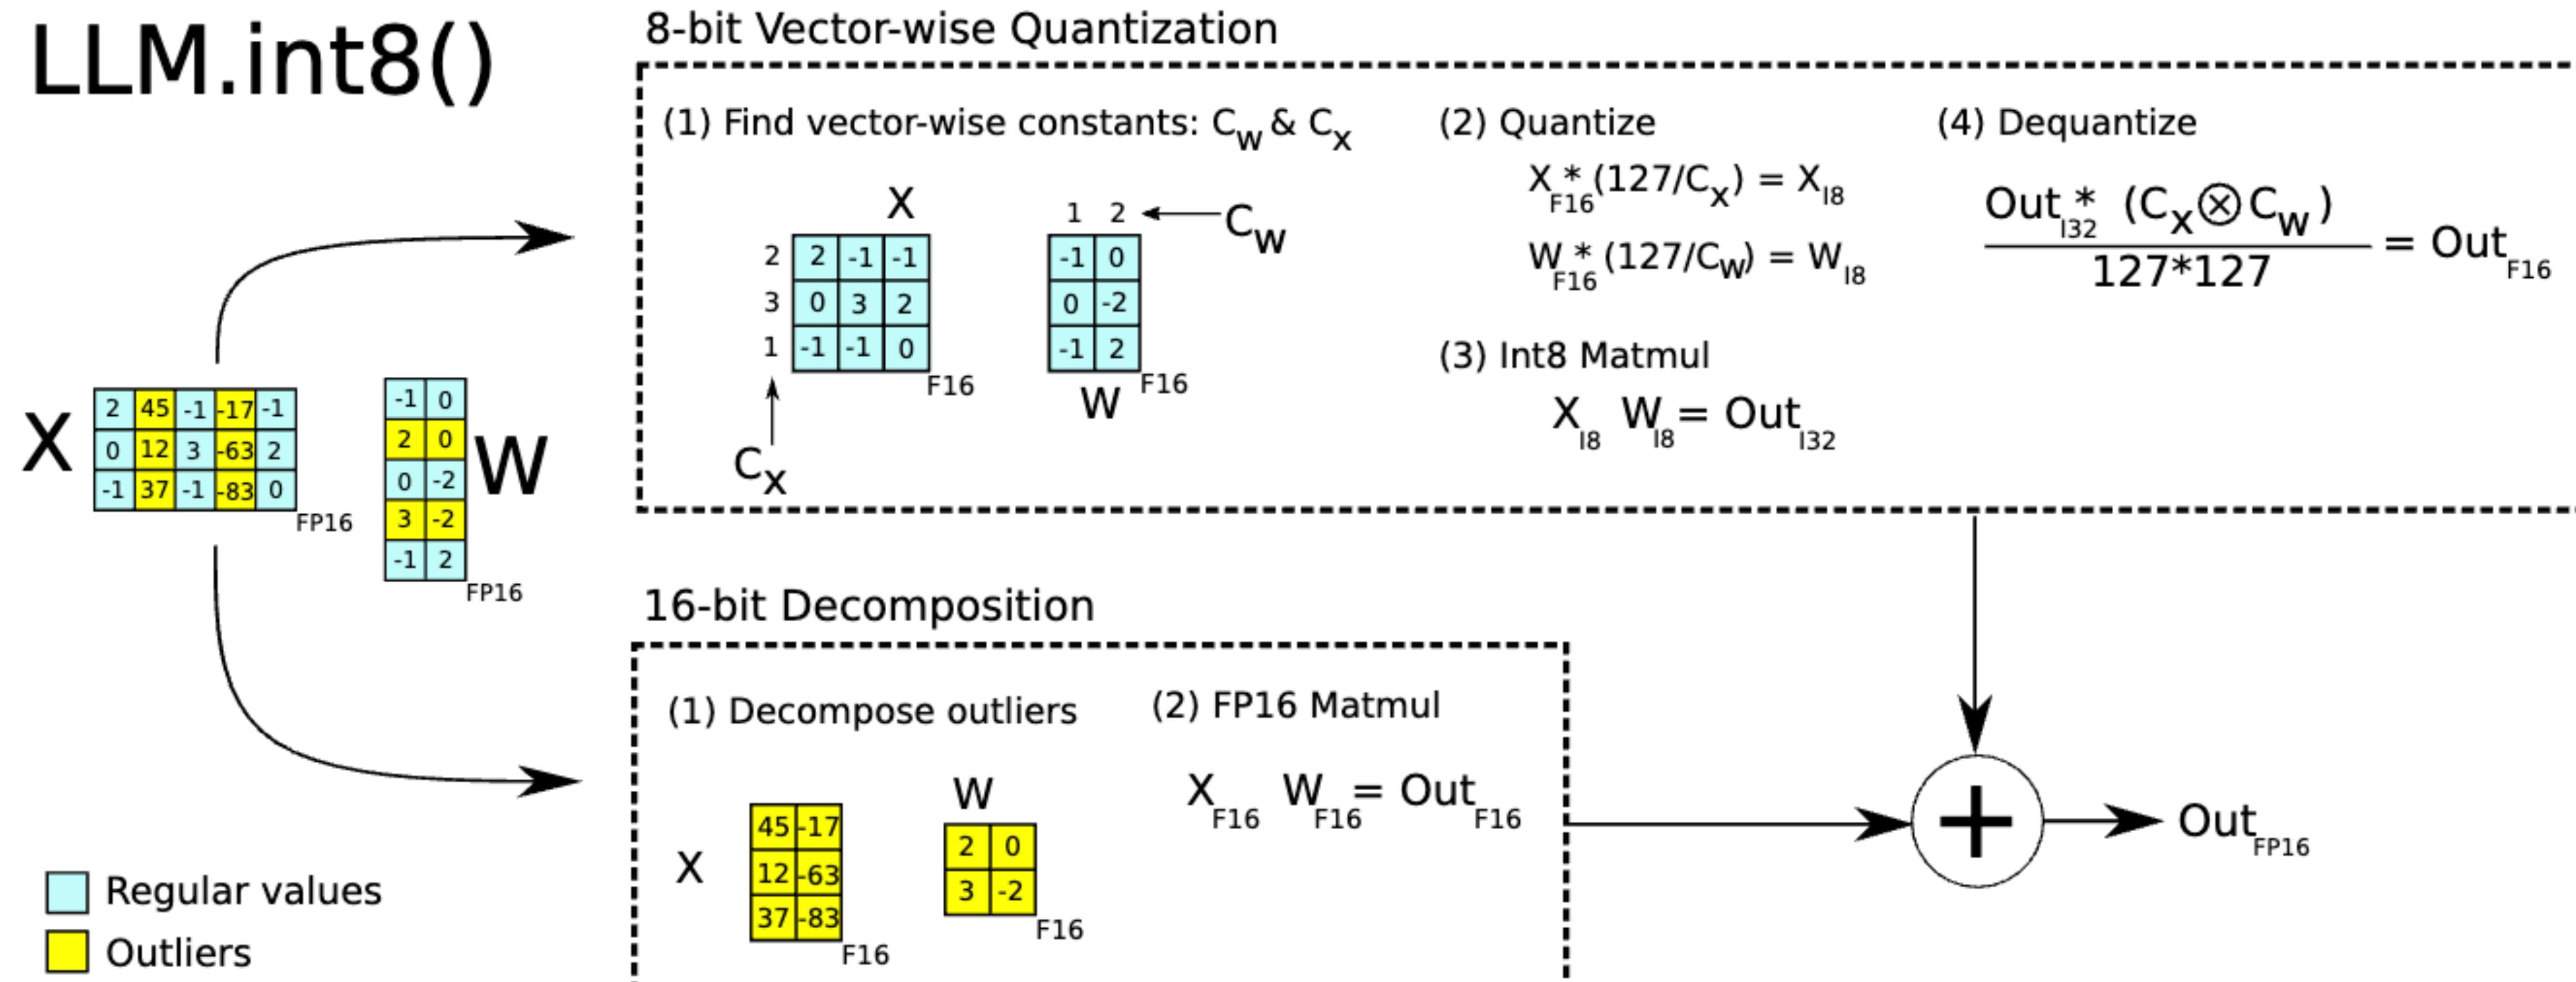


(c) Per-embedding-group

Figure 3: An overview for several choices of activation quantization granularity. The color indicates quantization parameter sharing. In all cases we assume per-tensor weight quantization.

Activation Quantization

- (1) Divide-and-Quantize
 - Assign **different precision** to different channels



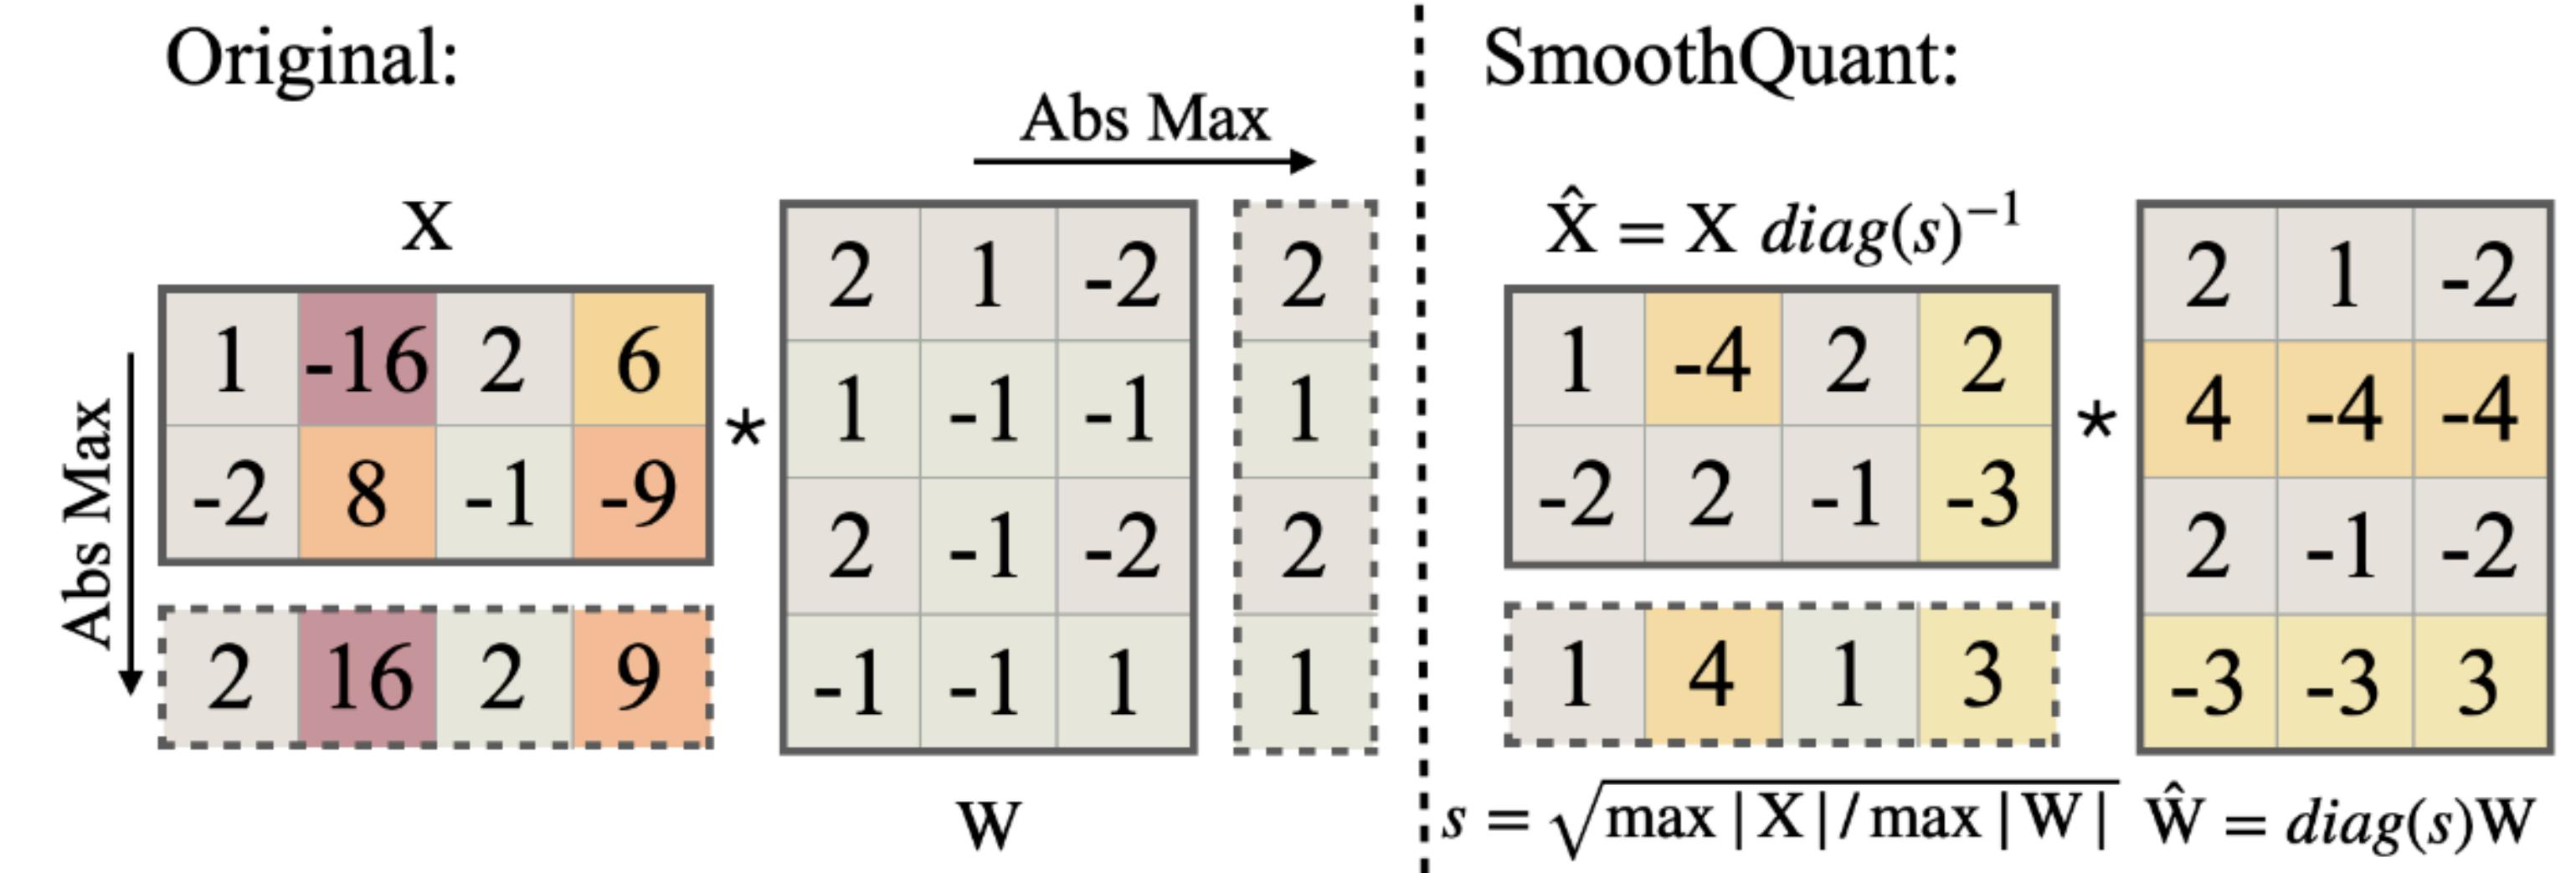
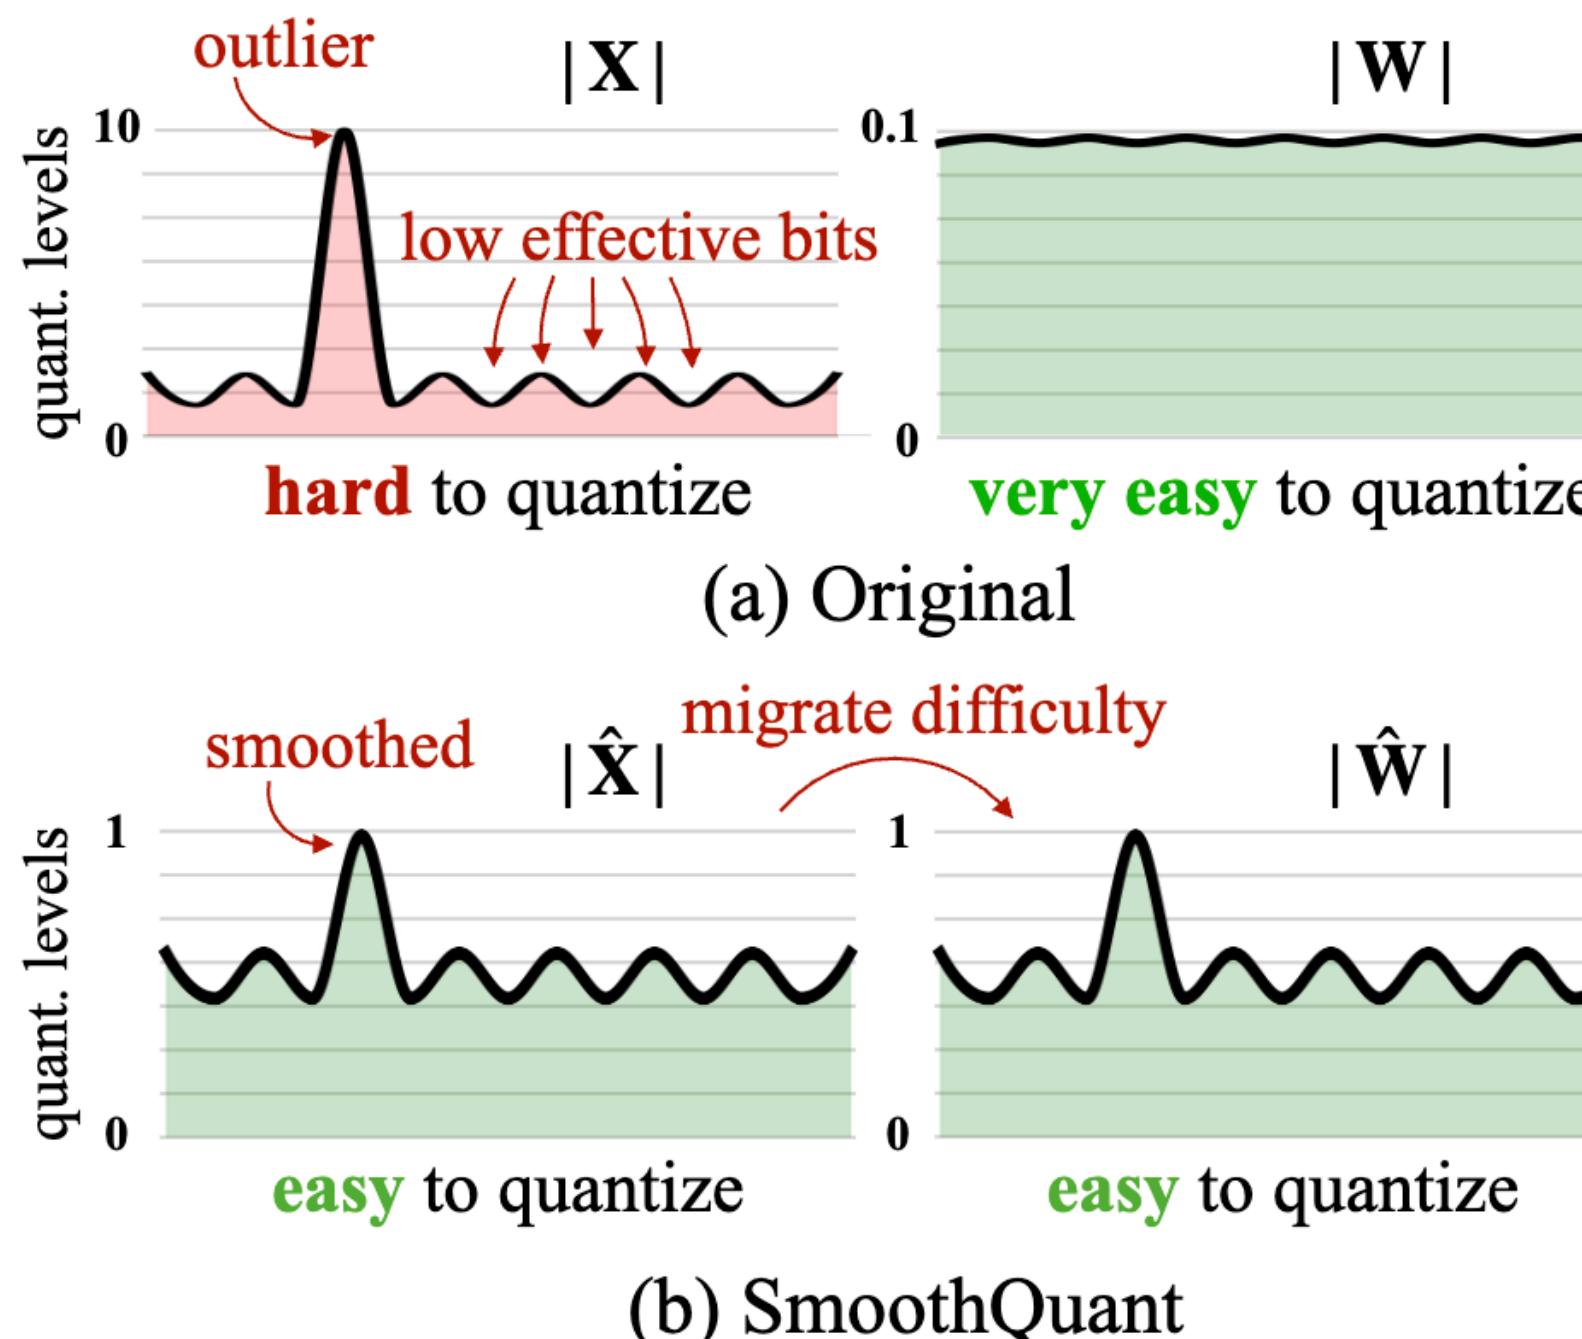
Activation Quantization

- (2) Migrate the difficulty to weight

- Defer scaling factors in the weights

$$\mathbf{WX} = (\mathbf{W}\Lambda^{-1})(\Lambda\mathbf{X})$$

- $\Lambda_i = (\max(|\mathbf{X}_i|)/\max(|\mathbf{W}_i|))^\alpha$



Activation Quantization

- (2) Migrate the difficulty to weight
 - Can be scaled using the scaling factors in the LayerNorm

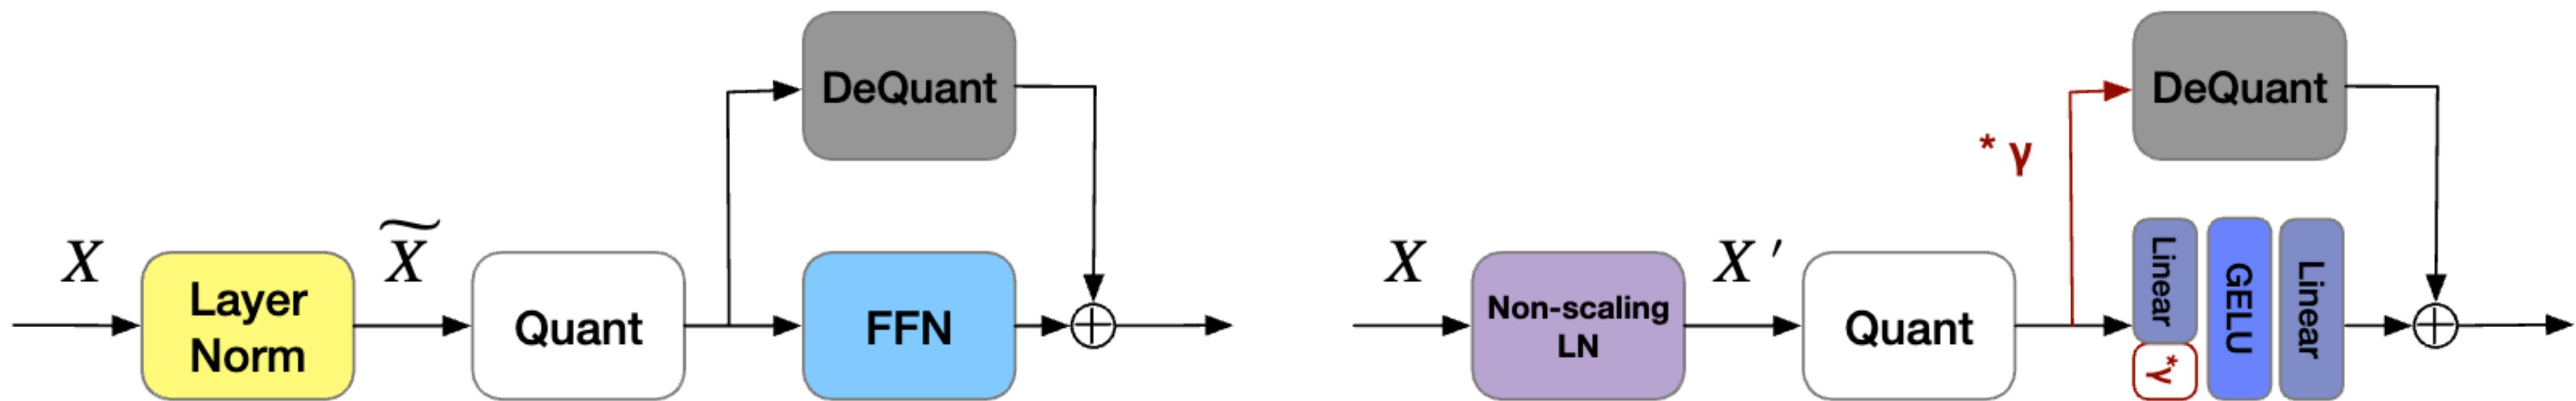


Figure 3: Comparison of the quantization flow before (left) and after (right) Gamma Migration. The original LayerNorm = the Non-scaling LayerNorm γ^* . For other detailed applications such as LayerNorm in encoder-decoder structure, see [Fig. 6](#), [Fig. 7](#).

Weight Quantization

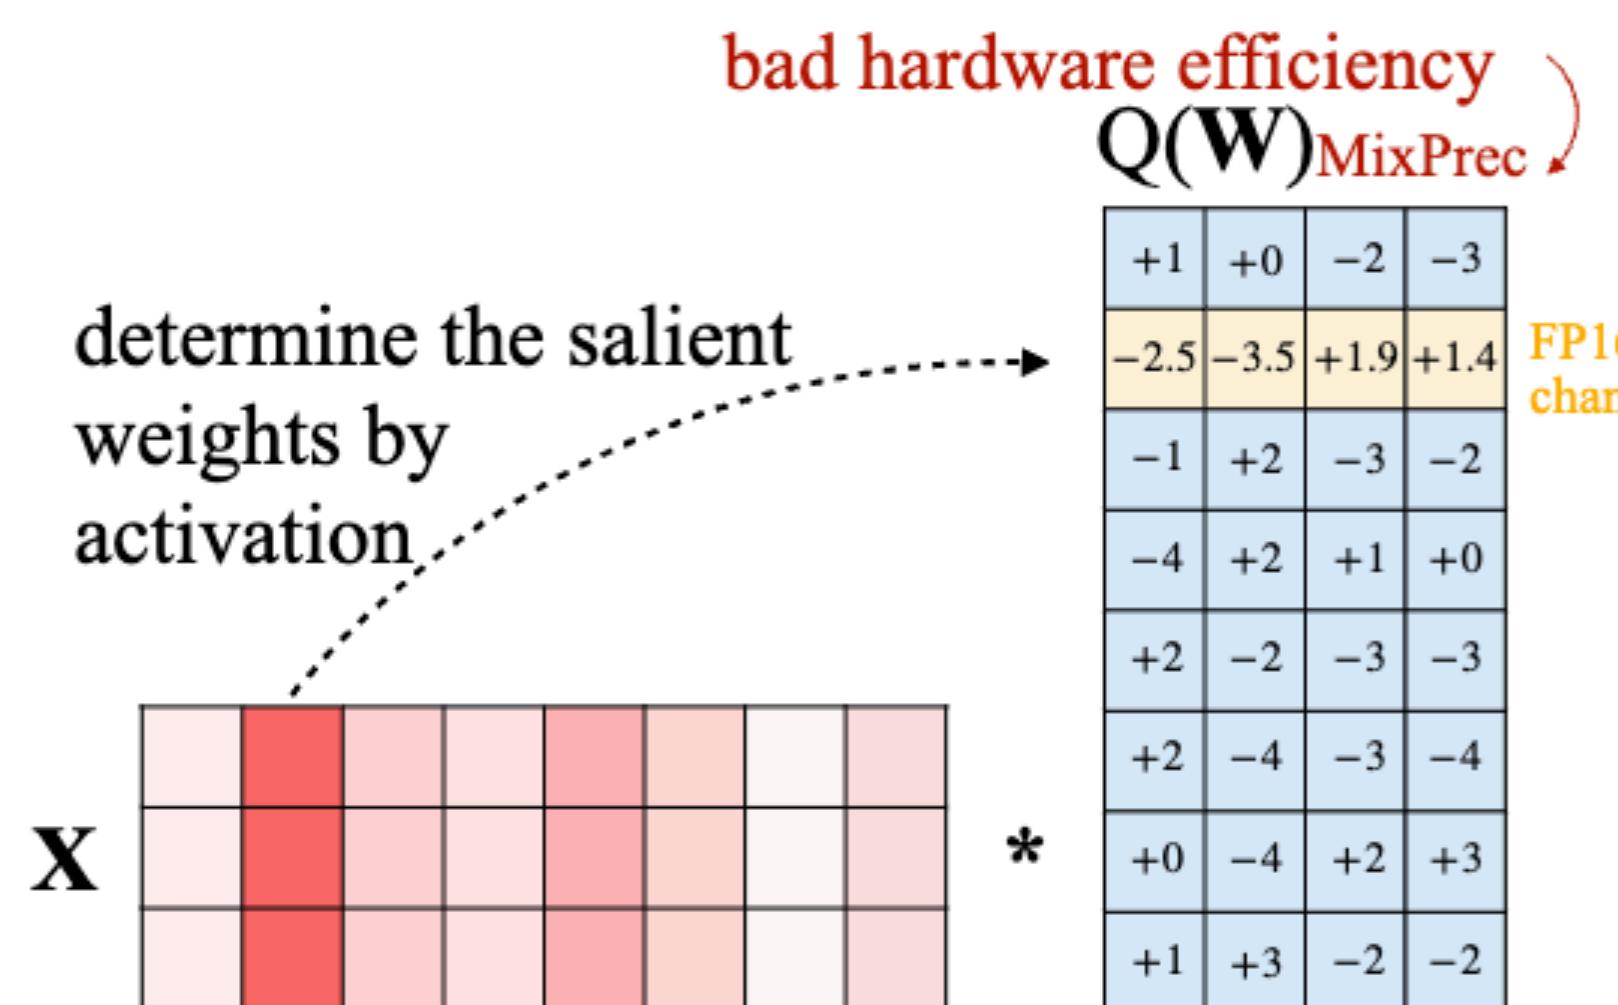
- Activation distribution helps discover important weights

\mathbf{W}_{FP16}			
+1.2	-0.2	-2.4	-3.4
-2.5	-3.5	+1.9	+1.4
-0.9	+1.6	-2.5	-1.9
-3.5	+1.5	+0.5	-0.1
+1.8	-1.6	-3.2	-3.4
+2.4	-3.5	-2.8	-3.9
+0.1	-3.8	+2.4	+3.4
+0.9	+3.3	-1.9	-2.3

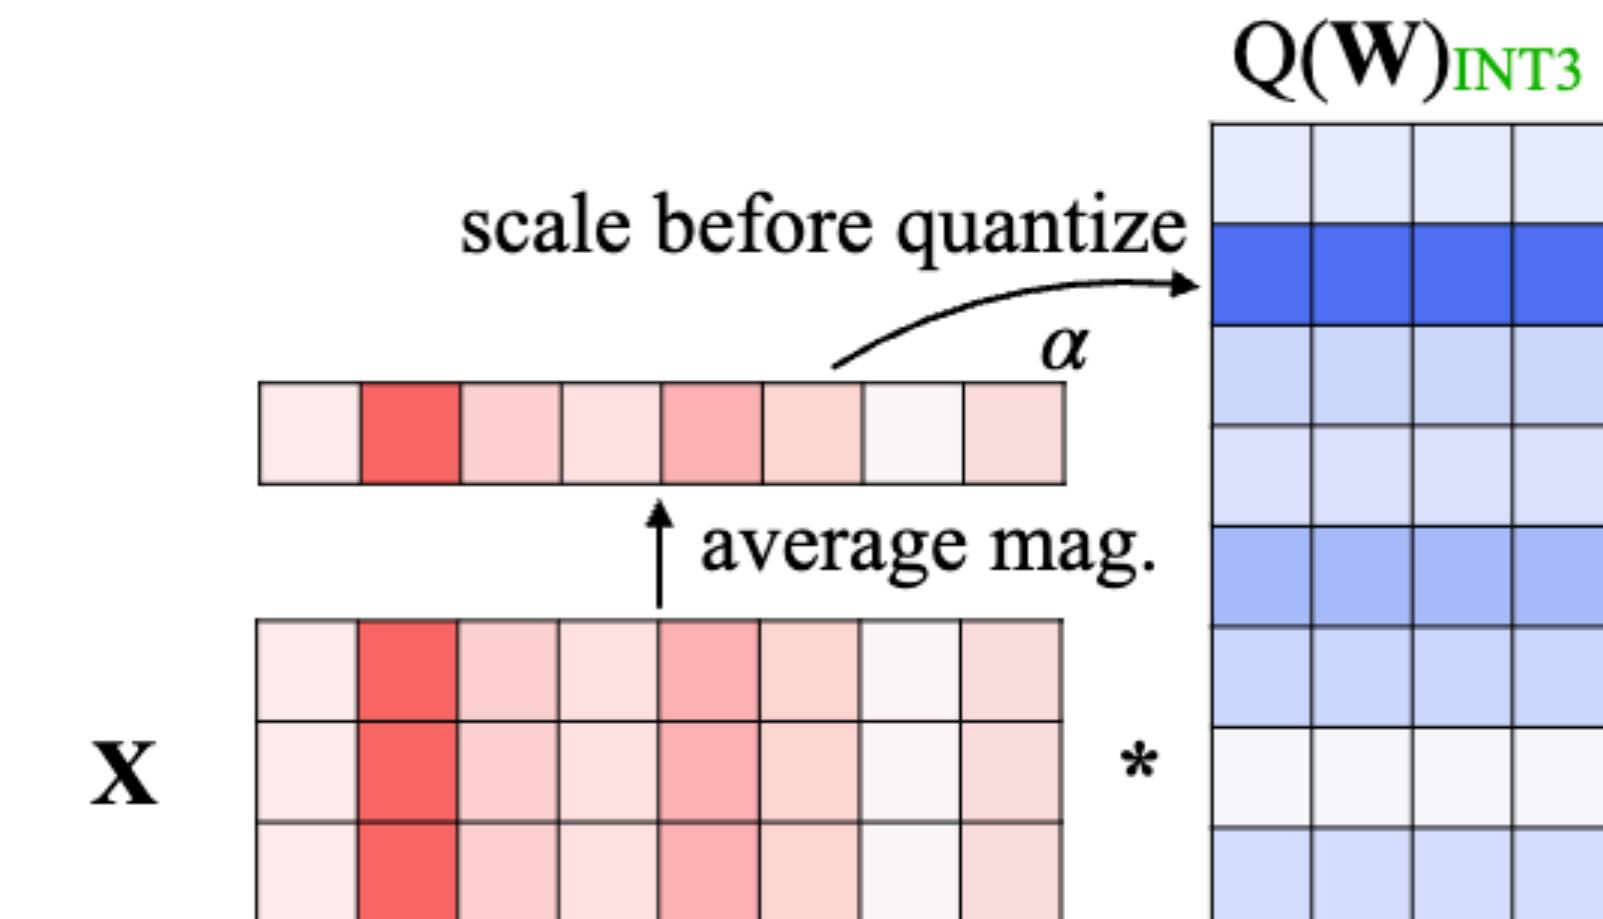
$Q(\mathbf{W})_{\text{INT3}}$			
+1	+0	-2	-3
-3	-4	+2	+1
-1	+2	-3	-2
-4	+2	+1	+0
+2	-2	-3	-3
+2	-4	-3	-4
+0	-4	+2	+3
+1	+3	-2	-2

$\rightarrow Q$

(a) RTN quantization (**PPL 43.2**)



(b) Keep 1% salient weights in FP16 (**PPL 13.0**)



(c) Scale the weights before quantization (**PPL 13.0**)

Weight Pruning

- Activation distribution helps discover important weights

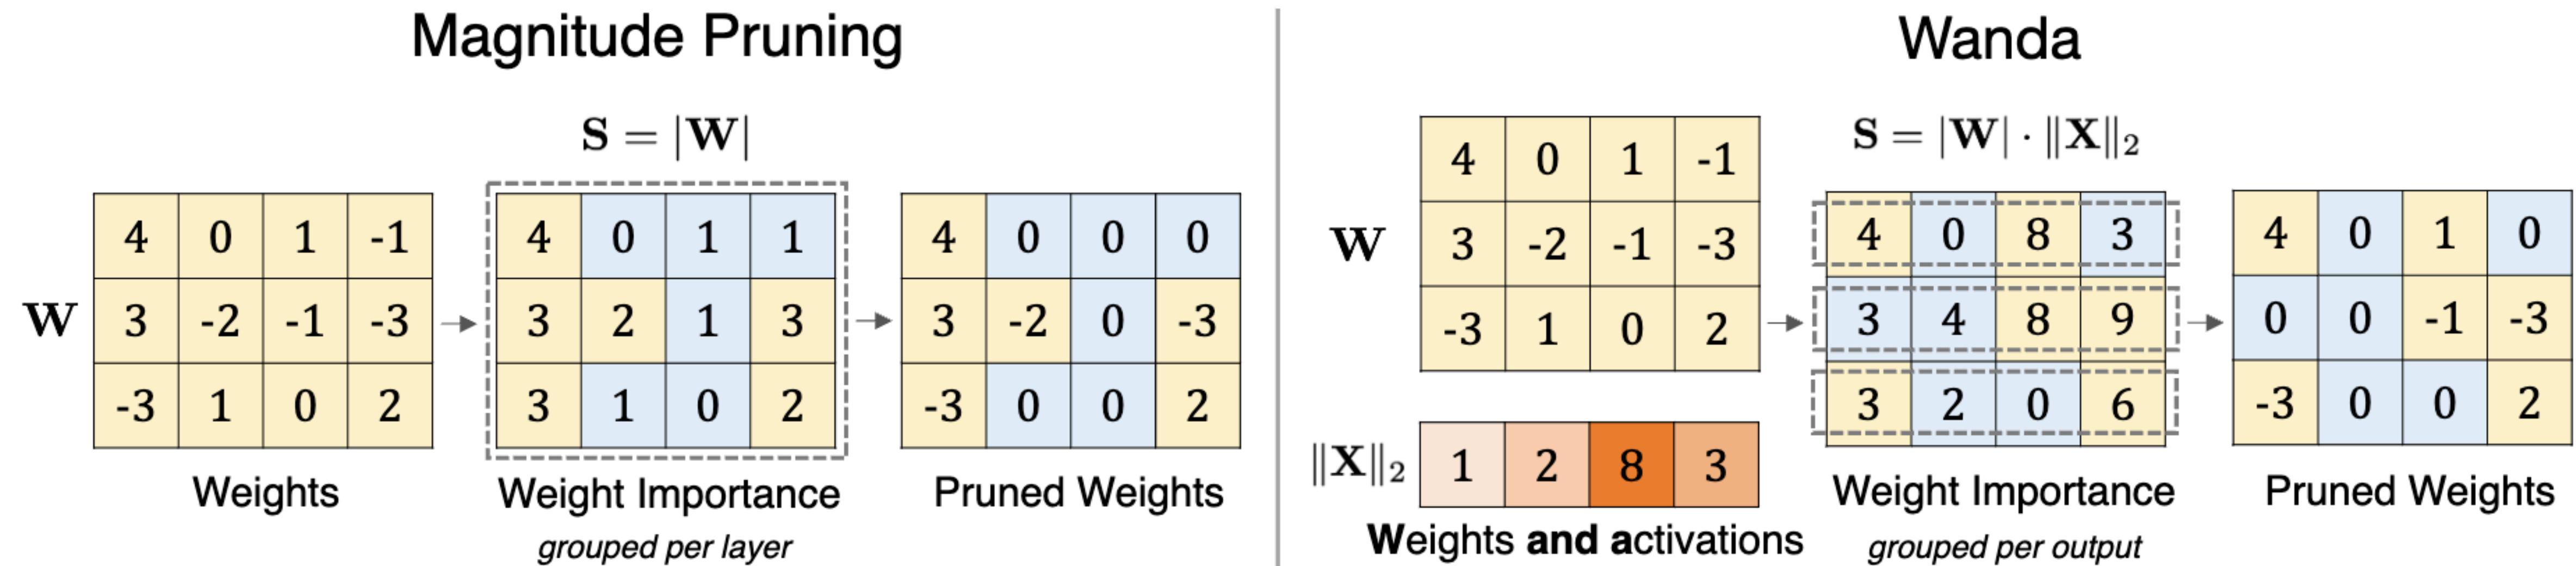
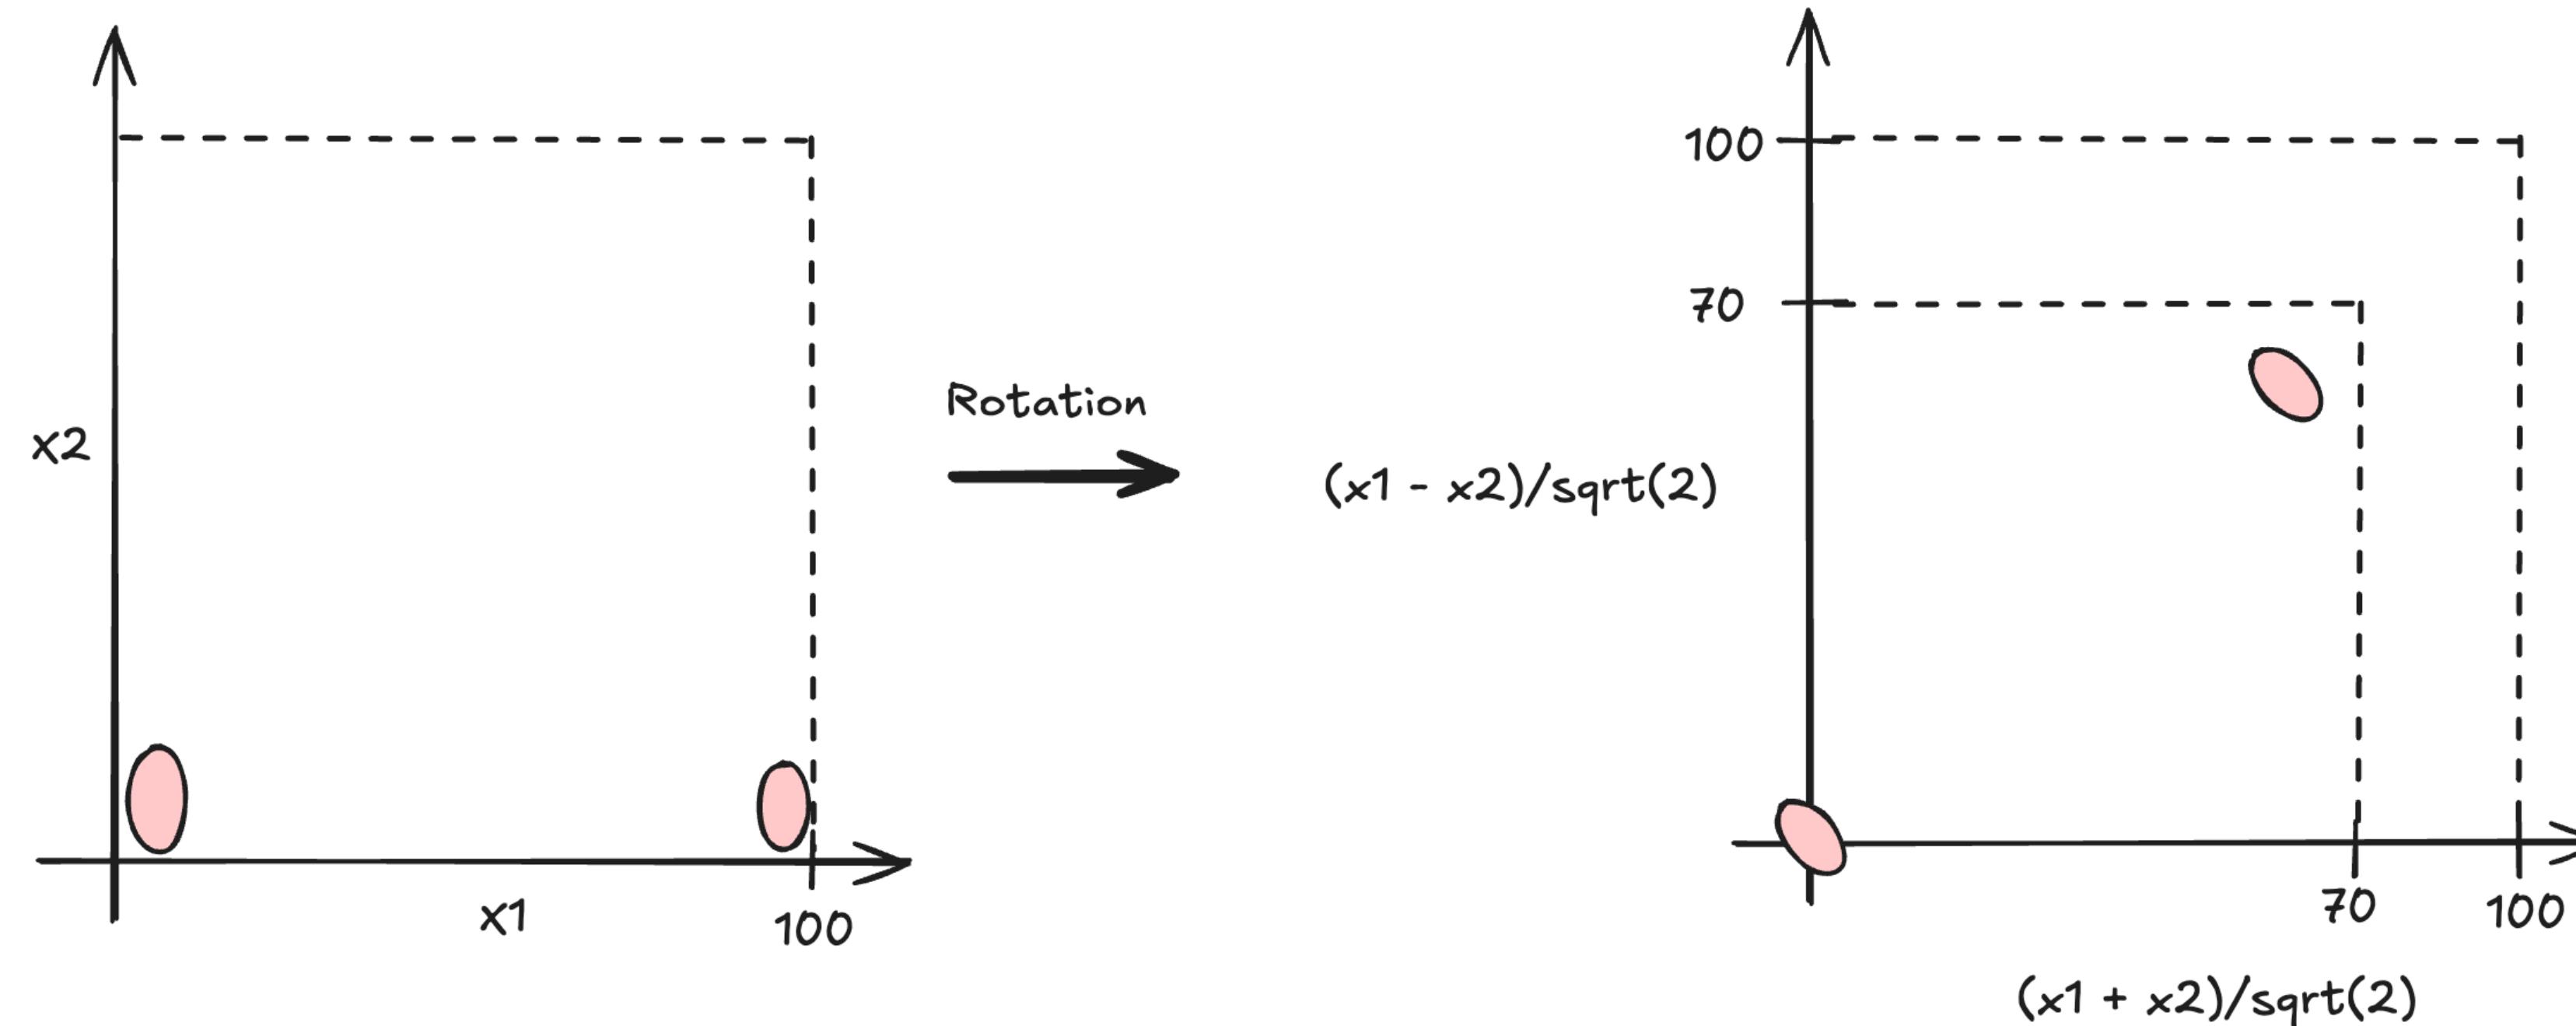


Figure 1: Illustration of our proposed method **Wanda** (Pruning by **Weights and activations**), compared with the magnitude pruning approach. Given a weight matrix \mathbf{W} and input feature activations \mathbf{X} , we compute the weight importance as the product between the weight magnitude and the norm of the corresponding input activations ($|\mathbf{W}| \cdot \|\mathbf{X}\|_2$). Weight importance scores are compared on a *per-output* basis (within each row in \mathbf{W}), rather than globally across the entire matrix.

Rotation

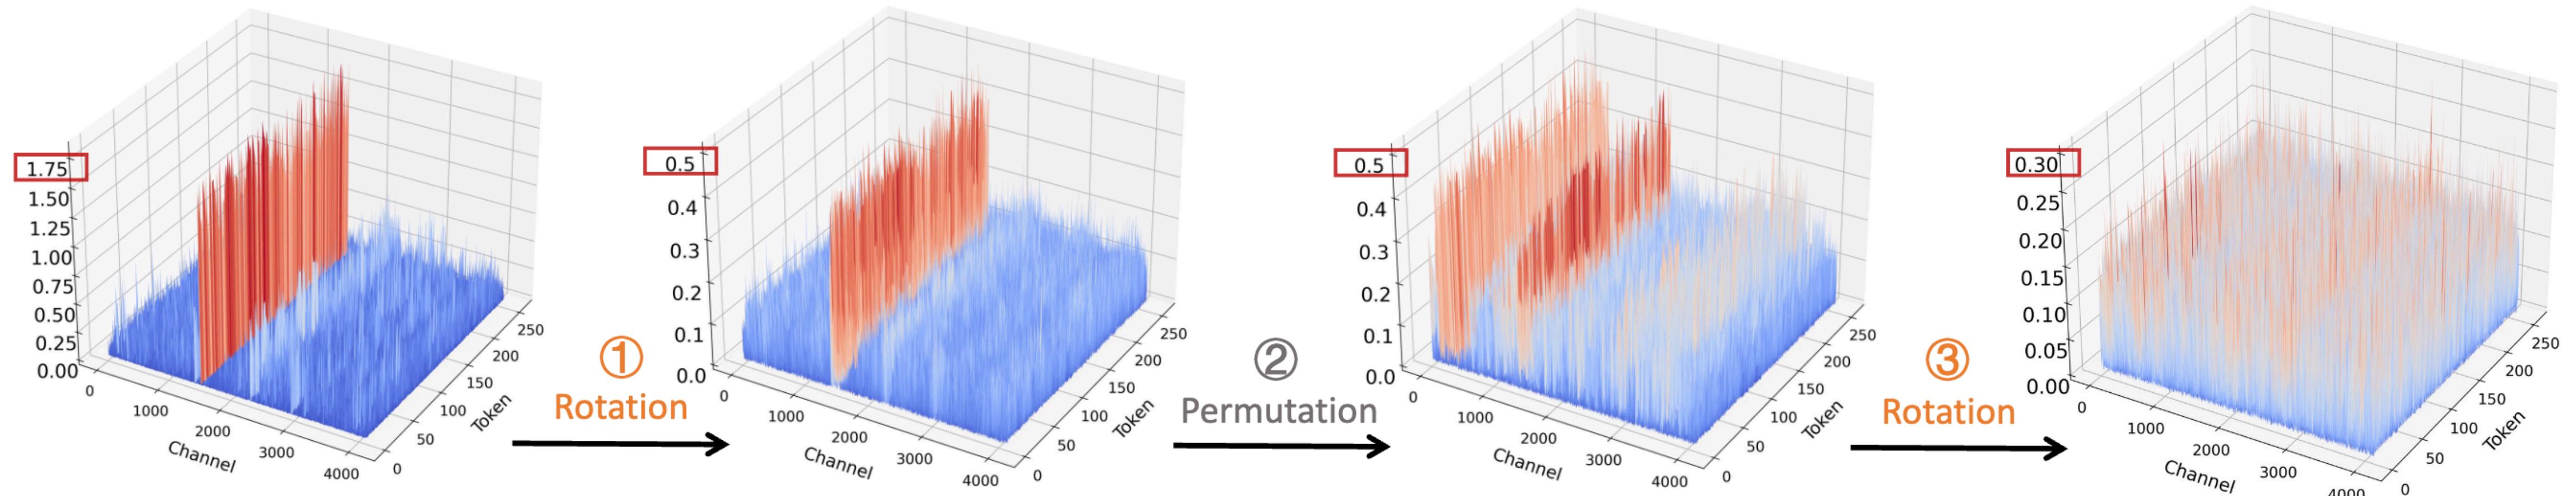
- **Idea.** Mitigate the outliers by multiplying the rotation matrix
 - For some orthogonal rotation matrix \mathbf{R} ($\mathbf{R}^\top \mathbf{R} = \mathbf{R}\mathbf{R}^\top = \mathbf{I}$, $R_{ij} \in \{\pm 1\}$)

$$\mathbf{WX} = (\mathbf{WR}^\top)(\mathbf{RX})$$



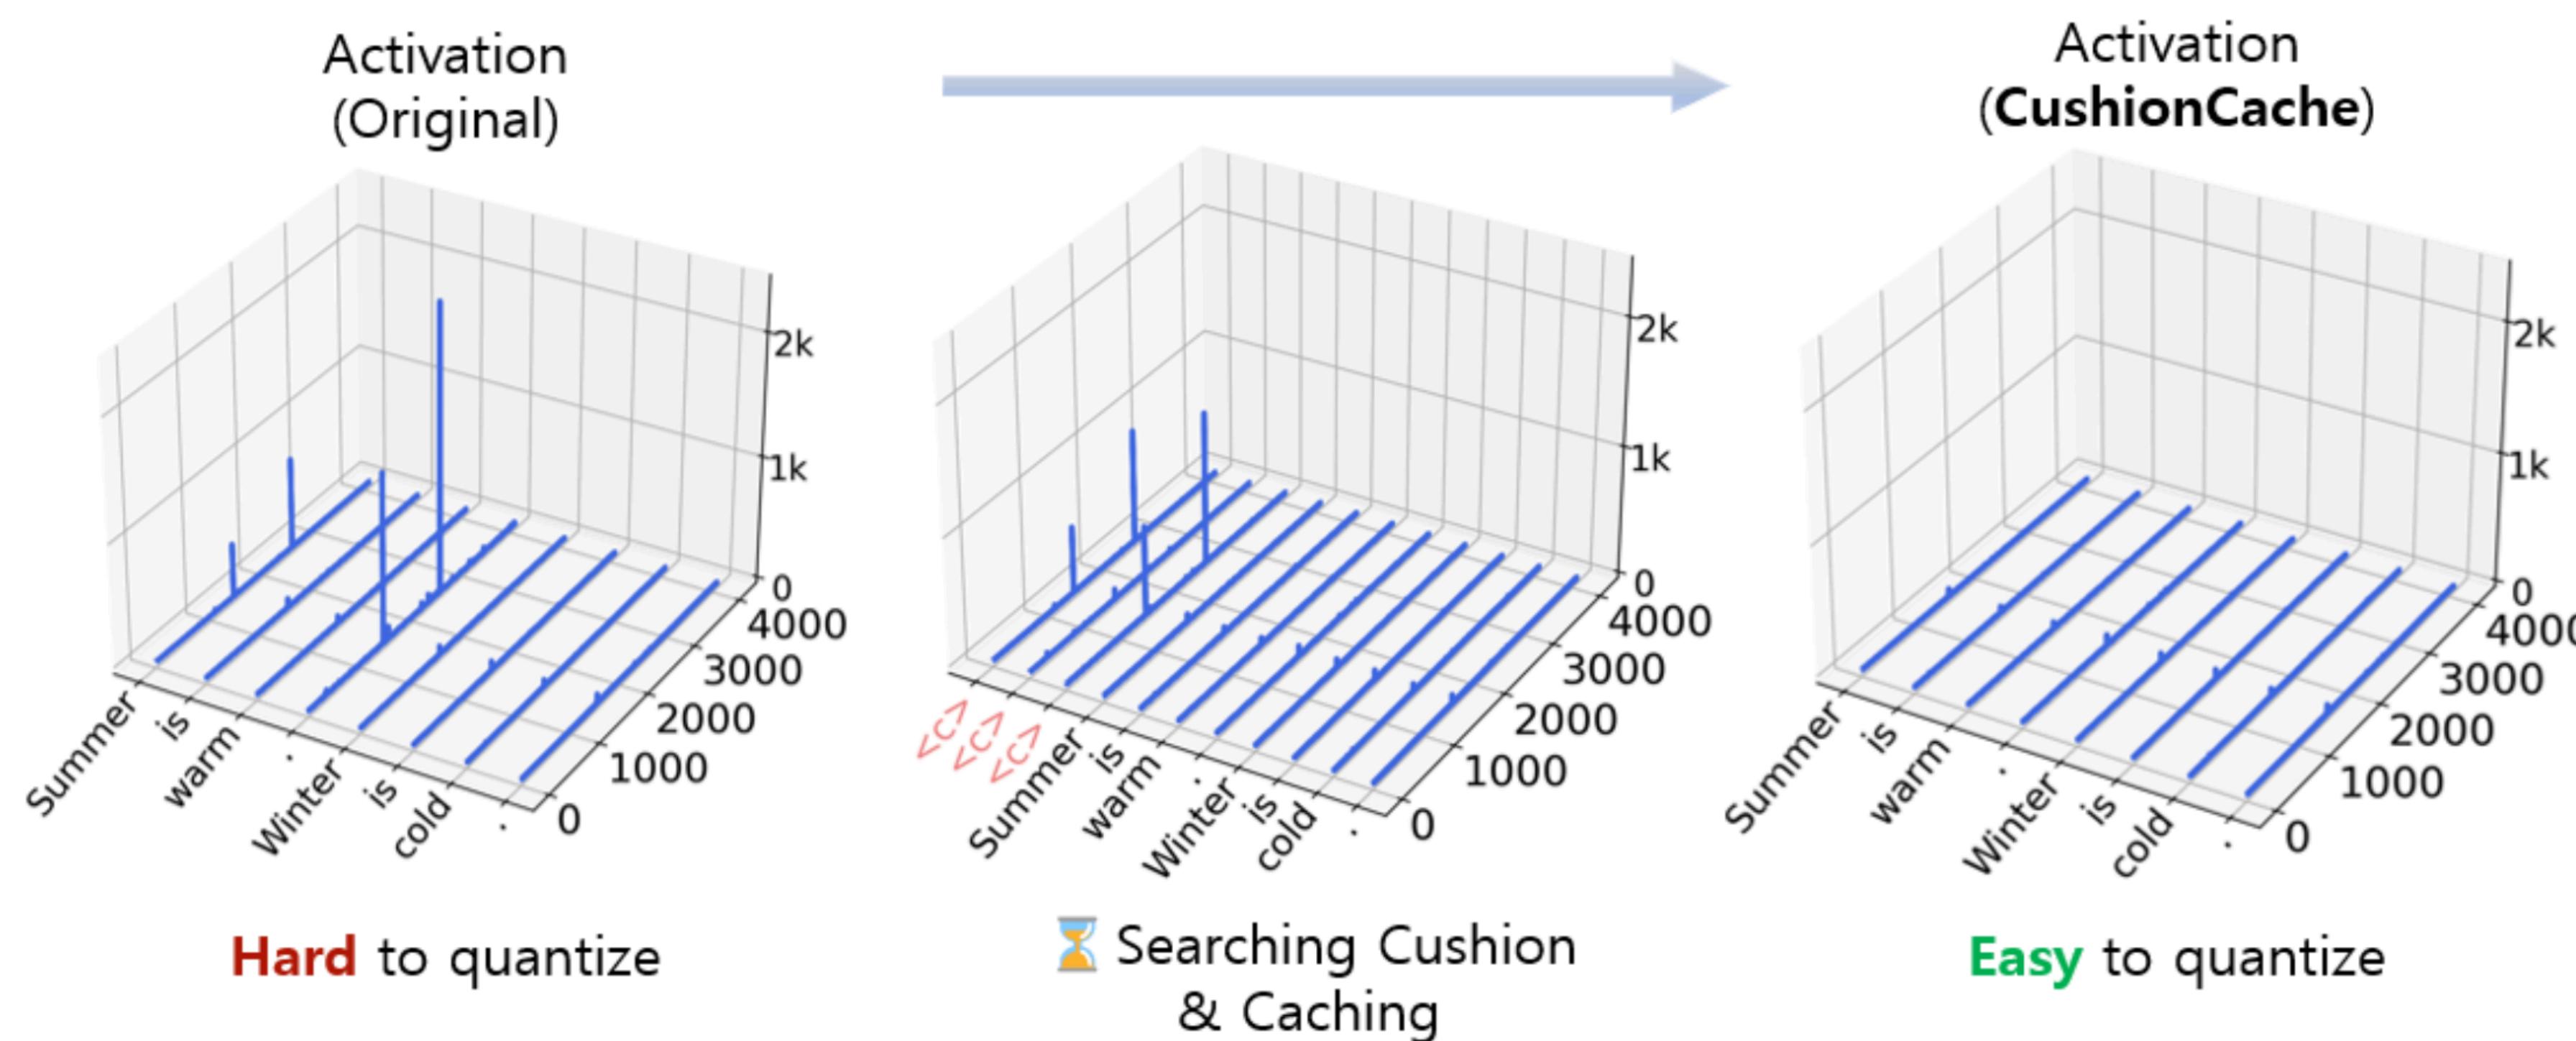
Rotation

- To get the rotation matrix:
 - Random Hadamard rotation
 - Learned rotation (e.g., SpinQuant, via Cayley SGD)
- Often represented as blockwise rotation + permutation



Prefix Tuning

- **Idea.** Add some special “attention sink” tokens as the prompt token
 - These suck up all attention, and mitigates outliers in later tokens
 - Apply further fine-tuning



Wrapping up

- LLM compression = Model compression + LLM-specific constraints
 - Much focus on practicality

That's it for today

