

Bits of Vision: Generative Modeling - 2

Recap


- **Last class.** Generative model for images
 - VAE (Variational Autonencoder)
 - GAN (Generative Adversarial Net)
- **Today.** Diffusion model

Recap: VAE

- In VAE, the decoder $p_{\theta}(\mathbf{x} | \mathbf{z})$ generates samples from a random latent code $\mathbf{z} \sim \mathcal{N}(0, I_k)$, such that


$$p_{\text{data}}(\mathbf{x}) \approx p_{\theta}(\mathbf{x})$$

- **Problem.** For training such a model, we need a good **inverse map**



Recap: VAE

- The solution was to **jointly train** an encoder which generates Gaussian from inputs
- **Problem.** As the “distribution of images” is **extremely complicated**, a single forward of neural network may not have a sufficient capacity to do this...



Diffusion model

- **Observation.** We can also generate Gaussian-like latent codes from the input (i.e., encode), via **gradually adding Gaussian noise** to the input


- i.e., sample the data \mathbf{x}_t from the distribution

$$q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N} \left(\mathbf{x}_t | \sqrt{\alpha_t} \mathbf{x}_{t-1}, (1 - \alpha_t)I \right)$$

- i.e., mix with the Gaussian noise:

$$\mathbf{x} \mapsto \sqrt{\alpha_t} \mathbf{x} + \sqrt{1 - \alpha_t} \epsilon, \quad \epsilon \sim \mathcal{N}(0, I)$$


(scaling factors for preserving the ℓ_2 norm)



Diffusion model

- Use this **noise-adding** as our probabilistic encoder!
 - How do we train the corresponding decoder?
- **Idea.** Train a **step-by-step model** $p_{\theta}(\mathbf{x}_{t-1} \mid \mathbf{x}_t)$ which approximates the posterior of the noise addition, i.e., $q(\mathbf{x}_{t-1} \mid \mathbf{x}_t)$
 - Parameterized as Gaussian, with trainable mean & variance

$$p_{\theta}(\mathbf{x}_{t-1} \mid \mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t-1} \mid \mu_{\theta,t}(\mathbf{x}_t), \Sigma_{\theta,t}(\mathbf{x}_t))$$



Training

- Draw a sample sequence $\mathbf{x}_0, \dots, \mathbf{x}_T$, using the **forward diffusion**

$$q(\mathbf{x}_{0:T}) = q(\mathbf{x}_0) \prod_{t=1}^T q(\mathbf{x}_t \mid \mathbf{x}_{t-1})$$

- Then, maximize the log probability of generating the real image

$$\mathbb{E}_{q(\mathbf{x}_0)} [\log p_\theta(\mathbf{x}_0)]$$

where the **reverse diffusion process** is given as

$$p_\theta(\mathbf{x}_{0:T}) = p_\theta(\mathbf{x}_T) \prod_{t=1}^T p_\theta(\mathbf{x}_{t-1} \mid \mathbf{x}_t)$$

Training

- To evaluate the log probability, we use ELBO (math alert!)
 - Use Jensen's inequality to get:

$$\begin{aligned}\mathbb{E}_{q(\mathbf{x}_0)} [\log p_\theta(\mathbf{x}_0)] &= \mathbb{E}_{q(\mathbf{x}_0)} \left[\log \left(\int p_\theta(\mathbf{x}_{0:T}) d\mathbf{x}_{1:T} \right) \right] \\ &= \mathbb{E}_{q(\mathbf{x}_0)} \left[\log \left(\int q(\mathbf{x}_{1:T} | \mathbf{x}_0) \frac{p_\theta(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T} | \mathbf{x}_0)} d\mathbf{x}_{1:T} \right) \right] \\ &= \mathbb{E}_{q(\mathbf{x}_0)} \left[\log \mathbb{E}_{q(\mathbf{x}_{1:T} | \mathbf{x}_0)} \left[\frac{p_\theta(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T} | \mathbf{x}_0)} \right] \right] \\ &\geq \mathbb{E}_{q(\mathbf{x}_{0:T})} \left[\log \frac{p_\theta(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T} | \mathbf{x}_0)} \right]\end{aligned}$$

Training

- Decomposing further, we get

$$\begin{aligned}\mathbb{E}_{q(\mathbf{x}_{0:T})} \left[\log \frac{p_\theta(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T} | \mathbf{x}_0)} \right] &= \mathbb{E}_q \left[\log \frac{p_\theta(\mathbf{x}_T) \prod_{t=1}^T p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)}{\prod_{t=1}^T q(\mathbf{x}_t | \mathbf{x}_{t-1})} \right] \\ &= \mathbb{E}_q \left[\log p_\theta(\mathbf{x}_T) + \sum_{t=1}^T \log \frac{p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)}{q(\mathbf{x}_t | \mathbf{x}_{t-1})} \right] \\ &= \mathbb{E}_q \left[\log p_\theta(\mathbf{x}_T) + \sum_{t=2}^T \log \frac{p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)}{q(\mathbf{x}_t | \mathbf{x}_{t-1})} + \log \frac{p_\theta(\mathbf{x}_0 | \mathbf{x}_1)}{q(\mathbf{x}_1 | \mathbf{x}_0)} \right]\end{aligned}$$

- Step-by-step operations, for each $\mathbf{x}_0, \dots, \mathbf{x}_T$

Training

- We apply an additional conditioning

$$\mathbb{E}_q \left[\log p_\theta(\mathbf{x}_T) + \sum_{t=2}^T \log \frac{p_\theta(\mathbf{x}_{t-1} \mid \mathbf{x}_t)}{q(\mathbf{x}_t \mid \mathbf{x}_{t-1})} + \log \frac{p_\theta(\mathbf{x}_0 \mid \mathbf{x}_1)}{q(\mathbf{x}_1 \mid \mathbf{x}_0)} \right]$$

$$= \mathbb{E}_q \left[\log p_\theta(\mathbf{x}_T) + \sum_{t=2}^T \log \left(\frac{p_\theta(\mathbf{x}_{t-1} \mid \mathbf{x}_t)}{q(\mathbf{x}_{t-1} \mid \mathbf{x}_t, \mathbf{x}_0)} \cdot \frac{q(\mathbf{x}_{t-1} \mid \mathbf{x}_0)}{q(\mathbf{x}_t \mid \mathbf{x}_0)} \right) + \log \frac{p_\theta(\mathbf{x}_0 \mid \mathbf{x}_1)}{q(\mathbf{x}_1 \mid \mathbf{x}_0)} \right]$$

$$q(\mathbf{x}_t \mid \mathbf{x}_{t-1}) = q(\mathbf{x}_t \mid \mathbf{x}_{t-1}, \mathbf{x}_0)$$

$$= \frac{q(\mathbf{x}_t, \mathbf{x}_{t-1} \mid \mathbf{x}_0)}{q(\mathbf{x}_{t-1} \mid \mathbf{x}_0)}$$

$$= \frac{q(\mathbf{x}_{t-1} \mid \mathbf{x}_t, \mathbf{x}_0)q(\mathbf{x}_t \mid \mathbf{x}_0)}{q(\mathbf{x}_{t-1} \mid \mathbf{x}_0)}$$

Training

- Then proceed with decomposition

$$\begin{aligned} & \mathbb{E}_q \left[\log p_\theta(\mathbf{x}_T) + \sum_{t=2}^T \log \left(\frac{p_\theta(\mathbf{x}_{t-1} \mid \mathbf{x}_t)}{q(\mathbf{x}_{t-1} \mid \mathbf{x}_t, \mathbf{x}_0)} \cdot \frac{q(\mathbf{x}_{t-1} \mid \mathbf{x}_0)}{q(\mathbf{x}_t \mid \mathbf{x}_0)} \right) + \log \frac{p_\theta(\mathbf{x}_0 \mid \mathbf{x}_1)}{q(\mathbf{x}_1 \mid \mathbf{x}_0)} \right] \\ &= \mathbb{E}_q \left[\log p_\theta(\mathbf{x}_T) + \sum_{t=2}^T \log \frac{p_\theta(\mathbf{x}_{t-1} \mid \mathbf{x}_t)}{q(\mathbf{x}_{t-1} \mid \mathbf{x}_t, \mathbf{x}_0)} + \sum_{t=2}^T \log \frac{q(\mathbf{x}_{t-1} \mid \mathbf{x}_0)}{q(\mathbf{x}_t \mid \mathbf{x}_0)} + \log \frac{p_\theta(\mathbf{x}_0 \mid \mathbf{x}_1)}{q(\mathbf{x}_1 \mid \mathbf{x}_0)} \right] \\ &= \mathbb{E}_q \left[\log \frac{p_\theta(\mathbf{x}_T)}{q(\mathbf{x}_T \mid \mathbf{x}_0)} + \sum_{t=2}^T \log \frac{p_\theta(\mathbf{x}_{t-1} \mid \mathbf{x}_t)}{q(\mathbf{x}_{t-1} \mid \mathbf{x}_t, \mathbf{x}_0)} + \log p_\theta(\mathbf{x}_0 \mid \mathbf{x}_1) \right] \\ &= \mathbb{E}_q[\log p_\theta(\mathbf{x}_0 \mid \mathbf{x}_1)] - \mathbb{E}_q D\left(q(\mathbf{x}_T \mid \mathbf{x}_0) \middle\| p(\mathbf{x}_T)\right) - \sum_{t=2}^T \mathbb{E}_q D\left(q(\mathbf{x}_{t-1} \mid \mathbf{x}_t, \mathbf{x}_0) \middle\| p_\theta(\mathbf{x}_{t-1} \mid \mathbf{x}_t)\right) \end{aligned}$$

Training

$$\mathbb{E}_q[\log p_\theta(\mathbf{x}_0 \mid \mathbf{x}_1)] - \mathbb{E}_q D\left(q(\mathbf{x}_T \mid \mathbf{x}_0) \middle\| p(\mathbf{x}_T)\right) - \sum_{t=2}^T \mathbb{E}_q D\left(q(\mathbf{x}_{t-1} \mid \mathbf{x}_t, \mathbf{x}_0) \middle\| p_\theta(\mathbf{x}_{t-1} \mid \mathbf{x}_t)\right)$$

- **First term.** We know that this is the squared loss of the mean predictor

- Assuming that $\Sigma = I$ for simplicity, we have:

$$\mathbb{E}_q[\log p_\theta(\mathbf{x}_0 \mid \mathbf{x}_1)] = -\frac{1}{2} \mathbb{E}_q \|\mathbf{x}_0 - \mu_{\theta,1}(\mathbf{x}_1)\|^2$$

- **Second term.** Does not have any learnable parameter
 - Thus, ignore

Training

$$-\frac{1}{2} \mathbb{E}_q \|\mathbf{x}_0 - \mu_{\theta,1}(\mathbf{x}_1)\|^2 - \sum_{t=2}^T \mathbb{E}_q D\left(q(\mathbf{x}_{t-1} \mid \mathbf{x}_t, \mathbf{x}_0) \middle\| p_{\theta}(\mathbf{x}_{t-1} \mid \mathbf{x}_t)\right)$$

- **Third term.** First, we look at the LHS of the KL divergence

- If we have the relationships $\# \bar{\alpha}_i := \alpha_1 \cdot \alpha_2 \cdot \dots \cdot \alpha_i$

$$\mathbf{x}_{t-1} = \sqrt{\bar{\alpha}_{t-1}} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_{t-1}} \epsilon, \quad \mathbf{x}_t = \sqrt{\alpha_t} \mathbf{x}_{t-1} + \sqrt{1 - \alpha_t} \epsilon'$$

- Then the following holds (exercise; use Bayes' theorem)

$$q(\mathbf{x}_{t-1} \mid \mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}\left(\frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \mathbf{x}_t + \frac{\sqrt{\bar{\alpha}_{t-1}}(1 - \alpha_t)}{1 - \bar{\alpha}_t} \mathbf{x}_0, \frac{(1 - \alpha_t)(1 - \sqrt{\bar{\alpha}_{t-1}})}{1 - \bar{\alpha}_t} I\right)$$

Training

$$-\frac{1}{2} \mathbb{E}_q \|\mathbf{x}_0 - \mu_{\theta,1}(\mathbf{x}_1)\|^2 - \sum_{t=2}^T \mathbb{E}_q D\left(q(\mathbf{x}_{t-1} \mid \mathbf{x}_t, \mathbf{x}_0) \middle\| p_{\theta}(\mathbf{x}_{t-1} \mid \mathbf{x}_t)\right)$$

- Now, the KL divergence between Gaussians can be written simply as

$$D\left(\mathcal{N}(\mu_1, \sigma_1^2 I) \middle\| \mathcal{N}(\mu_2, \sigma_2^2 I)\right) = \log \frac{\sigma_2}{\sigma_1} - \frac{d}{2} + \frac{d\sigma_1^2 + \|\mu_1 - \mu_2\|^2}{2\sigma_2^2}$$

- Plug this in, to get the loss (ignoring the variance terms)

$$\sum_{i=2}^T \left\| \mu_{\theta,t}(\mathbf{x}_t) - \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \mathbf{x}_t + \frac{\sqrt{\bar{\alpha}_{t-1}}(1 - \alpha_t)}{1 - \bar{\alpha}_t} \mathbf{x}_0 \right\|^2 =: \sum_{i=1}^T \|\mu_{\theta,t}(\mathbf{x}_t) - \mu_q(\mathbf{x}_t, \mathbf{x}_0)\|^2$$

In a nutshell

- Training the reverse diffusion process is simply:
 - Sample an image \mathbf{x}_0 from the dataset
 - Using $q(\cdot)$, sample $\mathbf{x}_1, \dots, \mathbf{x}_T$.
 - Pick a time t :
 - Train $\mu_{\theta,t}(\cdot)$ to minimize $\|\mu_{\theta,t}(\mathbf{x}_t) - \mu_q(\mathbf{x}_t, \mathbf{x}_0)\|^2$
 - Repeat

In a nutshell

- This is typically reparametrized as a **noise prediction**
 - i.e., predict the residual of the prediction
 - Recent works suggest that this reparametrization may be flawed...
<https://arxiv.org/abs/2511.13720>

Algorithm 1 Training

- 1: **repeat**
- 2: $\mathbf{x}_0 \sim q(\mathbf{x}_0)$
- 3: $t \sim \text{Uniform}(\{1, \dots, T\})$
- 4: $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 5: Take gradient descent step on
$$\nabla_{\theta} \|\boldsymbol{\epsilon} - \boldsymbol{\epsilon}_{\theta}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}, t)\|^2$$
- 6: **until** converged

In a nutshell


- The inference is done by starting from a Gaussian distribution
 - Then, keep denoising

Algorithm 2 Sampling


- 1: $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 2: **for** $t = T, \dots, 1$ **do**
- 3: $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ if $t > 1$, else $\mathbf{z} = \mathbf{0}$
- 4: $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t}} \boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$
- 5: **end for**
- 6: **return** \mathbf{x}_0

Latent diffusion

- We use diffusion in some latent space
 - Combine with the ideas of VAE
 - Plus, we do some conditioning



Pros & Cons



More references

- **Beginner**
 - <https://huggingface.co/blog/annotated-diffusion>
 - <https://lilianweng.github.io/posts/2021-07-11-diffusion-models/>
- **Advanced**
 - <https://arxiv.org/abs/2403.18103>
 - <https://arxiv.org/abs/2510.21890> <- highly recommended

</lecture 18>