

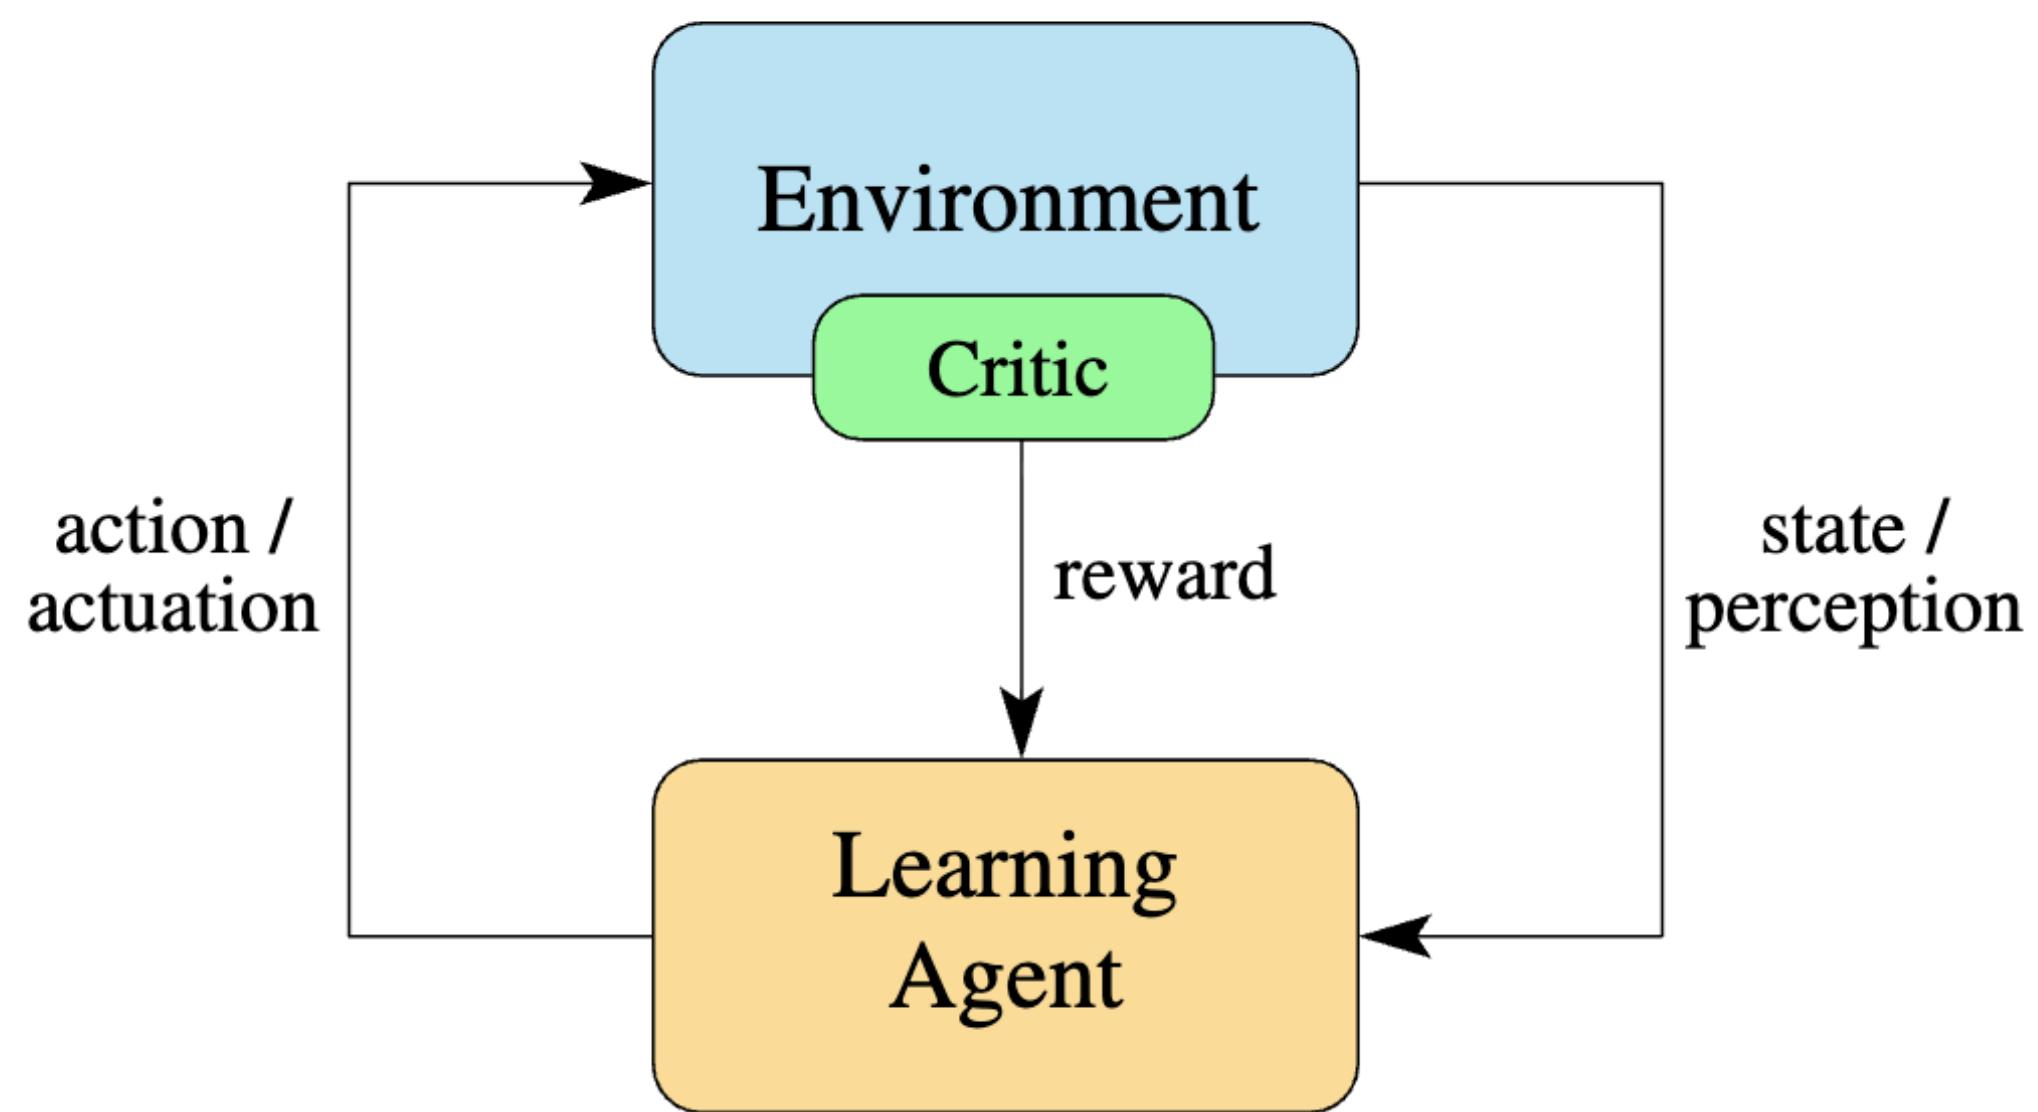
Reinforcement Learning - 3

Recap

- **Last class.**
 - Markov Process + Reward + Decision
 - Evaluating the policy
 - Value function & Q-function
 - Assumption. Full knowledge of MDP:
 - the state transition $p(s' | s, a)$
 - the reward $r(s, a)$
- **Today.**
 - Evaluating the policy without full knowledge

Recap: Reinforcement learning

- We were interested in a **sequential decision-making** scenario
 - State transitions as some Markov process $s_{t+1} \sim p(\cdot | s_t, a_t)$
 - Agent makes actions according to some policy $a_t \sim \pi(\cdot | s_t)$
 - Agent collects the reward $r_t \sim r(\cdot | s_t, a_t)$



```
for  $t = 1, \dots, n$  do
    The agent perceives state  $s_t$ 
    The agent performs action  $a_t$ 
    The environment evolves to  $s_{t+1}$ 
    The agent receives reward  $r_t$ 
end for
```

Recap: Evaluation of policy

- We were interested in the **return** – i.e., accumulated discounted reward

$$G_t = r_t + \gamma r_{t+1} + \cdots + \gamma^{H-1} r_{t+H-1}$$

- We discussed how we can evaluate the policy π by computing:
 - **Value function.** Expected return, beginning at some state

$$V^\pi(s) = \mathbb{E}[r_t + \gamma r_{t+1} + \cdots + \gamma^{H-1} r_{t+H-1} \mid s_t = s]$$

- **Q-function.** Expected return from making an action at some state, assuming that we'll continue using the policy π afterwards

$$Q^\pi(s, a) = \mathbb{E}[G_t \mid r_t + \gamma r_{t+1} + \cdots + \gamma^{H-1} r_{t+H-1} \mid s_t = s, a_t = a]$$

Recap: Evaluation of policy

- We can use **Bellman equation** to simplify this for infinite horizon

- **Value function**

$$V^\pi(s) = \sum_{a \in A} \pi(a | s) \left(R(a, s) + \gamma \sum_{s' \in S} p(s' | s, a) V^\pi(s') \right)$$

- **Q-function**

$$Q^\pi(s, a) = R(s, a) + \gamma \sum_{s' \in S} p(s' | s, a) V^\pi(s')$$

- Can be computed by iterative method

- Requires knowing $R(s, a)$ and $p(s' | s, a)$

Evaluation through experience

- Today, we assume that we don't know $R(s, a)$ or $p(s' | s, a)$
- Instead, we assume that we can **collect data from environment**
 - Deploy an agent that uses $\pi(a | s)$, and collect
$$e = (s_1, a_1, r_1, s_2, a_2, r_2 \dots)$$
 - Collect multiple **episodes** e_1, e_2, \dots, e_n
- **Goal.** Estimate $V^\pi(s)$ from e_1, \dots, e_n

Monte Carlo policy evaluation

- **Idea.** Simply measure the **mean return**, starting from each state
 - Look at each sub-trajectory
 - At each starting state s_t , compute the G_t
 - Average over all sub-trajectories with the same starting state
- There are two well-known variants
 - First-visit
 - Every-visit

First-visit Monte Carlo

- Initialize $N(s) = 0, G(s) = 0 \quad \forall s \in S$
- Loop:
 - Sample episode $i = s_{i,1}, a_{i,1}, r_{i,1}, s_{i,2}, a_{i,2}, r_{i,2}, \dots, s_{i,T_i}, a_{i,T_i}, r_{i,T_i}$
 - Define $G_{i,t} = r_{i,t} + \gamma r_{i,t+1} + \gamma^2 r_{i,t+2} + \dots + \gamma^{T_i-1} r_{i,T_i}$ as the return from time step t onwards in i -th episode
 - For each time step t until T_i – If this is the **first time t** that state s is visited in episode i
 - Increment the counter $N(s) \leftarrow N(s) + 1$
 - Increment the total return $G(s) \leftarrow G(s) + G_{i,t}$
 - Update the estimate $\hat{V}^\pi(s) = G(s)/N(s)$

Every-visit Monte Carlo

- Initialize $N(s) = 0, G(s) = 0 \quad \forall s \in S$
- Loop:
 - Sample episode $i = s_{i,1}, a_{i,1}, r_{i,1}, s_{i,2}, a_{i,2}, r_{i,2}, \dots, s_{i,T_i}, a_{i,T_i}, r_{i,T_i}$
 - Define $G_{i,t} = r_{i,t} + \gamma r_{i,t+1} + \gamma^2 r_{i,t+2} + \dots + \gamma^{T_i-1} r_{i,T_i}$ as the return from time step t onwards in i -th episode
 - For each time step t until T_i :
 - Increment the counter $N(s) \leftarrow N(s) + 1$
 - Increment the total return $G(s) \leftarrow G(s) + G_{i,t}$
 - Update the estimate $\hat{V}^\pi(s) = G(s)/N(s)$

Every-visit Monte Carlo

- Note that we can think of an equivalent, **sequential-update** version
 - Natural, as we usually collect data in sequence rather than in batch
- For state s visited at time step t in episode i ,
 - Increment the counter $N(s) \leftarrow N(s) + 1$
 - Update the estimate:
$$\hat{V}^\pi(s) \leftarrow \hat{V}^\pi(s) + \frac{1}{N(s)}(G_{i,t} - \hat{V}^\pi(s))$$
- Can be viewed as iterative updates
 - $1/N(s)$ is some **learning rate**

Properties

- **First-visit**

- $\hat{V}^\pi(s)$ is unbiased: $\mathbb{E}[\hat{V}^\pi(s)] = V^\pi(s)$
- $\hat{V}^\pi(s)$ is consistent: $\lim_{n \rightarrow \infty} \hat{V}^\pi(s) = V^\pi(s)$

- **Every-visit**

- $\hat{V}^\pi(s)$ is biased: $\mathbb{E}[\hat{V}^\pi(s)] \neq V^\pi(s)$
- $\hat{V}^\pi(s)$ is consistent: $\lim_{n \rightarrow \infty} \hat{V}^\pi(s) = V^\pi(s)$
- Often has a smaller MSE:

$$\mathbb{E}[\|\hat{V}_{\text{ev}}^\pi - V^\pi\|^2] \leq \mathbb{E}[\|\hat{V}_{\text{fv}}^\pi - V^\pi\|^2]$$

Properties

- Both first-visit and every-visit Monte Carlo estimate tend to have a large variance, i.e.,

$$\mathbb{E}[\|\hat{V}_{\text{ev}}^{\pi} - V^{\pi}\|^2] \gg 0$$

- Requires many samples for convergence
 - In RL, the data collection is usually very expensive
 - Example. Robot Learning
 - Example. Drug discovery
 - Requires the episode to terminate before updating

Temporal Difference Learning

- **Idea.** Use **recursive estimates**, just like Bellman's equation

$$V^\pi(s) = \sum_{a \in A} \pi(a | s) \left(R(a, s) + \gamma \sum_{s' \in S} p(s' | s, a) V^\pi(s') \right)$$

- We can use the current estimates of $\hat{V}^\pi(s_{t+1})$, instead of waiting for the episode to be fully unrolled
- **TD(0) Learning.** At every transition (s_t, a_t, r_t, s_{t+1}) , do:

$$\hat{V}^\pi(s) \leftarrow \hat{V}^\pi(s_t) + \alpha \left([r_t + \gamma \hat{V}^\pi(s_{t+1})] - \hat{V}^\pi(s_t) \right)$$

- α is some “learning rate”

Properties

- Biased, in general
- Usually has a smaller variance and is more sample-efficient
- Requires less storage, and updates online

Wrapping up

So far

- We have looked at various ML algorithms:
 - From small-scale to large-scale methods
- **Before mid-term: Classic ML**
 - Empirical Risk Minimization vs. Bayesian Approach
 - Supervised Learning
 - Linear models, SVM, Nonlinear methods (based on kernels)
 - Unsupervised Learning
 - Clustering, GMM, PCA

So far

- **After mid-term: Modern ML**
 - Deep Learning & Optimization techniques
 - SGD and Backpropagation
 - Vision
 - Convolutions, Self-supervised Learning VAE, GAN, Diffusion,
 - Language
 - Tokenization, Language models, Post-training, Multimodal AI
 - Reinforcement Learning
 - Markov Decision Process, Temporal Difference Learning

What is missing?

- We have focused on fundamental topics, mostly
 - Stayed away from **research frontiers**
- Here is a very brief look at some “emerging topics”
 - **Dynamic world model.** <https://deepmind.google/blog/genie-3-a-new-frontier-for-world-models/>
 - **Weather forecasting.** <https://blog.google/technology/google-deepmind/weathernext-2/>
- Common: Want to encode “physical understanding” to save computations and training data needed

</lecture 25>